Protocol v3.1: proofs appendix

Part A: cryptographic and DP guarantees · Part B: sensitivity and statistical guarantees
Part C: composition across runs and the hockey-stick ledger

Draft for review

This appendix collects the proofs for protocol v3.1 (honest-but-curious model). Part A (page 4) covers the encoding sensitivity, smudging, the worst-case error bound B_{tot}, output privacy (Theorem 1), input and cohort privacy, and correctness. Part B (page 40) covers the replace-one sensitivity, operator-norm concentration, the two-stage choice of (λ, Ψ), the bound on $\|\boldsymbol{\beta}\|$, coverage of both intervals, and the bootstrap RSS correction. Part C (page 66) proves composition across runs for an adaptive, possibly colluding coalition of researchers (the paper's Proposition `prop:compose`), with a self-contained Rényi filter, and the hockey-stick treatment of the smudging step that the ledger uses (three sums, the reserved share ξ_L of ε_L, and the rule for κ_L). Its result numbers are those the paper cites as "proofs, composition appendix". Sections 4, 7 and 9 of Part A and S7–S10 of Part B were added after a second reading of the companion book on the protocol, whose proofs of these results were checked step by step (with the corrections noted where they occur), and include new results (Part A, Lemma 7.5 and Proposition 7.7; Part B, Lemma S0(b), Lemma S8′, Propositions S10a and S10b, Theorem S10c). Each part labels every step as exact, first order, asymptotic, numerical or open, and lists its corrections to the specification.

Status at a glance

Claim	Status	Notes
Lemma 1: encoding sensitivity	proven	$\sqrt{2/N}$ and the $\sqrt{N}$ rounding term are correct
Lemma 2: smudging	proven	SD $\le \|e\|_2/(\sqrt{2\pi}\sigma_1) + 4\epsilon/(1-\epsilon)$ with $\sigma_1 = \sqrt{\sigma_c^2 - 2.1^2}$; the per-release form $2^{-\kappa}(1+2^{-80}) + 2^{-97}$ needs $\sigma_c \ge 2^{42}$ (A, Corollary 2.4, corrected from 2^{20}); the hockey-stick form needs no floor (C, Lemma 5.1)
Obligation O1: worst-case B_{tot}	proven (explicit)	$\approx 2^{25.3}$ for the circuit as specified; errors measured in OpenFHE runs stay 6–8 bits below it (no OpenFHE noise estimate is used); assumes (M5)
O1 for OpenFHE's gadget, decoupled scales	proven for model (M3′) (A, Lemma 4.1)	$\varsigma_\uparrow = 2\max_j k_j$, $\varsigma_\downarrow = 2k_P$; $B_{\text{tot},2} = 2^{26.59}$ at the evaluation setting; match of OpenFHE 1.5.1 with (M3′) by code reading; (M5) for OpenFHE's sampler numerical
Centre sensitivity, decoupled scales	proven (A, Lemma 4.3)	$\Delta_{x,2} = 2^{27.57}$, remainder $\epsilon_x = 2^{-36.5}$ covered by the "+1"; exact rounding in Ecd assumed
Decoding offset	proven (A, Proposition 4.5)	$\bar{d}_2$ between 10^{-8} and $5 \cdot 10^{-6}$; coefficient-domain $d_{\max} \le 2^{38.6}$

Claim	Status	Notes
zCDP of the ideal mechanism; tight conversion	proven	$\rho_{\text{run}} = 0.01496$ matches the specification
Theorem 1: output privacy	proven (hybrids)	one-sample RLWE for the re-randomizer; plain RLWE for O's public-key share, the only key material O sends (circular security only if its rotation-key shares were sent); O1; R's errors honestly sampled; slack γ statistical
Composition across runs	proven (C, Theorem 4.2)	any PPT coalition under the filter; runs fixed in advance under rule (N); the ledger charges a run's full cost at admission; RLWE for a finite menu of parameter sets, against non-uniform adversaries; rule (N) with adaptive requests not proved
Hockey-stick ledger and κ_L	proven (C, Theorem 5.2, Corollary 5.4, Proposition 5.5)	$\kappa_L = 22, 21, 19, 17, 15$ at $(10, 10^{-6}, 0.1, 10^{-3})$ for $\varepsilon_{\text{run}} = 0.5, \ldots, 8$
G1, G2: input and cohort privacy	proven	simulators, up to the cuckoo-hashing failure probability, now bounded (A, Lemma 7.5: $2^{-49.1}$ at $n_R^{\max} = 2{,}000$, $\leq 2^{-40}$ up to load 0.639; needs an insertion that finds an assignment whenever one exists) and R's observable aborts
G4: disclosure from the Link phase	proven, ROM and CDH (A, Theorem 7.3)	first layer, without circuit privacy; second layer (SanitizeBFV) proven for the textbook BFV model (A, Proposition 7.7: evaluation noise $\leq 2^{-31}$ after the switch); OpenFHE's BFVrns not checked against the model
Soundness of Calibrate; standard simulation with output; O's work	proven (A, Section 9)	the run computes its DP functionality against R when the functionality outputs the centred residue modulo q_0; $\Omega(n_{\text{Sc}})$ reads in a cell-probe model
Correctness (no wraparound)	proven	failure probability $\leq \delta_B + 2Ne^{-50}$ per release; no condition at the earlier levels (A, Proposition 8.1)
S1: replace-one sensitivity	proven, exact for $d_O \geq 2$	triu and Frobenius suprema are equal; regime 3 reduces to three cubics; $d_O = 1$ is valid but conservative
When $\mathbf{y}^\top \mathbf{y}$ is free	partly proven	exact criterion; free if $w \leq 2(r + o)$, never free if $w > 3(r + o)$; boundary in between only numerical
S2: operator-norm radius	proven (one-sided)	r_1 valid for the event Algorithm 6 uses; about $1.55\times$ conservative
S7: exact radius	proven	exact level with fresh draws; transfer to the lattice noise with the TV constant now proved (S0)
S8: switch-off	bracket proven; size of the bias partly	bracket tight among Gram matrices with a given $\lambda_{\min}$; spectral bound on the bias by Sudakov–Fernique; empirical factors 1.2/1.4 not proved
S9: mechanism bias	proven, explicit constants	leading term, ridge, decay in λ
S10: FWS $\hat{s}^2$; clipping; full-record Δ	proven	plug-in RSS biased by $-\sigma^2(\operatorname{tr}\mathbf{P} + \hat{\beta}^\top M_\chi \hat{\beta})$; interval valid under binding clipping with a public variance bound; closed form for blocks of dimension ≥ 2
S3: two-stage (λ, Ψ)	proven	independent of the stage-2 noise; closed form for λ_O
S4: bound on $\|\boldsymbol{\beta}\|$	proven	holds for the projection coefficient with no model assumed
S5: coverage	partly first order	Wald remainder bounded exactly; plug-in variance first order; normality of the sampling term asymptotic
S6: bootstrap RSS correction	proven to first order	removes the $O(\sigma^2)$ bias up to $O(\sigma^4)$; constants not explicit

Corrections to the specification

Both parts found claims in `protocol_v31.tex` that need to change. The most consequential:

1. **The CKKS scale is undersized.** Flooding the whole error in ℓ_2 needs $\log_2 \Omega \approx 61.5$ for the circuit as specified, not 52.5. A restructured circuit brings this down to ≈ 57.2; adding encoding at scale 2^{60} by R brings it to ≈ 53.7. Calibrate must use the data-independent worst-case bound, never OpenFHE's noise estimate.
2. **The centre of the mechanism.** The ideal mechanism must be centred at the noiseless circuit output, not at $\mathsf{Ecd}(\mathbf{z})$. Its sensitivity rises by a relative 2^{-33}. Ω must be the exact final scale of the circuit.
3. **δ accounting.** The per-release statistical distance is $2^{-\kappa}(1 + 2^{-80}) + 2^{-97}$, for $\sigma_c \geq \max\{2^{\kappa} B_{\mathrm{tot}}/\sqrt{2\pi}, 2^{42}\}$ (the floor 2^{20} stated earlier is too small for the factor $1 + 2^{-80}$). For a single (ε, δ) statement, the slack enters with a factor $(1 + e^{\varepsilon})$. Across runs, the ledger charges a run's full cost at admission and accounts the smudging step in hockey-stick divergence (Part C).
4. **Assumptions.** RLWE is needed in two forms. Disclosure G4 needs circuit privacy of O's BFV replies; the modulus switch alone is heuristic. G1 and G2 need explicit assumptions and the hashing slack.
5. **Algorithm 6.**
   - The radius comment should be one-sided, or use $\ln(2/\alpha_L)$.
   - U_{yy} spends α_U twice, and σ_{yy} is undefined.
   - $\boldsymbol{\beta}$ must be defined as a working-model or projection coefficient on the clipped data, and the homoskedastic variance fails when clipping binds.
   - The bootstrap needs truncation, since the moments do not exist otherwise.
   - The bias bound $\bar{b}_j$ is first order only.
   - The private $\|\boldsymbol{\beta}\|$ bound is usually infinite while the ridge is active.
   - The specification should state whether coverage is conditional on the gate.
6. **Smaller points.**
   - The sampler needs a rational σ_c^2 and a defined notion of constant time.
   - The modulus switch to q_0 is a no-op after two rescales.
   - The printed $\Delta = 86.18$ comes from unrounded bounds.
   - The earlier paper's "$\mathbf{y}^\top \mathbf{y}$ is free" and "Frobenius is conservative" are false in general.

Open items

- Full proofs for the two improved circuit variants; a concrete constant-time exact sampler. (Adaptive composition across runs is proved in Part C; rule (N) with adaptively chosen requests is not. No wraparound needs no condition at the earlier levels; the chain is sized by Calibrate, which is proved sound.) Whether OpenFHE's BFVrns satisfies the BFV model of the B_ψ bound. (The cuckoo insertion now falls back to augmenting paths, so Lemma 7.5 of Part A covers it.)
- The free-set boundary for $\mathbf{y}^\top \mathbf{y}$; a closed form for Δ with $\mathbf{y}^\top \mathbf{y}$ when $d_O = 1$; explicit constants for S6; a guarantee for the data-driven variance when clipping binds (a public bound gives a valid interval, S10); a spectrum-sharp bound on the downward bias of $\lambda_{\min}$ (S8).
- Sixteen citation details are flagged in the two parts for verification, among them theorem numbers in Canonne et al. (2020), Vershynin, and Boucheron, Lugosi and Massart.

Protocol v3.1: proofs of the cryptographic and differential-privacy claims

Technical appendix: encoding sensitivity, smudging, a worst-case error bound, output privacy (SIM-CDP), input and cohort privacy, correctness

Draft for review

Abstract

This appendix proves, corrects, or marks as open each cryptographic and differential-privacy claim of the specification *Protocol v3.1* (`protocol_v31.tex`). We work in the honest-but-curious model of that specification and reuse its notation. Section 1 proves the encoding-sensitivity lemma. Section 2 proves the smudging lemma for the exact discrete Gaussian, with explicit constants. Section 3 discharges Obligation O1: an explicit bound B_{tot} on the ℓ_2 norm of the whole error ring element, worst case over the clipped data domain, with probability only over keys and fresh randomness. Section 4 extends O1 to decoupled encoding scales and to OpenFHE's uncentred key switching, proves the sensitivity of the centre with decoupled scales, and bounds the decoding offset. Section 5 proves output privacy (Theorem 1) by hybrids. Section 6 proves input and cohort privacy (G1, G2). Section 7 proves the first layer of the disclosure claim G4 (random-oracle model, CDH), bounds the cuckoo-hashing slack, and bounds the reply noise of the PSI in a model of BFV. Section 8 treats correctness, at every level of the circuit. Section 9 proves that Calibrate meets the hypotheses of Theorem 1, that the run computes its DP functionality against R in the standard sense, and that O's work is linear in its scope. Section 10 lists the corrections to the specification.

Summary of results.

- **Lemma 1** holds with the constant $\sqrt{2/N}$ for the stated convention. One precision: Ω must be the *exact* final scale Ω_{fin} of the circuit, not the nominal scale (Correction 2).
- **Lemma 2** holds with the constant $1/\sqrt{2\pi}$ up to an explicit additive smoothing term ($\leq 2^{-98}$ here). An elementary proof gives the constant $1/2$ with no side condition.
- **O1.** The mechanism must be centred at the noiseless circuit output rather than at $\mathsf{Ecd}(\mathbf{z})$ (Correction 1). With that centre, and with every rounding treated adversarially, the circuit exactly as specified gives $B_{\text{tot}} \approx 2^{25.3}$ for stage 2 of the example, with $\delta_B = 2^{-50}$. This needs $\log_2 \Omega \approx 61.5$, not 52.5. The specification's 52.5 floods only the fresh-encryption error, in mean. With the specification's centre, the deterministic plaintext-arithmetic terms alone reach $2^{30.7}$ to $2^{34.5}$, which would need $\log_2 \Omega \approx 67$ to 71. Changes to the circuit bring B_{tot} down to $2^{17.6}$, so that $\log_2 \Omega \approx 53.7$ suffices (Section 3, Table 1):
  - evaluate EvalSum before the first rescale;
  - use stochastic rounding and a single final rescale;
  - make the special modulus larger;
  - give R's encoding a larger scale.
- **Theorem 1** holds under decision-RLWE for two secret distributions, a circular-security assumption for O's rotation-key shares, and O1, with the corrected centring. It holds up to an

explicit statistical slack γ. PSI security is *not* needed for G3 under the specification's adjacency. The (ε, δ) accounting needs the factor $(1+e^{\varepsilon})$ on γ when it is read as indistinguishability-based CDP (Correction 5).

- **G1 and G2** hold under RLWE in s_R (with circular security for R's key material), the random-oracle model for the seed expansion, IND-CPA of the PSI's BFV, and obliviousness of the OPRF. G2 holds up to the cuckoo-hashing failure probability.
- **Correctness.** The q_0 condition in CALIBRATE is sufficient, with a correctness-failure probability of at most $\delta_B + 2Ne^{-50}$ per release, which the specification should state. The exact condition is derived below; no condition is needed at the earlier levels.
- **Added after a second reading of the companion book** (Sections 4, 7 and 9): O1 for OpenFHE's uncentred gadget under model (M3′), with $\varsigma_{\uparrow} = 2\max_j k_j$ and $\varsigma_{\downarrow} = 2k_P$ (Lemma 4.1; the match of OpenFHE 1.5.1 with (M3′) is a code reading); the centre sensitivity with decoupled scales (Lemma 4.3); a public bound on the decoding offset (Proposition 4.5); G4's first layer under CDH in the random-oracle model, without circuit privacy (Theorem 7.3); the soundness of CALIBRATE (Proposition 9.1); standard simulation of R's view jointly with its output (Corollary 9.3); and O's $\Omega(n_{\mathrm{Sc}})$ work (Proposition 9.5). New here: a proved bound on the cuckoo-hashing slack (Lemma 7.5) and a worst-case bound on the PSI's evaluation noise in the textbook BFV model (Proposition 7.7), which covers it by the "+1" of the implementation's B_{err} in that model.

Steps that are not fully rigorous are marked **[Not fully rigorous: ...]**. Citation details we could not verify are marked *[citation detail unverified: ...]*.

Contents

1 Conventions and Lemma 1 (encoding sensitivity) 3

2 Lemma 2 (smudging with the exact discrete Gaussian) 5

3 Obligation O1: a worst-case bound on the whole error 7
3.1 Probability tools . . . 8
3.2 The circuit, made explicit . . . 9
3.3 The noiseless centre and the error decomposition . . . 10
3.4 Term-by-term bounds, worst case over data . . . 11
3.5 Numerical evaluation . . . 13
3.6 Empirical sanity check with OpenFHE . . . 14

4 Decoupled scales, OpenFHE's key switching and the decoding offset 15
4.1 OpenFHE's uncentred key switching . . . 15
4.2 The sensitivity of the centre with decoupled scales . . . 17
4.3 The decoding offset . . . 18

5 Output privacy (Theorem 1) 19
5.1 Differential-privacy tools . . . 19
5.2 SIM-CDP . . . 20
5.3 Assumptions . . . 21
5.4 The theorem . . . 22

6 Input privacy (G1) and cohort privacy (G2) 24

7 Disclosure (G4), the cuckoo-hashing slack and the reply noise **25**
7.1 The BFV evaluation noise of the replies 27

8 Correctness: no wraparound, and the exact sampler **29**

9 Calibrate, the run as a secure computation, and O's work **31**
9.1 Soundness of CALIBRATE 31
9.2 The run computes the DP functionality against R 32
9.3 O's work is linear in its scope 32

10 Corrections to the specification **33**

1 Conventions and Lemma 1 (encoding sensitivity)

Ring and embeddings. Let $N \geq 8$ be a power of two, $\mathcal{R} = \mathbb{Z}[X]/(X^N+1)$, $\mathcal{R}_q = \mathcal{R}/q\mathcal{R}$, and $\zeta = e^{i\pi/N}$, a primitive $2N$-th root of unity. The roots of X^N+1 are ζ^j for $j \in \mathbb{Z}_{2N}^*$, the odd residues. Let $\mathcal{H} = \langle 5 \rangle \leq \mathbb{Z}_{2N}^*$. Then:

- $\mathcal{H}$ has order $N/2$ and $\mathcal{H} = \{g \in \mathbb{Z}_{2N} : g \equiv 1 \bmod 4\}$;
- $\mathbb{Z}_{2N}^* = \mathcal{H} \sqcup (-\mathcal{H})$;
- the slots are ordered by $h_k = 5^k \bmod 2N$, for $k = 0, \ldots, N/2-1$.

For $f \in \mathbb{R}[X]/(X^N+1)$, the *canonical embedding* is $\sigma(f) = (f(\zeta^j))_{j \in \mathbb{Z}_{2N}^*} \in \mathbb{C}^N$, and the *slot map* is $\tau(f) = (f(\zeta^{h_k}))_{k<N/2} \in \mathbb{C}^{N/2}$. Since f is real, $f(\zeta^{-h}) = \overline{f(\zeta^h)}$. So $\sigma(f)$ consists of $\tau(f)$ and its complex conjugate. The decoding of the specification is $\mathsf{Dcd}(v) = \tau(v)/\Omega$.

We write $\|f\| = \|f\|_2$, $\|f\|_\infty$ and $\|f\|_1$ for the norms of the coefficient vector, and $\|f\|_\infty^{\mathrm{can}} = \|\sigma(f)\|_\infty$. Norms of vectors in $\mathbb{C}^{N/2}$ are complex ℓ_2 norms. For $g \in \mathbb{Z}_{2N}^*$, ψ_g is the automorphism $X \mapsto X^g$.

Lemma 1.1 (Parseval for the canonical embedding)**.** *For every $f \in \mathbb{R}[X]/(X^N+1)$, $\|\sigma(f)\|_2^2 = N\|f\|_2^2$ and $\|\tau(f)\|_2^2 = \frac{N}{2}\|f\|_2^2$. The map τ is an $\mathbb{R}$-linear bijection onto $\mathbb{C}^{N/2}$.*

Proof. We have $\sigma(f) = Wf$ with $W_{j,k} = \zeta^{jk}$, for $j \in \mathbb{Z}_{2N}^*$ and $0 \leq k < N$.

Orthogonality. For $k \neq l$, put $d = l-k$, so that $0 < |d| < N$. Then

$$(W^*W)_{k,l} = \sum_{j \text{ odd}} \zeta^{jd} = \zeta^d \sum_{t=0}^{N-1} \zeta^{2td} = \zeta^d \frac{1-\zeta^{2Nd}}{1-\zeta^{2d}} = 0,$$

because $\zeta^{2Nd} = 1 \neq \zeta^{2d}$. Also $(W^*W)_{k,k} = N$, so $W^*W = N\mathbf{I}$.

Norms. The first identity follows from $W^*W = N\mathbf{I}$. The second follows because $\sigma(f)$ is $\tau(f)$ together with its conjugate.

Bijectivity. Let $y \in \mathbb{C}^N$ satisfy $y_{-j} = \overline{y_j}$. Then $f = N^{-1}W^*y$ has coefficients

$$f_k = N^{-1}\sum_j \zeta^{-jk} y_j, \qquad \overline{f_k} = N^{-1}\sum_j \zeta^{jk}\,\overline{y_j} = N^{-1}\sum_j \zeta^{jk} y_{-j} = f_k,$$

so f is real. Given $z \in \mathbb{C}^{N/2}$, take $y = (z, \bar{z})$ arranged by $\mathcal{H} \sqcup (-\mathcal{H})$. This f satisfies $\tau(f) = z$, and it is unique because W is injective. □

Lemma 1.2 (Encoding sensitivity; Lemma 1 of the specification)**.** *Let $\tau^{-1} : \mathbb{C}^{N/2} \to \mathbb{R}[X]/(X^N+1)$ be the inverse of the slot map.*

(a) $\|\tau^{-1}(\mathbf{z})\|_2 = \sqrt{2/N}\,\|\mathbf{z}\|_2$ *for every* $\mathbf{z} \in \mathbb{C}^{N/2}$.
(b) Let $\mathsf{Ecd}_\Omega(\mathbf{z}) = \lfloor \Omega\,\tau^{-1}(\mathbf{z}) \rceil$, *rounded coefficient-wise with any fixed tie rule. For all* $\mathbf{z}, \mathbf{z}'$,

$$\|\mathsf{Ecd}_\Omega(\mathbf{z}) - \mathsf{Ecd}_\Omega(\mathbf{z}')\|_2 \le \sqrt{2/N}\,\Omega\,\|\mathbf{z} - \mathbf{z}'\|_2 + \sqrt{N}.$$

(c) Let $\Delta_s = \sup_{D \sim D'} \|\mathbf{z}_s(D) - \mathbf{z}_s(D')\|_2$, *where each released entry occupies exactly one slot and all other slots are zero. Then* $\sup_{D\sim D'} \|\mathsf{Ecd}_\Omega(\mathbf{z}_s(D)) - \mathsf{Ecd}_\Omega(\mathbf{z}_s(D'))\|_2 \le \Delta_{\mathrm{enc},s} = \sqrt{2/N}\,\Omega\,\Delta_s + \sqrt{N}$.

Proof. (a) Apply Lemma 1.1 to $f = \tau^{-1}(\mathbf{z})$.

(b) Write $\mathsf{Ecd}_\Omega(\mathbf{z}) = \Omega\tau^{-1}(\mathbf{z}) + \mathbf{r}$ and $\mathsf{Ecd}_\Omega(\mathbf{z}') = \Omega\tau^{-1}(\mathbf{z}') + \mathbf{r}'$, with $\mathbf{r}, \mathbf{r}' \in [-\frac{1}{2}, \frac{1}{2}]^N$. Then $\|\mathbf{r} - \mathbf{r}'\|_\infty \le 1$, so $\|\mathbf{r} - \mathbf{r}'\|_2 \le \sqrt{N}$. Apply (a) to $\mathbf{z} - \mathbf{z}'$, using the $\mathbb{R}$-linearity of τ^{-1}, and the triangle inequality.

(c) Follows from (b). □

Remark 1.3 (Scope of Lemma 1.2)**.** (i) The constant $\sqrt{2/N}$ belongs to the convention of $N/2$ complex slots measured in complex ℓ_2. Under the full N-dimensional embedding the constant would be $1/\sqrt{N}$. With real data the imaginary parts are zero, so the specification's Δ_s, a real ℓ_2 sensitivity, is the right input.

(ii) The rounding term $\sqrt{N}$ cannot be improved in the worst case. At $N = 2^{14}$ it is 2^7, against $\sqrt{2/N}\,\Omega\Delta_s \approx 2^{52}$.

(iii) The lemma is about $\mathsf{Ecd}_\Omega(\mathbf{z}_s)$. The protocol's decrypted value is centred instead at the output of a noiseless evaluation of the circuit. Section 3 shows that the two differ by a deterministic term of size up to about 2^{30}, and Corollary 3.13 gives the sensitivity of the correct centre. Here Ω must be the *exact* final scale Ω_{fin} (Correction 2).

Facts used in Section 3.

Lemma 1.4 (Norms and automorphisms)**.** *For all* $f, g \in \mathbb{R}[X]/(X^N + 1)$*:*

- $\|fg\|_2 \le \|f\|_\infty^{\mathrm{can}}\,\|g\|_2$ *and* $\|fg\|_2 \le \|f\|_1 \|g\|_2$*;*
- $\|f\|_2 \le \|f\|_\infty^{\mathrm{can}} \le \|f\|_1 \le \sqrt{N}\|f\|_2$.

Every ψ_g *acts as a signed permutation on coefficients and as a permutation on* σ*. It therefore preserves* $\|\cdot\|_2$, $\|\cdot\|_\infty$, $\|\cdot\|_1$ *and* $\|\cdot\|_\infty^{\mathrm{can}}$.

Proof. Products. We have $\sigma(fg) = \sigma(f) \odot \sigma(g)$. By Lemma 1.1,

$$\|fg\|_2 = N^{-1/2}\|\sigma(f) \odot \sigma(g)\|_2 \le \|f\|_\infty^{\mathrm{can}}\, N^{-1/2}\|\sigma(g)\|_2 = \|f\|_\infty^{\mathrm{can}}\,\|g\|_2.$$

The bound $\|f\|_\infty^{\mathrm{can}} \le \|f\|_1$ holds because $|f(\zeta^j)| \le \sum_k |f_k|$. Also $\|f\|_2 = N^{-1/2}\|\sigma(f)\|_2 \le \|f\|_\infty^{\mathrm{can}}$, and $\|f\|_1 \le \sqrt{N}\|f\|_2$ by Cauchy–Schwarz.

Automorphisms. Take g odd. Then $X^k \mapsto X^{gk \bmod 2N} = \pm X^{gk \bmod N}$, and $k \mapsto gk \bmod N$ is a bijection of $\mathbb{Z}_N$. Moreover $\sigma(\psi_g f)_j = f(\zeta^{jg}) = \sigma(f)_{jg}$. □

Lemma 1.5 (The slot sum is a trace)**.** *Let* $L = \log_2(N/2)$ *and* $g_i = 5^{2^i}$. *For* $0 \le i \le L$, *let* $T_i = \prod_{l=i}^{L-1}(1 + \psi_{g_l})$, *so that* $T_i = \sum_{g \in \mathcal{H}_i} \psi_g$ *with* $\mathcal{H}_i = \langle 5^{2^i} \rangle$ *and* $|\mathcal{H}_i| = 2^{L-i}$. *The log-step* EvalSum *evaluates* $\mathrm{A} := T_0 = \sum_{g \in \mathcal{H}} \psi_g$, *and*

$$\mathrm{A}(f) = \tfrac{N}{2}\big(f_0 + f_{N/2} X^{N/2}\big), \qquad \tau(\mathrm{A}f) = \tfrac{N}{2}\big(f_0 + i f_{N/2}\big)\mathbf{1} = \Big(\sum_{h \in \mathcal{H}} \sigma(f)_h\Big)\mathbf{1}.$$

Proof. The product formula. Every $g \in \mathcal{H}$ has a unique representation $5^{\sum_i b_i 2^i}$ with $b_i \in \{0,1\}$. This gives the formula for T_i.

The trace of a monomial. By linearity it suffices to treat $f = X^k$. For $j \in \mathcal{H}$ we have $j\mathcal{H} = \mathcal{H}$, so

$$\sigma_j(\mathrm{A}X^k) = S_k := \sum_{g\in\mathcal{H}} \zeta^{gk} = \zeta^k \sum_{t=0}^{N/2-1} \zeta^{4kt}.$$

If $\zeta^{4k} \neq 1$, the geometric sum equals $(1-\zeta^{2Nk})/(1-\zeta^{4k}) = 0$. Now $\zeta^{4k} = 1$ if and only if $k \equiv 0 \pmod{N/2}$. Hence $S_0 = N/2$, $S_{N/2} = \frac{N}{2}\zeta^{N/2} = \frac{N}{2}i$, and $S_k = 0$ otherwise. For $j \in -\mathcal{H}$ the values are the conjugates.

Identification. The polynomial $\frac{N}{2}([k{=}0] + [k{=}N/2]X^{N/2})$ has the same embedding, because $\zeta^{jN/2} = i^j$ equals i on $\mathcal{H}$ and $-i$ on $-\mathcal{H}$. The claim follows from the injectivity of σ. The last expression is $\sigma_j(\mathrm{A}f) = \sum_{g\in\mathcal{H}} \sigma(f)_{jg}$ evaluated at $j = 1$. □

Lemma 1.6 (Masked slot functionals)**.** *Let $f \in \mathbb{R}[X]/(X^N+1)$, $h \in \mathcal{H}$ and $\mathcal{H}' \subseteq \mathcal{H}$. The real-linear map $\xi \mapsto \sigma_h\big(\sum_{g\in\mathcal{H}'} \psi_g(f\xi)\big)$, for $\xi \in \mathbb{R}^N$, equals $\sum_k \xi_k\beta_k$, with*

$$\|\beta\|_2^2 = N \sum_{g\in\mathcal{H}'} |\sigma(f)_{hg}|^2 \le N \min\big\{N\|f\|_2^2,\ |\mathcal{H}'|\,(\|f\|_\infty^{\mathrm{can}})^2\big\}.$$

In particular $\|\beta\|_2^2 = N|\mathcal{H}'|$ when $f = 1$. The real and imaginary parts have coefficient vectors $\Re\beta$ and $\Im\beta$, with $\|\Re\beta\|^2 + \|\Im\beta\|^2 = \|\beta\|^2$.

Proof. We have $\sigma_h(\psi_g(f\xi)) = \sigma(f)_{hg}\sum_k \xi_k\zeta^{hgk}$, so $\beta_k = \sum_{g\in\mathcal{H}'} \sigma(f)_{hg}\zeta^{hgk}$. Expanding,

$$\|\beta\|^2 = \sum_{g,g'\in\mathcal{H}'} \sigma(f)_{hg}\overline{\sigma(f)_{hg'}} \sum_{k<N} \zeta^{hk(g-g')}.$$

If $g \neq g'$, then $m = h(g-g')$ is even and $m \not\equiv 0 \pmod{2N}$, so ζ^m is an N-th root of unity other than 1 and the inner sum vanishes. If $g = g'$, the inner sum is N. The bounds follow from Lemma 1.1 and from $|\sigma(f)_{hg}| \le \|f\|_\infty^{\mathrm{can}}$. □

Lemma 1.7 (Masking)**.** *Let $\mathbf{u}_k$ be the one-hot slot vector and $m_{u_k} = \mathsf{Ecd}_{\Omega_u}(\mathbf{u}_k) = \Omega_u\tau^{-1}(\mathbf{u}_k) + r_{u_k}$, with $\|r_{u_k}\|_\infty \le \frac{1}{2}$. For polynomials y_e and distinct slots k_e:*

$$\Big\|\sum_e y_e\,\Omega_u\tau^{-1}(\mathbf{u}_{k_e})\Big\|_2 = \Omega_u\sqrt{2/N}\Big(\sum_e |\tau(y_e)_{k_e}|^2\Big)^{1/2}, \qquad \|y_e\,r_{u_k}\|_2 \le \tfrac{N}{2}\|y_e\|_2.$$

Proof. The slot map is multiplicative, so $\tau(y_e\tau^{-1}(\mathbf{u}_{k_e})) = \tau(y_e)_{k_e}\mathbf{u}_{k_e}$. Apply Lemma 1.2(a) to the sum, whose slots are disjoint. The second bound is Lemma 1.4 with $\|r_{u_k}\|_1 \le N/2$. □

2 Lemma 2 (smudging with the exact discrete Gaussian)

Conventions. For $\sigma > 0$ and $x \in \mathbb{R}^N$, let $\rho_\sigma(x) = \exp(-\|x\|^2/(2\sigma^2))$. For a countable set S, write $\rho_\sigma(S) = \sum_{x\in S}\rho_\sigma(x)$. For $c \in \mathbb{R}^N$, $D_{\mathbb{Z}^N+c,\sigma}$ is the distribution on $\mathbb{Z}^N + c$ with probability mass function $\rho_\sigma(x)/\rho_\sigma(\mathbb{Z}^N + c)$. The specification's $\mathcal{N}_\mathbb{Z}(0,\sigma^2)^N$ is $D_{\mathbb{Z}^N,\sigma}$, the product of N independent copies of the one-dimensional discrete Gaussian $\mathcal{N}_\mathbb{Z}(0,\sigma^2)$ of Canonne, Kamath and Steinke [4]. SD is statistical (total-variation) distance.

Smoothing parameters are usually stated with the Gaussian width $s = \sqrt{2\pi}\sigma$, where $\rho_\sigma(x) = e^{-\pi\|x\|^2/s^2}$. For a lattice Λ and $\epsilon > 0$, $\eta_\epsilon(\Lambda)$ is the least s such that $\sum_{y\in\Lambda^*\setminus\{0\}} e^{-\pi s^2\|y\|^2} \le \epsilon$ [13].

Lemma 2.1 (Gaussian sums are maximal at the lattice)**.** *For all $\sigma > 0$ and $c \in \mathbb{R}^N$, $\rho_\sigma(\mathbb{Z}^N + c) \le \rho_\sigma(\mathbb{Z}^N)$. Equality holds if $c \in \mathbb{Z}^N$.*

Proof. The sum factorizes over coordinates, so it suffices to take $N = 1$. By Poisson summation for the Gaussian,

$$\sum_{z\in\mathbb{Z}} e^{-(z+c)^2/(2\sigma^2)} = \sqrt{2\pi}\,\sigma \sum_{k\in\mathbb{Z}} e^{-2\pi^2\sigma^2k^2}\cos(2\pi kc) \le \sqrt{2\pi}\,\sigma \sum_{k\in\mathbb{Z}} e^{-2\pi^2\sigma^2k^2} = \rho_\sigma(\mathbb{Z}).$$

For integer c the sum is a re-indexing of $\rho_\sigma(\mathbb{Z})$. □

Lemma 2.2 (Divergences of an integer shift)**.** *Let $\eta \sim D_{\mathbb{Z}^N,\sigma}$, $\mathbf{e} \in \mathbb{Z}^N$, and let $P_{\mathbf{e}}$ be the law of $\eta + \mathbf{e}$. For every $\alpha > 1$,*

$$D_\alpha(P_{\mathbf{e}}\|P_{\mathbf{0}}) \le \frac{\alpha\|\mathbf{e}\|^2}{2\sigma^2} \qquad \textit{and} \qquad \mathrm{KL}(P_{\mathbf{e}}\|P_{\mathbf{0}}) = \frac{\|\mathbf{e}\|^2}{2\sigma^2}.$$

Proof. Normalizers. Since $\mathbf{e}$ is integral, $P_{\mathbf{e}}(y) = \rho_\sigma(y-\mathbf{e})/Z$ and $P_{\mathbf{0}}(y) = \rho_\sigma(y)/Z$ with the same $Z = \rho_\sigma(\mathbb{Z}^N)$.

Rényi divergence. Use the identity $\alpha\|y-\mathbf{e}\|^2 + (1-\alpha)\|y\|^2 = \|y-\alpha\mathbf{e}\|^2 + \alpha(1-\alpha)\|\mathbf{e}\|^2$. With Lemma 2.1,

$$\sum_y P_{\mathbf{e}}(y)^\alpha P_{\mathbf{0}}(y)^{1-\alpha} = \frac{\rho_\sigma(\mathbb{Z}^N - \alpha\mathbf{e})}{Z}\, e^{\alpha(\alpha-1)\|\mathbf{e}\|^2/(2\sigma^2)} \le e^{\alpha(\alpha-1)\|\mathbf{e}\|^2/(2\sigma^2)}.$$

Take $\frac{1}{\alpha-1}\ln$ of both sides.

Kullback–Leibler divergence. By symmetry $\mathbb{E}\eta = \mathbf{0}$, so

$$\mathrm{KL}(P_{\mathbf{e}}\|P_{\mathbf{0}}) = \mathbb{E}\big[\|\eta+\mathbf{e}\|^2 - \|\eta\|^2\big]/(2\sigma^2) = \big(2\mathbb{E}\langle\eta,\mathbf{e}\rangle + \|\mathbf{e}\|^2\big)/(2\sigma^2) = \|\mathbf{e}\|^2/(2\sigma^2).$$

□

Lemma 2.3 (Smudging; Lemma 2 of the specification, with explicit constants)**.** *Let $\eta \sim D_{\mathbb{Z}^N,\sigma}$ and $\mathbf{e} \in \mathbb{Z}^N$.*

(a) Elementary bound, valid for all $\sigma > 0$: $\mathrm{SD}(\eta+\mathbf{e},\eta) \le \|\mathbf{e}\|_2/(2\sigma)$.

(b) Sharp constant: *let $\epsilon \in (0,1)$, $\sigma_2 \ge \eta_\epsilon(\mathbb{Z}^N)/\sqrt{2\pi}$, $\sigma > \sigma_2$ and $\sigma_1 = \sqrt{\sigma^2 - \sigma_2^2}$. Then*

$$\mathrm{SD}(\eta+\mathbf{e},\eta) \le \frac{\|\mathbf{e}\|_2}{\sqrt{2\pi}\,\sigma_1} + \frac{4\epsilon}{1-\epsilon}.$$

By [13, Lemma 3.3], $\eta_\epsilon(\mathbb{Z}^N) \le \sqrt{\ln(2N(1+1/\epsilon))/\pi}$. For $N = 2^{14}$ and $\epsilon = 2^{-100}$ this gives $\eta_\epsilon(\mathbb{Z}^N)/\sqrt{2\pi} \le 2.01$, so $\sigma_2 = 2.1$ is admissible.

Proof. (a) Apply Pinsker's inequality, $\mathrm{SD} \le \sqrt{\mathrm{KL}/2}$, to Lemma 2.2.

(b) The proof constructs a shift-equivariant post-processing F of a continuous Gaussian whose output is close to η.

The post-processing. Let $x \sim \mathcal{N}(0,\sigma_1^2\mathbf{I}_N)$ be continuous. Let $F(x) = x + z'$, where $z' \sim D_{\mathbb{Z}^N - x,\sigma_2}$ is sampled given x. Then $F(x) \in \mathbb{Z}^N$. For $\mathbf{e} \in \mathbb{Z}^N$ we have $\mathbb{Z}^N - x - \mathbf{e} = \mathbb{Z}^N - x$, so $F(x+\mathbf{e})$ and $F(x) + \mathbf{e}$ have the same law.

$F(x)$ is close to η. The law of $F(x)$ is

$$P(y) = \int \phi_{\sigma_1}(x)\,\frac{\rho_{\sigma_2}(y-x)}{\rho_{\sigma_2}(\mathbb{Z}^N - x)}\,dx, \qquad y \in \mathbb{Z}^N.$$

Since $\sqrt{2\pi}\sigma_2 \geq \eta_\epsilon(\mathbb{Z}^N)$, the smoothing lemma [13, Lemma 4.4] gives $\rho_{\sigma_2}(\mathbb{Z}^N - x) \in [\frac{1-\epsilon}{1+\epsilon}, 1]\,\rho_{\sigma_2}(\mathbb{Z}^N)$ for every x. The Gaussian convolution identity gives $\int \phi_{\sigma_1}(x)\rho_{\sigma_2}(y-x)\,dx = (\sigma_2/\sigma)^N \rho_\sigma(y)$. Hence $P(y) = C\,\kappa(y)\rho_\sigma(y)$, with $\kappa(y) \in [1, r]$ and $r = \frac{1+\epsilon}{1-\epsilon}$. Normalizing, $P(y)/D_{\mathbb{Z}^N,\sigma}(y) \in [1/r, r]$, and therefore $\mathrm{SD}(F(x), \eta) \leq r - 1 = 2\epsilon/(1-\epsilon)$.

The shift. By the triangle inequality and the data-processing inequality,

$$\mathrm{SD}(\eta + \mathbf{e}, \eta) \leq \mathrm{SD}(F(x+\mathbf{e}), F(x)) + \frac{4\epsilon}{1-\epsilon} \leq \mathrm{SD}(x+\mathbf{e}, x) + \frac{4\epsilon}{1-\epsilon}.$$

By rotation invariance, $\mathrm{SD}(x+\mathbf{e}, x) = 2\Phi(\|\mathbf{e}\|/(2\sigma_1)) - 1 \leq \|\mathbf{e}\|/(\sqrt{2\pi}\sigma_1)$, since $\Phi' \leq 1/\sqrt{2\pi}$. □

Corollary 2.4 (Flooding with side information)**.** *Let $(Z, \mathbf{e})$ have any joint law, with $\mathbf{e} \in \mathbb{Z}^N$. Let $\eta \sim D_{\mathbb{Z}^N,\sigma}$ be independent of $(Z, \mathbf{e})$, and let $B > 0$. Then*

$$\mathrm{SD}\big((Z, \mathbf{e}+\eta), (Z, \eta)\big) \leq \Pr[\|\mathbf{e}\|_2 > B] + \sup_{\|\mathbf{e}'\|_2 \leq B} \mathrm{SD}(\mathbf{e}'+\eta, \eta).$$

The same bound holds after reducing $\mathbf{e}+\eta$ and η modulo q_0, by data processing. In particular, suppose $\sigma \geq \max\{2^\kappa B/\sqrt{2\pi},\, 2^{42}\}$, with $\epsilon = 2^{-100}$ and $\sigma_2 = 2.1$ as above. Then the right-hand side is at most

$$\Pr[\|\mathbf{e}\| > B] + 2^{-\kappa}\sigma/\sigma_1 + 2^{-97.9} \leq \Pr[\|\mathbf{e}\| > B] + 2^{-\kappa}(1 + 2^{-80}) + 2^{-97}.$$

Proof. Condition on $(Z, \mathbf{e}) = (z, \mathbf{e}')$. Because η is independent of $(Z, \mathbf{e})$, the conditional distance is $\mathrm{SD}(\mathbf{e}'+\eta, \eta)$. This is at most the supremum when $\|\mathbf{e}'\| \leq B$, and at most 1 otherwise. Averaging over $(z, \mathbf{e}')$ gives the bound. For the numerical form, Lemma 2.3(b) gives $\mathrm{SD}(\mathbf{e}'+\eta, \eta) \leq \|\mathbf{e}'\|/(\sqrt{2\pi}\sigma_1) + 4\epsilon/(1-\epsilon) \leq 2^{-\kappa}\sigma/\sigma_1 + 2^{-97.9}$. Moreover $\sigma/\sigma_1 = (1-\sigma_2^2/\sigma^2)^{-1/2} \leq 1 + \sigma_2^2/\sigma^2$ whenever $\sigma \geq 2\sigma_2$, and for $\sigma \geq 2^{42}$ this is at most $1 + 4.41 \cdot 2^{-84} < 1 + 2^{-81}$. Hence $2^{-\kappa}\sigma/\sigma_1 + 2^{-97.9} \leq 2^{-\kappa}(1+2^{-80}) + 2^{-97}$. □

Remark 2.5 (The floor 2^{42}; corrected)**.** An earlier version of this corollary stated the final inequality, in the form $\leq \Pr[\|\mathbf{e}\| > B] + 2^{-\kappa} + 2^{-97}$, "for every $\sigma \geq 2^{20}$". That is false for $\kappa < 60$: at $\sigma = 2^{20}$ the factor σ/σ_1 is $1 + 2^{-38.9}$, and the proof's bound $1 + \sigma_2^2/\sigma^2$ is $1 + 2^{-37.9}$, so the excess $2^{-\kappa} \cdot 2^{-37.9}$ exceeds the room $2^{-97} - 2^{-97.9}$ unless $\kappa \geq 60$. The factor $1 + 2^{-80}$ needs $\sigma \geq 2.1 \cdot 2^{40} \approx 2^{41.1}$; the floor is now 2^{42} (composition appendix, Remark 2.6). Nothing numerical changes: at every evaluation setting $\sigma_c \geq 2^\kappa B_{\mathrm{tot}}/\sqrt{2\pi} \geq 2^{55}$. The hockey-stick form of this step (composition appendix, Lemma 5.1) needs no floor at all.

Remark 2.6 (Comparison with the specification)**.** The specification states Lemma 2 with "$\sigma_c \gg 1$" and "+negl". Lemma 2.3(b) makes both precise:

- the condition is $\sigma_c > \sigma_2 \approx 2.1$;
- the negligible term is $4\epsilon/(1-\epsilon) \leq 2^{-97}$;
- the per-release distance is $2^{-\kappa}(1 + 2^{-80}) + 2^{-97}$ rather than exactly $2^{-\kappa}$.

The difference is immaterial, but the δ accounting should use the corrected value (Correction 5). An implementation may prefer the elementary bound (a), which needs no smoothing argument. The condition then becomes $\sigma_c \geq 2^\kappa B_{\mathrm{tot}}/2$, which costs $\log_2(\sqrt{2\pi}/2) \approx 0.33$ bits of Ω.

3 Obligation O1: a worst-case bound on the whole error

The review requires the bound to be worst case over the data, for two reasons. CKKS decryptions leak the secret through their error [12]. And flooding calibrated to average-case noise estimates has been broken by adversarially chosen inputs [9]. Below, every polynomial determined by the data, or by the public randomness, is treated adversarially. Probabilities are taken only over secrets, errors and fresh coins, which are independent of those polynomials (Lemma 3.6).

3.1 Probability tools

A real random variable X is *v-subgaussian* if $\mathbb{E}e^{\lambda X} \le e^{v\lambda^2/2}$ for all $\lambda \in \mathbb{R}$.

Lemma 3.1 (Variance proxies)**.** *(a) The uniform distribution on $\{-1, 0, 1\}$ is $\frac{2}{3}$-subgaussian. Sums of independent variables add proxies, so each coefficient of $s = s_R + s_O$ is $\frac{4}{3}$-subgaussian.*
(b) $\mathcal{N}_{\mathbb{Z}}(0, \sigma^2)$ is σ^2-subgaussian.
(c) A mean-zero variable supported on an interval of length 1 is $\frac{1}{4}$-subgaussian.

Proof. (a) Expand in powers of λ: $\frac{1}{3} + \frac{2}{3}\cosh\lambda = 1 + \sum_{k\ge1} \frac{2}{3(2k)!}\lambda^{2k}$ and $e^{\lambda^2/3} = 1 + \sum_{k\ge1} \frac{\lambda^{2k}}{3^k k!}$. Compare the series term by term, using $\frac{2}{3}\,3^k k! \le (2k)!$ for all $k \ge 1$.

(b) Complete the square:

$$\mathbb{E}e^{\lambda X} = \frac{\sum_z e^{-z^2/(2\sigma^2)+\lambda z}}{\rho_\sigma(\mathbb{Z})} = e^{\lambda^2\sigma^2/2}\,\frac{\rho_\sigma(\mathbb{Z} - \lambda\sigma^2)}{\rho_\sigma(\mathbb{Z})} \le e^{\lambda^2\sigma^2/2},$$

by Lemma 2.1.

(c) This is Hoeffding's lemma. □

Lemma 3.2 (Hsu, Kakade and Zhang [11], Theorem 1)**.** *Let $x \in \mathbb{R}^n$ satisfy $\mathbb{E}e^{\alpha^\top x} \le e^{\|\alpha\|^2/2}$ for all α. This holds, for example, when the coordinates are independent and 1-subgaussian. Let $A \in \mathbb{R}^{m\times n}$ and $\Sigma = A^\top A$. Then, for all $t > 0$,*

$$\Pr\big[\|Ax\|^2 > \operatorname{tr}\Sigma + 2\sqrt{\operatorname{tr}(\Sigma^2)t} + 2\|\Sigma\|t\big] \le e^{-t}.$$

Since $\operatorname{tr}(\Sigma^2) \le \|\Sigma\|\operatorname{tr}\Sigma$, it follows that $\|Ax\|_2 \le \|A\|_F + \sqrt{2t}\,\|A\|_{\mathrm{op}}$ with probability at least $1 - e^{-t}$.

Corollary 3.3 (Products with fixed polynomials)**.** *Let $\xi_1, \dots, \xi_m \in \mathbb{Z}^N$ be independent random vectors with independent coordinates, where those of ξ_j are mean-zero and v_j-subgaussian. Let $f_1, \dots, f_m$ be fixed polynomials, or random but independent of the ξ_j. Then, with probability at least $1 - e^{-t}$,*

$$\Big\|\sum_j f_j\xi_j\Big\|_2 \le \Big(N\sum_j v_j\|f_j\|_2^2\Big)^{1/2} + \Big(2t\sum_j v_j(\|f_j\|_\infty^{\mathrm{can}})^2\Big)^{1/2}.$$

More generally, for any fixed real-linear map $\ell : (\mathbb{R}^N)^m \to \mathbb{R}^d$ with matrix A in the scaled variables $\xi_j/\sqrt{v_j}$, $\|\ell(\xi)\| \le \|A\|_F(1 + \sqrt{2t})$ with probability at least $1 - e^{-t}$.

Proof. Multiplication by f has the matrix $M_f = N^{-1}W^*\mathrm{diag}(\sigma(f))W$. This matrix is normal, so $\|M_f\|_F^2 = \|\sigma(f)\|^2 = N\|f\|^2$ and $\|M_f\|_{\mathrm{op}} = \|f\|_\infty^{\mathrm{can}}$. Apply Lemma 3.2 to the block matrix $[\sqrt{v_1}M_{f_1} \mid \cdots \mid \sqrt{v_m}M_{f_m}]$. Its operator norm satisfies $\|A\|_{\mathrm{op}}^2 \le \sum_j v_j\|M_{f_j}\|_{\mathrm{op}}^2$. For the general form, use $\|A\|_{\mathrm{op}} \le \|A\|_F$. When the f_j are random and independent of the ξ_j, condition on them. □

Lemma 3.4 (Norms of random polynomials)**.** *Let $\xi \in \mathbb{Z}^N$ have independent, mean-zero, v-subgaussian coordinates. Then:*

- $\|\xi\|_2 \le \sqrt{v}(\sqrt{N} + \sqrt{2t})$ *with probability at least $1 - e^{-t}$;*
- $\|\xi\|_\infty^{\mathrm{can}} \le \sqrt{2vN\ln(2N/\delta)}$ *with probability at least $1 - \delta$.*

Proof. The first bound is Lemma 3.2 with $A = \mathbf{I}$.

For the second, take h odd. The real part $\Re\sigma(\xi)_h = \sum_k \xi_k\cos(\pi hk/N)$ is $vN/2$-subgaussian, because $\sum_{k<N}\cos^2(\pi hk/N) = N/2$; the same holds for the imaginary part. So each part exceeds a in absolute value with probability at most $2e^{-a^2/(vN)}$. There are $N/2$ conjugate pairs of slots and two parts each, so a union bound gives $\|\xi\|_\infty^{\mathrm{can}} \le \sqrt{2}\,a$ except with probability $2Ne^{-a^2/(vN)}$. □

3.2 The circuit, made explicit

The specification fixes the circuit only up to standard choices for CKKS [7], its RNS variant [8] and hybrid key switching [10]. We make the following choices explicit; Section 10 discusses deviations from them.

(M1) *Moduli.* Ciphertexts start at $Q_2 = q_0q_1q_2$, where q_0 may be a product of RNS primes. P is the special modulus. $\textsc{Rescale}_q$ maps $c \mapsto (c - [c]_q)/q$ component-wise, as RNS-CKKS does. Its rounding error on each component lies in $[-\frac{1}{2}, \frac{1}{2}]^N$.

(M2) *Scales.* R encodes its vectors at scale Ω_U. O encodes V_e at Ω_V and the masks at Ω_u. The final scale $\Omega_{\text{fin}} := \Omega_U\Omega_V\Omega_u/(q_1q_2)$ is public, and R decodes with Ω_{fin}. The example uses $\Omega_U = q_2$ and $\Omega_V = \Omega_u = q_1 = \Omega_{\text{fin}} = \Omega$.

(M3) *Key switching*, in the hybrid form.

- *Gadget.* The primes of $Q_1 = q_0q_1$ are partitioned into d_{num} digits with products $\tilde{Q}_j$. Let g_j be the CRT idempotent of digit j: it is $\equiv 1$ modulo the primes of digit j and $\equiv 0$ modulo the others.
- *Decomposition.* $\tilde{D}_j(c) \in \mathcal{R}$ satisfies $\tilde{D}_j(c) \equiv c \pmod{\tilde{Q}_j}$ and $\|\tilde{D}_j(c)\|_\infty \le \varsigma_\uparrow \tilde{Q}_j/2$. Here $\varsigma_\uparrow = 1$ for exact ModUp, and $\varsigma_\uparrow \le k_j$, the number of primes in digit j, for fast base extension.
- *Joint rotation keys.* The key for ψ_g is $\text{rk}_j = (b_j, a_j)$, with $b_j = -a_js + P\mathsf{g}_j\psi_g(s) + e_j^{\text{rk}} \bmod PQ_2$. Here a_j is a common random string and $e_j^{\text{rk}} = e_{R,j}^{\text{rk}} + e_{O,j}^{\text{rk}}$. This is the sum of the two parties' shares $-a_js_P + P\mathsf{g}_j\psi_g(s_P) + e_{P,j}^{\text{rk}}$ in the protocol of Mouchet et al. [17]. It is used at level 1 after reduction modulo PQ_1.
- *The switch.* $\text{KS}(c_0, c_1) = \big(c_0 + \lfloor P^{-1}\sum_j \tilde{D}_j(c_1)b_j\rceil,\ \lfloor P^{-1}\sum_j \tilde{D}_j(c_1)a_j\rceil\big)$. The ModDown rounding errors rd_0, rd_1 lie in $[-\frac{1}{2}, \frac{1}{2}]^N$.

(M4) *Release* is Algorithm 5 at level 0, with $p_0 = -as + e^{\text{pk}} \bmod q_0$ and $e^{\text{pk}} = e_R^{\text{pk}} + e_O^{\text{pk}}$. After two rescales the ciphertext is at q_0. The "modulus switch to q_0" of Algorithm 4 is therefore a no-op, or an exact limb drop.

(M5) *Distributions.* $s_R, s_O, u \leftarrow \chi_{\text{key}}$, uniform ternary. Every error is drawn from $\chi_{\text{err}} = \mathcal{N}_{\mathbb{Z}}(0, \sigma_e^2)^N$, with $\sigma_e = 3.2$.

Lemma 3.5 (Key switching)**.** *Suppose* $c_0 + c_1\psi_g(s) \equiv \mu \pmod{Q_1}$. *Then* $\text{KS}(c_0, c_1) = (c_0', c_1')$ *satisfies* $c_0' + c_1's \equiv \mu + e_{\text{ks}} \pmod{Q_1}$, *where*

$$e_{\text{ks}} = P^{-1}\sum_j \tilde{D}_j(c_1)\, e_j^{\text{rk}} + \text{rd}_0 + \text{rd}_1 s.$$

Proof. Modulo PQ_1, $\sum_j \tilde{D}_j(b_j + a_js) = P\sum_j \tilde{D}_j\mathsf{g}_j\psi_g(s) + \sum_j \tilde{D}_j e_j^{\text{rk}}$. By the CRT, $\sum_j \tilde{D}_j(c_1)\mathsf{g}_j \equiv c_1 \pmod{Q_1}$. Divide by P, which is exact modulo Q_1 for multiples of P, and add the rounding errors. □

Lemma 3.6 (The c_1-components are independent of all secrets and errors)**.** *At every step of Algorithm 4, the* c_1*-component of every ciphertext is a deterministic function of three things: the seeds' expansions* a_h*, the common random strings* a_j *of the rotation keys, and the data. Consequently* $\tilde{D}_j(\cdot)$, rd_1 *and the* c_1*-rounding errors of every rescale are independent of* $(s_R, s_O, u,$ *all errors*$)$.

Proof. By induction over the operations:

- the conversion leaves $c_1 = a_h$;
- a ct–pt product maps $c_1 \mapsto m_V c_1$, and additions add c_1-components;
- a rescale maps $c_1 \mapsto \lfloor c_1/q\rceil$;
- an automorphism maps $c_1 \mapsto \psi_g(c_1)$;
- key switching outputs $\lfloor P^{-1}\sum_j \tilde{D}_j(c_1)a_j\rceil$.

None of these uses c_0, a secret or an error. □

3.3 The noiseless centre and the error decomposition

For entry e, let U_e be R's vector and V_e O's vector. Define the plaintexts $m_{U_e,h} = \mathsf{Ecd}_{\Omega_U}(U_e[h])$, $m_{V_e,h} = \mathsf{Ecd}_{\Omega_V}(V_e[h])$ and $m_{u_e} = \mathsf{Ecd}_{\Omega_u}(\mathbf{u}_{k_e}) = \Omega_u \tau^{-1}(\mathbf{u}_{k_e}) + r_{u_e}$.

Definition 3.7 (Noiseless evaluation)**.** For release s,

$$\widehat{Z}_s(D) = \sum_{e \in \mathcal{E}_s} \frac{1}{q_1 q_2}\, \mathrm{A}\Big(\sum_{h \le m} m_{U_e,h} m_{V_e,h}\Big) m_{u_e} \ \in \mathbb{Q}[X]/(X^N+1), \qquad f_s(D) = \lfloor \widehat{Z}_s(D) \rceil \in \mathbb{Z}^N.$$

This is a deterministic function of O's data, R's data and the public parameters. It does not depend on any key, error or random coin.

Proposition 3.8 (Error decomposition)**.** *Let $v_s = h_{O,s} + c'_{1,s} s_R \bmod q_0$. Under (M1)–(M5),*

$$v_s \equiv f_s(D) + e_{\mathrm{tot},s} + \eta_s \pmod{q_0}, \qquad e_{\mathrm{tot},s} = \sum_{i=1}^{8} T_i \in \mathbb{Z}^N,$$

where, writing $\Pi_e(y) := q_1^{-1} y\, \Omega_u \tau^{-1}(\mathbf{u}_{k_e})$ for the main part of masking and rescaling,

$$\begin{aligned}
T_1 &= \sum_e \Pi_e\big(\mathrm{A}(x_e^R)\big), \quad x_e^R = q_2^{-1} \sum_h e_{R,h}^{(U_e)} m_{V_e,h} && \text{(R's encryption errors)}\\
T_2 &= \sum_e \Pi_e\big(\mathrm{A}(\varepsilon_{0,e} + \varepsilon_{1,e} s)\big) && \text{(first rescale)}\\
T_3 &= \sum_e \Pi_e\Big(\sum_{i<L} T_{i+1}\big(P^{-1} \textstyle\sum_j \tilde{D}_{e,i,j} e_j^{\mathrm{rk},(i)}\big)\Big) && \text{(key-switching gadget)}\\
T_4 &= \sum_e \Pi_e\Big(\sum_{i<L} T_{i+1}\big(\mathrm{rd}_{0,e,i} + \mathrm{rd}_{1,e,i} s\big)\Big) && \text{(ModDown roundings)}\\
T_5 &= \sum_e (\varepsilon'_{0,e} + \varepsilon'_{1,e} s) && \text{(second rescale)}\\
T_6 &= e^{\mathrm{pk}} u + e_0 + e_1 s && \text{(fresh encryption)}\\
T_7 &= \widehat{Z}_s(D) - f_s(D), \quad \|T_7\|_\infty \le \tfrac{1}{2} && \text{(rounding of the centre)}\\
T_8 &= \sum_e q_1^{-1} Y_e\, r_{u_e} && \text{(mask spill of the errors).}
\end{aligned}$$

Here $Y_e = \mathrm{A}(x_e^R + \varepsilon_{0,e} + \varepsilon_{1,e} s) + \sum_i T_{i+1}(e_{\mathrm{ks},e,i})$. All ε and rd *vectors lie in $[-\frac{1}{2}, \frac{1}{2}]^N$, and $e_j^{\mathrm{rk},(i)}$ is the error of digit j of the key for ψ_{g_i}.*

Proof. Track the decryption invariant $c_0 + c_1 s$, as integer polynomials, through Algorithm 4.

Conversion and products. The conversion gives $c_0 + a_h s = m_{U,h} + e_{R,h}$ exactly. After the products and the chunk sum, the value is $\mu_{\mathbb{Z}} + KQ_2$ for some $K \in \mathcal{R}$, with $\mu_{\mathbb{Z}} = \sum_h m_{U,h} m_{V,h} + \sum_h e_{R,h} m_{V,h}$.

First rescale. Write the rounding of c_j/q_2 as $c_j/q_2 + \varepsilon_j$. Then $c'_0 + c'_1 s = \mu_{\mathbb{Z}}/q_2 + \varepsilon_0 + \varepsilon_1 s + KQ_1$.

EvalSum. By Lemma 3.5 and $\psi_g(c_0) + \psi_g(c_1)\psi_g(s) = \psi_g(c_0 + c_1 s)$, induction over the L steps gives $\mathrm{A}(\mu^{(2)}) + \sum_i T_{i+1}(e_{\mathrm{ks},i}) \pmod{Q_1}$, by Lemma 1.5.

Masking and second rescale. Multiplying by m_{u_e} and rescaling by q_1 adds $\varepsilon'_0 + \varepsilon'_1 s$ modulo q_0.

Release. Summing over e, the release adds $(-as + e^{\mathrm{pk}})u + e_0 + (au + e_1)s + \eta = e^{\mathrm{pk}} u + e_0 + e_1 s + \eta$.

Collecting terms. Split $m_{u_e} = \Omega_u \tau^{-1}(\mathbf{u}_{k_e}) + r_{u_e}$. The terms not involving any key, error or rounding make up exactly $\widehat{Z}_s(D)$; the rest are $T_1, \dots, T_6$ and T_8. Since v_s, f_s and η_s are integral, so is $e_{\mathrm{tot},s}$. □

Remark 3.9 (Why the centre matters)**.** The specification centres the mechanism at $\mathsf{Ecd}_\Omega(\mathbf{z}_s(D))$ and counts the encoding roundings of V_e and $\mathbf{u}_k$ as error. The deterministic difference $\widehat{Z}_s(D) - \Omega\tau^{-1}(\mathbf{z}_s(D))$ contains the following terms:

- the *mask spill* of the statistic itself, $\sum_e(\Omega_U\Omega_V/(q_1q_2))S_e\, r_{u_e}$, where S_e is the entry's value. Its ℓ_2 norm is about $\sqrt{N/12}\,(\sum_e S_e^2)^{1/2}$. This is up to $2^{30.7}$ for $\bar{n}_O = 10^6$, and up to $2^{34.5}$ by the triangle inequality;
- the cross terms $\mathrm{A}(r_U \cdot \Omega_V\tau^{-1}(V))$ between R's rounding and O's data. These are up to $\frac{\sqrt{N}}{2}\sqrt{m}\,(\bar{n}_O\Gamma_\mathcal{E})^{1/2} \approx 2^{25}$, in the notation of Section 3.4;
- a nominal-versus-exact scale mismatch of relative size about 2^{-30} when the implementation decodes with a nominal scale, as OpenFHE's FIXEDMANUAL does. This gives about 2^{23}; see Section 3.6.

All three are deterministic functions of D. They are part of the *query*, not hints about secrets. They need no flooding, provided the mechanism is centred at $f_s(D)$ and the sensitivity of f_s is used (Corollary 3.13). Flooding them instead would require $B_{\mathrm{tot}} \approx 2^{30.7}$ to $2^{34.5}$, and hence $\log_2\Omega \approx 67$ to 71 for stage 2 of the example (script `scratch/ecd_centre.py`).

3.4 Term-by-term bounds, worst case over data

Let $|\mathcal{E}|$ be the number of entries in the release and $m = \lceil \bar{n}_O/N_s \rceil$. Fix δ_B and split it into $n_{\mathrm{ev}} = 12$ equal parts. Put

$$t = \ln(n_{\mathrm{ev}}/\delta_B), \qquad \tau_t = 1 + \sqrt{2t}, \qquad \ell_N = \ln(2Nn_{\mathrm{ev}}/\delta_B), \qquad v_s = \tfrac{4}{3}, \quad v_u = \tfrac{2}{3}.$$

Lemma 3.10 (Data bounds)**.** *Every dataset in the clipped domain satisfies* $\sum_{e\in\mathcal{E}} V_{e,t}^2 \le \Gamma_\mathcal{E}$ *for every row* t*, where*

$$\Gamma_{\mathcal{E}_G} = o^2 + (d_R+1)o, \qquad \Gamma_{\mathcal{E}_2} = o^2 + (d_R+1)(o+w) + ow + w^2.$$

Consequently $\sum_e\sum_h \|m_{V_e,h}\|_2^2 \le M_V^2$*, with* $M_V := \Omega_V\sqrt{2/N}\,(\bar{n}_O\Gamma_\mathcal{E})^{1/2} + \frac{1}{2}(|\mathcal{E}|mN)^{1/2}$.

Proof. Bound the entries one block at a time:

- for $\mathrm{triu}\,\mathbf{B}$: $\sum_{j\le k}(x_jx_k)^2 = \frac{1}{2}(\|\mathbf{x}\|^4 + \sum_j x_j^4) \le \|\mathbf{x}\|^4 \le o^2$;
- for $\mathbf{C}_{jk}$, where $V = x_{O,k}$ for d_R+1 values of j: $(d_R+1)\|\mathbf{x}\|^2$;
- for $\mathbf{c}_R$: $(d_R+1)y^2$;
- for $\mathbf{c}_O$: $\|\mathbf{x}\|^2y^2$;
- for $\mathbf{y}^\top\mathbf{y}$: y^4.

Dummy and unmatched rows have $V = 0$, or are bounded in the same way. For the second claim, $\|m_{V,h}\| \le \Omega_V\sqrt{2/N}\|V_h\| + \sqrt{N}/2$ by Lemma 1.2. Then apply Minkowski's inequality over the family (e,h), and use $\sum_{e,h}\|V_{e,h}\|^2 \le \bar{n}_O\Gamma_\mathcal{E}$. □

Proposition 3.11 (Bounds on the terms)**.** *Assume (M1)–(M5), and let* $\tilde{Q}_{\max} = \max_j \tilde{Q}_j$*. Fix all public randomness: the expansions* a_h*, the common random strings* a, a_j*, and the dataset* D *in the clipped domain. Then each of the following bounds fails with probability at most* δ_B/n_{ev} *over*

$(s_R, s_O, u, \textit{all errors})$*; the bound on* T_6 *uses five of these events.*

$$\begin{aligned}
\|T_1\| &\le B_1 = \sqrt{\tfrac{2}{N}}\,\tfrac{\Omega_u}{q_1}\,\tfrac{N}{2q_2}\,\sqrt{2}\,\sigma_e\,M_V\,\tau_t,\\
\|T_2\| &\le B_2 = \sqrt{\tfrac{2}{N}}\,\tfrac{\Omega_u}{q_1}\,\tfrac{N}{2}\Big(\sqrt{|\mathcal{E}|/2} + \sqrt{v_s}\sqrt{2|\mathcal{E}|}\,\tfrac{\sqrt{N}}{2}\,\tau_t\Big),\\
\|T_3\| &\le B_3 = \sqrt{\tfrac{2}{N}}\,\tfrac{\Omega_u}{q_1}\,\sqrt{2}\,\sigma_e\,\big(|\mathcal{E}|\,L\,d_{\mathrm{num}}\big)^{1/2}\,\frac{\varsigma_\uparrow N^{3/2}\tilde{Q}_{\max}}{2P}\,\tau_t,\\
\|T_4\| &\le B_4 = \sqrt{\tfrac{2}{N}}\,\tfrac{\Omega_u}{q_1}\,\sqrt{|\mathcal{E}|}\Big(\tfrac{N}{2}\sum_{i<L} 2^{(L-1-i)/2} + \sqrt{v_s}\,L\,\tfrac{N^{3/2}}{2}\,\tau_t\Big),\\
\|T_5\| &\le B_5 = |\mathcal{E}|\tfrac{\sqrt{N}}{2} + \sqrt{v_s}\,|\mathcal{E}|\,\tfrac{N}{2}\,\tau_t,\\
\|T_6\| &\le B_6 = \big(N(v_u\beta_{\mathrm{pk}}^2 + \sigma_e^2 + \sigma_e^2\beta_s^2)\big)^{1/2} + \big(2t(v_u\gamma_{\mathrm{pk}}^2 + \sigma_e^2 + \sigma_e^2\gamma_s^2)\big)^{1/2},\\
\|T_7\| &\le B_7 = \sqrt{N}/2, \qquad \|T_8\| \le B_8 \quad (\textit{defined in the proof}).
\end{aligned}$$

Here $\beta_{\mathrm{pk}} = \sqrt{2}\sigma_e(\sqrt{N}+\sqrt{2t})$, $\gamma_{\mathrm{pk}} = (4\sigma_e^2 N\ell_N)^{1/2}$, $\beta_s = \sqrt{v_s}(\sqrt{N}+\sqrt{2t})$ *and* $\gamma_s = (2v_s N\ell_N)^{1/2}$.

Proof. Three facts are used throughout.

- By Lemma 1.7, the main part of $\sum_e \Pi_e(y_e)$ has norm $\sqrt{2/N}(\Omega_u/q_1)(\sum_e |\tau(y_e)_{k_e}|^2)^{1/2}$, because the entries occupy distinct slots.
- By Lemma 3.6, the polynomials $\tilde{D}$, rd_1, ε_1 and ε_1' are independent of the random variables they multiply. Corollary 3.3 therefore applies with them treated as fixed, and with no assumption on their distribution. This is the worst case over data required by the review.
- The deterministic roundings ε_0, ε_0' and rd_0 are bounded only through $\|\cdot\|_\infty \le \frac{1}{2}$.

T_1. By Lemma 1.5, $\tau(\mathrm{A}x)_k = \frac{N}{2}(x_0 + ix_{N/2})$. The coefficients $(fg)_0$ and $(fg)_{N/2}$ are inner products of f with signed permutations of g. So $e_R \mapsto (\frac{N}{2}x^R_{e,0}, \frac{N}{2}x^R_{e,N/2})_e$ is linear with Frobenius norm $\frac{N}{2q_2}\sqrt{2}(\sum_{e,h}\|m_{V_e,h}\|^2)^{1/2} \le \frac{N}{2q_2}\sqrt{2}M_V$, by Lemma 3.10. Entries that share R's ciphertext share e_R; the general form of Corollary 3.3 allows this. The coordinates of e_R are σ_e^2-subgaussian (Lemma 3.1).

T_2. The deterministic part satisfies $\frac{N}{2}|\varepsilon_{0,0} + i\varepsilon_{0,N/2}| \le \frac{N}{2}\cdot\frac{1}{\sqrt{2}}$ per entry. The random part is linear in s. Its $2|\mathcal{E}|$ rows have norm at most $\frac{N}{2}\|\varepsilon_{1,e}\| \le \frac{N}{2}\cdot\frac{\sqrt{N}}{2}$.

T_3. Apply Lemma 1.6 with $f = P^{-1}\tilde{D}_{e,i,j}$ and $\mathcal{H}' = \mathcal{H}_{i+1}$. This gives coefficient vectors with $\|\beta\| \le N\|f\|_2 \le N\cdot\sqrt{N}\varsigma_\uparrow\tilde{Q}_j/(2P)$. The coordinates of $e^{\mathrm{rk}} = e_R^{\mathrm{rk}} + e_O^{\mathrm{rk}}$ are $2\sigma_e^2$-subgaussian, and different (i,j) have independent errors. Sum the squared norms over (e,i,j).

T_4. For rd_0, Lemma 1.6 with $f=1$ gives $|\sum_{g\in\mathcal{H}'}\sigma(\mathrm{rd}_0)_{hg}| \le \|\mathrm{rd}_0\|_2\sqrt{N|\mathcal{H}'|} \le \frac{N}{2}\sqrt{|\mathcal{H}_{i+1}|}$. For $\mathrm{rd}_1 s$, take $f = \mathrm{rd}_1$: then $\|\beta\| \le N\|\mathrm{rd}_1\| \le N^{3/2}/2$ for each i, and at most $LN^{3/2}/2$ after summing over i for a fixed entry.

T_5. The deterministic part has $\|\sum_e \varepsilon'_{0,e}\| \le |\mathcal{E}|\sqrt{N}/2$. For the random part, $F = \sum_e \varepsilon'_{1,e}$ is independent of s, with $\|F\|_2 \le |\mathcal{E}|\sqrt{N}/2$ and $\|F\|_\infty^{\mathrm{can}} \le \|F\|_1 \le |\mathcal{E}|N/2$. Apply Corollary 3.3.

T_6. Condition on (e^{pk}, s), which is independent of the fresh (u, e_0, e_1). Apply Corollary 3.3 with $(f_j) = (e^{\mathrm{pk}}, 1, s)$. Then bound $\|e^{\mathrm{pk}}\|, \|e^{\mathrm{pk}}\|_\infty^{\mathrm{can}}, \|s\|, \|s\|_\infty^{\mathrm{can}}$ by Lemma 3.4, which costs four further events.

T_8. By Lemma 1.7, $\|T_8\| \le \frac{N}{2q_1}\sum_e\|Y_e\|$. We have $\|\mathrm{A}y\|_2 = \frac{N}{2}(y_0^2 + y_{N/2}^2)^{1/2}$, which is at most $\sqrt{N/2}(q_1/\Omega_u)(B_1+B_2)$ for each entry, and $\|T_{i+1}w\| \le |\mathcal{H}_{i+1}|\|w\|$. A union bound over the $|\mathcal{E}|L$ key switchings with $t' = t + \ln(|\mathcal{E}|L)$ bounds each $\|e_{\mathrm{ks}}\|$ by $\epsilon_{\mathrm{ks}} = \sqrt{2}\sigma_e d_{\mathrm{num}}\frac{N\tilde{Q}_{\max}}{2P}\tau_{t'} + \frac{\sqrt{N}}{2} + \sqrt{v_s}\frac{N}{2}\tau_{t'}$. This gives

$$B_8 = \tfrac{N}{2q_1}\big(|\mathcal{E}|\sqrt{N/2}\,(q_1/\Omega_u)(B_1+B_2) + |\mathcal{E}|\tfrac{N}{2}\epsilon_{\mathrm{ks}}\big).$$

□

Theorem 3.12 (Obligation O1 discharged)**.** *Under (M1)–(M5), for every dataset D in the clipped domain and every value of the public randomness,*

$$\Pr\big[\|e_{\mathrm{tot},s}\|_2 \le B_{\mathrm{tot},s} := B_1 + \cdots + B_8\big] \ge 1 - \delta_B,$$

with the probability taken over the key shares, all errors and the fresh randomness only. The centre is $f_s(D)$ of Definition 3.7. The bound covers the whole ring element, in the coefficient ℓ_2 norm.

Proof. Combine Proposition 3.8, Proposition 3.11, the triangle inequality and a union bound over the n_{ev} events. □

Corollary 3.13 (Sensitivity of the centre)**.** *Let $U_{\max,e}$ bound $|U_e|$: 1 for $\mathbf{b}$, and B_R for $\dot{\mathbf{x}}_j$. Put $\nu_U = (\sum_e U_{\max,e}^2)^{1/2}$ and $\nu_V = \sup_{D\sim D'}(\sum_e (V_e - V_e')_t^2)^{1/2} \le 2\Gamma_{\mathcal{E}}^{1/2}$. Then, with $\Omega_U = \Omega_V = \Omega_u = q_1 = q_2$,*

$$\sup_{D\sim D'} \|f_s(D) - f_s(D')\|_2 \le \widehat{\Delta}_{\mathrm{enc},s} := \sqrt{2/N}\,\Omega_{\mathrm{fin}}\Delta_s + \Delta_{x,s} + \sqrt{N},$$

$$\Delta_{x,s} = \sqrt{N/2}\,(\nu_V + \sqrt{N}\nu_U) + \tfrac{\sqrt{N}}{2}\sqrt{|\mathcal{E}|}\,\Delta_s + \tfrac{N^{3/2}}{4q_1}\sqrt{|\mathcal{E}|}(\nu_V + \sqrt{N}\nu_U).$$

Terms of order $N^{3/2}\sqrt{|\mathcal{E}|}/\Omega$, which come from the $\sqrt{N}/2$ parts of $\|m_U\|$, are omitted; they are below 2^{-20} for the example and are covered by adding 1 to $\Delta_{x,s}$. For the example, $\Delta_{x,2} \approx 2^{18.8}$ against $\sqrt{2/N}\,\Omega_{\mathrm{fin}}\Delta_2 \approx 2^{52.4}$. So $\widehat{\Delta}_{\mathrm{enc}}$ exceeds the specification's Δ_{enc} by a relative $2^{-33.7}$, and ρ_s by a relative $2^{-32.7}$.

Proof. Neighbours differ in one row, which sits in chunk h^* and slot l^*. Let $\Delta m_e = \mathsf{Ecd}(V_e[h^*]) - \mathsf{Ecd}(V_e'[h^*]) = \Omega_V\tau^{-1}(\Delta V_e \mathbf{u}_{l^*}) + \delta r_e$, with $\|\delta r_e\|_\infty \le 1$.

The slot value. By Lemma 1.5, $\mathrm{A}(m_{U_e,h^*}\Delta m_e)$ has all slots equal to

$$c_e = \sum_l \tau(m_U)_l \tau(\Delta m_e)_l = \Omega_U\Omega_V U_{e,l^*}\Delta V_e + \Omega_V\tau(r_U)_{l^*}\Delta V_e + \sum_l \tau(m_U)_l\tau(\delta r_e)_l.$$

The first term is $\Omega_U\Omega_V$ times the change of entry e, and these changes have ℓ_2 norm at most Δ_s. Bound $|\tau(r_U)_{l^*}| \le \|r_U\|_1 \le N/2$. By Cauchy–Schwarz and Lemma 1.1, the last sum is at most $\frac{N}{2}\|m_U\|\|\delta r\| \le \frac{N}{2}(\Omega_U U_{\max} + \frac{\sqrt{N}}{2})\sqrt{N}$.

Masking. Apply Lemma 1.7. The main part contributes $\sqrt{2/N}\,(\Omega_u/(q_1q_2))\|(c_e)_e\|$, and the spill contributes $\frac{\sqrt{N}}{2q_1q_2}\sum_e |c_e|$.

Rounding. Finally, $\|\lfloor x\rceil - \lfloor x'\rceil\| \le \|x - x'\| + \sqrt{N}$. □

3.5 Numerical evaluation

Take the example parameters: $N = 2^{14}$, $L = 13$, $d_R = 4$, $d_O = 6$, $\bar{n}_O = 10^6$ (so $m = 123$), $B_R = 5.59$, $B_O = 6.24$, $B_y = 2.89$, $\sigma_e = 3.2$, ternary s_R, s_O, u, $\delta_B = 2^{-50}$ (so $t = 37.1$ and $\tau_t = 9.62$), $d_{\mathrm{num}} = 2$ and $\varsigma_\uparrow = 1$. Then $|\mathcal{E}_G| = 51$, $|\mathcal{E}_2| = 63$ (joint mode), $\Gamma_{\mathcal{E}_G} = 1711$ and $\Gamma_{\mathcal{E}_2} = 2147$.

The budget is $\varepsilon = 1$ and $\delta_{\mathrm{run}} = 10^{-9}$, with three releases, $\kappa = 40$ and $f_1 = 0.1$. The conversion of Proposition 5.4 gives $\rho_{\mathrm{run}} = 0.01496$ and a stage-2 slot s.d. of 525.

The column $\log_2\Omega$ is the least scale meeting CALIBRATE's condition $\sigma_{c,s} \ge 2^\kappa B_{\mathrm{tot},s}/\sqrt{2\pi}$, with $\sigma_{c,s} = \Delta_{\mathrm{enc},s}/\sqrt{2\rho_s}$. The condition is binding for stage 2. Script: `scratch/btot.py`.

Configuration (stage 2)	B_1	B_2	B_3	B_4	B_5	B_6	$\log_2 B_{\mathrm{tot}}$	$\log_2 \Omega$
(A) Algorithm 4 as specified, $\tilde{Q}_{\max}/P = 1$	20.9	19.5	24.3	23.7	22.5	17.2	25.34	61.5
(A′) as (A), $P \geq 2^{20}\tilde{Q}_{\max}$	20.9	19.5	4.3	23.7	22.5	17.2	24.39	60.5
(C) EvalSum before the first rescale; stochastic rounding; one final rescale; $P \geq 2^{10}\tilde{Q}_{\max}$	20.9	14.1	< 0	< 0	14.0	17.2	21.07	57.2
(C′) as (C), with $\Omega_U = q_2 = 2^{60}$	13.4	14.1	< 0	< 0	14.0	17.2	17.56	53.7
Fresh-encryption error only (what v3.1 assumed)	–	–	–	–	–	17.2	17.17	53.3

Table 1: $\log_2$ of the bounds of Proposition 3.11 for stage 2 of the example. In all rows $B_7 = 2^6$ and $B_8 < 1$. For stage 1 the required scale is 1.4 to 1.6 bits lower; for (A) it is 59.9. Using Lemma 2.3(a) instead of (b) adds 0.33 bits.

Reading the table.

- The specification's $\log_2 \Omega \approx 52.5$ floods only the fresh-encryption error, measured by its *mean* ℓ_2 norm of about $2^{16.4}$. A high-probability bound on that term alone already needs 53.3.
- For the circuit exactly as specified, the worst-case analysis is dominated by three terms:
  - the gadget term B_3, unless $P \gg \tilde{Q}_{\max}$;
  - the ModDown roundings B_4;
  - the $|\mathcal{E}|$ separate second rescales B_5.

  The latter two are large because roundings are treated as adversarially aligned with a slot. That is the price of a bound with no randomness assumption on rounding errors.
- Variant (C) removes these terms in three ways:
  - EvalSum at level 2, so that key-switching errors are divided by q_2 and the first-rescale rounding is not traced;
  - stochastic rounding, so that rounding errors are mean-zero and $\frac{1}{4}$-subgaussian given the ciphertext, by Lemma 3.1(c);
  - a single rescale after summing the masked products.

  (C′) also shrinks T_1 by encoding R's vectors at a larger scale.

The bounds for (C) and (C′) use the same lemmas. We give the formulas in the script rather than repeating the proofs. [**Not fully rigorous:** The proofs for variants (C) and (C′) are only sketched. The T_2 and T_5 formulas for (C) apply Corollary 3.3 conditionally on s, to the independent stochastic-rounding coins.]

3.6 Empirical sanity check with OpenFHE

Setup. We ran the circuit shape of Algorithm 4 in OpenFHE [1], through its Python bindings, (`scratch/ofhe_check.py`). The run used:

- a single key, FIXEDMANUAL, $N = 2^{14}$, scale 2^{50}, depth 2;
- one entry, with $m \in \{4, 16\}$ chunks;
- 0/1 vectors U and $|V| \leq o/2$.

Two runs on the same data with independent keys and randomness isolate the randomness-dependent error, because the deterministic part cancels in the difference.

Results.

- The measured coefficient ℓ_2 norm of the random error was $2^{14.0}$ for $m = 4$ and $2^{15.7}$ for $m = 16$. Theorem 3.12 at the same setting gives $2^{21.8}$ and $2^{21.9}$. The bound holds, with a margin of 6

to 8 bits, as expected for a worst-case, high-probability bound.

- The deviation of the decryption from $\mathsf{Ecd}(\mathbf{z})$ was $2^{23.0}$. It is dominated by a deterministic relative error of $8.7 \cdot 10^{-10} \approx 2^{-30}$ at the target slot, from decoding with the nominal scale 2^{50} rather than the exact Ω_{fin}. This confirms Remark 3.9 and Correction 2.

4 Decoupled scales, OpenFHE's key switching and the decoding offset

The evaluation runs circuit (A) with *decoupled* encoding scales ($\Omega_U = 2^{68}$, $\Omega_u = q_1$, Ω_V chosen by CALIBRATE) and with the hybrid key switching of OpenFHE 1.5.1. Proposition 3.11 and Theorem 3.12 are already stated for general $\Omega_U, \Omega_V, \Omega_u, q_1, q_2$: no step of their proofs uses $\Omega_U = q_2$ or $\Omega_V = \Omega_u = q_1$. Three things are not covered by Sections 3.2–3.5: OpenFHE's uncentred gadget (Lemma 4.1), the sensitivity of the centre with decoupled scales (Lemma 4.3; Corollary 3.13 assumes equal scales), and a public bound on the decoding offset (Proposition 4.5), which the statistical part needs. This section proves them. The proofs were first written in the companion book on the protocol (its Appendix B, Lemmas `btot.gadget` and `btot.sensZ`, Proposition `btot.offset`, and Chapter 9, Corollary `agg.sensZ` and Lemma `agg.offset`); they were checked here step by step, with the corrections noted, and the numbers recomputed (script `experiments/proofs_check/check_btot_offset.py` of the repository).

4.1 OpenFHE's uncentred key switching

(M3′) As (M3), with three differences. (i) The digit $\tilde{D}_j(c)$ is the integer polynomial $\sum_{i \in j} [c_i (\tilde{Q}_j/q_i)^{-1}]_{q_i} (\tilde{Q}_j/q_i)$, with residues $[\cdot]_{q_i} \in [0, q_i)$; so $\tilde{D}_j(c) \equiv c \pmod{\tilde{Q}_j}$ and $\tilde{D}_j(c) \in [0, k_j \tilde{Q}_j)^N$ (fast base extension without centring). (ii) ModDown is $\mathrm{MD}(x) = (x - y_P)/P$, with $y_P = \sum_i [x_i (P/p_i)^{-1}]_{p_i} (P/p_i) \in [0, k_P P)^N$ and $y_P \equiv x \pmod{P}$, where $p_1, \ldots, p_{k_P}$ are the special primes; so $\mathrm{MD}(x) = x/P + \mathrm{rd}$ with $\mathrm{rd} = -y_P/P \in (-k_P, 0]^N$, the division being exact. (iii) The key switch from s to $\psi_g^{-1}(s)$ is applied *before* the automorphism: the key is $b_j = -a_j \psi_g^{-1}(s) + P \mathsf{g}_j s + e_j^{\mathrm{rk}}$, and the rotation outputs $(\psi_g(c_0 + \mathrm{MD}(\sum_j \tilde{D}_j(c_1) b_j)),\ \psi_g(\mathrm{MD}(\sum_j \tilde{D}_j(c_1) a_j)))$.

In the notation of Proposition 3.11, (M3′) gives $\|\tilde{D}_j\|_\infty < k_j \tilde{Q}_j \le \varsigma_\uparrow \tilde{Q}_{\max}/2$ with $\varsigma_\uparrow = 2 \max_j k_j$, and $\|\mathrm{rd}\|_\infty < k_P = \varsigma_\downarrow/2$ with $\varsigma_\downarrow = 2k_P$, against $\varsigma_\uparrow \le \max_j k_j$ and $\varsigma_\downarrow = 1$ for the centred gadget of (M3).

Lemma 4.1 (Uncentred digits and ModDown rounding)**.** *Assume (M1), (M2), (M3′), (M4), (M5). Then Proposition 3.8 holds verbatim with every $\tilde{D}_{e,i,j}$, $e_j^{\mathrm{rk},(i)}$, $\mathrm{rd}_{0,e,i}$, $\mathrm{rd}_{1,e,i}$ replaced by its image under the automorphism ψ_{g_i} of rotation i. Proposition 3.11 and Theorem 3.12 hold with $\varsigma_\uparrow = 2 \max_j k_j$ in B_3 and ϵ_{ks}, and with the ModDown terms (all of B_4, and the terms $\sqrt{N}/2$ and $\sqrt{v_s} \frac{N}{2} \tau_{t'}$ of ϵ_{ks}) multiplied by $\varsigma_\downarrow = 2k_P$. B_1, B_2, B_5, B_6, B_7 are unchanged. No mean-zero or symmetry property of the digits or of the roundings is used.*

Proof. The switch. Let $c_0 + c_1 s \equiv \mu \pmod{Q_1}$. Modulo PQ_1, $\sum_j \tilde{D}_j(c_1)(b_j + a_j \psi_g^{-1}(s)) = P \sum_j \tilde{D}_j(c_1) \mathsf{g}_j s + \sum_j \tilde{D}_j(c_1) e_j^{\mathrm{rk}}$, and $\sum_j \tilde{D}_j(c_1) \mathsf{g}_j \equiv c_1 \pmod{Q_1}$ by the CRT, because $\tilde{D}_j(c_1) \equiv c_1 \pmod{\tilde{Q}_j}$ holds for uncentred digits as well. $\mathrm{MD}(x) \bmod Q_1$ depends only on $x \bmod PQ_1$ (adding $PQ_1 K$ to x adds $Q_1 K$ to $\mathrm{MD}(x)$). Dividing by P and adding the roundings, the switched pair

(c_0', c_1') satisfies

$$c_0' + c_1'\psi_g^{-1}(s) \equiv \mu + e'_{\mathrm{ks}} \pmod{Q_1}, \qquad e'_{\mathrm{ks}} = P^{-1}\sum_j \tilde{D}_j(c_1)e_j^{\mathrm{rk}} + \mathrm{rd}_0 + \mathrm{rd}_1\psi_g^{-1}(s).$$

The automorphism afterwards. $\psi_g(c_0') + \psi_g(c_1')s = \psi_g(c_0' + c_1'\psi_g^{-1}(s))$, so the rotated pair decrypts to $\psi_g(\mu) + \psi_g(e'_{\mathrm{ks}})$ under s, and

$$\psi_g(e'_{\mathrm{ks}}) = P^{-1}\sum_j \tilde{D}_j^\star e_j^\star + \mathrm{rd}_0^\star + \mathrm{rd}_1^\star s, \qquad \tilde{D}_j^\star = \psi_g(\tilde{D}_j(c_1)),\ e_j^\star = \psi_g(e_j^{\mathrm{rk}}),\ \mathrm{rd}_b^\star = \psi_g(\mathrm{rd}_b).$$

This is the form of Lemma 3.5. The proof of Proposition 3.8 only adds such errors and applies the traces T_{i+1} to them, so the decomposition holds with the starred quantities.

The hypotheses of the bounds. (a) ψ_g is a signed permutation of coefficients (Lemma 1.4), so the starred digits and roundings have the same ℓ_∞, ℓ_2 and ℓ_1 norms as the unstarred ones. (b) The coordinates of $e_j^\star$ are those of e_j^{rk} up to order and sign; a sign flip preserves v-subgaussianity, so $e_j^\star$ has independent $2\sigma_e^2$-subgaussian coordinates, independent across (i, j). (c) $\tilde{D}_j^\star$ and $\mathrm{rd}_1^\star$ are functions of the c_1-component, which by the induction of Lemma 3.6 (the new step outputs $\psi_g(\mathrm{MD}(\sum_j \tilde{D}_j(c_1)a_j))$ and uses neither c_0, nor a secret, nor an error) is a function of the public randomness and the data only. (d) The rescales are those of (M1). The proofs of B_3, B_4 and B_8 use only $\|\tilde{D}\|_\infty$, $\|\mathrm{rd}\|_\infty$, the proxy of the key errors, (c), and Lemma 1.6, which holds for every polynomial f; so they hold with the stated factors. □

Remark 4.2 (Conformance of OpenFHE 1.5.1 to (M3′); a code reading, not a proof)**.** We re-read the relevant code at tag `v1.5.1` (the headers installed with the implementation are identical to the tagged ones):

- `core/include/lattice/hal/default/dcrtpoly-impl.h`, `ApproxSwitchCRTBasis`: residues are stored in $[0, q_i)$; both code paths compute $\sum_i [x_i\hat{Q}_i^{-1}]_{q_i}[\hat{Q}_i]_p \bmod p$, the residue of the integer $\sum_i [x_i\hat{Q}_i^{-1}]_{q_i}\hat{Q}_i \in [0, k\tilde{Q})$. `ApproxModDown`: `ans` $= (x - \mathrm{partPSwitchedToQ}) \cdot P^{-1}$, with no centring; the t-multiplications are skipped for CKKS, whose noise scale is 1 (`gen-cryptocontext-ckksrns-internal.h`).
- `pke/lib/keyswitch/keyswitch-hybrid.cpp`: `KeySwitchGenInternal` sets $b = -a\,s_{\mathrm{new}} + e + [P]_{q_i}s_{\mathrm{old}}$ on the primes of digit j and $-a\,s_{\mathrm{new}} + e$ elsewhere, which is $P\mathbf{g}_j s_{\mathrm{old}}$; `EvalKeySwitchPrecomputeCore` keeps the digit's own residues and extends them by `ApproxSwitchCRTBasis`; `EvalFastKeySwitchCore` applies `ApproxModDown` to both components; `KeySwitchInPlace` adds the first to c_0 and replaces c_1 by the second.
- `pke/lib/schemebase/base-leveledshe.cpp`, `EvalAutomorphism`: `KeySwitchInPlace`, then `AutomorphismTransform` of both components; CKKS does not override it. `base-multiparty.cpp`, `MultiEvalAutomorphismKeyGen`: the new key is $\psi_{g^{-1}}(s_P)$.
- The implementation (`implementation/native/vfr_native.cpp`) builds the rotation-key shares in this layout itself (`rotkey_share`: $b_j = -a_j\psi_g^{-1}(s) + e + P\mathbf{g}_j s$), sums them, calls the scheme's `EvalAutomorphism` inside its own EvalSum, and uses its own centred deterministic rescale (`rescale_poly`, remainder in $(-q/2, q/2]$), its own encodings and its own release. So (M1) holds as stated and only key switching is OpenFHE's.

These are statements about a program and are not proved. (M5) for OpenFHE's table sampler, which draws the errors, remains a numerical check (Section 3.6; the paper's Check (T4)).

4.2 The sensitivity of the centre with decoupled scales

For release s let $\nu_U = (\sum_{e\in\mathcal{E}} U_{\max,e}^2)^{1/2}$, with $U_{\max,e} = 1$ for $\mathbf{b}$ and B_R for a column of $\dot{\mathbf{x}}$, and $\nu_V = \sup_{D\sim D'} \max_t (\sum_{e\in\mathcal{E}} (V_{e,t} - V'_{e,t})^2)^{1/2} \le 2\Gamma_{\mathcal{E}}^{1/2}$ (Lemma 3.10).

Lemma 4.3 (Sensitivity of the centre, decoupled scales)**.** *Assume (M1), (M2) with arbitrary $\Omega_U, \Omega_V, \Omega_u, q_1, q_2$, and that* Ecd *rounds to the nearest integer coefficient-wise. Put*

$$\begin{aligned}
\epsilon_x &= \frac{\Omega_u}{q_1 q_2}\,\frac{N^{3/2}\sqrt{|\mathcal{E}|}}{2\sqrt{2}} + \frac{N^{5/2}|\mathcal{E}|}{8 q_1 q_2},\\
\Delta_{x,s} &= \sqrt{\tfrac{N}{2}}\,\frac{\Omega_u}{q_1 q_2}\big(\Omega_V \nu_V + \sqrt{N}\,\Omega_U \nu_U\big) + \frac{\sqrt{N}}{2}\,\frac{\Omega_U \Omega_V}{q_1 q_2}\sqrt{|\mathcal{E}|}\,\Delta_s\\
&\quad + \frac{N^{3/2}}{4 q_1 q_2}\sqrt{|\mathcal{E}|}\big(\Omega_V \nu_V + \sqrt{N}\,\Omega_U \nu_U\big) + 1.
\end{aligned}$$

For every pair of neighbours $D \sim D'$ (differing in one row of O*, matched or not) and every input and alignment of* R*, $\|\widehat{Z}_s(D) - \widehat{Z}_s(D')\|_2 \le \sqrt{2/N}\,\Omega_{\mathrm{fin}}\Delta_s + \Delta_{x,s} - 1 + \epsilon_x$; hence, if $\epsilon_x \le 1$,*

$$\|f_s(D) - f_s(D')\|_2 \le \widehat{\Delta}_{\mathrm{enc},s} := \sqrt{2/N}\,\Omega_{\mathrm{fin}}\Delta_s + \Delta_{x,s} + \sqrt{N}.$$

With equal scales this reduces to Corollary 3.13.

Proof. The changed row sits in chunk h^* and slot l^*; only the encodings m_{V_e,h^*} change. Write $\Delta V_e = V_{e,l^*} - V'_{e,l^*}$ and $\Delta m_e = m_{V_e,h^*} - m'_{V_e,h^*} = \Omega_V \tau^{-1}(\Delta V_e \mathbf{u}_{l^*}) + \delta r_e$ with $\|\delta r_e\|_\infty \le 1$. Then $\widehat{Z}_s(D) - \widehat{Z}_s(D') = \sum_e (q_1 q_2)^{-1} y_e m_{u_e}$ with $y_e = \mathrm{A}(m_{U_e,h^*}\Delta m_e)$.

Slot value. By Lemma 1.5 and multiplicativity of τ, every slot of y_e equals $c_e = \sum_l \tau(m_{U_e,h^*})_l \tau(\Delta m_e)_l$. With $m_{U_e,h^*} = \Omega_U \tau^{-1}(U_e[h^*]) + r_{U_e}$,

$$c_e = \Omega_U\Omega_V\, U_{e,l^*}\Delta V_e + \Omega_V\, \tau(r_{U_e})_{l^*}\Delta V_e + \sum_l \tau(m_{U_e,h^*})_l \tau(\delta r_e)_l.$$

The vector $(U_{e,l^*}\Delta V_e)_e$ is the change of the statistic, of norm at most Δ_s; it is zero if the row is unmatched, because R's vectors vanish there. $|\tau(r_{U_e})_{l^*}| \le \|r_{U_e}\|_1 \le N/2$, so the second terms have norm at most $\Omega_V \frac{N}{2}\nu_V$. By Cauchy–Schwarz and Lemma 1.1 the third term is at most $\frac{N}{2}\|m_{U_e,h^*}\|\|\delta r_e\| \le \frac{N}{2}(\Omega_U U_{\max,e} + \frac{\sqrt{N}}{2})\sqrt{N}$, using $\|m_U\| \le \Omega_U\sqrt{2/N}\|U_e[h^*]\| + \sqrt{N}/2$ and $\|U_e[h^*]\| \le \sqrt{N/2}\,U_{\max,e}$. So

$$\|(c_e)_e\| \le \Omega_U\Omega_V\Delta_s + \tfrac{N}{2}\big(\Omega_V\nu_V + \sqrt{N}\,\Omega_U\nu_U\big) + \tfrac{N^2}{4}\sqrt{|\mathcal{E}|}.$$

Masking. Split $m_{u_e} = \Omega_u \tau^{-1}(\mathbf{u}_{k_e}) + r_{u_e}$. The main part has norm $\sqrt{2/N}\,(\Omega_u/(q_1q_2))\|(c_e)_e\|$ (Lemma 1.7). For the spill, y_e has only the coefficients of 1 and $X^{N/2}$ (Lemma 1.5): $y_e = a + bX^{N/2}$ with $a^2 + b^2 = |c_e|^2$. Multiplication by $X^{N/2}$ is an orthogonal map J with $J^\top = -J$, so $\langle r, Jr\rangle = 0$ and $\|y_e r_{u_e}\|^2 = |c_e|^2\|r_{u_e}\|^2 \le |c_e|^2 N/4$. Hence the spill is at most $\frac{\sqrt{N}}{2q_1q_2}\sum_e |c_e| \le \frac{\sqrt{N}}{2q_1q_2}\sqrt{|\mathcal{E}|}\|(c_e)_e\|$.

Collecting. With $\Omega_U\Omega_V\Omega_u/(q_1q_2) = \Omega_{\mathrm{fin}}$, the main part gives $\sqrt{2/N}\,\Omega_{\mathrm{fin}}\Delta_s$, the first term of $\Delta_{x,s}$, and $\sqrt{2/N}(\Omega_u/(q_1q_2))\frac{N^2}{4}\sqrt{|\mathcal{E}|}$; the spill gives the second and third terms and $\frac{\sqrt{N}}{2q_1q_2}\frac{N^2}{4}|\mathcal{E}|$. The two leftovers sum to ϵ_x. Finally $\|\lfloor x\rceil - \lfloor x'\rceil\| \le \|x - x'\| + \sqrt{N}$. □

Remark 4.4 (Second reading)**.** The book states the lemma for a changed row that is matched; the proof covers any row, as shown (an unmatched row contributes ΔV_e only through the rounding

cross terms), and the paper's adjacency (Definition `def:adj`) changes a row "matched or not". A brute-force check at $N = 16$ with decoupled scales $(\Omega_U, \Omega_V, \Omega_u, q_1, q_2) = (2^{10}, 2^9, 2^{10}, 2^{10}, 2^9)$ and exact rational arithmetic gives a largest ratio 0.93 of the change of f_s to the bound over 30 random instances. If the encoder's output differs from exact rounding by at most δ per coefficient, every rounding bound $\frac{1}{2}$ becomes $\frac{1}{2} + \delta$ and the lemma holds with the first three terms of $\Delta_{x,s}$ multiplied by $(1 + 2\delta)^2$; the implementation's double-double encoder is not analysed here.

4.3 The decoding offset

Let $\zeta_{\mathcal{E}}^2 = \sup_{\text{rows}} \sum_e (U_{e,t} V_{e,t})^2$, the largest squared contribution of one row to the statistic: with $r = B_R^2$, $\zeta_{\mathcal{E}_G}^2 \le o^2 + ro$ and $\zeta_{\mathcal{E}_2}^2 \le o^2 + r(o + w) + ow + w^2$ (joint mode), by the block computation of Lemma 3.10. Let $n_{\max}$ bound the number of matched rows ($n_R^{\max}$ in the protocol).

Proposition 4.5 (A public bound on the decoding offset)**.** *Assume (M1), (M2) and exact rounding in* Ecd*. Put*

$$A_s = \frac{N\sqrt{m}}{2\sqrt{2}} \Big(\frac{\sqrt{n_{\max}}\, \nu_U}{\Omega_V} + \frac{(\bar{n}_O \Gamma_{\mathcal{E}})^{1/2}}{\Omega_U} \Big) + \frac{\sqrt{|\mathcal{E}|}\, mN^2}{8\, \Omega_U \Omega_V}.$$

For every dataset in the clipped domain with at most $n_{\max}$ matched rows, the decoded centre satisfies, on the slots of the entries,

$$\Big\| \big(\mathsf{Dcd}(f_s(D))_{k_e} - \mathbf{z}_e(D) \big)_{e \in \mathcal{E}} \Big\|_2 \le \bar{d}_s, \qquad \bar{d}_s := A_s + \frac{N\sqrt{|\mathcal{E}|}}{2\sqrt{2}\, \Omega_u} \big(n_{\max} \zeta_{\mathcal{E}} + A_s \big) + \frac{N}{2\sqrt{2}\, \Omega_{\text{fin}}},$$

with $\mathsf{Dcd}(v) = \tau(v)/\Omega_{\text{fin}}$*, and, in the coefficient domain,*

$$\|\widehat{Z}_s(D) - \Omega_{\text{fin}} \tau^{-1}(\mathbf{z}_s(D))\|_2 \le d_{\max},$$

$$d_{\max} := \sqrt{\tfrac{2}{N}}\, \frac{\Omega_u}{q_1 q_2}\, \bar{r}_s + \frac{\sqrt{N}}{2 q_1 q_2} \sqrt{|\mathcal{E}|} \big(\Omega_U \Omega_V n_{\max} \zeta_{\mathcal{E}} + \bar{r}_s \big), \qquad \bar{r}_s = \Omega_U \Omega_V A_s.$$

Proof. Write $m_{U_e,h} = \Omega_U \tau^{-1}(U_e[h]) + r_{U,e,h}$ and $m_{V_e,h} = \Omega_V \tau^{-1}(V_e[h]) + r_{V,e,h}$, with roundings of sup-norm at most $\frac{1}{2}$, so $\|\tau(r)\| = \sqrt{N/2}\|r\| \le N/(2\sqrt{2})$. As in the previous proof, every slot of $y_e = \mathrm{A}(\sum_h m_{U_e,h} m_{V_e,h})$ equals

$$c_e = \Omega_U \Omega_V \mathbf{z}_e + \bar{r}_e, \qquad \bar{r}_e = \Omega_U a_e + \Omega_V b_e + d_e,$$

with $a_e = \sum_h \langle U_e[h], \tau(r_{V,e,h}) \rangle$, $b_e = \sum_h \langle \tau(r_{U,e,h}), V_e[h] \rangle$, $d_e = \sum_h \langle \tau(r_{U,e,h}), \tau(r_{V,e,h}) \rangle$ (bilinear sums over the slots). By Cauchy–Schwarz over slots and chunks, $|a_e| \le \frac{N}{2\sqrt{2}} \sqrt{m} \|U_e\|$, $|b_e| \le \frac{N}{2\sqrt{2}} \sqrt{m} \|V_e\|$, $|d_e| \le mN^2/8$. R's vectors vanish off the matched rows, so $\sum_e \|U_e\|^2 \le n_{\max} \nu_U^2$, and $\sum_e \|V_e\|^2 \le \bar{n}_O \Gamma_{\mathcal{E}}$ (Lemma 3.10); hence $\|(\bar{r}_e)_e\| \le \Omega_U \Omega_V A_s = \bar{r}_s$.

Slots. Since $\tau(y_{e'} r_{u_{e'}}) = c_{e'} \tau(r_{u_{e'}})$, $\tau(\widehat{Z}_s)_k = (q_1 q_2)^{-1} \sum_{e'} c_{e'} (\Omega_u [k = k_{e'}] + \tau(r_{u_{e'}})_k)$. Dividing by Ω_{fin},

$$\mathsf{Dcd}(\widehat{Z}_s)_{k_e} - \mathbf{z}_e = \frac{\bar{r}_e}{\Omega_U \Omega_V} + \sum_{e'} \frac{c_{e'}\, \tau(r_{u_{e'}})_{k_e}}{\Omega_U \Omega_V \Omega_u}.$$

The first term has norm at most A_s over the entries. For the second, $\|(\tau(r_{u_{e'}})_{k_e})_e\| \le \|\tau(r_{u_{e'}})\| \le N/(2\sqrt{2})$, so it is at most

$$\tfrac{N}{2\sqrt{2}\Omega_u} \textstyle\sum_{e'} |c_{e'}|/(\Omega_U \Omega_V) \le \tfrac{N\sqrt{|\mathcal{E}|}}{2\sqrt{2}\Omega_u} (\|\mathbf{z}\| + A_s),$$

and $\|\mathbf{z}\| \le n_{\max}\zeta_{\mathcal{E}}$ by the triangle inequality over the matched rows. Finally $f_s - \widehat{Z}_s$ has sup-norm at most $\frac{1}{2}$, so its slot vector has norm at most $N/(2\sqrt{2})$, which contributes $N/(2\sqrt{2}\,\Omega_{\text{fin}})$. Taking real parts does not increase norms.

Coefficients. The main part of $\widehat{Z}_s$ is

$$\textstyle\sum_e (\Omega_u/(q_1q_2))y_e\tau^{-1}(\mathbf{u}_{k_e}) = \Omega_{\text{fin}}\tau^{-1}(\mathbf{z}_s) + (\Omega_u/(q_1q_2))\tau^{-1}(\sum_e \bar{r}_e\mathbf{u}_{k_e}),$$

whose second piece has norm $\sqrt{2/N}(\Omega_u/(q_1q_2))\|(\bar{r}_e)\|$ (Lemma 1.2(a)). The spill is at most $\frac{\sqrt{N}}{2q_1q_2}\sum_e |c_e|$, as in the previous proof. □

The numbers at the evaluation setting. Chain $[60, 59, 59, 59]$ with $q_0 = p_0p_1 \approx 2^{119}$, $q_1, q_2 \approx 2^{59}$, two special primes of 60 bits, $d_{\text{num}} = 2$ digits at the EvalSum level ($\max_j k_j = 2$, $\tilde{Q}_{\max}/P \approx 2^{-1}$), $\Omega_U = 2^{68}$, $\Omega_u = q_1$, simulation population ($d_R = 4$, $d_O = 6$, $(B_R, B_O, B_y) = (5.59, 6.24, 2.89)$), $\varepsilon = 1$, $\delta = 10^{-9}$, $\kappa = 40$, previous setting ($f_1 = 0.1$), recomputed with the implementation's Calibrate:

- $\log_2 \Omega_{\text{fin}} = 62.75$; $B_{\text{tot},1} = 2^{26.43}$ and $B_{\text{tot},2} = 2^{26.59}$ by Lemma 4.1 ($\varsigma_\uparrow = 4$, $\varsigma_\downarrow = 4$), the same at $\bar{n}_O = 4{,}096$, 16,384 and 10^6;
- $\Delta_{x,2} = 2^{27.57}$ and $\Delta_{x,1} = 2^{27.46}$; $\epsilon_x = 2^{-36.5}$ and $2^{-36.7}$, so the "+1" of $\Delta_{x,s}$ covers it;
- $\bar{d}_2 = 9.9 \cdot 10^{-9}$ at $\bar{n}_O = 4{,}096$ with $n_{\max} = n_R^{\max} = 2{,}000$; $2.0 \cdot 10^{-8}$ with $n_{\max} = \bar{n}_O = 4{,}096$; $1.6 \cdot 10^{-8}$ and $4.8 \cdot 10^{-6}$ at $\bar{n}_O = 10^6$ with $n_{\max} = 2{,}000$ and 10^6; against $\sigma_2 = 525$;
- $d_{\max} = 2^{30.3}$ at $\bar{n}_O = 10^6$ with $n_{\max} = 2{,}000$, and $2^{38.6}$ with $n_{\max} = 10^6$.

A brute-force check at $N = 16$ with decoupled scales gives a largest ratio 0.07 of the offset to $\bar{d}_s$.

5 Output privacy (Theorem 1)

5.1 Differential-privacy tools

Definition 5.1 (zCDP, Bun and Steinke [3])**.** A randomized mechanism $\mathcal{M}$ is ρ-zCDP with respect to $\sim$ if $D_\alpha(\mathcal{M}(D)\|\mathcal{M}(D')) \le \rho\alpha$ for all $D \sim D'$ and all $\alpha > 1$. Here D_α is the Rényi divergence of order α.

Lemma 5.2 ([3], composition and post-processing)**.**

- Composition. *Suppose $\mathcal{M}_1$ is ρ_1-zCDP, and $\mathcal{M}_2(\cdot, y)$ is ρ_2-zCDP for every y. Then $D \mapsto (y, \mathcal{M}_2(D, y))$ with $y \leftarrow \mathcal{M}_1(D)$ is $(\rho_1 + \rho_2)$-zCDP.*
- Post-processing. *$g \circ \mathcal{M}$ is ρ-zCDP for every randomized g that does not depend on D.*

Proposition 5.3 (The ideal mechanism; multivariate discrete Gaussian)**.** *Let $f : \mathcal{D} \to \mathbb{Z}^N$ satisfy $\sup_{D\sim D'}\|f(D) - f(D')\|_2 \le \widehat{\Delta}$, and let $\mathcal{M}(D) = f(D) + \eta$ with $\eta \sim D_{\mathbb{Z}^N,\sigma}$. Then $\mathcal{M}$ is $\widehat{\Delta}^2/(2\sigma^2)$-zCDP.*

Hence $\mathcal{M}^{\text{id}}(D) = (f_s(D) + \eta_s)_s$, with independent $\eta_s \sim D_{\mathbb{Z}^N,\sigma_{c,s}}$ and f_s from Definition 3.7, is $\sum_s \rho_s$-zCDP, with $\rho_s = \widehat{\Delta}^2_{\text{enc},s}/(2\sigma^2_{c,s})$.

Proof. Single release. $\mathcal{M}(D)$ and $\mathcal{M}(D')$ are integer shifts of one another, by $\mathbf{e} = f(D) - f(D') \in \mathbb{Z}^N$. Lemma 2.2 gives $D_\alpha \le \alpha\|\mathbf{e}\|^2/(2\sigma^2)$. This is the multivariate statement of [4] [*citation detail unverified:* numbered Theorem 14 in the arXiv version], proved here in a self-contained way.

Several releases. The releases are computed from fixed ciphertexts M_1, M_2, M_y. R's choices between releases do not change what O releases, so non-adaptive composition (Lemma 5.2) applies.

The centre. Corollary 3.13 bounds the sensitivity of f_s. The specification's $\Delta_{\text{enc},s}$ is the sensitivity of $\mathsf{Ecd}_\Omega(\mathbf{z}_s)$ (Lemma 1.2); it must be replaced by $\widehat{\Delta}_{\text{enc},s}$, a relative change of 2^{-33}. □

Proposition 5.4 (Conversion from zCDP, Canonne, Kamath and Steinke [4])**.** *If $\mathcal{M}$ is ρ-zCDP, then it is $(\varepsilon, \delta(\rho,\varepsilon))$-DP for every $\varepsilon \geq 0$, where*

$$\delta(\rho,\varepsilon) = \inf_{\alpha>1} \frac{e^{(\alpha-1)(\alpha\rho-\varepsilon)}}{\alpha-1}\Big(1-\frac{1}{\alpha}\Big)^{\alpha}.$$

The map $\rho \mapsto \delta(\rho,\varepsilon)$ is continuous and nondecreasing, and $\delta(\rho,\varepsilon) \to 0$ as $\rho \to 0$. Hence $\rho_{\mathrm{run}} = \max\{\rho : \delta(\rho,\varepsilon_{\mathrm{run}}) \leq \delta_{\mathrm{DP}}\}$ in Calibrate *is well defined, and the ρ_{run}-zCDP mechanism is $(\varepsilon_{\mathrm{run}}, \delta_{\mathrm{DP}})$-DP.*

Proof. Let $P = \mathcal{M}(D)$ and $Q = \mathcal{M}(D')$, with $D \sim D'$.

Hockey-stick form. (ε,δ)-indistinguishability is equivalent to $\mathbb{E}_Q[(P/Q - e^{\varepsilon})_+] \leq \delta$ (with the usual convention when $P \not\ll Q$; here both have full support on $\mathbb{Z}^N$).

A pointwise bound. For $x > 0$ and $\alpha > 1$, $(x - e^{\varepsilon})_+ \leq x^{\alpha} c_{\alpha}$, where

$$c_{\alpha} = \sup_{y>0} \frac{(y-e^{\varepsilon})_+}{y^{\alpha}} = \frac{e^{-(\alpha-1)\varepsilon}}{\alpha-1}\Big(1-\frac{1}{\alpha}\Big)^{\alpha}.$$

The supremum is attained at $y = \alpha e^{\varepsilon}/(\alpha-1)$.

Conclusion. Hence $\mathbb{E}_Q[(P/Q - e^{\varepsilon})_+] \leq c_{\alpha}\mathbb{E}_Q[(P/Q)^{\alpha}] = c_{\alpha} e^{(\alpha-1)D_{\alpha}(P\|Q)} \leq c_{\alpha} e^{(\alpha-1)\alpha\rho}$. Take the infimum over α. Each term is continuous and increasing in ρ, so the infimum is nondecreasing and upper semicontinuous. It is also continuous here, because the infimum may be restricted to a compact set of α for ρ in a compact range. As $\rho \to 0$, taking $\alpha = 1 + \sqrt{\varepsilon/\rho}$ shows $\delta \to 0$ for $\varepsilon > 0$. □

Remark 5.5. The conversion is valid for every ρ-zCDP mechanism, but it is not the tightest possible for the Gaussian mechanism. For the example, $\varepsilon = 1$ and $\delta_{\mathrm{DP}} \approx 10^{-9}$ give $\rho_{\mathrm{run}} = 0.01496$. This agrees with the specification.

5.2 SIM-CDP

Let λ be the computational security parameter. $\mathrm{view}_{\mathsf{R}}(D)$ comprises three things: R's input (cohort identifiers and $\mathbf{X}_R$), its random tape, and every message it receives in one run on dataset D.

Definition 5.6 (SIM-CDP for R, after Mironov et al. [16])**.** The protocol is (ε,δ)-*SIM-CDP for* R *with statistical slack* γ if two objects exist:

- a randomized mechanism $\mathcal{M}$ that is (ε,δ)-DP with respect to $\sim$;
- a probabilistic polynomial-time simulator $\mathcal{S}$.

They must satisfy, for every dataset D and every PPT distinguisher $\mathcal{A}$,

$$\big|\Pr[\mathcal{A}(\mathrm{view}_{\mathsf{R}}(D)) = 1] - \Pr[\mathcal{A}(\mathcal{S}(\mathcal{M}(D))) = 1]\big| \leq \gamma + \mathrm{negl}(\lambda).$$

With $\gamma = 0$ this is the simulation-based notion of [16], which those authors state for pure ε [*citation detail unverified:* whether MPRV09 also define the (ε,δ) variant]. The extension to (ε,δ) is the standard one. The slack γ makes explicit the statistical terms that are not negligible in λ at concrete κ.

Lemma 5.7 (From SIM-CDP to indistinguishability)**.** *Suppose the protocol is (ε,δ)-SIM-CDP with slack γ. Then for all $D \sim D'$ and every PPT $\mathcal{A}$,*

$$\Pr[\mathcal{A}(\mathrm{view}(D)) = 1] \leq e^{\varepsilon}\Pr[\mathcal{A}(\mathrm{view}(D')) = 1] + \delta + (1+e^{\varepsilon})\gamma + \mathrm{negl}(\lambda).$$

That is, the protocol is $(\varepsilon, \delta + (1+e^{\varepsilon})\gamma)$-IND-CDP in the sense of [16].

Proof. Chain three steps: $\mathrm{view}(D) \approx \mathcal{S}(\mathcal{M}(D))$; then DP of $\mathcal{S} \circ \mathcal{M}$, which is post-processing; then $\mathcal{S}(\mathcal{M}(D')) \approx \mathrm{view}(D')$. □

Lemma 5.8 (Post-processing and composition of SIM-CDP)**.** *(a) If the protocol is (ε, δ)-SIM-CDP with slack γ, then so is $g(\mathrm{view}_{\mathsf{R}}(D))$ for every PPT g, with simulator $g \circ \mathcal{S}$.*

(b) Consider K runs on the same dataset, with independent keys and randomness and with R's inputs fixed. Suppose run k is simulated by $\mathcal{S}_k$ from a ρ_k-zCDP mechanism $\mathcal{M}_k$ with slack γ_k. Then the joint view is simulated by $(\mathcal{S}_k)_k$ from $(\mathcal{M}_k)_k$, which is $\sum_k \rho_k$-zCDP, with slack $\sum_k \gamma_k$.

In particular, the per-base-register ledger of the specification bounds the privacy loss of all runs, even if researchers collude. This holds when runs use fresh keys, which Algorithm 1 provides.

Proof. (a) is immediate. For (b), use a hybrid over runs: replace run k by $\mathcal{S}_k(\mathcal{M}_k(D))$ one run at a time. This is possible because run k uses keys and randomness independent of the other runs, so the reduction can simulate those runs itself. Then apply Lemma 5.2. If the researchers' inputs and the runs' parameters depend adaptively on earlier outputs, the hybrid is written out in Part C, the composition appendix (Theorem 4.2 there, `thm:main`): all smudging steps are done first in one statistical hybrid, the computational steps follow run by run, and a Rényi filter replaces Lemma 5.2. Part C also requires the ledger to charge a run's full cost at admission, since the δ_B part of an aborted run's unmade releases must stay charged, and gives the ledger with a hockey-stick slack (Corollary 5.4 there). □

5.3 Assumptions

Assumption 5.9 (RLWE for the re-randomizer)**.** Let $a \leftarrow \mathcal{R}_{q_0}$ be uniform, $u \leftarrow \chi_{\mathrm{key}}$ (uniform ternary) and $e_1 \leftarrow \chi_{\mathrm{err}} = \mathcal{N}_{\mathbb{Z}}(0, \sigma_e^2)^N$. Then $(a, au + e_1) \approx_c (a, \mathrm{U}(\mathcal{R}_{q_0}))$. This is decision-RLWE with one sample, a ternary secret and discrete Gaussian error with $\sigma_e = 3.2$. Let Adv_1 denote the best advantage.

Assumption 5.10 (RLWE with circular security for O's key shares)**.** Let $s_O \leftarrow \chi_{\mathrm{key}}$, let every error be drawn from χ_{err}, and let all a's be uniform common random strings. The tuple

$$\left(a,\ -as_O + e_O^{\mathrm{pk}};\ \{a_{g,j},\ -a_{g,j}s_O + P\mathsf{g}_j\psi_g(s_O) + e_{O,g,j}^{\mathrm{rk}}\}_{g\in\{g_0,\ldots,g_{L-1}\},\, j\le d_{\mathrm{num}}}\right)$$

is computationally indistinguishable from uniform over $\mathcal{R}_{Q_2} \times \mathcal{R}_{PQ_2}^{Ld_{\mathrm{num}}}$; let Adv_2 be the advantage. The first component is plain RLWE in s_O. The rotation-key components encrypt key-dependent messages $P\mathsf{g}_j\psi_g(s_O)$, so they need a circular (key-dependent-message) security assumption, as do all CKKS and BFV rotation keys.

These shares are exactly what R receives from O in the one-round public-key and rotation-key generation protocols of Mouchet et al. [17]. Under this assumption a semi-honest R's view of key generation is simulated by uniform elements; this is the "security of threshold key generation" used in the specification. [*citation detail unverified:* Mouchet et al. state that the rotation-key protocol relies on the same circular-security assumption as single-key RLWE schemes; we did not locate the exact statement number.]

The proof also uses three things proved above: Theorem 3.12 (Obligation O1), Lemma 2.3 with Corollary 2.4, and Proposition 5.3 with Proposition 5.4. It also uses O's honest behaviour in Algorithm 5:

- (u, e_0, e_1, η) are fresh and independent for each release;
- η is *exactly* $D_{\mathbb{Z}^N, \sigma_{c,s}}$;

- c'_0 is neither sent nor reused.

The proof does *not* use PSI security, the random-oracle model, or any assumption on the distribution of rounding errors.

5.4 The theorem

Theorem 5.11 (Output privacy against an honest-but-curious R; Theorem 1 of the specification, corrected)**.** *Assume (M1)–(M5), Assumptions 5.9 and 5.10, and that for every release s:*

- $\sigma_{c,s} \geq \max\{2^{\kappa} B_{\mathrm{tot},s}/\sqrt{2\pi},\, 2^{42}\}$, *with $B_{\mathrm{tot},s}$ from Theorem 3.12, or from Lemma 4.1 for OpenFHE's key switching (M3′) (the floor was 2^{20} in earlier versions; Remark 2.5);*
- $\rho_s = \widehat{\Delta}^2_{\mathrm{enc},s}/(2\sigma^2_{c,s})$, *with $\widehat{\Delta}_{\mathrm{enc},s}$ from Corollary 3.13 (equal scales) or Lemma 4.3 (decoupled scales, as in the evaluation);*
- $\sum_s \rho_s \leq \rho_{\mathrm{run}}$, *with ρ_{run} from Proposition 5.4 at $(\varepsilon_{\mathrm{run}}, \delta_{\mathrm{DP}})$.*

Then R's view of one run is $(\varepsilon_{\mathrm{run}}, \delta_{\mathrm{DP}})$-SIM-CDP for $\sim$, with statistical slack

$$\gamma = n_{\mathrm{rel}}\big(\delta_B + 2^{-\kappa}(1 + 2^{-80}) + 2^{-97}\big)$$

and computational loss $n_{\mathrm{rel}}\mathrm{Adv}_1 + \mathrm{Adv}_2$. The mechanism is

$$\mathcal{M}(D) = \big(\omega_{\mathsf{R}},\ \Lambda,\ (f_s(D;\pi) + \eta_s)_s\big),$$

where ω_{R} is a fresh random tape for R, and Λ is R's view of the Link phase (Algorithm 2) generated with fresh coins for O. Those coins include the OPRF key, the dummy identifiers' labels and the permutation τ, hence π; the same τ fixes the arrangement used in $f_s(\cdot;\pi)$. By Lemma 5.7, the view is $(\varepsilon_{\mathrm{run}}, \delta_{\mathrm{DP}} + (1 + e^{\varepsilon_{\mathrm{run}}})\gamma)$-IND-CDP.

Proof. *$\mathcal{M}$ is ρ_{run}-zCDP, and hence $(\varepsilon_{\mathrm{run}}, \delta_{\mathrm{DP}})$-DP.* The pair $(\omega_{\mathsf{R}}, \Lambda)$ has the same law under D and D'. Neighbours differ only in $(\mathbf{x}_O, y)$ of one row, and neither the Link phase nor R's coins depend on $(\mathbf{x}_O, y)$: the labels are positions $\tau(j)$, and the identifiers, including the dummy ones, are part of the fixed part of the dataset. Conditionally on $(\omega_{\mathsf{R}}, \Lambda)$, the last component is the composition of the mechanisms $f_s(\cdot;\pi) + \eta_s$. Their sensitivity is at most $\widehat{\Delta}_{\mathrm{enc},s}$ for every arrangement π, by Corollary 3.13 or Lemma 4.3. Apply Proposition 5.3, Lemma 5.2 and Proposition 5.4.

The simulator. On input $(\omega_{\mathsf{R}}, \Lambda, (w_s)_s)$, $\mathcal{S}$ proceeds as follows:

1. output ω_{R} and Λ;
2. replace O's key-generation shares by uniform elements of the right rings;
3. for each release, sample $c^*_{1,s} \leftarrow \mathrm{U}(\mathcal{R}_{q_0})$ and output $(c^*_{1,s},\ w_s - c^*_{1,s}s_R \bmod q_0)$, reading s_R from ω_{R};
4. output the public parameters, including the outputs of Calibrate. These depend on public inputs only, because $B_{\mathrm{tot},s}$ is a worst-case bound.

Hybrids.

H0 The real view.

H1 Compute $h_{O,s}$ as $v_s - c'_{1,s}s_R$, with $v_s = c'_{0,s} + c'_{1,s}s + \eta_s \bmod q_0$. This is identical to H0, since $h_{O,s} = c'_{0,s} + c'_{1,s}s_O + \eta_s$.

H2$^{(s)}$ For $s = 1, \ldots, n_{\mathrm{rel}}$ in turn, replace v_s by $f_s(D;\pi) + \eta_s \bmod q_0$. By Proposition 3.8, $v_s \equiv f_s + e_{\mathrm{tot},s} + \eta_s$. Let Z be everything else in the hybrid. Then η_s is independent of $(Z, e_{\mathrm{tot},s})$, because it is fresh and used only in v_s. Corollary 2.4 with $B = B_{\mathrm{tot},s}$ and Theorem 3.12, which holds for every D and every value of the public randomness, bound the statistical distance between consecutive hybrids by $\delta_B + 2^{-\kappa}(1 + 2^{-80}) + 2^{-97}$. Summing over s gives γ. Three precisions:

– *The side information includes* R*'s random tape.* Z contains ω_{R}, hence s_R, R's encryption errors and its key-share errors, all of which enter $e_{\mathrm{tot},s}$ (through the term of R's errors times O's data, the shares of e^{pk} and e^{rk}, and the s-dependent terms). Corollary 2.4 allows any joint law of $(Z, e_{\mathrm{tot},s})$, so the simulator needs only $\mathcal{M}(D)$, not the exact $f_s(D)$: this is simulation in the standard sense, not the relaxed sense of Hwang et al. [15] in which the simulator also receives the exact output.

– R*'s errors must be honestly sampled.* Theorem 3.12 bounds $e_{\mathrm{tot},s}$ under (M5), with every error drawn from χ_{err}. A researcher who inflated its encryption errors would push the term of R's errors times O's data beyond $B_{\mathrm{tot},s}$, and that term is correlated with O's data. This is part of the honest-but-curious model, and it must be stated as such.

– γ *is statistical.* The slack γ is a concrete statistical distance at the chosen κ and δ_B, not a quantity negligible in λ; it is accounted for explicitly in δ (Lemma 5.7).

H3$^{(s)}$ Replace $c'_{1,s}$ by a uniform element of $\mathcal{R}_{q_0}$. After H2, u_s and $e_{1,s}$ enter the view only through $c'_{1,s} = c_{1,s} + au_s + e_{1,s}$. The reason is that $c'_{0,s}$ is not sent, and neither v_s nor anything else uses $c'_{0,s}$, $e_{0,s}$, u_s or $e_{1,s}$; in particular the hint $e^{\mathrm{pk}}u_s + e_{1,s}s$ was removed with $e_{\mathrm{tot},s}$.

The reduction to Assumption 5.9 works as follows. On a challenge (a^*, b^*), it sets a to a uniform CRT lift of a^*, so that a is uniform at the full modulus. It samples s_O, all keys, all errors and all other releases itself; it knows D. It sets $c'_{1,s} = c_{1,s} + b^*$. Releases $s' < s$ are already uniform and releases $s' > s$ are real, and both are computable because the reduction knows a. The loss per release is Adv_1. Fresh u_s per release is essential here. If u were reused, then $c'_{1,s} - c'_{1,s'}$ would equal $c_{1,s} - c_{1,s'} + e_{1,s} - e_{1,s'}$. That is a short perturbation of a deterministic function of the data and the public randomness, and the hybrid would fail.

H4 Replace O's key-generation shares by uniform elements. After H3, the view depends on $(s_O, e_O^{\mathrm{pk}}, e_O^{\mathrm{rk}})$ only through these shares:

– $h_{O,s}$ no longer contains s_O;

– the ciphertexts M_s, and therefore the key-switching and conversion steps, no longer influence the view;

– $e_{\mathrm{tot},s}$ has been removed.

The reduction to Assumption 5.10 embeds its challenge as O's shares. It computes the joint public key and R's own shares from the challenge and ω_{R}, and runs the rest honestly. The loss is Adv_2.

Only $p_{0,O}$ is sent. In the protocol as implemented and as stated in the paper, O sends R only its public-key share $p_{0,O} = -as_O + e_O^{\mathrm{pk}}$; O sums the rotation-key shares itself and never sends its own. After H3 the rotation keys enter R's view only through the ciphertexts M_s, which no longer enter it. So H4 needs only the first component of Assumption 5.10: plain decision-RLWE with one sample in s_O, with no circular-security assumption. The rotation-key components are needed only if O's rotation-key shares are sent.

H4 is distributed exactly as $\mathcal{S}(\mathcal{M}(D))$, with the coins of $\mathcal{M}$ coupled to those of the hybrid: the Link part of H4 is generated exactly as Λ, and every other component matches the simulator. □

Remark 5.12 (Comparison with the specification's proof sketch)**.** The sketch's hybrids H0–H3 match ours, with these differences:

(i) the centre is $f_s(D)$, not $\mathsf{Ecd}(\mathbf{z}_s(D))$, which changes Δ_{enc} to $\widehat{\Delta}_{\mathrm{enc}}$;

(ii) the per-release distance is $\delta_B + 2^{-\kappa}(1 + 2^{-80}) + 2^{-97}$;

(iii) the statistical slack is kept separate from δ, and it enters δ with the factor $1 + e^{\varepsilon}$ under the indistinguishability reading;

(iv) the sketch's H4 invokes PSI security, which is unnecessary under the specification's adjacency.

If one wants the simulator to use only the leakage G4, that is, $(I, \pi|_I)$, rather than a transcript

generated from the identifiers, then one needs the semi-honest receiver security of the labeled PSI [5]. That in turn needs circuit privacy of O's BFV replies (Correction 7).

Remark 5.13 (Why exactness of the sampler matters)**.** Proposition 5.3 needs η_s to be exactly $D_{\mathbb{Z}^N,\sigma}$, through Lemmas 2.1 and 2.2. A sampler at statistical distance ν from it adds ν per coordinate, so $N\nu$ per release, to γ. That is acceptable only if $N\nu \ll 2^{-\kappa}$. A floating-point sampler at $\sigma \approx 2^{52}$ is not of this kind. Float64 values of magnitude at least 2^{53} are spaced at least 2 apart, so the rounded output misses odd integers in that range. More generally, the low-order bits of its output are not smoothed. So ν is not small, and the Rényi divergence between neighbouring outputs can be infinite. This is the review's point, and it is why SampleDG must be exact (Section 8).

6 Input privacy (G1) and cohort privacy (G2)

O's view. Besides its own input D and random tape, O receives the following from R:

(i) *Key-generation shares:* $p_{0,R} = -as_R + e_R^{\mathrm{pk}}$, and the rotation-key shares $-a_{g,j}s_R + P\mathsf{g}_j\psi_g(s_R) + e_{R,g,j}^{\mathrm{rk}}$.

(ii) *Link messages:* the blinded OPRF queries, exactly $n_R^{\max}$ of them; R's BFV public key and any evaluation keys used by the PSI; and the BFV ciphertexts of windowed powers. Their number is fixed by public PSI parameters and $n_R^{\max}$.

(iii) *Uploads:* $(c_{0,h}, \varsigma_h)$ for each of the $(d_R+2)m$ chunks, with $c_{0,h} = -a_h s_R + e_h + m_{U,h}$ and $a_h = H(\varsigma_h)$.

(iv) *Control messages:* the requests for the two or three releases. Their schedule is fixed by Algorithm 6 and does not depend on R's data. R's refusal at line 7 happens after the last release and is not sent to O.

O receives no decryption share and no ciphertext component that involves s_R other than these.

Assumption 6.1 (For G1 and G2)**.** (a) *RLWE with circular security in s_R.* Let $s_R \leftarrow \chi_{\mathrm{key}}$, let every error be drawn from χ_{err}, and let all a's be uniform. Then the tuple consisting of

- R's public-key share,
- its Ld_{num} rotation-key shares, and
- $(d_R+2)m$ samples $(a_h, -a_h s_R + e_h)$ at modulus Q_2

is pseudorandom. This is multi-sample RLWE plus the circular assumption for the rotation-key shares. Samples at modulus Q_2 are reductions of samples at modulus PQ_2, so a single modulus suffices.

(b) *Random oracle.* The seed expansion H is a random oracle, and seeds are λ-bit strings drawn uniformly.

(c) *The PSI's encryption.* The PSI's BFV is IND-CPA under RLWE, with circular security for any relinearization key R sends.

(d) *The OPRF.* The blinded OPRF queries of [5] are distributed independently of the queried identifiers. For the Diffie–Hellman-type OPRF this holds perfectly in a prime-order group. [*citation detail unverified:* the concrete OPRF instantiation in [5] and whether its query blinding is perfectly or computationally hiding]

Theorem 6.2 (G1: R-input privacy)**.** *Under Assumption 6.1 there is a PPT simulator $\mathcal{S}_O$ with the following property. Given O's input and random tape, $(n_R^{\max}, \bar{n}_O)$ and the public parameters, $\mathcal{S}_O$ outputs a view that is indistinguishable from O's real view, for every input of R. The distance is at most*

$$\mathrm{Adv}_{\mathrm{RLWE},R} + \mathrm{Adv}_{\mathrm{BFV}} + q_H^2/2^{\lambda} + p_{\mathrm{cuckoo}},$$

where q_H is the number of oracle queries and p_{cuckoo} is the probability that R*'s cuckoo hashing fails in a way that* O *can observe, for example as an abort or a retry (bounded in Lemma 7.5).* R*'s abort on a detected false match is observable too; its probability, at most p_{false} (the paper's Remark on collisions), is added to the slack in the same way.*

Proof. The simulator. $\mathcal{S}_O$ outputs:

- uniform key-generation shares;
- $n_R^{\max}$ uniform group elements as the OPRF queries;
- a fresh BFV public key, any evaluation keys, and encryptions of zero in the public number;
- $(d_R + 2)m$ pairs $(c_{0,h}^*, \varsigma_h^*)$, with $c_{0,h}^*$ uniform in $\mathcal{R}_{Q_2}$ and ς_h^* uniform seeds;
- the fixed control messages.

Hybrids.

1. Replace the OPRF queries by uniform group elements. The distance is zero for a perfectly hiding blinding, and the OPRF advantage otherwise.
2. Replace R's BFV ciphertexts of powers by encryptions of zero, at cost Adv_{BFV}.
3. In the random-oracle model, the $a_h = H(\varsigma_h)$ are uniform and independent unless two seeds collide, which happens with probability at most $q_H^2/2^\lambda$. The reduction programs $H(\varsigma_h)$ to the a-components of its RLWE challenge. The tuple of R's key shares and the values $-a_h s_R + e_h$ is then replaced by uniform elements, at cost $\text{Adv}_{\text{RLWE},R}$ by Assumption 6.1(a). Adding the known plaintext $m_{U,h}$ to a uniform element leaves it uniform.
4. The real protocol leaks the cuckoo-hashing outcome. Removing that leak costs at most p_{cuckoo}.

The final hybrid is $\mathcal{S}_O$'s output. O's own messages, such as the OPRF evaluations, PSI replies and releases, are functions of its view, so they need no simulation. □

Corollary 6.3 (G2: cohort privacy)**.** *$\mathcal{S}_O$ does not use* R*'s cohort, n_R or $\mathbf{X}_R$. Hence, for any two cohorts of size at most $n_R^{\max}$,* O*'s views are indistinguishable, up to twice the distance of Theorem 6.2. Padding the OPRF queries to exactly $n_R^{\max}$ is what makes the number of OPRF calls cohort-independent. The per-run fresh OPRF key and the quota are not needed against an honest-but-curious* O*; they matter only for linkage across runs.*

Remark 6.4.
(i) The statement "O never receives a decryption share from R" is what makes G1 a pure encryption argument. If any future version sent R's share $c_1 s_R + e$, G1 would need its own flooding.
(ii) The random-oracle model is used only to treat the expanded a_h as uniform. A standard-model alternative is to send a_h in full, at twice the upload size.
(iii) G1 and G2 hold against an honest-but-curious O only. A deviating O could, for example, choose its PSI replies so as to learn the outcome of R's decryption through R's later behaviour. This is outside the model.

7 Disclosure (G4), the cuckoo-hashing slack and the reply noise

This section proves the first layer of G4 (what R learns from the Link phase) in the random-oracle model under CDH, bounds the cuckoo-hashing failure probability that enters the slack of G1 and G2, and bounds the BFV evaluation noise of the PSI replies in an explicit model of BFV. The first result was first written in the companion book (Chapter 8, Lemma `aln.keyed` and Theorem `aln.R`) and was checked here; the other two are new. They replace Correction 7 as far as the first layer is concerned.

The Link phase. O draws a fresh OPRF key $k \leftarrow \mathbb{Z}_{q_\mathbb{G}}^*$ and a uniform permutation τ of its $\bar{n}_O$ padded rows. Row j gets $y_j = F_k(\mathrm{id}_j) = H_2(\mathrm{id}_j, H_1(\mathrm{id}_j)^k)$; its item, bins and label key are domain-separated hashes of the two halves of y_j, and its label is $\tau(j)$ padded with the label key modulo t. R pads its identifiers to a list $\bar{X}$ of exactly $n_R^{\max}$, obtains F_k on $\bar{X}$ through the blinded protocol ($\alpha = H_1(x)^r$, $\beta = \alpha^k$), cuckoo-hashes its items, and sends BFV encryptions of the windowed powers under its own key. O evaluates the masked membership and label polynomials of each partition and replies with SanitizeBFV of the results. R's output is $(I, \pi|_I)$, or an abort.

Assumption 7.1 (CDH in the OPRF group)**.** Given (g, g^k, h) with k and h uniform, no PPT algorithm (with non-uniform advice) outputs h^k with probability above $\mathrm{Adv}_{\mathrm{CDH}}$.

Lemma 7.2 (OPRF-keyed items and labels are pseudorandom outside the match)**.** *Model H_1, H_2 and the derived hash functions as random oracles, let R be semi-honest, and let the distinguisher make at most q_{H_2} queries to H_2. Let J be the set of O's rows whose identifiers are not in $\bar{X}$. Then R's view of the OPRF phase together with $(F_k(\mathrm{id}_j))_{j\in J}$ changes by at most $q_{H_2}\mathrm{Adv}_{\mathrm{CDH}} + \epsilon_{\mathrm{coll}}$ in distinguishing advantage when these outputs are replaced by independent uniform strings, for every input; $\epsilon_{\mathrm{coll}} \le \bar{n}_O n_R^{\max} 2^{-\ell_{\mathrm{id}}} + (q_{H_1} + \bar{n}_O)/q_\mathbb{G}$ accounts for a collision of a fresh dummy identifier of length ℓ_{id} with the other list and for H_1 hitting the identity.*

Proof. Identical until a bad event. In the real experiment $F_k(\mathrm{id}_j) = H_2(\mathrm{id}_j, H_1(\mathrm{id}_j)^k)$; in the modified one it is a fresh uniform string, and H_2 is answered uniformly elsewhere. An honest R queries H_2 only at points $(x, H_1(x)^k)$ with $x \in \bar{X}$, so the two experiments are identical unless the distinguisher queries H_2 at $(\mathrm{id}_j, H_1(\mathrm{id}_j)^k)$ for some $j \in J$ (event Bad), which has the same probability in both.

A CDH solver. Given (g, g^k, h), run the modified experiment with O's input fixed (non-uniform advice), programming $H_1(x) = g^{a_x}$ at every queried point other than O's identifiers in J, and $H_1(\mathrm{id}_j) = h\, g^{b_j}$ for $j \in J$, with fresh uniform exponents. Answer a blinded query $\alpha = g^{a_x r}$ with $(g^k)^{a_x r} = \alpha^k$; compute the outputs of the matched rows as $H_2(x, (g^k)^{a_x})$ and those of the rows in J as fresh strings. Preprocessing, evaluation and sanitization use no other secret, and R's own coins are sampled by the solver. At the end, pick a uniformly random H_2 query (id_j, Z) with $j \in J$ and output $Z\,(g^k)^{-b_j}$. On Bad some query has $Z = h^k (g^k)^{b_j}$, picked with probability at least $1/q_{H_2}$. So $\Pr[\mathsf{Bad}] \le q_{H_2}\mathrm{Adv}_{\mathrm{CDH}}$; the programmed values are uniform except on the events counted in ϵ_{coll}. □

Theorem 7.3 (G4, first layer: R learns only the match and its positions)**.** *Under the hypotheses of Lemma 7.2 there is a PPT simulator $\mathcal{S}_\Lambda$ that, given R's input, its output $(I, \pi|_I)$ and the public sizes, outputs R's view of the Link phase up to distinguishing advantage $q_{H_2}\mathrm{Adv}_{\mathrm{CDH}} + \epsilon_{\mathrm{coll}}$, for every input. Jointly with the output, the distance grows by the probability of the correctness failures on which R's output differs from $(I, \pi|_I)$ (overflow of a bin, a cuckoo failure, a false match; Lemma 7.5 and the paper's Remark on collisions). The statement needs neither circuit privacy of the BFV evaluation nor* SanitizeBFV*, nor a bound on the reply noise.*

Proof. $\mathcal{S}_\Lambda$ samples R's tape, runs R's code to produce the padded list, the blinded elements, the BFV keys and the encrypted windowed powers; draws a key k^* and answers the blinded elements with it; creates one row per $i \in I$ with output $F_{k^*}(i)$ (computed by R's code) and position $\pi(i)$, and $\bar{n}_O - |I|$ rows with uniform outputs and positions forming a uniform bijection onto $[\bar{n}_O] \setminus \pi(I)$; and runs O's preprocessing (partition placement included), evaluation, masking and sanitization with fresh coins. By Lemma 7.2 the real view is within the stated distance of the experiment in which O's rows outside the match carry uniform outputs. In that experiment the key enters only through R's own OPRF answers, so k^* may replace k; positions outside the match are a uniform bijection

onto the complement of $\pi(I)$ because τ is uniform; the labels outside the match are padded by uniform keys. Every reply, its noise included, is a fixed randomized function of R's ciphertexts and of these inputs, so the two experiments coincide. Overflow aborts, cuckoo failures, false matches and the partition index of a match arise from identically distributed items and bins and are reproduced exactly. □

Remark 7.4 (Second reading; what remains of Correction 7)**.** The book's lemma omits the collision term ϵ_{coll}, which is negligible; otherwise its proof is correct as written. The reduction is non-uniform (it embeds O's input), so Adv_{CDH} is against non-uniform adversaries. Circuit privacy is not needed for the first layer: whatever the noise of a reply reveals is a function of pseudorandom items and labels. The second layer, SanitizeBFV, makes the reply noise independent of O's data up to a statistical distance, given a worst-case bound B_ψ (Proposition 7.7 below bounds its evaluation part in a model of BFV).

Lemma 7.5 (The cuckoo-hashing slack)**.** *Let n items each have h candidate bins drawn independently and uniformly from m bins (in the random-oracle model for the bin hashes; with $m = 2^{14}$ dividing 2^{256} the reduction modulo m of a uniform hash is exactly uniform). The probability that no assignment of the items to distinct candidate bins exists is at most*

$$p_{\text{noassign}}(n, m, h) \le \sum_{s=2}^{n} \binom{n}{s} \binom{m}{s-1} \Big(\frac{s-1}{m}\Big)^{hs}.$$

An insertion procedure that finds an assignment whenever one exists (augmenting paths, or the implementation's random walk followed, on failure, by an augmenting-path search) fails with exactly this probability, so $p_{\text{cuckoo}} \le p_{\text{noassign}}$. At $m = 2^{14}$ and $h = 3$ the bound is $1.7 \cdot 10^{-15} = 2^{-49.1}$ for $n = 2{,}000$ (the evaluation's $n_R^{\max}$, one table), $2^{-47.0}$ for $n = 4{,}096$, $2^{-45.0}$ for $n = 8{,}192$, and at most 2^{-40} for every $n \le 10{,}472$ (load 0.639).

Proof. By Hall's theorem an assignment exists unless some set S of $s \ge 2$ items has fewer than s candidate bins in total (a single item always has a bin). Fix S and a set T of $s-1$ bins; all hs candidates of S lie in T with probability $((s-1)/m)^{hs}$. A union bound over S and T gives the sum (a set with fewer than $s-1$ bins is contained in some T). The numbers are from `experiments/proofs_check/tv1_cuckoo.py`, which sums every term. □

Remark 7.6 (What is not covered)**.** The implementation's insertion is a random walk with at most 2,000 evictions per item. Until 2026-10-25 it had no fallback, and its failure probability when an assignment exists was not bounded. It now completes the walk's partial assignment by augmenting paths when the walk exceeds its step limit (`implementation/vfr/psi.py`, `cuckoo` and `_augment`), so it finds an assignment whenever one exists and $p_{\text{cuckoo}} = p_{\text{noassign}}$. The union bound is useless above load about 0.64 at $m = 2^{14}$ (the implementation allows tables up to load 0.9, used only when $n_R^{\max} > 10{,}472$).

7.1 The BFV evaluation noise of the replies

We bound the evaluation part $w_{\text{eval}} = (q'/q_B)v$ of the switched reply error in the following model of BFV (with R's ternary key s_B, plaintext modulus t, $\Omega_{\text{B}} = \lfloor q_B/t \rfloor$, $r_t = q_B - t\Omega_{\text{B}} \in [0, t)$, ring dimension $N = N_B$, and every integer polynomial represented with centred coefficients).

(B1) A ciphertext (c_0, c_1) with plaintext m, $\|m\|_\infty \le t/2$, satisfies $c_0 + c_1 s_B = \Omega_{\text{B}} m + v + q_B k$ as integer polynomials; $\|c_i\|_\infty \le q_B/2$ and $\|v\|_\infty \le q_B/2$ give $\|k\|_\infty \le \kappa_k := (N+3)/2$.

(B2) Every error coefficient has absolute value at most B_e (OpenFHE's table sampler at $\sigma = 3.19$ saturates at $B_e = 27$), so R's fresh encryptions have $\|v\|_\infty \le E_0 := (2N+1)B_e$: $v = e_B^{\mathrm{pk}} u' + e_0 + e_1 s_B$ for a public-key encryption, and $v = e$, with $\|v\|_\infty \le B_e$, for the secret-key encryption that the implementation uses.

(B3) A ciphertext product computes $\tilde{d}_i = \lfloor t\, d_i / q_B \rceil$ from the tensor $(d_0, d_1, d_2) = (c_0 c_0',\ c_0 c_1' + c_1 c_0',\ c_1 c_1')$, with rounding errors of sup-norm at most $\frac{1}{2}$ up to an additive error of sup-norm at most E_{rns} (an allowance for an RNS implementation), and relinearizes $\tilde{d}_2$, adding an error of sup-norm at most E_{rl}.

(B4) A plaintext multiplication multiplies both components by a plaintext polynomial a with $\|a\|_\infty \le t/2$; a plaintext addition adds $\Omega_{\mathrm{B}} a$ to c_0.

Proposition 7.7 (The reply's evaluation noise in the BFV model)**.** *Under (B1)–(B4), let* $F(E,E') = \frac{Nt^2}{4} + t(\frac{Nt}{4} + \frac{1}{2}) + \frac{Nt}{2}(E+E') + \frac{tN}{q_B}EE' + Nt^2\kappa_k + tN\kappa_k(E+E') + \frac{1+N+N^2}{2} + E_{\mathrm{rl}} + E_{\mathrm{rns}}$*. A product of two ciphertexts with noise sup-norms at most* E *and* E' *(both below* $q_B/(2t)$*) has noise sup-norm at most* $F(E,E')$*. Hence, if every power* $\mathbf{w}^d$*,* $d \le D$*, is a product tree of depth at most* $L_\times$ *over the fresh windows, its noise is at most* $E_{L_\times}$ *with* $E_{l+1} = F(E_l, E_l)$*, and a reply* $\sum_{d=0}^{D} a_d \odot \mathrm{ct}(\mathbf{w}^d)$ *(with* $\mathrm{ct}(\mathbf{w}^0)$ *the plaintext* 1*) has noise at most*

$$E_{\mathrm{eval}} = D\Big(\tfrac{Nt}{2} E_{L_\times} + t\big(\tfrac{Nt}{4} + \tfrac{1}{2}\big)\Big) + \tfrac{t(D+1)}{2},$$

so that $\|w_{\mathrm{eval}}\|_2 \le \sqrt{N}\,(q'/q_B)\,E_{\mathrm{eval}}$*, provided* $E_{\mathrm{eval}} < q_B/(2t)$ *(which also gives correct decryption before the switch).*

Proof. Product. Multiply the two invariants of (B1) and scale by t/q_B. With $t\Omega_{\mathrm{B}} = q_B - r_t$ and $mm' = [mm']_t + t\mu$, $\|\mu\|_\infty \le Nt/4 + \frac{1}{2}$, the identity

$$\begin{aligned}\tfrac{t}{q_B}(\Omega_{\mathrm{B}} m + v + q_B k)(\Omega_{\mathrm{B}} m' + v' + q_B k') \equiv \Omega_{\mathrm{B}}[mm']_t - \tfrac{r_t}{q_B}\Omega_{\mathrm{B}} mm' - r_t\mu \\ + \tfrac{t\Omega_{\mathrm{B}}}{q_B}(mv' + m'v) + \tfrac{t}{q_B}vv' - r_t(mk' + m'k) + t(vk' + v'k)\end{aligned}$$

holds modulo q_B (the terms $t\Omega_{\mathrm{B}}(mk' + m'k)$ and $tq_B kk'$ reduce to $-r_t(mk' + m'k)$ and 0). The rounded tensor adds $r_0 + r_1 s_B + r_2 s_B^2$, of sup-norm at most $\frac{1}{2}(1 + \|s_B\|_1 + \|s_B^2\|_1) \le \frac{1}{2}(1 + N + N^2)$, plus E_{rns}; relinearization adds E_{rl}. Bound each term with $\|ab\|_\infty \le N\|a\|_\infty\|b\|_\infty$: $\frac{r_t\Omega_{\mathrm{B}}}{q_B} \le 1$ and $\|mm'\|_\infty \le Nt^2/4$; $r_t < t$; $\|mv' + m'v\|_\infty \le \frac{Nt}{2}(E+E')$; $\|vv'\|_\infty \le NEE'$; $\|mk' + m'k\|_\infty \le Nt\kappa_k$; $\|vk' + v'k\|_\infty \le N\kappa_k(E+E')$. The sum is $F(E,E')$. The bound is increasing in E and E', so a tree of depth l has noise at most E_l.

Evaluation. $a \odot (c_0, c_1)$ has invariant $a(\Omega_{\mathrm{B}} m + v + q_B k) \equiv \Omega_{\mathrm{B}}[am]_t - r_t\mu_a + av$, with $\|\mu_a\|_\infty \le Nt/4 + \frac{1}{2}$ and $\|av\|_\infty \le \frac{Nt}{2}\|v\|_\infty$. Summing D such terms and the plaintext term a_0, the messages add up to at most $(D+1)t/2$ in sup-norm; reducing the sum modulo t changes the noise by $r_t\nu$ with $\|\nu\|_\infty \le (D+1)/2$, at most $t(D+1)/2$. The switch multiplies the evaluation part of the noise by q'/q_B, and $\|\cdot\|_2 \le \sqrt{N}\|\cdot\|_\infty$. □

Corollary 7.8 (A proved B_ψ in the model)**.** *At the evaluation's PSI parameters (*$N_B = 2^{14}$*,* $t =$ 4,293,918,721*,* $D = 64$*,* $W = 7$ *windows and product trees of depth* $L_\times = 3$*, as in the implementation's balanced split of* d *by popcount,* q_B *of seven* 60*-bit limbs and* q' *of two), Proposition 7.7 gives* $E_0 = 2^{19.8}$*,* $E_3 = 2^{211}$*,* $E_{\mathrm{eval}} = 2^{262} < q_B/(2t) \approx 2^{387}$ *and* $\|w_{\mathrm{eval}}\|_2 \le 2^{-31}$ *for any* $E_{\mathrm{rl}}, E_{\mathrm{rns}} \le 2^{26}$*, and* 2^{-30} *for any* $E_{\mathrm{rl}}, E_{\mathrm{rns}} \le 2^{90}$ *(script* `experiments/proofs_check/bpsi_model.py`*). So, in the model, the evaluation part is covered by the "+1" of the implementation's* $B_{\mathrm{err}} = 2(\sqrt{N_B}/2 + N_B^{3/2}/2) + 1$*, and* $B_\psi := \sqrt{2\pi}B_{\mathrm{err}}$ *bounds* $\|w\|_2$ *with the rounding part (exact) and the fresh-encryption part (whose high-probability bound* $(q'/q_B)\beta_{\mathrm{fresh}}$ *is below* 2^{-250}*).*

Remark 7.9 (Status of B_ψ)**.** Proposition 7.7 is a theorem about the textbook BFV model (B1)–(B4); it is worst case over O's inputs and needs only the honest sampling of R's errors. That OpenFHE 1.5.1's BFVrns (its RNS multiplication and scaling, and its relinearization) satisfies (B3) with small E_{rns} and E_{rl} was not checked; the measured evaluation noise after the switch is 2^{-72}, far below the model's 2^{-31}. G4's second layer is therefore proved for the model and remains conditional on this conformance for the implementation.

8 Correctness: no wraparound, and the exact sampler

Proposition 8.1 (Decryption without wraparound)**.** *Let $S_{\max} = \max_e \sup_D |\mathbf{z}_e(D)|$. Each entry is a sum over at most $\bar{n}_O$ rows of a product bounded by one of o, $\sqrt{ro}$, $\sqrt{rw}$, $\sqrt{ow}$ or w; so $S_{\max} \le \bar{n}_O \max(o, \sqrt{ro}, \sqrt{rw}, \sqrt{ow}, w)$. Let $d_{\max}$ bound $\|\widehat{Z}_s(D) - \Omega_{\mathrm{fin}}\tau^{-1}(\mathbf{z}_s(D))\|_2$ over all D; Proposition 4.5 gives an explicit, proved $d_{\max}$ ($2^{30.3}$ at the evaluation setting with $\bar{n}_O = 10^6$ and $n_R^{\max} = 2{,}000$, and $2^{38.6}$ if every row may match). Suppose*

$$q_0 > 2\Big(\Omega_{\mathrm{fin}}\,\tfrac{2|\mathcal{E}_s|}{N}\,S_{\max} + d_{\max} + \tfrac{1}{2} + B_{\mathrm{tot},s} + t_\eta\,\sigma_{c,s}\Big). \tag{Q}$$

Then v_s, reduced to its centred representative modulo q_0, equals the integer polynomial $f_s(D) + e_{\mathrm{tot},s} + \eta_s$. This holds except with probability at most

$$\delta_B + 2Ne^{-t_\eta^2/2}$$

over the keys and the fresh randomness. Moreover $\mathsf{Dcd}(v_s) = \tau(\widehat{Z}_s)/\Omega_{\mathrm{fin}} + \tau(e_{\mathrm{tot},s} + T_7' + \eta_s)/\Omega_{\mathrm{fin}}$, where $T_7' = f_s - \widehat{Z}_s$.

Proof. Size of the centre. From the inversion formula in the proof of Lemma 1.1, $\|\tau^{-1}(\mathbf{z})\|_\infty \le \frac{2}{N}\|\mathbf{z}\|_1 \le \frac{2|\mathcal{E}|}{N}S_{\max}$. So $\|f_s\|_\infty \le \Omega_{\mathrm{fin}}\frac{2|\mathcal{E}|}{N}S_{\max} + d_{\max} + \frac{1}{2}$.

Size of the error. $\|e_{\mathrm{tot},s}\|_\infty \le \|e_{\mathrm{tot},s}\|_2 \le B_{\mathrm{tot},s}$, except with probability δ_B (Theorem 3.12).

Size of the noise. Each η_k is σ^2-subgaussian (Lemma 3.1(b)), so $\Pr[|\eta_k| > t_\eta\sigma] \le 2e^{-t_\eta^2/2}$. A union bound over the N coefficients gives the second term.

Under (Q), the integer polynomial $f_s + e_{\mathrm{tot},s} + \eta_s$ lies strictly inside $(-q_0/2, q_0/2)^N$. So the centred reduction of v_s returns it exactly, and the decoding identity follows.

Earlier levels: no condition is needed. Proposition 3.8 is an identity of congruences that holds whatever the size of the intermediate plaintexts. A rescale satisfies $c_0' + c_1's = (c_0 + c_1 s)/q + \varepsilon_0 + \varepsilon_1 s$ exactly in $\mathbb{Q}[X]/(X^N + 1)$; if $c_0 + c_1 s = \mu + KQ$ with K integral, whatever the size of μ, then $(c_0 + c_1 s)/q = \mu/q + K(Q/q)$, so the invariant modulo Q/q is $\mu/q + \varepsilon_0 + \varepsilon_1 s$. Plaintext products, additions and automorphisms are congruences modulo the current modulus, and so is key switching (Lemmas 3.5 and 4.1), again with no size condition; the rounding errors are bounded through their sup-norms only. Hence only the final reduction modulo q_0 can wrap around, and (Q) suffices. (This replaces the argued step of an earlier version, which bounded the intermediate plaintexts; the argument is that of the companion book, Chapter 9, Proposition `agg.wrap`, checked here.) □

Remark 8.2 (The specification's condition)**.** CALIBRATE requires

$$q_0 > 2\big(\Omega S_{\max} + 2^\kappa B_{\mathrm{tot}} + 10\,\sigma_{c,\max}\big).$$

This implies (Q) with $t_\eta = 10$ when $\Omega = \Omega_{\mathrm{fin}}$. The reasons:

- $2|\mathcal{E}|/N \le 1$, and the slack $\Omega_{\mathrm{fin}}S_{\max}(1 - 2|\mathcal{E}|/N)$ is far larger than $d_{\max} + \frac{1}{2}$;
- $2^\kappa B_{\mathrm{tot}} \ge B_{\mathrm{tot}}$.

The specification should still state the correctness-failure probability $2Ne^{-50} \approx 2^{-57.1}$ per release. This is a correctness event, not a privacy event, because the view is the same whether or not wraparound occurs. The term $\Omega S_{\max}$ overestimates by a factor $N/(2|\mathcal{E}|) \approx 2^7$, so q_0 could be about 7 bits smaller.

The exact sampler. SAMPLEDG is Algorithm 3 of Canonne, Kamath and Steinke [4]. It draws a discrete Laplace candidate Y with integer scale $t = \lfloor \sigma \rfloor + 1$, accepts it with probability $\exp(-(|Y| - \sigma^2/t)^2/(2\sigma^2))$, and repeats otherwise. The acceptance coins come from an exact Bernoulli($e^{-\gamma}$) sampler for rational γ. CKS20 prove the following, for rational σ^2:

- the output is *exactly* $\mathcal{N}_{\mathbb{Z}}(0, \sigma^2)$;
- only integer arithmetic and uniform random bits are used;
- the expected number of iterations is bounded by a constant.

[*citation detail unverified:* the exact constant bounding the expected iterations in CKS20]

Two implementation requirements follow:

(i) $\sigma^2_{c,s}$ *must be rational.* Set $\sigma^2_{c,s} := \lceil \widehat{\Delta}^2_{\mathrm{enc},s}/(2\rho_s) \rceil$. Increasing σ lowers ρ_s and only helps flooding.

(ii) *Big integers must be arbitrary-precision.* At $\sigma \approx 2^{55}$, t and σ^2 exceed 64 bits, so all arithmetic must be exact arbitrary-precision integer or rational arithmetic, and the coins must come from a CSPRNG.

What "constant-time" must mean. The model of the specification abstracts side channels away. Its remaining obligation O2 nevertheless asks for a constant-time sampler, and R can at least observe O's response time. The requirement is therefore:

> (CT) The joint distribution of everything R can observe about the execution of the release, in particular its duration, is independent of (η, s_O, u, e_0, e_1) given the public parameters.

The rejection structure helps only in part:

- In rejection sampling, the number of *rejected* trials is independent of the accepted value.
- But the time of the accepting trial depends on the accepted Y. The discrete Laplace and Bernoulli($e^{-\gamma}$) subroutines loop for value-dependent numbers of iterations, and big-integer operations take time that depends on operand length.

So a direct implementation of CKS20 does not satisfy (CT).

A sufficient implementation has four parts:

(a) every inner loop runs for a fixed public number of iterations $T_{\max}$, with the result selected branch-free;

(b) the probability that some loop would have needed more than $T_{\max}$ iterations is at most $2^{-\kappa'}$ per sample, and the total $Nn_{\mathrm{rel}}2^{-\kappa'}$ is added to the slack γ;

(c) all arithmetic is done in fixed-width limbs;

(d) the N samples of a release are produced in one batch whose duration is padded to a public bound.

[**Not fully rigorous:** We have not analysed a concrete constant-time variant of CKS20. Generic constant-time samplers by convolution, such as that of Micciancio and Walter [14], are Rényi-close rather than exact; using one would require adding its divergence to the privacy analysis.]

9 Calibrate, the run as a secure computation, and O's work

The three results of this section were first written in the companion book (Chapter 19, Proposition `par.calibrate`; Chapter 4, Proposition `sc.secureDP` with Chapter 17, Theorem `run.secure`; Chapter 20, Proposition `cost.lower`). They were checked here, with one correction to the second.

9.1 Soundness of Calibrate

CALIBRATE (the paper's Section on calibration) takes the parameter vector $\mathcal{P}$ and the public ledger parameters and runs seven steps: the security test $\log_2(QP) \le L_{\max}(N, \lambda_{\sec})$; the sensitivities and the mode; $\kappa = \max\{\kappa_{\min}, \kappa_{\run}, \kappa_L\}$ (unless $\mathcal{P}$ fixes κ), with the test of the run's share of the ledger's slack capacity; γ and $\delta_{\DP}$; $\rho_{\run}$ and its split, the stage-2 share rounded down; the smallest grid scale at which $\sigma_{c,s} = (\lceil \widehat{\Delta}^2_{\enc,s}/(2\rho_s) \rceil)^{1/2} \ge \max\{2^{\kappa} B_{\tot,s}/\sqrt{2\pi}, 2^{42}\}$ for every release; and the test on q_0. It refuses when a step has no solution.

Proposition 9.1 (CALIBRATE is sound)**.** *Suppose* CALIBRATE *does not refuse, and that it evaluates $B_{\tot,s}$ and $\Delta_{x,s}$ by the formulas of Theorem 3.12 with Lemma 4.1, and of Lemma 4.3, for the circuit model in use. Then:*

(a) every output is a deterministic function of $\mathcal{P}$ and the public ledger parameters;

(b) for every release, $\sigma_{c,s} \ge \max\{2^{\kappa} B_{\tot,s}/\sqrt{2\pi}, 2^{42}\}$, $\widehat{\Delta}^2_{\enc,s}/(2\sigma^2_{c,s}) \le \rho_s$ with $\sigma^2_{c,s}$ an integer, $\sum_s \rho_s \le \rho_{\run}$, $\rho_{\run} = \max\{\rho : \delta(\rho, \varepsilon_{\run}) \le \delta_{\DP}\}$ with $\delta_{\DP} = \delta_{\run} - (1 + e^{\varepsilon_{\run}})\gamma > 0$; these are the hypotheses of Theorem 5.11;

(c) condition (Q) of Proposition 8.1 holds for every release, with $d_{\max}$ of Proposition 4.5, and the chain meets the security bound;

(d) if κ is derived, $(1 + e^{\varepsilon_{\run}})\gamma \le \zeta \delta_{\run}$ and $\rho_{\run} \ge \tilde{\rho}_{\run}$; with a ledger cap, $n_{\rel}\varepsilon'(\kappa) \le (\rho_{\run}/\rho_L)E'_L$ and $n_{\rel}(\delta_B + \delta') \le (\rho_{\run}/\rho_L)\Gamma_L$, and the costs it returns are the per-run costs of the composition appendix (Part C, Section 5).

Proof. (a) No step reads data, keys, errors or coins: $B_{\tot,s}$ and $\Delta_{x,s}$ are functions of public parameters only (Theorem 3.12, Lemmas 4.1 and 4.3), the grid and its order are fixed, and every step is deterministic. (b) The scale loop stops only where the flooding condition holds for every release. $\sigma^2_{c,s} = \lceil \widehat{\Delta}^2_{\enc,s}/(2\rho_s) \rceil \ge \widehat{\Delta}^2_{\enc,s}/(2\rho_s)$. The split sums to at most $\rho_{\run}$ because the stage-2 share is rounded down. $\delta_{\DP} > 0$ by the test, and $\rho_{\run}$ exists and is positive because $\delta(\cdot, \varepsilon)$ is continuous, increasing and 0 at 0. (c) The q_0 test is $q_0 > 2(\Omega_{\fin} S_{\max} + 2^{\kappa} B_{\tot,\max} + t_\eta \sigma_{c,\max})$ with $S_{\max} = \bar{n}_O \max(r, o, w, \sqrt{ro}, \sqrt{rw}, \sqrt{ow})$. Since $\Omega_{\fin} S_{\max} = \Omega_{\fin} \frac{2|\mathcal{E}_s|}{N} S_{\max} + (1 - \frac{2|\mathcal{E}_s|}{N})\Omega_{\fin} S_{\max}$ and the second part exceeds $d_{\max} + \frac{1}{2}$ whenever $d_{\max} + \frac{1}{2} \le (1 - 2|\mathcal{E}_s|/N)\Omega_{\fin} S_{\max}$ (at the evaluation setting $2^{38.6}$ against more than 2^{74}), and $2^{\kappa} B_{\tot,\max} \ge B_{\tot,s}$, $\sigma_{c,\max} \ge \sigma_{c,s}$, the test implies (Q). (d) $\gamma(k)$ is strictly decreasing, so $\kappa \ge \kappa_{\run}$ gives the share and $\delta_{\DP} \ge (1 - \zeta)\delta_{\run}$, hence $\rho_{\run} \ge \tilde{\rho}_{\run}$ by monotonicity; $\varepsilon'(k)$ is decreasing in k, so $\kappa \ge \kappa_L$ and $\tilde{\rho}_{\run} \le \rho_{\run}$ give the smudging inequality, and the share test gives the slack inequality. □

Remark 9.2 (Second reading)**.** The proof is in exact arithmetic. The implementation computes $\sigma^2_{c,s}$ as an exact fraction, rounds the split and the ledger's capacities in the safe direction, and stores the ledger's sums exactly; it evaluates $B_{\tot,s}$ in double precision (relative error about 10^{-15}) and, since 2026-10-25, makes the flooding test conservative: $B_{\tot,s}$ and the needed $\sigma_{c,s}$ are rounded up by a relative 2^{-40}, and the integer $\sigma^2_{c,s}$ is compared with the square of the rounded value in exact rational arithmetic, so floating-point error cannot pass a case that exact arithmetic would refuse. Its formula for the term T_8 uses $d_{\num}$ where $\sqrt{d_{\num}}$ suffices, which is conservative. Its diagnostic value `q0_need_Q` used $d_{\max} = 2^{35}$ until 2026-10-25, which Proposition 4.5 proves when at most

$n_R^{\max} = 2{,}000$ rows match ($2^{30.3}$) but not when every row of a register of 10^6 may match ($2^{38.6}$). It now uses the $d_{\max}$ of Proposition 4.5 for $n_{\max} = \min(n_R^{\max}, \bar{n}_O)$; the condition (Q) still needs 74.1 bits at $\bar{n}_O = 4{,}096$ and 81.9 bits at 10^6 ($\varepsilon = 1$), because its dominant term is about 2^{81}. The refusal test itself uses the condition of (c), which needs no $d_{\max}$.

9.2 The run computes the DP functionality against R

Let $\mathcal{M}(D) = (\omega_{\mathsf{R}}, \Lambda, ([f_s(D;\pi) + \eta_s]_{q_0})_s)$ be the mechanism of Theorem 5.11, where $[\cdot]_{q_0}$ is the centred residue modulo q_0 (a post-processing, so $\mathcal{M}$ is still ρ_{run}-zCDP), and let Out be R's output map: its decryption of the Link replies and the centred residues $v_s = h_{O,s} + c'_{1,s} s_R \bmod q_0$.

Corollary 9.3 (Standard simulation for R, jointly with its output)**.** *Under the hypotheses of Theorem 5.11, the pair (R's view, R's output) is within statistical distance γ plus computational loss $n_{\text{rel}}\text{Adv}_1 + \text{Adv}_2$ of $(\mathcal{S}(\mathcal{M}(D)), \mathcal{M}(D))$: the run computes the functionality that hands $\mathcal{M}(D)$ to R, against a semi-honest R, in the standard (not the liberal) sense. With Theorem 7.3 in place of the transcript Λ, the functionality may output $(I, \pi|_I)$ instead of Λ, at the additional cost of that theorem.*

Proof. Out is an efficient deterministic function of the view, and the simulator of Theorem 5.11 is output-consistent: on input w_s it outputs $(c^*_{1,s}, w_s - c^*_{1,s} s_R \bmod q_0)$, whose reconstruction $h^* + c^* s_R \bmod q_0$, taken as a centred residue, is w_s exactly because $w_s = [f_s + \eta_s]_{q_0}$ is already reduced. Applying $v \mapsto (v, \text{Out}(v))$ to the two views of Theorem 5.11 preserves both the statistical and the computational distance, and maps the simulated pair to $(\mathcal{S}(\mathcal{M}(D)), \mathcal{M}(D))$. □

Remark 9.4 (Correction to the book)**.** The book's mechanism outputs $f_s + \eta_s$ unreduced, and its claim that the reconstruction returns it "exactly" fails on the wraparound event $\|f_s + \eta_s\|_\infty \geq q_0/2$, of probability at most $2Ne^{-t_\eta^2/2}$ per release under (Q). Either the functionality outputs the centred residue, as above (nothing changes for privacy), or that probability is added to the slack. The book's general statement (SIM-CDP with an output-consistent simulator implies security jointly with the outputs, and the converse fails without output-consistency) is correct as written.

9.3 O's work is linear in its scope

Model O's computation as reading cells of its register (one cell per row), its final state as a function of its coins and the cells read, and the decrypted statistic as a function of the final states; the protocol is $(\tau_{\text{dec}}, p_{\text{err}})$-correct if, for every input in the clipped domain, the decrypted $\mathbf{z}$ is within τ_{dec} of the true one in every entry except with probability p_{err}.

Proposition 9.5 (Work linear in the scope)**.** *Let the protocol be $(\tau_{\text{dec}}, p_{\text{err}})$-correct with $B_O^2 > 2\tau_{\text{dec}}$, and let every efficient distinguisher of O's views for two cohorts of the same size have advantage at most* adv*. If n_{Sc} rows of the scope register are non-dummy, each matchable by some cohort, then O reads at least $(1 - 2p_{\text{err}} - \text{adv})\, n_{\text{Sc}}$ rows in expectation, for every input.*

Proof. Fix a non-dummy row t and let E_t be the event that cell t is never read; it is a function of O's view. If t is matched, couple two executions with the same coins and inputs except $\mathbf{x}_{O,t} = \mathbf{0}$ in one and $B_O \mathbf{e}_1$ in the other. They coincide until cell t is read, so E_t has the same probability in both, and on E_t the decrypted statistics coincide while the true entries $\mathbf{B}_{11}$ differ by $B_O^2 > 2\tau_{\text{dec}}$; so on E_t at least one execution errs, and $\Pr[E_t] \leq 2p_{\text{err}}$, for every content of the other cells. If t is unmatched, replace one identifier of the cohort by that of row t; the indicator of E_t is an efficient function of O's view, so its probability changes by at most adv. Summing $\Pr[t \text{ read}] \geq 1 - 2p_{\text{err}} - \text{adv}$ over the n_{Sc} rows gives the bound. □

In v3.1, adv is the slack of G2 (Corollary 6.3), so O's work is $\Omega(n_{\mathrm{Sc}})$; O may skip its dummy rows, which it knows, so the bound is in n_{Sc}, not $\bar{n}_O$. The argument is that of the classical lower bound on the server's work in single-server private information retrieval.

10 Corrections to the specification

Items are ordered by importance. "Spec" refers to `protocol_v31.tex`.

Correction 1 (The centre of the mechanism; affects Lemma 1, Theorem 1 and Calibrate)**.** *Spec:* the ideal mechanism is $\bar{\mathbf{z}}_s(D) + \eta_s$ with $\bar{\mathbf{z}}_s = \mathsf{Ecd}(\mathbf{z}_s(D))$, and $e_{M,s}$ includes O's encoding roundings of V_e and $\mathbf{u}_k$.

Problem: the decrypted value differs from $\mathsf{Ecd}(\mathbf{z}_s(D))$ by deterministic, data-dependent terms. Chief among them is the mask spill $\sum_e S_e\, r_{u_e}$, of norm about $2^{30.7}$ to $2^{34.5}$ for $\bar{n}_O = 10^6$ (Remark 3.9). Flooding them would require $\log_2 \Omega \approx 67$ to 71.

Corrected statement: centre the mechanism at the noiseless circuit output $f_s(D) = \lfloor \widehat{Z}_s(D) \rceil$ (Definition 3.7), and use the sensitivity

$$\widehat{\Delta}_{\mathrm{enc},s} = \sqrt{2/N}\,\Omega_{\mathrm{fin}}\Delta_s + \Delta_{x,s} + \sqrt{N}$$

(Corollary 3.13), where $\Delta_{x,2} \approx 2^{18.8}$. Then $e_{\mathrm{tot},s}$ consists only of randomness-dependent terms (Proposition 3.8), and Theorem 5.11 holds with $\rho_s = \widehat{\Delta}^2_{\mathrm{enc},s}/(2\sigma^2_{c,s})$. Numerically this changes ρ_s by a relative $2^{-32.7}$. Lemma 1.2 itself is correct as stated.

Correction 2 (Exact final scale)**.** Throughout Lemma 1, Calibrate, σ_s and Dcd, Ω must be the exact final scale $\Omega_{\mathrm{fin}} = \Omega_U\Omega_V\Omega_u/(q_1q_2)$. It must be tracked exactly, as in FLEXIBLEAUTO, or computed from the primes. Decoding with a nominal power-of-two scale, as FIXEDMANUAL does, causes a deterministic relative bias of about 2^{-30}; we measured $8.7 \cdot 10^{-10}$ (Section 3.6). This bias is not a privacy problem once the mechanism is centred at f_s, but it is an avoidable estimation bias.

Correction 3 (Obligation O1: the value of B_{tot} and the required scale)**.** *Spec:* $\log_2 \Omega \approx 52.5$ at $\varepsilon = 1$.

Problem: that figure floods only the fresh-encryption error, and only at its mean ℓ_2 norm. For the circuit of Algorithm 4 as written, Theorem 3.12 gives a rigorous worst-case bound of $B_{\mathrm{tot}} \approx 2^{25.3}$ for stage 2, with $\delta_B = 2^{-50}$. That requires $\log_2 \Omega \approx 61.5$, or 60.5 if $P \geq 2^{20}\tilde{Q}_{\max}$ (Table 1).

Recommended changes: the following give $B_{\mathrm{tot}} \approx 2^{17.6}$ and $\log_2 \Omega \approx 53.7$:

(i) evaluate EvalSum before the first rescale, at level 2;
(ii) sum the masked products of a release *before* a single second rescale, instead of rescaling each P_e;
(iii) use stochastic rounding in Rescale and in ModDown;
(iv) choose $P \geq 2^{10}\tilde{Q}_{\max}$;
(v) encode R's vectors at $\Omega_U = q_2 \approx 2^{60}$.

By a rough count, $\log_2(QP) \approx 320$ for (A) and ≈ 285 for (C′). Both are below the 128-bit limit of about 438 for $N = 2^{14}$ [2]. [**Not fully rigorous:** This count is approximate; the modulus chain should be re-derived with the chosen d_{num}.]

Further requirement: Calibrate must use the *worst-case*, data-independent $B_{\mathrm{tot},s}(\Omega)$ of Theorem 3.12. A value estimated from the actual ciphertexts, for example by OpenFHE's noise estimation, would depend on the data. It would then leak through $\sigma_{c,s}$, and it would not be a worst-case bound. Such an estimate may serve as a sanity check only.

Correction 4 (Lemma 2: constants)**.** *Corrected statement:* for $e \in \mathbb{Z}^N$ and $\eta \sim D_{\mathbb{Z}^N,\sigma_c}$ with $\sigma_c > \sigma_2 \geq \eta_\epsilon(\mathbb{Z}^N)/\sqrt{2\pi}$,

$$\mathrm{SD}(\eta + e, \eta) \leq \frac{\|e\|_2}{\sqrt{2\pi}\sqrt{\sigma_c^2 - \sigma_2^2}} + \frac{4\epsilon}{1-\epsilon}.$$

Alternatively, $\mathrm{SD}(\eta + e, \eta) \leq \|e\|_2/(2\sigma_c)$ for every $\sigma_c > 0$. With CALIBRATE's condition, the per-release distance is at most $2^{-\kappa}(1 + 2^{-80}) + 2^{-97}$ (Lemma 2.3, Corollary 2.4).

Correction 5 (δ accounting)**.** *Spec:* $\delta_{\mathrm{DP}} = \delta_{\mathrm{run}} - n_{\mathrm{rel}}(\delta_B + 2^{-\kappa})$.

Corrected statement: Theorem 5.11 gives $(\varepsilon_{\mathrm{run}}, \delta_{\mathrm{DP}})$-SIM-CDP with statistical slack

$$\gamma = n_{\mathrm{rel}}\big(\delta_B + 2^{-\kappa}(1 + 2^{-80}) + 2^{-97}\big),$$

plus the sampler-cap failures of Section 8. A single indistinguishability guarantee $(\varepsilon_{\mathrm{run}}, \delta_{\mathrm{run}})$ requires

$$\delta_{\mathrm{DP}} = \delta_{\mathrm{run}} - (1 + e^{\varepsilon_{\mathrm{run}}})\gamma$$

(Lemma 5.7). The factor $1 + e^{\varepsilon}$ is missing in the spec. For the example ($\gamma \approx 2.7 \cdot 10^{-12}$) the effect on ρ_{run} is negligible.

Correction 6 (Assumptions of Theorem 1)**.**

- *RLWE in two forms.* RLWE is needed both for the re-randomizer (one sample per release, at modulus q_0, with ternary u and $\mathcal{N}_{\mathbb{Z}}(0, 3.2^2)$ errors) and for O's key shares (multi-sample, at modulus PQ_2, with circular security for the rotation-key shares). These are Assumptions 5.9 and 5.10.
- *What "threshold key-generation security" means here.* For the one-round protocols of [17] it is exactly Assumption 5.10.
- *Unneeded assumptions.* PSI security is *not* needed for G3 under the spec's adjacency, because the identifiers are fixed and the PSI transcript is independent of $(\mathbf{x}_O, y)$. Nor is the random-oracle model.
- *An implicit requirement.* CALIBRATE's outputs must be data-independent.

Correction 7 (PSI and the disclosure claim G4)**.** G4 says that R learns only I and $\pi|_I$ from the Link phase. This requires the labeled PSI to be secure against a semi-honest receiver. For the FHE-based PSI of [6, 5], that requires *circuit privacy* of O's BFV replies, for example by noise flooding or re-randomization before the reply. Without it, the reply noise depends on O's identifier set. A modulus switch alone (Algorithm 2, line 4) is a heuristic measure. *Update:* with OPRF-keyed items and labels, the first layer of G4 holds without circuit privacy, in the random-oracle model under CDH (Section 7, Theorem 7.3); SANITIZEBFV adds a statistical second layer, proved for a model of BFV (Proposition 7.7). [*citation detail unverified:* how CLR17 and CHLR18 phrase the noise-flooding requirement]

Correction 8 (G1 and G2: assumptions and slack)**.** *Missing assumptions:* G1 needs Assumption 6.1:

- RLWE with circular security in s_R, covering R's rotation-key shares;
- the random oracle for the seed expansion;
- IND-CPA of the PSI's BFV, with circular security for any relinearization key;
- obliviousness of the OPRF.

Simulator inputs: the simulator also takes the public parameters and O's own input.

Slack: G1 and G2 hold up to the probability p_{cuckoo} that R's cuckoo hashing fails observably.

Correction 9 (Correctness)**.** The q_0 condition in CALIBRATE is sufficient once $\Omega = \Omega_{\mathrm{fin}}$ (Proposition 8.1). The spec should state the per-release correctness-failure probability, $\delta_B + 2Ne^{-50} \approx 2^{-50}$.

The sharper condition (Q) saves about 7 bits of q_0. The "modulus switch to q_0" in Algorithm 4 is a no-op, or an exact limb drop, for a chain $q_0q_1q_2$. A scaling modulus switch would add a rounding term bounded like T_5.

Correction 10 (Sampler)**.**

- Set $\sigma_{c,s}^2$ to a rational, for example $\lceil \widehat{\Delta}_{\mathrm{enc},s}^2/(2\rho_s) \rceil$, as CKS20's exact sampler requires.
- "Constant-time" must mean (CT) of Section 8. A direct implementation of CKS20 does not satisfy (CT).
- The slot-domain noise has variance at most $\sigma_{c,s}^2 N/(2\Omega_{\mathrm{fin}}^2)$ per real slot component. The variance of $\mathcal{N}_{\mathbb{Z}}(0,\sigma^2)$ is below σ^2, and at $\sigma \gg 1$ the difference is negligible. So $\sigma_s = \sigma_{c,s}\sqrt{N/2}/\Omega_{\mathrm{fin}}$ is an upper bound on the noise s.d., not an exact value.

Remark 10.1 (Statements verified as correct)**.** The following claims of the spec were checked and need no correction:

- the identity $\|\tau^{-1}(\mathbf{z})\|_2 = \sqrt{2/N}\|\mathbf{z}\|_2$ and the rounding term $\sqrt{N}$ of Lemma 1;
- the zCDP cost $\rho_s = \Delta_{\mathrm{enc},s}^2/(2\sigma_{c,s}^2)$ of the lattice discrete Gaussian, once Δ_{enc} is replaced by $\widehat{\Delta}_{\mathrm{enc}}$;
- the validity of the conversion used in CALIBRATE;
- the release format (c_1', h_O) with c_0' withheld, which removes the C3 leak;
- the need for fresh (u, e_0, e_1, η) in every release;
- the absence of any decryption share from R.

The spec's remark that stage 1's error "cannot depend on the entries released only in stage 2" is true but not needed: flooding covers any correlation.

References

[1] A. Al Badawi, J. Bates, F. Bergamaschi, D. B. Cousins, S. Erabelli, N. Genise, S. Halevi, H. Hunt, A. Kim, Y. Lee, Z. Liu, D. Micciancio, I. Quah, Y. Polyakov, R. V. Saraswathy, K. Rohloff, J. Saylor, D. Suponitsky, M. Triplett, V. Vaikuntanathan, V. Zucca. OpenFHE: Open-source fully homomorphic encryption library. In *WAHC '22: Workshop on Encrypted Computing & Applied Homomorphic Cryptography*, ACM, 2022. [*citation detail unverified:* full author list]

[2] M. Albrecht, M. Chase, H. Chen, J. Ding, S. Goldwasser, S. Gorbunov, S. Halevi, J. Hoffstein, K. Laine, K. Lauter, S. Lokam, D. Micciancio, D. Moody, T. Morrison, A. Sahai, V. Vaikuntanathan. Homomorphic Encryption Security Standard. HomomorphicEncryption.org, Toronto, Canada, November 2018.

[3] M. Bun and T. Steinke. Concentrated differential privacy: Simplifications, extensions, and lower bounds. In *Theory of Cryptography (TCC 2016-B)*, LNCS 9985, pages 635–658. Springer, 2016.

[4] C. L. Canonne, G. Kamath, and T. Steinke. The discrete Gaussian for differential privacy. In *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*, 2020. Extended version in *Journal of Privacy and Confidentiality* 12(1), 2022.

[5] H. Chen, Z. Huang, K. Laine, and P. Rindal. Labeled PSI from fully homomorphic encryption with malicious security. In *ACM Conference on Computer and Communications Security (CCS)*, pages 1223–1237, 2018.

[6] H. Chen, K. Laine, and P. Rindal. Fast private set intersection from homomorphic encryption. In *ACM Conference on Computer and Communications Security (CCS)*, pages 1243–1255, 2017.

[7] J. H. Cheon, A. Kim, M. Kim, and Y. Song. Homomorphic encryption for arithmetic of approximate numbers. In *Advances in Cryptology – ASIACRYPT 2017*, LNCS 10624, pages 409–437. Springer, 2017.

[8] J. H. Cheon, K. Han, A. Kim, M. Kim, and Y. Song. A full RNS variant of approximate homomorphic encryption. In *Selected Areas in Cryptography (SAC 2018)*, LNCS 11349, pages 347–368. Springer, 2019.

[9] Q. Guo, D. Nabokov, E. Suvanto, and T. Johansson. Key recovery attacks on approximate homomorphic encryption with non-worst-case noise flooding countermeasures. In *33rd USENIX Security Symposium*, 2024. [*citation detail unverified:* author list and exact title]

[10] K. Han and D. Ki. Better bootstrapping for approximate homomorphic encryption. In *Topics in Cryptology – CT-RSA 2020*, LNCS 12006, pages 364–390. Springer, 2020. [*citation detail unverified:* page numbers]

[11] D. Hsu, S. M. Kakade, and T. Zhang. A tail inequality for quadratic forms of subgaussian random vectors. *Electronic Communications in Probability*, 17(52):1–6, 2012.

[12] B. Li and D. Micciancio. On the security of homomorphic encryption on approximate numbers. In *Advances in Cryptology – EUROCRYPT 2021*, LNCS 12696, pages 648–677. Springer, 2021.

[13] D. Micciancio and O. Regev. Worst-case to average-case reductions based on Gaussian measures. *SIAM Journal on Computing*, 37(1):267–302, 2007.

[14] D. Micciancio and M. Walter. Gaussian sampling over the integers: Efficient, generic, constant-time. In *Advances in Cryptology – CRYPTO 2017*, LNCS 10402, pages 455–485. Springer, 2017.

[15] I. Hwang, Y. Hwang, M. Kim, D. Lee, and Y. Song. On the (in)security of approximate computation protocols from CKKS. Cryptology ePrint Archive, Paper 2025/395, 2025. Earlier title: Provably secure approximate computation protocols from CKKS. [*citation detail unverified:* later title and venue (ASIACRYPT 2026) from search-engine extracts only]

[16] I. Mironov, O. Pandey, O. Reingold, and S. Vadhan. Computational differential privacy. In *Advances in Cryptology – CRYPTO 2009*, LNCS 5677, pages 126–142. Springer, 2009.

[17] C. Mouchet, J. Troncoso-Pastoriza, J.-P. Bossuat, and J.-P. Hubaux. Multiparty homomorphic encryption from ring-learning-with-errors. *Proceedings on Privacy Enhancing Technologies*, 2021(4):291–311, 2021.

Protocol v3.1 — statistical appendix

Proofs of the sensitivity, concentration and inference claims (S0–S6)

Technical appendix to `protocol_v31.tex`

Abstract

We prove the statistical claims of Protocol v3.1. Proposition S1 gives the exact replace-one ℓ_2 sensitivity of the released vector $(\operatorname{triu}\mathbf{B}, \mathbf{C}, \mathbf{c}_R, \mathbf{c}_O, \mathbf{y}^\top\mathbf{y})$ in all three regimes, including a proof that the norms sit at their bounds (and where they do not), the triu-versus-Frobenius question, the case $d_O = 1$, the sensitivities $\Delta_{\neg yy}$, Δ_G, Δ_y, and the set of bounds on which $\mathbf{y}^\top\mathbf{y}$ is free. Lemma S2 bounds the operator norm of the masked symmetric noise matrix; Lemma S3 analyses the two-stage choice of the ridge; Lemma S4 bounds $\|\boldsymbol{\beta}\|$; Proposition S5 states the coverage of both intervals of Algorithm 6, separating exact, first-order and asymptotic parts; Lemma S6 analyses the parametric-bootstrap RSS correction. Sections S7–S10, added after a second reading of the companion book on the protocol, prove the exact stage-1 radius and its transfer to the lattice noise, a bracket for the switch-off of the ridge and a bound on the downward bias of $\lambda_{\min}$, the mechanism bias with explicit constants, the bias of the plug-in residual variance used in the FWS comparison, a valid bias-aware interval under binding clipping, and a closed form for the full-record sensitivity. Section 14 lists every place where the specification is false, imprecise or needs an extra condition, with corrected statements. Every step that is not fully rigorous is marked [first order], [asymptotic], [numerical] or [open].

Contents

1 Summary **2**

2 Setting, notation and the noise model **2**

3 Proposition S1: replace-one sensitivity **5**
3.1 The change vector . . . 5
3.2 Statement . . . 5
3.3 Two lemmas . . . 6
3.4 Proof of Theorem 3.1 . . . 7
3.5 The case $d_O = 1$ . . . 8
3.6 Numerical verification . . . 9

4 Lemma S2: operator norm of the masked noise **9**

5 Lemma S3: the two-stage choice of the ridge **11**

6 Lemma S4: the bound on $\|\boldsymbol{\beta}\|$ **12**

7 Proposition S5: coverage **13**

8 Lemma S6: the bootstrap RSS correction and covariance **15**

9 S7: the exact stage-1 radius and the transfer to the protocol **17**

10 S8: where the ridge switches off **18**

11 S9: mechanism bias with explicit constants **19**

12 S10: the FWS residual variance, binding clipping, and the full-record sensitivity **20**
12.1 The plug-in residual variance of the FWS comparison 20
12.2 Binding clipping and the bias-aware interval 20
12.3 The full-record sensitivity in closed form 21

13 Simulation of Algorithm 6 end to end **22**

14 Corrections to the specification **23**

15 Steps not made fully rigorous **25**

1 Summary

2 Setting, notation and the noise model

Notation. We follow `protocol_v31.tex`. R's block has $d_A := d_R + 1$ coordinates (intercept included) and O's block d_O, so $p = d_A + d_O$. For a matched record, $\mathbf{a} := \dot{\mathbf{x}} \in \mathbb{R}^{d_A}$ is R's full row ($\|\mathbf{a}\| \le B_R$, $a_1 = 1$), $\mathbf{x} := \mathbf{x}_O \in \mathbb{R}^{d_O}$ ($\|\mathbf{x}\| \le B_O$) and y ($|y| \le B_y$). We write

$$r = B_R^2, \quad o = B_O^2, \quad w = B_y^2, \quad b = B_y.$$

Π_O is the coordinate projection on O's block. The mask $\chi \in \{0,1\}^{p\times p}$ has $\chi_{ij} = 0$ iff $i, j \le d_A$. For the ridge chosen by Algorithm 6 we write

$$\mathbf{G}_\lambda = \mathbf{G} + \lambda\Psi, \qquad \mathbf{P}_\lambda = \mathbf{G}_\lambda^{-1}, \qquad \boldsymbol{\theta}_\lambda = \mathbf{P}_\lambda \mathbf{c}, \qquad \vartheta_0 = \frac{\lambda_{\min}(\mathbf{G}_\lambda)}{2\sigma_2\sqrt{p}}, \qquad \hat{\vartheta} = \frac{\lambda_{\min}(\hat{\mathbf{G}}_2 + \lambda\Psi)}{2\sigma_2\sqrt{p}}.$$

ϑ_0 is the (unobservable) validity ratio of the ridge at the true Gram matrix; $\hat{\vartheta}$ is the gate statistic of Algorithm 6, line 7. The notation table of the specification calls $\hat{\vartheta}$ "the validity ratio"; the guarantee of Lemma S3 is about ϑ_0. σ_{yy} is the noise s.d. of $\widehat{yy}$ (σ_2 in joint mode, σ_y in separate mode). $z_a = \Phi^{-1}(a)$, ϕ and Φ are the standard normal density and distribution function, and $Q = 1 - \Phi$.

The masked noise model. A *masked symmetric noise matrix* with scale σ is $\mathbf{E} = \mathbf{E}^\top \in \mathbb{R}^{p\times p}$ with $E_{ij} = 0$ when $\chi_{ij} = 0$ and $\{E_{ij} : i \le j,\ \chi_{ij} = 1\}$ independent, each $\mathcal{N}(0, \sigma^2)$ (the diagonal of the O block included, since triu $\mathbf{B}$ is released entry by entry). Its covariance is

$$\mathrm{Cov}(E_{ij}, E_{kl}) = \sigma^2\chi_{ij}\big(\delta_{ik}\delta_{jl} + \delta_{il}\delta_{jk} - \delta_{ij}\delta_{ik}\delta_{il}\big). \tag{1}$$

Stage s releases $\hat{\mathbf{G}}_s = \mathbf{G} + \mathbf{E}_s$ and (stage 2) $\hat{\mathbf{c}}_2 = \mathbf{c} + \mathbf{f}_2$ with $\mathbf{f}_2 \sim \mathcal{N}(0, \sigma_2^2\mathbf{I}_p)$ (every entry of $\mathbf{c}_R$ and $\mathbf{c}_O$ depends on $\mathbf{y}$), and $\widehat{yy} = \mathbf{y}^\top\mathbf{y} + g$ with $g \sim \mathcal{N}(0, \sigma_{yy}^2)$; the noises of different releases are independent and independent of the data. For $\boldsymbol{\theta} \in \mathbb{R}^p$ we use

$$\mathbf{K}(\boldsymbol{\theta}) := \sigma^{-2}\,\mathrm{Cov}(\mathbf{E}\boldsymbol{\theta}), \qquad K(\boldsymbol{\theta})_{ik} = \delta_{ik}\Big(\sum_j \chi_{ij}\theta_j^2 - \chi_{ii}\theta_i^2\Big) + \chi_{ik}\theta_i\theta_k, \tag{2}$$

Item	What is proved	Status
S0	The slot noise is the image of the lattice discrete Gaussian under a scaled orthogonal map: exactly uncorrelated, σ_s^2-subgaussian, and within total variation $N(0.0202\,\sigma_c^{-2} + 0.0315\,\sigma_c^{-3} + e^{-2\pi^2\sigma_c^2})$ of i.i.d. $\mathcal{N}(0,\sigma_s^2)$ plus a perturbation of norm $\le N/(2\sqrt{2}\,\Omega)$.	exact (TV constant now proved)
S1	Exact Δ in regimes 1–2 (closed form), exact finite algorithm in regime 3; norms at bounds; triu = Frobenius supremum for $d_O \ge 2$; $d_O = 1$ exact reduction and closed-form Δ_G; $\Delta_{\neg yy}$ in closed form in all regimes; $\mathbf{y}^\top\mathbf{y}$ free if $w \le 2(r+o)$, costly if $w > 3(r+o)$.	exact; boundary of the free set [numerical]
S2	$\mathbb{E}\lambda_{\max}(\mathbf{E}) \le 2\sigma\sqrt{p}$; $\sqrt{2}\sigma$-Lipschitz concentration; $\Pr[\lambda_{\max}(\mathbf{E}) \ge r_1/\sigma_1 \cdot \sigma] \le \alpha_L$.	exact (one-sided); $\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}} \le 2\sigma\sqrt{p}$ [open]
S3	Measurability and independence; on Ω_L: $\lambda_{\min}(\mathbf{G}+\lambda\Psi) \ge 2\vartheta^\star\sigma_2\sqrt{p}$, $\hat{L} \le \lambda_{\min}(\mathbf{G})$; closed-form λ_O; gate-failure bound.	exact
S4	$\|\boldsymbol{\beta}\|^2 \le \mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}]/\lambda_{\min}(\mathbf{G})$, also for the projection estimand; U_{yy} valid at level $2\alpha_U$ as written, α_U after correction.	exact (needs conditional independence)
S5	Exact Gaussian first-order term, exact resolvent remainder identity, rigorous finite-sample coverage inequality for the oracle Wald interval; AKK inequality and union bound for the β_j interval.	plug-in [first order]; normality of sampling term [asymptotic]
S6	Second-order RSS bias in closed form; Corr removes it up to a relative $O(\vartheta_0^{-2})$ remainder; bootstrap covariance consistent to relative $O_p(\vartheta_0^{-1})$.	[first order]; constants not explicit
S7	Exact stage-1 radius: exact level $(M-j)/(M+1)$ with fresh Monte Carlo draws, confidence c over the draws; transfer to the protocol with the proved TV constant and offset bound.	exact
S8	Bracket for the switch-off, tight among Gram matrices with a given $\lambda_{\min}$; spectral lower bound on $\mathbb{E}\lambda_{\min}(\mathbf{G}+\mathbf{E})$ by Sudakov–Fernique.	exact; a spectrum-sharp tail bound [open]
S9	Leading-order mechanism bias, ridge version and decay in λ, with explicit remainder constants.	exact remainder bounds
S10	Bias of FWS's plug-in RSS with explicit remainder; a valid bias-aware interval under binding clipping with a public variance bound; closed form of the full-record sensitivity for blocks of dimension ≥ 2.	exact; data-driven variance under clipping [open]

Table 1: What is proved. Corrections to the specification are collected in Section 14; open steps in Section 15.

and $\mathbf{M}_\chi(\mathbf{P}) := \sigma^{-2}\mathbb{E}[\mathbf{EPE}] = \chi \circ \mathbf{P} + \mathrm{diag}_i\big(\sum_j \chi_{ij}P_{jj} - \chi_{ii}P_{ii}\big)$, both direct consequences of (1). $0 \preceq \mathbf{K}(\boldsymbol{\theta})$ and $\|\mathbf{K}(\boldsymbol{\theta})\| \le 2\|\boldsymbol{\theta}\|^2$.

Assumption 2.1 (Gaussian slot model, N)**.** Every released entry equals its exact value plus an independent $\mathcal{N}(0,\sigma_s^2)$ variable, independently across releases and of the data.

All statistical statements below are proved under Assumption 2.1. Lemma S0 shows how they transfer to the actual mechanism, whose noise is the linear image of a lattice discrete Gaussian.

Lemma 2.2 (S0: the slot noise)**.** *Let $\eta \sim \mathcal{N}_{\mathbb{Z}}(0,\sigma_c^2)^N$ and let $\boldsymbol{\xi}^{\mathrm{raw}} \in \mathbb{R}^m$ be the noise in m released slot values (real parts of distinct slots, or any m distinct coordinates of $(\mathrm{Re}\,\tau(\eta), \mathrm{Im}\,\tau(\eta))/\Omega$). Write $\sigma_s = \sigma_c\sqrt{N/2}/\Omega$.*

(a) $\boldsymbol{\xi}^{\mathrm{raw}} = (\sigma_s/\sigma_c)\,\mathbf{U}\eta$ for a matrix $\mathbf{U} \in \mathbb{R}^{m\times N}$ with orthonormal rows. Hence the coordinates of $\boldsymbol{\xi}^{\mathrm{raw}}$ are exactly *uncorrelated, with common variance $\sigma_s^2\,v(\sigma_c)/\sigma_c^2$, where $v(\sigma) = \mathrm{Var}\,\mathcal{N}_{\mathbb{Z}}(0,\sigma^2)$ satisfies $\sigma^2\big(1 - O(\sigma^2 e^{-2\pi^2\sigma^2})\big) \le v(\sigma) \le \sigma^2$; and $\boldsymbol{\xi}^{\mathrm{raw}}$ is σ_s^2-subgaussian: $\mathbb{E}e^{\langle \mathbf{u}, \boldsymbol{\xi}^{\mathrm{raw}}\rangle} \le e^{\sigma_s^2\|\mathbf{u}\|^2/2}$ for all $\mathbf{u}$.*

(b) There is a coupling of η with $\mathbf{Z} \sim \mathcal{N}(0, \sigma_c^2 \mathbf{I}_N)$ such that $\Pr[\eta \neq \lfloor \mathbf{Z} \rceil] \leq N \,\mathrm{TV}_1(\sigma_c)$, where $\mathrm{TV}_1(\sigma) = d_{\mathrm{TV}}\big(\mathcal{N}_{\mathbb{Z}}(0, \sigma^2), \lfloor \mathcal{N}(0, \sigma^2) \rceil\big)$ satisfies, for every $\sigma > 0$,

$$\mathrm{TV}_1(\sigma) \leq \frac{\phi(1)}{12\,\sigma^2} + \frac{V}{48\,\sigma^3} + \sum_{j \geq 1} e^{-2\pi^2\sigma^2 j^2}, \qquad V = 2\phi(0) + 8\phi(\sqrt{3}) = 1.5100,$$

that is, $\mathrm{TV}_1(\sigma) \leq 0.020165\,\sigma^{-2} + 0.031459\,\sigma^{-3} + 1.0001\,e^{-2\pi^2\sigma^2}$ for $\sigma \geq 1$. On that event $\boldsymbol{\xi}^{\mathrm{raw}} = \boldsymbol{\xi}^G + \boldsymbol{\nu}$ with $\boldsymbol{\xi}^G = (\sigma_s/\sigma_c)\mathbf{U}\mathbf{Z} \sim \mathcal{N}(0, \sigma_s^2 \mathbf{I}_m)$ exactly and $\|\boldsymbol{\nu}\| \leq N/(2\sqrt{2}\,\Omega)$.

(c) The decrypted value also carries the flooded error $e_{\mathrm{tot},s}$ and the encoding rounding. On $\{\|e_{\mathrm{tot},s}\| \leq B_{\mathrm{tot},s}\}$ (probability $\geq 1 - \delta_B$) their slot-domain contribution has ℓ_2 norm at most $N/(2\sqrt{2}\,\Omega) + \sqrt{2\pi}\,2^{-\kappa}\sigma_s$, using $\sigma_c \geq 2^{\kappa} B_{\mathrm{tot}}/\sqrt{2\pi}$.

Consequently every statement proved under Assumption 2.1 holds for Protocol v3.1 after adding $\delta_B + N\,\mathrm{TV}_1(\sigma_{c,s})$ (and the wraparound probability excluded by Calibrate*) per release to the failure probabilities, and enlarging every radius, in operator or Frobenius norm, by $\sqrt{2}\,\bar{\nu}$ with $\bar{\nu} := N/(\sqrt{2}\,\Omega) + \sqrt{2\pi}\,2^{-\kappa}\sigma_s$ (the factor $\sqrt{2}$ accounts for the mirrored off-diagonal entries of $\hat{\mathbf{B}}$). With the indicative parameters ($N = 2^{14}$, $\log_2 \Omega \approx 52.5$, $\kappa = 40$, $\sigma_c \approx 2^{50}$) both additions are below 10^{-11} in relative terms.*

Proof. (a) For $m(X) = \sum_k \eta_k X^k$, slot j is $\tau(m)_j = \sum_k \eta_k \zeta^{e_j k}$ with $\zeta = e^{i\pi/N}$ and $e_j = 5^j \bmod 2N$ odd. For $j \neq l$, $\sum_{k=0}^{N-1} \cos(\pi k e_j/N)\cos(\pi k e_l/N) = \frac{1}{2}\sum_k [\cos(\pi k(e_j - e_l)/N) + \cos(\pi k(e_j + e_l)/N)] = 0$, because $e_j \pm e_l$ is even and non-zero modulo $2N$ (the subgroup generated by 5 does not contain -1), and $\sum_{k<N} e^{i\pi k m/N} = (1 - e^{i\pi m})/(1 - e^{i\pi m/N}) = 0$ for even $m \not\equiv 0$. The same computation gives $\sum_k \cos^2 = \sum_k \sin^2 = N/2$ and $\sum_k \cos \cdot \sin = 0$. Hence the N real numbers $(\mathrm{Re}\,\tau(m), \mathrm{Im}\,\tau(m))$ equal $\sqrt{N/2}\,\mathbf{U}_0\eta$ for an orthogonal $\mathbf{U}_0$ (this is the identity $\|\tau^{-1}(\mathbf{z})\| = \sqrt{2/N}\|\mathbf{z}\|$ of Lemma 1 of the specification), and $\mathbf{U}$ is a set of m rows of $\mathbf{U}_0$. The variance and subgaussian bounds for $\mathcal{N}_{\mathbb{Z}}(0, \sigma^2)$ are in Canonne, Kamath and Steinke [6] (Section 2); the lower bound on v follows from Poisson summation, $v(\sigma) = \sigma^2 \frac{1 + \sum_{m \neq 0}(1 - 4\pi^2\sigma^2 m^2) e^{-2\pi^2\sigma^2 m^2}}{1 + \sum_{m \neq 0} e^{-2\pi^2\sigma^2 m^2}}$. Subgaussianity of $\boldsymbol{\xi}^{\mathrm{raw}}$ follows because $\langle \mathbf{u}, \boldsymbol{\xi}^{\mathrm{raw}} \rangle = (\sigma_s/\sigma_c)\langle \mathbf{U}^\top \mathbf{u}, \eta \rangle$ with $\|\mathbf{U}^\top \mathbf{u}\| = \|\mathbf{u}\|$ and independent coordinates.

(b) Couple coordinatewise with a maximal coupling; a union bound gives $N\,\mathrm{TV}_1$. Write $p_k \propto e^{-k^2/2\sigma^2}$ and $q_k = \int_{k-1/2}^{k+1/2} \phi_\sigma$, so $\mathrm{TV}_1 = \frac{1}{2}\sum_k |p_k - q_k|$. First, $p_k = \phi_\sigma(k)/\sum_l \phi_\sigma(l)$ and, by Poisson summation, $\sum_l \phi_\sigma(l) = 1 + 2\epsilon$ with $\epsilon = \sum_{j \geq 1} e^{-2\pi^2\sigma^2 j^2}$; so $\sum_k |p_k - \phi_\sigma(k)| = |1 - \sum_l \phi_\sigma(l)| = 2\epsilon$. Second, $q_k - \phi_\sigma(k) = \int_{-1/2}^{1/2} (\phi_\sigma(k + u) - \phi_\sigma(k) - u\phi'_\sigma(k))\,du$ (the linear term integrates to zero), so $|q_k - \phi_\sigma(k)| \leq \frac{1}{24}\sup_{I_k} |\phi''_\sigma|$ with $I_k = [k - \frac{1}{2}, k + \frac{1}{2}]$. For any function f of bounded variation, $\sup_{I_k} |f| \leq \int_{I_k} |f| + \mathrm{TV}_{I_k}(|f|)$, so $\sum_k \sup_{I_k} |\phi''_\sigma| \leq \int |\phi''_\sigma| + \mathrm{TV}(|\phi''_\sigma|) = 4\phi(1)\sigma^{-2} + V\sigma^{-3}$: $|\phi''(x)| = \phi(x)|x^2 - 1|$ for the standard density has zeros at ± 1 and local maxima $\phi(0)$ at 0 and $2\phi(\sqrt{3})$ at $\pm\sqrt{3}$, so its total variation is $2\phi(0) + 8\phi(\sqrt{3})$, and $\int |\phi''| = 4\phi(1)$. Halving the sum of the two parts gives the bound. It is within a factor 1.02 of the exact value at $\sigma = 100$ and within a factor 3.2 at $\sigma = 1$ (exact high-precision evaluation, `experiments/proofs_check/tv1_cuckoo.py`); at the evaluation setting ($N = 2^{14}$, $\sigma_c \approx 2^{65.3}$) $N\,\mathrm{TV}_1(\sigma_c) \leq 2^{-122.2}$. (An earlier version gave only the leading constant, with the $O(\sigma^{-3})$ term checked numerically; so does the companion book's Lemma `rel.slot`.) Finally $\|\lfloor \mathbf{Z} \rceil - \mathbf{Z}\| \leq \sqrt{N}/2$ and $\|(\sigma_s/\sigma_c)\mathbf{U}\|_{\mathrm{op}} = \sqrt{N/2}/\Omega$.

(c) The rounding of Ecd has sup-norm $\leq 1/2$; the slot image of a ring element e has norm $\sqrt{N/2}\|e\|/\Omega$, and $\sqrt{N/2}B_{\mathrm{tot}}/\Omega \leq \sqrt{N/2}\sqrt{2\pi}\sigma_c 2^{-\kappa}/\Omega = \sqrt{2\pi}2^{-\kappa}\sigma_s$. □

3 Proposition S1: replace-one sensitivity

3.1 The change vector

Replacing $(\mathbf{x}, y)$ by $(\mathbf{x}', y')$ in a matched record with R-row $\mathbf{a}$ changes $(\operatorname{triu}\mathbf{B}, \mathbf{C}, \mathbf{c}_R, \mathbf{c}_O, \mathbf{y}^\top\mathbf{y})$ by

$$\big(\operatorname{triu}\mathbf{D},\ \mathbf{a}(\mathbf{x}-\mathbf{x}')^\top,\ \mathbf{a}(y-y'),\ \mathbf{x}y-\mathbf{x}'y',\ y^2-y'^2\big), \qquad \mathbf{D} := \mathbf{x}\mathbf{x}^\top - \mathbf{x}'\mathbf{x}'^\top.$$

A change in an unmatched (or dummy) row changes nothing, because its O-side entries are multiplied by $b_t = 0$ or by $\dot{\mathbf{x}}_t = \mathbf{0}$ (Algorithm 4). Hence $\Delta^2 = \max_{\mathcal{D}} \Phi_{\text{triu}}$ with

$$\Phi_{\text{triu}}(\mathbf{a}, \mathbf{x}, \mathbf{x}', y, y') = \|\operatorname{triu}\mathbf{D}\|^2 + \|\mathbf{a}\|^2\big(\|\mathbf{x}-\mathbf{x}'\|^2 + (y-y')^2\big) + \|\mathbf{x}y - \mathbf{x}'y'\|^2 + (y^2-y'^2)^2,$$
$$\mathcal{D} = \{\|\mathbf{a}\| \le B_R,\ \|\mathbf{x}\|, \|\mathbf{x}'\| \le B_O,\ |y|, |y'| \le B_y\}$$

(compact, so the supremum is a maximum). Φ_F denotes the same expression with $\|\operatorname{triu}\mathbf{D}\|^2$ replaced by $\|\mathbf{D}\|_F^2$, and $\Phi_{\text{triu}}^{\neg yy}$, $\Phi_F^{\neg yy}$ the expressions without the last term. The intercept constraint $a_1 = 1$ is irrelevant: the maximum has $\|\mathbf{a}\| = B_R \ge 1$, which is attained by $\mathbf{a} = (1, \mathbf{z})$ with $\|\mathbf{z}\|^2 = r - 1$ whenever $d_R \ge 1$ (if $d_R = 0$ then $B_R = 1$). Δ is the supremum over all cohorts; for a given cohort the sensitivity can only be smaller.

Define on $[-o, o] \times [-b, b]$

$$F(s,t) = 2o^2 - 2s^2 + 2r(o-s) + r(b-t)^2 + o(w+t^2) - 2bts + (w-t^2)^2, \tag{3}$$
$$F^{\neg yy}(s,t) = F(s,t) - (w-t^2)^2, \qquad s^*(t) = \operatorname{clip}\big(-(r+bt)/2;\ -o,\ o\big), \qquad G(t) = F(s^*(t), t).$$

F is the function of the previous paper (with $s = oc$), so the reduction below is the missing justification of its use.

3.2 Statement

Theorem 3.1 (S1)**.** *Let $d_O \ge 2$ and $b > 0$ ($b = 0$ is covered by setting $w = 0$).*

(i) ***Reduction.*** $\Delta^2 = \max_{|s|\le o, |t|\le b} F(s,t) = \max_{|t|\le b} G(t)$.

(ii) ***Regime 1****, $|r - w| \le 2o$: $\Delta^2 = \frac{1}{2}(2o + r + w)^2 + 2rw$, attained exactly at $t = -b$, $s = (w - r)/2$.*

(iii) ***Regime 2****, $r - w \ge 2o$: $\Delta^2 = 4r(o + w)$, attained exactly at $t = -b$, $s = -o$.*

(iv) ***Regime 3****, $w - r > 2o$: with $t_1 = -(r + 2o)/b$ and $t_2 = (2o - r)/b$ one has $-b < t_1 < t_2 < b$ and G is C^1 and piecewise quartic:*

$$G_A(t) = t^4 + (r + o - 2w)t^2 - 2b(r + o)t + (r + o)w + w^2 \qquad t \in [-b, t_1]\ (s^* = o),$$
$$G_B(t) = t^4 + (r + o - \tfrac{3}{2}w)t^2 - rbt + 2o^2 + 2ro + (r + o)w + w^2 + \tfrac{1}{2}r^2 \qquad t \in [t_1, t_2]\ (|s^*| < o),$$
$$G_C(t) = t^4 + (r + o - 2w)t^2 + 2b(o - r)t + 4ro + (r + o)w + w^2 \qquad t \in [t_2, b]\ (s^* = -o).$$

$\Delta^2 = \max_{t\in\mathcal{T}} G(t)$ for the finite set $\mathcal{T} = \{-b, t_1, t_2, b\} \cup \{$real roots of G_A', G_B', G_C' in their intervals$\}$; this is exact (three cubics). $t = -b$ is a strict local maximizer with $G(-b) = 4w(r + o)$; if $\Delta^2 > 4w(r + o)$, every maximizer has $|t^*| < b$.

(v) ***Norms.*** *At every maximizer $\|\mathbf{a}\| = B_R$, $\|\mathbf{x}\| = \|\mathbf{x}'\| = B_O$ and $\max(|y|, |y'|) = B_y$. Both $|y| = |y'| = B_y$ in regimes 1–2; in regime 3 this fails whenever $\mathbf{y}^\top\mathbf{y}$ is not free ($|y'| = |t^*| < B_y$).*

(vi) triu ***versus Frobenius.*** $\max_{\mathcal{D}} \Phi_{\text{triu}} = \max_{\mathcal{D}} \Phi_F$ *(and the same without $\mathbf{y}^\top\mathbf{y}$); a configuration attains the* triu *maximum iff it attains the Frobenius maximum and $\mathbf{D}$ is diagonal.*

(vii) ***Worst case.*** *Let $s^\dagger$, $t^\dagger$ be a maximizer of F and $c = s^\dagger/o$, $\alpha_\pm = \frac{1}{2}(\sqrt{1+c} \pm \sqrt{1-c})$. A worst-case pair is $\|\mathbf{a}\| = B_R$, $y = B_y$, $y' = t^\dagger$,*

$$\mathbf{x} = B_O(\alpha_+\mathbf{e}_1 + \alpha_-\mathbf{e}_2), \qquad \mathbf{x}' = B_O(\alpha_-\mathbf{e}_1 + \alpha_+\mathbf{e}_2),$$

for which $\mathbf{D} = o\sqrt{1-c^2}\,(\mathbf{e}_1\mathbf{e}_1^\top - \mathbf{e}_2\mathbf{e}_2^\top)$. *In regime 1:* $y' = -y$, $\cos\angle(\mathbf{x},\mathbf{x}') = -(r-w)/(2o)$; *in regime 2:* $y' = -y$, $\mathbf{x}' = -\mathbf{x}$ *(*$\mathbf{D} = 0$*).*

(viii) ***Other sensitivities.*** $\Delta^2_{\neg yy} = \frac{1}{2}(2o + r + w)^2 + 2rw$ *if* $|r-w| \le 2o$ *and* $\Delta^2_{\neg yy} = 4M(o+m)$ *otherwise, with* $M = \max(r,w)$ *and* $m = \min(r,w)$. $\Delta_G = \Delta(B_R, B_O, 0)$ *equals* $(2o+r)/\sqrt{2}$ *if* $r \le 2o$ *and* $2\sqrt{ro}$ *otherwise.* $\Delta_y = w$.

(ix) ***When*** $\mathbf{y}^\top\mathbf{y}$ ***is free.*** $\Delta = \Delta_{\neg yy}$ *iff* $\max_{|t|\le b} G(t) \le 4w(r+o)$ *or* $w \le r + 2o$. *In particular it is free whenever* $w \le 2(r+o)$ *(this covers regimes 1, 2 and part of regime 3), and costly whenever* $w > 3(r+o)$. *For* $o = 0$ *the exact threshold is* $w \le \frac{11+5\sqrt{5}}{8}\, r \approx 2.7725\, r$.

The specification's DeltaRep ("exact for $d_O \ge 2$") is therefore correct in regimes 1–2 and, after the reduction (i), in regime 3, where (iv) replaces the grid search by an exact finite computation. For $d_O = 1$ see Proposition 3.5.

3.3 Two lemmas

Lemma 3.2 (triu versus Frobenius)**.** *For all* $\mathbf{x}, \mathbf{x}'$: $\|\operatorname{triu}\mathbf{D}\|^2 = \frac{1}{2}\big(\|\mathbf{D}\|_F^2 + \sum_j D_{jj}^2\big) \le \|\mathbf{D}\|_F^2$, *with equality iff* $\mathbf{D}$ *is diagonal. If* $d_O \ge 2$, *there is an orthogonal* $\mathbf{Q}$ *such that* $\mathbf{Q}\mathbf{D}\mathbf{Q}^\top$ *is diagonal; replacing* $(\mathbf{x},\mathbf{x}')$ *by* $(\mathbf{Q}\mathbf{x},\mathbf{Q}\mathbf{x}')$ *changes* $\mathbf{D}$ *into* $\mathbf{Q}\mathbf{D}\mathbf{Q}^\top$ *and leaves every other term of* Φ *and the constraints unchanged. Hence* $\max\Phi_{\text{triu}} = \max\Phi_F$. *For* $d_O = 1$, $\operatorname{triu}\mathbf{D} = \mathbf{D}$.

Proof. The identity holds because $\mathbf{D}$ is symmetric. $\mathbf{D}$ has rank ≤ 2 and range in $\operatorname{span}\{\mathbf{x},\mathbf{x}'\}$, so $\mathbf{D} = \mu_1\mathbf{q}_1\mathbf{q}_1^\top + \mu_2\mathbf{q}_2\mathbf{q}_2^\top$ with orthonormal $\mathbf{q}_1, \mathbf{q}_2$; any orthogonal $\mathbf{Q}$ with $\mathbf{Q}\mathbf{q}_i = \mathbf{e}_i$ diagonalizes $\mathbf{D}$ (this needs $d_O \ge 2$). The other terms depend on $(\mathbf{x},\mathbf{x}')$ only through $\|\mathbf{x}\|, \|\mathbf{x}'\|, \mathbf{x}^\top\mathbf{x}'$, which are invariant. Thus $\Phi_{\text{triu}}(\cdot, \mathbf{Q}\mathbf{x}, \mathbf{Q}\mathbf{x}', \cdot) = \Phi_F(\cdot, \mathbf{x}, \mathbf{x}', \cdot) \ge \Phi_{\text{triu}}(\cdot, \mathbf{x}, \mathbf{x}', \cdot)$. □

Lemma 3.3 (Reduction)**.** *Let* $d_O \ge 2$. *With* $\rho = \|\mathbf{a}\|^2$, $u = \|\mathbf{x}\|^2$, $v = \|\mathbf{x}'\|^2$, $s = \mathbf{x}^\top\mathbf{x}'$,

$$\Phi_F = \Upsilon := u^2 + v^2 - 2s^2 + \rho\big(u + v - 2s + (y-y')^2\big) + y^2u + y'^2v - 2yy's + (y^2 - y'^2)^2,$$

and the image of $\mathcal{D}$ *is* $\{0 \le \rho \le r,\ 0 \le u, v \le o,\ s^2 \le uv,\ |y|, |y'| \le b\}$.

(a) Υ *is non-decreasing in* ρ, *strictly unless* $(\mathbf{x}, y) = (\mathbf{x}', y')$.

(b) For fixed (ρ, v, s, y, y'), Υ *is strictly increasing in* u *on the feasible set* $\{u \in [0, o] : s^2 \le uv\}$, *which is an interval with right end* o. *Hence every maximizer has* $u = o$, *and likewise* $v = o$.

(c) For fixed (ρ, u, v, s), $\Upsilon = \mathcal{Q}(y, y') + (y^2 - y'^2)^2 + \text{const}$ *with* $\mathcal{Q}$ *positive semidefinite; hence* $\Upsilon(\kappa y, \kappa y') \ge \Upsilon(y, y')$ *for* $\kappa \ge 1$, *and every maximizer has* $\max(|y|, |y'|) = b$.

(d) Υ *is invariant under* $(u, y) \leftrightarrow (v, y')$ *and under* $(y, y') \mapsto (-y, -y')$.

Consequently $\Delta^2 = \max F$ *over* $[-o, o] \times [-b, b]$: *set* $\rho = r$, $u = v = o$, $y = b$, $y' = t$.

Proof. For $d_O \ge 2$, every (u, v, s) with $s^2 \le uv$ is realized by some $(\mathbf{x}, \mathbf{x}')$. (a) The coefficient of ρ is $\|\mathbf{x} - \mathbf{x}'\|^2 + (y - y')^2$. (b) $\partial_u\Upsilon = 2u + \rho + y^2 > 0$ (as $\rho \ge 1$ on $\mathcal{D}$); the feasible set in u is $[s^2/v, o]$ (or $[0, o]$ if $v = 0$), which is upward closed. This is the step the earlier proof sketch asserted without proof: monotonicity holds at fixed inner product s, *not* at fixed cosine (at fixed cosine c, $\partial_u(-2s^2) = -2vc^2$ can dominate). (c) $\mathcal{Q}(y, y') = \rho(y - y')^2 + uy^2 + vy'^2 - 2syy'$ has matrix $\left(\begin{smallmatrix} \rho+u & -(\rho+s) \\ -(\rho+s) & \rho+v \end{smallmatrix}\right)$ with non-negative diagonal and determinant $\rho(u + v - 2s) + uv - s^2 \ge 0$. So $\Upsilon(\kappa y, \kappa y') - \Upsilon(y, y') = (\kappa^2 - 1)\mathcal{Q} + (\kappa^4 - 1)(y^2 - y'^2)^2 \ge 0$. If a maximizer had $\max(|y|, |y'|) < b$, scaling would not increase Υ only if $\mathcal{Q}(y, y') = 0$ and $y^2 = y'^2$. Either $y = y' = 0$, and then moving to $(b, -b)$ increases Υ by $\mathcal{Q}(b, -b) = w(4\rho + u + v + 2s) > 0$; or $\det = 0$, which with $\rho > 0$ forces $u = v = s$, i.e. $\mathbf{x} = \mathbf{x}'$, and then $\mathcal{Q} = (\rho + u)(y - y')^2 = 0$ forces $y = y'$ and $\Upsilon = 0$, not a maximum. (d) is immediate. By (d) we may take $|y| = b$ and then $y = b$. □

3.4 Proof of Theorem 3.1

(i) Lemmas 3.2 and 3.3. For fixed t, F is a strictly concave quadratic in s, $F(s,t) = -2s^2 - 2s(r + bt) + H(t)$ with $H(t) = 2o^2 + 2ro + r(b-t)^2 + o(w+t^2) + (w-t^2)^2$, maximized over $[-o, o]$ at $s^*(t)$.

(ii) Relax $s \in \mathbb{R}$: $G_u(t) := \max_{s\in\mathbb{R}} F = H(t) + \frac{1}{2}(r+bt)^2 = t^4 + (r + o - \frac{3}{2}w)t^2 - rbt + K$, with K the constant term of G_B. Using $w - t^2 = (b-t)(b+t)$ and $w^2 - t^4 = (w-t^2)(w+t^2)$,

$$G_u(-b) - G_u(t) = (b+t)\big[rb + (b-t)\big(r + o - \tfrac{w}{2} + t^2\big)\big]. \tag{4}$$

If $w \le r + 2o$ then $r + o - w/2 \ge r/2 > 0$, so (4) is > 0 for $t > -b$. If also $r - w \le 2o$, the relaxed maximizer at $t = -b$, $s = -(r-w)/2$, is feasible. Hence $\Delta^2 = G_u(-b) = 2o^2 + 2o(r+w) + 4rw + \frac{1}{2}(r-w)^2 = \frac{1}{2}(2o + r + w)^2 + 2rw$, with a unique maximizer (strict concavity in s).

(iii) If $r - w \ge 2o$ then $-(r+bt)/2 \le -(r-w)/2 \le -o$ for all $|t| \le b$, so $s^* \equiv -o$ and $G(t) = F(-o,t) = 4ro + r(b-t)^2 + o(b+t)^2 + (w-t^2)^2$. Then

$$G(-b) - G(t) = (b+t)\big[r(3b-t) - o(b+t) - (b-t)^2(b+t)\big].$$

With $r \ge 2o + w$ and $t = b\tau$, the bracket is at least $o(5b-3t) + b^3\big[(3-\tau) - (1-\tau)^2(1+\tau)\big] = o(5b-3t) + b^3(2 + \tau^2 - \tau^3) \ge 2ob + 2b^3 > 0$. So the maximum is $G(-b) = 4r(o+w)$, attained only at $t = -b$.

(iv) $s^*(t) = o$ iff $r + bt \le -2o$ iff $t \le t_1$, and $s^*(t) = -o$ iff $t \ge t_2$; in regime 3, $t_1 > -b$ and $t_2 < b$. Substituting $s = \pm o$ or $s = -(r+bt)/2$ in (3) gives G_A, G_B, G_C. By Danskin's theorem [8] (unique maximizer $s^*(t)$, continuous in t), $G'(t) = \partial_t F(s^*(t), t)$ is continuous, so a maximizer of G on $[-b, b]$ is an endpoint or a zero of G', i.e. a root of the cubic G'_k of the piece that contains it. Further $G_A(-b) = 4w(r+o)$, $G'_A(-b) = -4b(r+o) < 0$, and $G(b) = G_C(b) = 4o(r+w) < 4w(r+o)$ because $w > r$ in regime 3. So if $\max G > 4w(r+o)$ every maximizer lies in $(-b, b)$.

(v) From Lemma 3.3 and the uniqueness statements in (ii)–(iv).

(vi) Lemma 3.2. If $\mathbf{D}$ is not diagonal at a Frobenius maximizer, the triu value is strictly smaller.

(vii) With $\|\mathbf{x}\| = \|\mathbf{x}'\| = B_O$ and $\mathbf{x}^\top\mathbf{x}' = oc$, $\mathbf{D}$ has eigenvalues $\pm o\sqrt{1-c^2}$ and triu attains $\|\mathbf{D}\|_F$ iff the eigenvectors are coordinate vectors; the stated pair has $\alpha_+^2 + \alpha_-^2 = 1$, $2\alpha_+\alpha_- = c$ and $\alpha_+^2 - \alpha_-^2 = \sqrt{1-c^2}$, which gives the displayed diagonal $\mathbf{D}$. A pair with the right norms and angle but non-diagonal $\mathbf{D}$ is *not* worst case: for $\mathbf{x} = B_O\mathbf{e}_1$, $\mathbf{x}' = B_O(c\mathbf{e}_1 + \sqrt{1-c^2}\mathbf{e}_2)$ one loses $o^2c^2(1-c^2)$ in $\|\operatorname{triu}\mathbf{D}\|^2$.

(viii) Lemma 3.3(a),(b) hold verbatim for $\Phi^{\neg yy}$. Without the quartic term, $\Upsilon^{\neg yy}$ is a positive semidefinite quadratic form in (y, y') plus a constant, hence convex, and its maximum over the square is at a vertex; by (d), $y = b$, $t = y' \in \{b, -b\}$. Now $F^{\neg yy}(s,-b) = 2o^2 + 2ro + 2ow + 4rw - 2s^2 - 2(r-w)s$ and $F^{\neg yy}(s,b) = 2o^2 - 2s^2 + 2(r+w)(o-s)$. Maximizing each over $|s| \le o$:

- $t = -b$: $\frac{1}{2}(2o + r + w)^2 + 2rw$ if $|r - w| \le 2o$; $4r(o+w)$ if $r - w > 2o$; $4w(o+r)$ if $w - r > 2o$.
- $t = b$: $\frac{1}{2}(2o + r + w)^2$ if $r + w \le 2o$, otherwise $4o(r+w)$.

Since $\frac{1}{2}(2o + r + w)^2 \ge 4o(r+w)$ (AM–GM), $4r(o+w) - 4o(r+w) = 4w(r-o) \ge 0$ in regime 2, and $4w(o+r) - 4o(r+w) = 4r(w-o) > 0$ in regime 3, the $t = -b$ values win, which gives the stated $\Delta_{\neg yy}$. Δ_G is the case $w = 0$: regime 1 if $r \le 2o$, regime 2 otherwise. $\Delta_y = \sup|y^2 - y'^2| = w$, attained at $(y, y') = (b, 0)$.

(ix) In regimes 1–2 the maximizer has $y' = -y$, so the $\mathbf{y}^\top\mathbf{y}$ term vanishes there and $\Delta = \Delta_{\neg yy}$ by the explicit formulas. In regime 3, $\Delta^2_{\neg yy} = 4w(r+o) = G(-b) \le \Delta^2$, so freeness is equivalent to $\max G \le 4w(r+o)$. By homogeneity, $F(s,t;r,o,w) = w^2F(s/w, t/b; r/w, o/w, 1)$, so take $w = b = 1$. Put $s = o - \varsigma$ with $\varsigma \in [0, 2o]$, $t = u - 1$ with $u \in [0, 2]$, $S = r + o$ and $\delta = 1 - r - 2o > 0$ (regime 3). A direct expansion (verified exactly in rational arithmetic by `s1_symbolic.py`, as are the other polynomial identities of this proof) gives

$$F(s,t) - 4S = 2\varsigma(u - \delta) - 2\varsigma^2 - S\,u(4-u) + u^2(2-u)^2. \tag{5}$$

Maximizing over $\varsigma \geq 0$ gives at most $\frac{1}{2}(u-\delta)_+^2 \leq \frac{1}{2}u^2$. If $w \leq 2(r+o)$, i.e. $S \geq \frac{1}{2}$, then $-Su(4-u) \leq -\frac{1}{2}u(4-u)$, and

$$F(s,t) - 4S \leq \tfrac{1}{2}u^2 - \tfrac{1}{2}u(4-u) + u^2(2-u)^2 = u(u-2)(u-1)^2 \leq 0 \qquad (0 \leq u \leq 2).$$

Hence $\Delta^2 \leq 4w(r+o) = \Delta^2_{\neg yy}$. If $w > 3(r+o)$, then $G(0) \geq F(o,0) = w(r+o+w) > 4w(r+o)$. For $o = 0$, $G(t) = r(b-t)^2 + (w-t^2)^2$, which is the function h of Proposition 3.5(d) with o replaced by w; the threshold follows from that computation. □

Remark 3.4 (The free set, numerically)**.** For fixed (r, o) the free set in regime 3 is an interval $r + 2o < w \leq w^*(r,o)$ [numerical: checked on a fine grid in w for $o/r \in [0, 100]$]; Theorem 3.1(ix) proves $2(r+o) \leq w^* \leq 3(r+o)$ and $w^*(r,0) = \frac{11+5\sqrt{5}}{8}r$. The ratio $w^*/(r+o)$ decreases from 2.7725 at $o = 0$ to 2 as $o/r \to \infty$:

o/r	0	0.01	0.1	0.25	0.5	1	2	5	10	100
$w^*/(r+o)$	2.773	2.767	2.725	2.677	2.623	2.534	2.414	2.250	2.153	2.019

Example: $(B_R, B_O, B_y) = (1, 1, 5)$ has $w = 25 > 3(r+o) = 6$, so $\mathbf{y}^\top\mathbf{y}$ costs: $\Delta = 26.0706$ (maximizer $t^* = -0.0704$, $s^* = -0.324$), against $\Delta_{\neg yy} = \sqrt{200} = 14.142$. For $(\sqrt{5}, \sqrt{5}, 5)$, $w = 25 \leq w^* = 25.33$, so $\mathbf{y}^\top\mathbf{y}$ is free ($\Delta = \Delta_{\neg yy} = 31.62$), which is why the earlier paper saw no difference there. The indicative bounds $(5.59, 6.24, 2.89)$ are in regime 1 (free), with $\Delta = 86.15$ and $\Delta_G = 77.16$.[1]

3.5 The case $d_O = 1$

Proposition 3.5 ($d_O = 1$)**.** *Let $d_O = 1$ and $\Phi_1(x, x', y, y') = (x^2 - x'^2)^2 + r(x-x')^2 + r(y-y')^2 + (xy - x'y')^2 + (y^2 - y'^2)^2$ (here* triu $\mathbf{D} = \mathbf{D}$*, so there is no* triu *issue).*

- *(a) $\Delta^{(1)} \leq \Delta$, $\Delta^{(1)}_{\neg yy} \leq \Delta_{\neg yy}$ and $\Delta^{(1)}_G \leq \Delta_G$; so the $d_O \geq 2$ formulas are valid (conservative) for $d_O = 1$.*
- *(b) $\Delta^{(1)2} = \max\{\max_{x',t} \Phi_1(B_O, x', b, t),\ \max_{x',t} \Phi_1(B_O, x', t, b)\}$ over $(x', t) \in [-B_O, B_O] \times [-b, b]$.*
- *(c) $\Delta^{(1)} = \Delta$ in regime 2 and whenever some maximizer of Theorem 3.1 has $|s^\dagger| = o$ (for instance in the free part of regime 3, where the maximizer is $\mathbf{x}' = \mathbf{x}$, $y' = -y$). $\Delta^{(1)} < \Delta$ in the interior of regime 1.*
- *(d) $\Delta^{(1)2}_G = 4ro$ if $o \leq \frac{11+5\sqrt{5}}{8}r$; otherwise $\Delta^{(1)2}_G = \frac{o^2}{16}(3-q)^3(1+q)$ with $q = \sqrt{1-2r/o}$.*

Proof. (a) A $d_O = 1$ configuration embedded in $\mathbb{R}^2$ (second coordinates zero) has diagonal $\mathbf{D}$ and the same value. (b) Φ_1 is non-decreasing under $(x, x') \mapsto \kappa(x, x')$, $\kappa \geq 1$ (each x-dependent term is a non-negative form of degree 4 or 2 in (x, x')), and under $(y, y') \mapsto \kappa(y, y')$ (as in Lemma 3.3(c), with $\mathcal{Q} = r(y-y')^2 + (xy - x'y')^2$, a sum of squares). So $\max(|x|, |x'|) = B_O$ and $\max(|y|, |y'|) = b$. Using the symmetries $(x, y) \leftrightarrow (x', y')$, $(x, x') \mapsto -(x, x')$ and $(y, y') \mapsto -(y, y')$, take $x = B_O$; then either $y = b$ or $y' = b$. (c) In regime 2 the maximizer $\mathbf{x}' = -\mathbf{x}$ is collinear. In the interior of regime 1 every maximizer of the $d_O \geq 2$ problem has $u = v = o$ and $s = (w-r)/2 \in (-o, o)$, so it is not collinear. The $d_O = 1$ feasible set is a compact subset of the $d_O = 2$ set that contains no maximizer, so its maximum is strictly smaller. (d) With $w = 0$, $x = B_O$, $x' = B_O\tau$: $\Phi_1 = o\,h(\tau)$ with $h(\tau) = o(1-\tau^2)^2 + r(1-\tau)^2$, and $h'(\tau) = -2(1-\tau)\big[2o\tau(1+\tau) + r\big]$. If $o \leq 2r$ the bracket is $\geq r - o/2 \geq 0$, so h decreases and the maximum is $h(-1) = 4r$. If $o > 2r$, h has a local maximum at $\tau_+ = (q-1)/2$, where $h(\tau_+) = (1-\tau_+)^2[o(1+\tau_+)^2 + r] = \frac{o}{16}(3-q)^3(1+q)$, while

[1] The specification prints $\Delta = 86.18$ and $\Delta_G = 77.20$, presumably computed from unrounded bounds; with the bounds as printed the formulas give 86.152 and 77.162.

$h(-1) = 4r = 2o(1-q^2)$. The interior value wins iff $(3-q)^3 > 32(1-q)$, i.e. $q > q_0 = 5 - 2\sqrt{5}$, i.e. $o/r > 2/(1-q_0^2) = \frac{11+5\sqrt{5}}{8}$ (and $o \le 2r$ falls in the first case). □

For $d_O = 1$ we give no closed form for $\Delta^{(1)}$ in regimes 1 and 3. (b) is an exact two-variable polynomial maximization over a rectangle, solvable to machine precision; for calibration the $d_O \ge 2$ formula is a valid upper bound. The losses are substantial in regime 1: $\Delta^{(1)}/\Delta = 0.894$ at $(1,1,1)$, 0.784 at $(1,2,1)$ and 0.912 at the indicative bounds (Table 2).

3.6 Numerical verification

Script `scratch_stats/s1_check.py` (output `s1_check.out`) compares, for 14 bound settings covering all regimes: the exact value of Theorem 3.1 (closed form, or candidate enumeration in regime 3); `sim_core.delta_rep` (read only); a brute-force multistart BFGS maximization of the triu objective Φ_{triu} over raw configurations $(\mathbf{x}, \mathbf{x}', y, y')$ for $d_O = 2$ and $d_O = 3$; and the explicit worst-case configuration of (vii) evaluated with the triu norm. All agree to 10^{-6} (relative difference to `delta_rep` is $\le 10^{-9}$, which is its deliberate rounding up). $\Delta_{\neg yy}$ agrees with brute force, and the $d_O = 1$ reduction (b) agrees with brute force over $d_O = 1$ configurations. Script `s1_freeset.py` produced the table of Remark 3.4, and `s1_band.py` confirms $\max_t G - 4(r+o) \le 4.4 \times 10^{-16}$ over a 300×120 grid of the band $r + 2o < w \le 2(r+o)$.

B_R	B_O	B_y	reg.	Δ exact	brute d_O=2	$\Delta_{\neg yy}$	$\Delta^{(1)}$ (d_O=1)	Δ_G	$\Delta_G^{(1)}$
1	1	1	1	3.162278	3.162278	3.16228	2.82843	2.12132	2.00000
1	1	5	3	26.070598	26.070598	14.14214	26.05763	2.12132	2.00000
3	1	1	2	8.485281	8.485281	8.48528	8.48528	6.00000	6.00000
2.236	2.236	5	3	31.621815	31.621815	31.62182	31.62182	10.60596	9.99939
5.59	6.24	2.89	1	86.152245	86.152245	86.15224	78.53500	77.16183	69.76320
1	2	1	1	7.211103	7.211103	7.21110	5.65685	6.36396	4.53635
4	1	0.5	2	8.944272	8.944272	8.94427	8.94427	8.00000	8.00000
1	0.5	3	3	9.672258	9.672258	6.70820	9.67226	1.00000	1.00000
1.5	0.7	2.2	3	7.283296	7.283296	7.28330	7.28330	2.10000	2.10000
1	1	1.8	3	5.091169	5.091169	5.09117	5.09117	2.12132	2.00000
1	0.2	2	3	4.566806	4.566806	4.07922	4.56681	0.40000	0.40000
1	3	4	1	25.387005	25.387005	25.38700	25.29822	13.43503	9.51475

Table 2: Proposition S1: the exact value equals `delta_rep`, the brute-force triu maximum for $d_O = 2$ and $d_O = 3$ (the $d_O = 3$ column is identical and omitted), and the explicit configuration of Theorem 3.1(vii), in every case.

4 Lemma S2: operator norm of the masked noise

Lemma 4.1 (S2). *Let* $\mathbf{E}$ *be a masked symmetric noise matrix with scale* σ *(Section 2); the proof allows any mask and any variances* $\le \sigma^2$ *off the diagonal and* $\le 2\sigma^2$ *on it, so it covers the full GOE-type matrix too. Let* $\mathbf{g} \sim \mathcal{N}(0, \mathbf{I}_p)$.

(a) $\mathbb{E}\lambda_{\max}(\mathbf{E}) \le 2\sigma\,\mathbb{E}\|\mathbf{g}\| < 2\sigma\sqrt{p}$, *and* $\lambda_{\max}(-\mathbf{E})$ *has the same law as* $\lambda_{\max}(\mathbf{E})$.

(b) $\mathbf{E} \mapsto \lambda_{\max}(\mathbf{E})$ *and* $\mathbf{E} \mapsto \|\mathbf{E}\|_{\text{op}}$ *are* $\sqrt{2}\sigma$*-Lipschitz functions of the standard Gaussian vector of independent entries. Hence, for* $t \ge 0$,

$$\Pr[\lambda_{\max}(\mathbf{E}) \ge \mathbb{E}\lambda_{\max}(\mathbf{E}) + t] \le e^{-t^2/(4\sigma^2)}, \qquad \Pr[\|\mathbf{E}\|_{\text{op}} \ge \mathbb{E}\|\mathbf{E}\|_{\text{op}} + t] \le e^{-t^2/(4\sigma^2)}.$$

(c) $\Pr\big[\lambda_{\max}(\mathbf{E}) \ge \sigma\big(2\sqrt{p} + 2\sqrt{\ln(1/\alpha)}\big)\big] \le \alpha$ *and* $\Pr\big[\|\mathbf{E}\|_{\mathrm{op}} \ge \sigma\big(2\sqrt{p} + 2\sqrt{\ln(2/\alpha)}\big)\big] \le \alpha$.
(d) $\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}} \le 2\sigma\sqrt{p} + c_0\sigma$, *with* $c_0 = 2\sqrt{\ln 2} + 4\sqrt{\pi}\,Q(\sqrt{2\ln 2}) \approx 2.512$, *and* $(\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}}^4)^{1/4} \le \sigma(2\sqrt{p} + c_0 + 2\sqrt{2})$.

In Algorithm 6, $r_1 = \sigma_1(2\sqrt{p} + 2\sqrt{\ln(1/\alpha_L)})$ *therefore satisfies* $\Pr[\lambda_{\max}(\hat{\mathbf{G}}_1 - \mathbf{G}) > r_1] \le \alpha_L$, *and this one-sided event is all the algorithm uses (Lemma S3).*

Proof. (a) For a unit vector $\mathbf{u}$ let $X_{\mathbf{u}} = \mathbf{u}^\top\mathbf{E}\mathbf{u} = \sum_{i\le j,\chi_{ij}=1} E_{ij}\, w_{ij}(\mathbf{u})$, with $w_{ii} = u_i^2$ and $w_{ij} = 2u_iu_j$. Then

$$\mathbb{E}(X_{\mathbf{u}} - X_{\mathbf{v}})^2 \le \sigma^2\Big[2\sum_i (u_i^2 - v_i^2)^2 + 4\sum_{i<j}(u_iu_j - v_iv_j)^2\Big] = 2\sigma^2\|\mathbf{u}\mathbf{u}^\top - \mathbf{v}\mathbf{v}^\top\|_F^2$$
$$= 4\sigma^2\big(1 - (\mathbf{u}^\top\mathbf{v})^2\big) \le 4\sigma^2\|\mathbf{u} - \mathbf{v}\|^2,$$

using $1 - (\mathbf{u}^\top\mathbf{v})^2 = (1 - \mathbf{u}^\top\mathbf{v})(1 + \mathbf{u}^\top\mathbf{v}) \le 2(1 - \mathbf{u}^\top\mathbf{v}) = \|\mathbf{u} - \mathbf{v}\|^2$. The process $Y_{\mathbf{u}} = 2\sigma\,\mathbf{g}^\top\mathbf{u}$ has $\mathbb{E}(Y_{\mathbf{u}} - Y_{\mathbf{v}})^2 = 4\sigma^2\|\mathbf{u} - \mathbf{v}\|^2$. By the Sudakov–Fernique inequality [12, 9] (see [14, Thm. 7.2.11]; both processes are centred, Gaussian and continuous on the sphere), $\mathbb{E}\lambda_{\max}(\mathbf{E}) = \mathbb{E}\sup_{\mathbf{u}} X_{\mathbf{u}} \le \mathbb{E}\sup_{\mathbf{u}} Y_{\mathbf{u}} = 2\sigma\mathbb{E}\|\mathbf{g}\| \le 2\sigma\sqrt{p}$. $\mathbf{E}$ and $-\mathbf{E}$ have the same law.

(b) Write $E_{ij} = \sigma_{ij}\zeta_{ij}$ with independent standard ζ_{ij}. For a unit $\mathbf{u}$, $|\mathbf{u}^\top(\mathbf{E}(\zeta) - \mathbf{E}(\zeta'))\mathbf{u}| \le \|\zeta - \zeta'\|\big(\sum \sigma_{ij}^2 w_{ij}(\mathbf{u})^2\big)^{1/2} \le \sqrt{2}\sigma\|\zeta - \zeta'\|$, since $\sum_i u_i^4 \cdot 2 + \sum_{i<j} 4u_i^2u_j^2 = 2$ (and with diagonal variance σ^2 the bound is $\sigma^2(2 - \sum_i u_i^4)$). Both $\lambda_{\max}$ and $\|\cdot\|_{\mathrm{op}} = \sup_{\mathbf{u}}|\mathbf{u}^\top\cdot\mathbf{u}|$ are therefore $\sqrt{2}\sigma$-Lipschitz, and Gaussian concentration [13] (see [5, Thm. 5.6]) gives $\Pr[f - \mathbb{E}f \ge t] \le e^{-t^2/(2L^2)}$ with $L = \sqrt{2}\sigma$.

(c) Combine (a) and (b) with $t = 2\sigma\sqrt{\ln(1/\alpha)}$. For the norm, $\|\mathbf{E}\|_{\mathrm{op}} = \max(\lambda_{\max}(\mathbf{E}), \lambda_{\max}(-\mathbf{E}))$ and a union bound gives the factor 2.

(d) With $m = 2\sigma\sqrt{p} \ge \mathbb{E}\lambda_{\max}(\pm\mathbf{E})$, $\|\mathbf{E}\|_{\mathrm{op}} \le m + \max\big((\lambda_{\max}(\mathbf{E}) - m)_+, (\lambda_{\max}(-\mathbf{E}) - m)_+\big)$, and the last maximum exceeds t with probability at most $\min(1, 2e^{-t^2/4\sigma^2})$; integrating gives c_0. For the fourth moment, $\|\mathbf{E}\|_{\mathrm{op}} = \mathbb{E}\|\mathbf{E}\|_{\mathrm{op}} + X$ with $\Pr[|X| \ge t] \le 2e^{-t^2/(4\sigma^2)}$ by (b), so $\mathbb{E}X^4 \le 8\int_0^\infty t^3 e^{-t^2/(4\sigma^2)}dt = 64\sigma^4$, and Minkowski's inequality gives the claim. □

Remark 4.2 (What is not proved)**.** The claim $\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}} \le 2\sigma\sqrt{p}$, which would give the two-sided bound with $\ln(1/\alpha)$, is [open] here. Sudakov–Fernique applied to $\sup_{\mathbf{u},\pm} \pm\mathbf{u}^\top\mathbf{E}\mathbf{u}$ fails, because the cross-sign increments $\mathbb{E}(\mathbf{u}^\top\mathbf{E}\mathbf{u} + \mathbf{v}^\top\mathbf{E}\mathbf{v})^2$ reach $8\sigma^2$ at $\mathbf{u} = \mathbf{v}$. The suggested Gaussian comparison does reduce the masked case to the full one: if $\mathbf{W} = \mathbf{E} + \mathbf{E}'$ with $\mathbf{E}'$ an independent centred Gaussian symmetric matrix supported on the masked entries (and, if desired, extra diagonal variance up to $2\sigma^2$), then $\mathbb{E}\|\mathbf{W}\|_{\mathrm{op}} = \mathbb{E}_{\mathbf{E}}\mathbb{E}_{\mathbf{E}'}\|\mathbf{E} + \mathbf{E}'\|_{\mathrm{op}} \ge \mathbb{E}\|\mathbf{E}\|_{\mathrm{op}}$ by Jensen's inequality, since $\|\cdot\|_{\mathrm{op}}$ is convex and $\mathbb{E}\mathbf{E}' = 0$. So the mask can only help, and it would suffice to prove $\mathbb{E}\|\mathbf{W}\|_{\mathrm{op}} \le 2\sigma\sqrt{p}$ for the full GOE-type matrix. For the GOE this holds asymptotically (the largest eigenvalue has mean $2\sigma\sqrt{p} - O(\sigma p^{-1/6})$ by Tracy–Widom asymptotics), but I know of no non-asymptotic proof with the constant exactly 2. Known sharp non-asymptotic bounds, such as that of Bandeira and van Handel [4], carry an additional $O(\sigma\sqrt{\log p})$ term. Monte Carlo (`s2_goe_ratio.py`) gives $\mathbb{E}\|\mathbf{W}\|_{\mathrm{op}}/(2\sigma\sqrt{p}) = 0.567, 0.724, 0.799, 0.864, 0.926, 0.954, 0.966$ for $p = 1, 2, 3, 5, 11, 20, 30$ (GOE normalization), increasing towards 1 from below, and 0.81 for the protocol's masked $p = 11$ (Table 3). Algorithm 6 does not need the claim.

Subgaussian (discrete-Gaussian) noise. By Lemma 2.2(b),(c), (c) holds for the protocol's noise with α replaced by $\alpha + N\,\mathrm{TV}_1(\sigma_{c,1}) + \delta_B$ and r_1 by $r_1 + \sqrt{2}\,\bar{\nu}$. Without the coupling, using only that the noise vector is σ^2-subgaussian (Lemma 2.2(a)), a $\frac{1}{4}$-net $\mathcal{N}$ of the sphere ($|\mathcal{N}| \le 9^p$)

and $\|\mathbf{E}\|_{\mathrm{op}} \le 2\max_{\mathbf{v}\in\mathcal{N}} |\mathbf{v}^\top \mathbf{E}\mathbf{v}|$ (see [14, Cor. 4.2.13, Ex. 4.4.3]), together with the fact that $\mathbf{v}^\top\mathbf{E}\mathbf{v} = \langle \mathbf{w}(\mathbf{v}), \boldsymbol{\xi}\rangle$ with $\|\mathbf{w}(\mathbf{v})\|^2 \le 2$ is $2\sigma^2$-subgaussian, give

$$\Pr\Big[\|\mathbf{E}\|_{\mathrm{op}} \ge 4\sigma\sqrt{p\ln 9 + \ln(2/\alpha)}\Big] \le \alpha.$$

For $p = 11$ and $\alpha = 0.0025$ this is 22.2σ, against 11.5σ for the Gaussian r_1. The coupling route is therefore the one to use.

Monte Carlo (`s2_check.py`, 2×10^5 draws). For the protocol shape ($p = 11$, 5×5 exact block): $\mathbb{E}\lambda_{\max} = 4.98\sigma$, $\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}} = 5.38\sigma$, against $2\sigma\sqrt{p} = 6.63\sigma$, a ratio of 0.81. The 0.9975-quantile of $\lambda_{\max}$ is 7.43σ (of $\|\mathbf{E}\|_{\mathrm{op}}$, 7.62σ), while $r_1/\sigma_1 = 11.53$ at $\alpha_L = 0.0025$, so r_1 is conservative by a factor of about 1.55. Table 3 gives other shapes. The discrete-Gaussian version (ring $N = 128$, canonical embedding, $\sigma_c \in \{0.8, 3\}$) has exactly uncorrelated slots and the same $\lambda_{\max}$ distribution to Monte Carlo accuracy (mean 4.97, 0.9975-quantile 7.38).

shape	$2\sqrt{p}$	$\mathbb{E}\lambda_{\max}$	$\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}}$	$q_{.9975}(\lambda_{\max})$	r_1/σ_1	$\Pr[\lambda_{\max} \ge 2\sqrt{p}+2]$
$p = 11$, $d_A = 5$ (protocol)	6.633	4.976	5.383	7.43	11.53	2.5×10^{-5}
$p = 11$, no mask, diag var 1	6.633	5.431	5.858	7.89	11.53	1.8×10^{-4}
$p = 11$, GOE (diag var 2)	6.633	5.645	6.148	8.29	11.53	9.0×10^{-4}
$p = 2$, $d_A = 1$	2.828	0.965	1.368	3.60	7.72	9×10^{-5}
$p = 20$, $d_A = 5$	8.944	7.702	8.096	10.00	13.84	4×10^{-5}
$p = 50$, $d_A = 10$	14.142	13.095	13.446	15.09	19.04	0

Table 3: Lemma S2, Monte Carlo in units of σ; r_1/σ_1 at $\alpha_L = 0.0025$. The bound $e^{-t^2/4}$ at $t = 2$ is 0.37; the observed tails are far lighter, because the Lipschitz constant $\sqrt{2}\sigma$ is attained only by spread-out eigenvectors (empirical s.d. of $\lambda_{\max}$: 0.77σ).

5 Lemma S3: the two-stage choice of the ridge

Let $\mathbf{E}_1 = \hat{\mathbf{G}}_1 - \mathbf{G}$ (scale σ_1), $\gamma := g^\star + r_1$ with $g^\star = 2\vartheta^\star\sigma_2\sqrt{p}$, and

$$\Omega_L := \{\lambda_{\max}(\mathbf{E}_1) \le r_1\}.$$

Lemma 5.1 (S3). *Under Assumption 2.1:*

(a) ***Measurability and independence.*** $(\lambda, \Psi) = \Lambda(\mathbf{A}, \hat{\mathbf{B}}_1, \hat{\mathbf{C}}_1)$ *for a Borel map* Λ *that depends otherwise only on public constants* $(\sigma_1, \sigma_2, p, \vartheta^\star, \alpha_L, \kappa_\Psi)$. *Hence, conditionally on the data* $(\mathbf{X}, \mathbf{y})$, $(\lambda, \Psi, \hat{L})$ *is independent of the stage-2 noise* $(\mathbf{E}_2, \mathbf{f}_2, g)$ *and of the separate* $\mathbf{y}^\top\mathbf{y}$ *noise.*

(b) ***Radius event.*** $\Pr[\Omega_L \mid \mathbf{X}, \mathbf{y}] \ge 1 - \alpha_L$. *On* Ω_L: $\lambda_{\min}(\mathbf{G} + \lambda\Psi) \ge g^\star$, *i.e.* $\vartheta_0 \ge \vartheta^\star$; *and* $\hat{L} \le \lambda_{\min}(\mathbf{G})$.

(c) ***The*** Π_O ***ridge.*** $\lambda_{\min}(\hat{\mathbf{G}}_1 + \lambda\Pi_O) \le \lambda_{\min}(\mathbf{A})$ *for every* $\lambda \ge 0$ *(Cauchy interlacing), and* $\lambda_{\min}(\hat{\mathbf{G}}_1 + \lambda\Pi_O) \uparrow \lambda_{\min}(\mathbf{A})$ *as* $\lambda \to \infty$. *If* $\lambda_{\min}(\mathbf{A}) > \gamma$,

$$\lambda_O = \max\big\{0,\ \gamma - \lambda_{\min}\big(\hat{\mathbf{B}}_1 - \hat{\mathbf{C}}_1^\top(\mathbf{A} - \gamma\mathbf{I})^{-1}\hat{\mathbf{C}}_1\big)\big\} < \infty;$$

if $\lambda_{\min}(\mathbf{A}) < \gamma$, $\lambda_O = \infty$; *the case* $\lambda_{\min}(\mathbf{A}) = \gamma$ *has probability zero (set* $\lambda_O = \infty$*). No bisection is needed.*

(d) ***Gate.*** *Conditionally on the data and stage 1,* $\Pr[\hat{\vartheta} < 1] \le \Pr\big[\lambda_{\max}(\mathbf{E}_2) > 2(\vartheta_0 - 1)\sigma_2\sqrt{p}\big] \le \exp\big(-p(\vartheta_0 - 2)^2\big)$ *for* $\vartheta_0 > 2$. *Hence* $\Pr[\textsf{refused} \mid \mathbf{X}, \mathbf{y}] \le \alpha_L + \exp(-p(\vartheta^\star - 2)^2)$ *when* $\vartheta^\star > 2$.

Proof. (a) $\lambda_{\min}$ is continuous; $\hat{L}$, λ_I and the closed form in (c) are continuous on their domains, and the comparison $\lambda_O \le \kappa_\Psi \lambda_I$ is a Borel event. Equivalently, $\{\lambda_O \le a\} = \{\lambda_{\min}(\hat{\mathbf{G}}_1 + a\Pi_O) \ge \gamma\}$ is closed, because $\lambda \mapsto \lambda_{\min}(\hat{\mathbf{G}}_1 + \lambda\Pi_O)$ is continuous and non-decreasing; so λ_O is lower semicontinuous. A bisection with a fixed number of steps is also Borel. The stage-2 noise is fresh: η and (u, e_0, e_1) are drawn anew in Algorithm 5, and (λ, Ψ) is a function of the data and $\mathbf{E}_1$ only. For the real protocol this independence holds exactly for the ideal mechanism $\mathcal{M}$ and up to the perturbation of norm $\le \sqrt{2}\,\bar{\nu}$ of Lemma 2.2.

(b) $\Pr[\Omega_L] \ge 1 - \alpha_L$ is Lemma 4.1(c). For any symmetric $\mathbf{M}$, $\lambda_{\min}(\mathbf{G}+\mathbf{M}) = \lambda_{\min}(\hat{\mathbf{G}}_1 + \mathbf{M} - \mathbf{E}_1) \ge \lambda_{\min}(\hat{\mathbf{G}}_1 + \mathbf{M}) - \lambda_{\max}(\mathbf{E}_1)$ (Weyl). With $\mathbf{M} = \lambda\Psi$: if $\Psi = \mathbf{I}$, $\lambda_{\min}(\hat{\mathbf{G}}_1 + \lambda_I \mathbf{I}) = \lambda_{\min}(\hat{\mathbf{G}}_1) + \lambda_I \ge \gamma$; if $\Psi = \Pi_O$ (then $\lambda_O < \infty$), $\lambda_{\min}(\hat{\mathbf{G}}_1 + \lambda_O \Pi_O) \ge \gamma$ by definition. So on Ω_L, $\lambda_{\min}(\mathbf{G}_\lambda) \ge \gamma - r_1 = g^\star$. With $\mathbf{M} = 0$: $\lambda_{\min}(\mathbf{G}) \ge \lambda_{\min}(\hat{\mathbf{G}}_1) - r_1$, and $\lambda_{\min}(\mathbf{G}) \ge 0$, so $\lambda_{\min}(\mathbf{G}) \ge \hat{L}$.

(c) Interlacing: $\mathbf{A}$ is a principal submatrix of $\hat{\mathbf{G}}_1 + \lambda\Pi_O$ (the $\mathbf{A}$ block is exact and unaffected by $\lambda\Pi_O$). For $\gamma < \lambda_{\min}(\mathbf{A})$, $\lambda_{\min}(\hat{\mathbf{G}}_1 + \lambda\Pi_O) \ge \gamma$ iff $\left(\begin{smallmatrix} \mathbf{A}-\gamma\mathbf{I} & \hat{\mathbf{C}}_1 \\ \hat{\mathbf{C}}_1^\top & \hat{\mathbf{B}}_1 + (\lambda-\gamma)\mathbf{I} \end{smallmatrix}\right) \succeq 0$ iff (Schur complement, $\mathbf{A} - \gamma\mathbf{I} \succ 0$) $\hat{\mathbf{B}}_1 - \hat{\mathbf{C}}_1^\top (\mathbf{A} - \gamma\mathbf{I})^{-1} \hat{\mathbf{C}}_1 + (\lambda - \gamma)\mathbf{I} \succeq 0$. The limit follows from the Schur complement as $\lambda \to \infty$. If $\gamma = \lambda_{\min}(\mathbf{A})$, feasibility requires $\hat{\mathbf{C}}_1^\top \mathbf{v} = 0$ for an eigenvector $\mathbf{v}$ of $\mathbf{A}$, a null event under continuous noise.

(d) $\lambda_{\min}(\hat{\mathbf{G}}_2 + \lambda\Psi) \ge \lambda_{\min}(\mathbf{G}_\lambda) - \lambda_{\max}(-\mathbf{E}_2)$, so $\hat{\vartheta} < 1$ implies $\lambda_{\max}(-\mathbf{E}_2) > 2(\vartheta_0 - 1)\sigma_2\sqrt{p} = 2\sigma_2\sqrt{p} + 2(\vartheta_0 - 2)\sigma_2\sqrt{p}$. Apply Lemma 4.1(a),(b) with $t = 2(\vartheta_0 - 2)\sigma_2\sqrt{p}$ and (b) of this lemma. □

Remark 5.2 (Size of the guarantees)**.** With $\vartheta^\star = 2.5$ and $p = 11$ the rigorous refusal bound is $\alpha_L + e^{-2.75} = \alpha_L + 0.064$. Monte Carlo at the worst case $\vartheta_0 = \vartheta^\star$ found no exceedance of $\lambda_{\max}(\mathbf{E}) > 9.95\sigma$ in 2×10^5 draws, and Algorithm 6 never refused in our simulations (Table 5). Because r_1 is conservative (Lemma S2), the realized ϑ_0 is typically far above $\vartheta^\star$: it was 7.1–8.2 in the ridge scenarios of Table 4. The private bound $\sqrt{U_{yy}/\hat{L}}$ is finite only if $\lambda_{\min}(\hat{\mathbf{G}}_1) > r_1$; when the ridge is active this often fails, and the β interval then rests on B_β^{pub} (Correction 10).

6 Lemma S4: the bound on $\|\boldsymbol{\beta}\|$

Lemma 6.1 (S4)**.** *Let* $\mathbf{X}$ *be the matched (clipped) design,* $\mathbf{G} = \mathbf{X}^\top\mathbf{X}$ *with* $\lambda_{\min}(\mathbf{G}) > 0$*, and* $\mathbf{y}$ *the clipped response,* $|y_i| \le B_y$*.*

(a) If $\mathbb{E}[\mathbf{y} \mid \mathbf{X}] = \mathbf{X}\boldsymbol{\beta}$*, then* $\|\boldsymbol{\beta}\|^2 \le \boldsymbol{\beta}^\top\mathbf{G}\boldsymbol{\beta}/\lambda_{\min}(\mathbf{G}) \le \mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}]/\lambda_{\min}(\mathbf{G})$*. Without any model, the projection coefficient* $\boldsymbol{\beta}^* := \mathbf{G}^{-1}\mathbf{X}^\top\mathbb{E}[\mathbf{y} \mid \mathbf{X}]$ *satisfies the same bound.*

(b) If the y_i*,* $i \in I$*, are conditionally independent given* $\mathbf{X}$*, then for* $s \ge 0$

$$\Pr\big[\widehat{yy} \le \mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}] - s \;\big|\; \mathbf{X}\big] \le \exp\Big(-\frac{s^2}{2(nB_y^4/4 + \sigma_{yy}^2)}\Big).$$

Hence $U_{yy}^{\text{c}} := \widehat{yy} + \sqrt{2\ln(1/\alpha_U)\,(nB_y^4/4 + \sigma_{yy}^2)}$ *satisfies* $\Pr[U_{yy}^{\text{c}} \ge \mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}] \mid \mathbf{X}] \ge 1 - \alpha_U$*. The specification's* $U_{yy} = \widehat{yy} + z_{1-\alpha_U}\sigma_{yy} + B_y^2\sqrt{n\ln(1/\alpha_U)/2}$ *satisfies this with probability* $\ge 1 - 2\alpha_U$ *only; with* $\alpha_U/2$ *in both terms it satisfies it with probability* $\ge 1 - \alpha_U$*. In all cases one may replace* U *by* $\min\{U, nB_y^2\}$*, since* $\mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}] \le nB_y^2$ *deterministically.*

(c) Let $\Omega_U = \{U \ge \mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}]\}$*. On* $\Omega_L \cap \Omega_U$*, and if* $\|\boldsymbol{\beta}\| \le B_\beta^{\text{pub}}$*, we have* $\|\boldsymbol{\beta}\| \le \hat{B}_\beta = \min\{B_\beta^{\text{pub}}, \sqrt{U/\hat{L}}\}$*.*

Proof. (a) $\mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}] = \|\mathbb{E}[\mathbf{y} \mid \mathbf{X}]\|^2 + \sum_i \mathrm{Var}(y_i \mid \mathbf{X}) \ge \|\mathbb{E}[\mathbf{y} \mid \mathbf{X}]\|^2$. Under the model, $\|\mathbb{E}[\mathbf{y} \mid \mathbf{X}]\|^2 = \boldsymbol{\beta}^\top\mathbf{G}\boldsymbol{\beta} \ge \lambda_{\min}(\mathbf{G})\|\boldsymbol{\beta}\|^2$. In general, $\boldsymbol{\beta}^{*\top}\mathbf{G}\boldsymbol{\beta}^* = \|\mathbf{H}\mathbb{E}[\mathbf{y} \mid \mathbf{X}]\|^2 \le \|\mathbb{E}[\mathbf{y} \mid \mathbf{X}]\|^2$ with $\mathbf{H}$ the hat matrix.

(b) $y_i^2 \in [0, w] = [0, B_y^2]$, so by Hoeffding's lemma [10] each centred y_i^2 is $(w^2/4) = (B_y^4/4)$-subgaussian given $\mathbf{X}$; independence gives $nw^2/4$ for the sum, and adding the independent $\mathcal{N}(0, \sigma_{yy}^2)$ release noise gives $nw^2/4 + \sigma_{yy}^2$. The Chernoff bound gives the display. For the specification's form, Hoeffding gives $\Pr[\mathbf{y}^\top\mathbf{y} \le \mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}] - w\sqrt{n\ln(1/\alpha)/2}] \le \alpha$ (this term is correct as stated) and $\Pr[\widehat{yy} \le \mathbf{y}^\top\mathbf{y} - z_{1-\alpha}\sigma_{yy}] = \alpha$; the union bound gives 2α, and $\alpha/2$ in each restores α. Without conditional independence, Hoeffding does not apply; a martingale version controls $\sum_i (y_i^2 - \mathbb{E}[y_i^2 \mid \mathcal{F}_{i-1}])$, which is not the target, so only the deterministic bound nB_y^2 remains.

(c) On Ω_L, $\hat{L} \le \lambda_{\min}(\mathbf{G})$ (Lemma 5.1(b)); on Ω_U, $U \ge \mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}]$; combine with (a). If $\hat{L} = 0$ the private bound is $+\infty$ and the claim reduces to the public one. □

Clipping and the estimand. R clips $\mathbf{x}_R$, O clips $\mathbf{x}_O$ and y. The clipped data need not follow a linear model even if the raw data do. Clipping y is a non-linear map of $\mathbf{x}^\top\boldsymbol{\beta}+\varepsilon$, and clipping $\mathbf{x}$ breaks linearity in the original features. Clipping also makes the errors heteroskedastic: $\mathrm{Var}(\mathrm{clip}(y) \mid \mathbf{x})$ depends on $\mathbf{x}$ through the probability of clipping. Precisely:

- The Wald interval targets $\boldsymbol{\theta}_\lambda = (\mathbf{G} + \lambda\Psi)^{-1}\mathbf{c}$ computed on the clipped matched data. It is well defined without any model; its randomness is the mechanism only, given the data and stage 1.
- The bias-aware interval targets $\boldsymbol{\beta}$ in the *working model on the clipped scale*: $\mathbb{E}[\tilde{y}_i \mid \tilde{\mathbf{X}}] = \tilde{\mathbf{x}}_i^\top\boldsymbol{\beta}$ with conditionally independent, homoskedastic errors, where tildes denote clipped values. Without that model the natural target is the conditional projection $\boldsymbol{\beta}^* = \mathbf{G}^{-1}\tilde{\mathbf{X}}^\top\mathbb{E}[\tilde{\mathbf{y}} \mid \tilde{\mathbf{X}}]$. For $\boldsymbol{\beta}^*$ the bias identity $\mathbb{E}[\boldsymbol{\theta}_\lambda \mid \mathbf{X}] - \boldsymbol{\beta}^* = -\lambda\mathbf{P}_\lambda\Psi\boldsymbol{\beta}^*$ (because $\tilde{\mathbf{X}}^\top\mathbb{E}[\tilde{\mathbf{y}} \mid \tilde{\mathbf{X}}] = \mathbf{G}\boldsymbol{\beta}^*$) and Lemma 6.1(a) still hold, but the homoskedastic variance formula does not. The interval is then approximate, and a sandwich variance (a further release) would be needed for validity.
- The raw-scale coefficient equals the clipped-scale one only if clipping never binds on the support, i.e. if the bounds are true bounds. Exactly Gaussian errors are incompatible with $|y| \le B_y$, so normality of the sampling term can only be asymptotic (Section 7).

7 Proposition S5: coverage

Condition on the data and stage 1, $\mathcal{F}_1 = \sigma(\mathbf{X}, \mathbf{y}, \mathbf{E}_1)$, so $\lambda, \Psi, \mathbf{G}_\lambda, \mathbf{P}_\lambda, \boldsymbol{\theta}_\lambda, \vartheta_0$ are fixed. Let

$$\boldsymbol{\xi} := \mathbf{P}_\lambda(\mathbf{f}_2 - \mathbf{E}_2\boldsymbol{\theta}_\lambda), \qquad \mathbf{V}_0 := \mathrm{Var}(\boldsymbol{\xi} \mid \mathcal{F}_1) = \sigma_2^2\,\mathbf{P}_\lambda\big(\mathbf{I} + \mathbf{K}(\boldsymbol{\theta}_\lambda)\big)\mathbf{P}_\lambda, \qquad q := \|\mathbf{G}_\lambda^{-1/2}\mathbf{E}_2\mathbf{G}_\lambda^{-1/2}\|_{\mathrm{op}}.$$

Theorem 7.1 (S5(a): Wald interval for $\boldsymbol{\theta}_\lambda$)**.** *Under Assumption 2.1:*

(i) *[exact]* $\boldsymbol{\xi} \mid \mathcal{F}_1 \sim \mathcal{N}(0, \mathbf{V}_0)$*; and on* $\{\mathbf{G}_\lambda + \mathbf{E}_2$ *invertible*$\}$, $\tilde{\boldsymbol{\theta}} - \boldsymbol{\theta}_\lambda = (\mathbf{I} + \mathbf{P}_\lambda\mathbf{E}_2)^{-1}\boldsymbol{\xi}$.

(ii) *[exact] If* $q < 1$*, then* $|\tilde{\theta}_j - \theta_{\lambda,j} - \xi_j| \le \frac{q}{1-q}\sqrt{(\mathbf{P}_\lambda)_{jj}}\,\|\mathbf{G}_\lambda^{1/2}\boldsymbol{\xi}\|$.

(iii) *[exact, finite sample] Let* $s_j^2 = (\mathbf{V}_0)_{jj}$, $z = z_{1-\alpha/2}$, $\varpi \in (0, z)$ *and* $\eta \in (0, 1)$ *with* $\eta\vartheta_0 \ge 1$*. The oracle interval* $\tilde{\theta}_j \pm z\,s_j$ *satisfies*

$$\Pr\big[|\tilde{\theta}_j - \theta_{\lambda,j}| > z s_j \;\big|\; \mathcal{F}_1\big] \le \alpha + 2\varpi\,\phi(z - \varpi) + 2e^{-p(\eta\vartheta_0 - 1)^2} + \Big(\frac{\eta}{1-\eta}\Big)^2\frac{\omega_j}{\varpi^2},$$

with $\omega_j := (\mathbf{P}_\lambda)_{jj}\,\sigma_2^2\,\mathrm{tr}\big((\mathbf{I} + \mathbf{K}(\boldsymbol{\theta}_\lambda))\mathbf{P}_\lambda\big)/s_j^2 \le (1 + 2\|\boldsymbol{\theta}_\lambda\|^2)\,\mathrm{tr}(\mathbf{P}_\lambda)/(\mathbf{P}_\lambda)_{jj}$.

(iv) *[first order] The reported interval uses* $\mathbf{V}_{\mathrm{mech}}$ *in place of* $\mathbf{V}_0$*. By Lemma S6,* $(\mathbf{V}_{\mathrm{mech}})_{jj} = s_j^2(1 + O_p(\vartheta_0^{-1}) + O(\vartheta_0^{-2}) + O_p(B^{-1/2}))$.

(v) *[asymptotic] Unconditionally,* $\Pr[\theta_{\lambda,j} \notin \mathrm{CI}_j,$ *not refused*$] \le \alpha_{\mathrm{CI}} + \alpha_L + \epsilon_j$*, where* $\epsilon_j \to 0$ *as* $\vartheta^\star \to \infty$ *and* $B \to \infty$ *with* p, $\|\boldsymbol{\theta}_\lambda\|$ *and* $\mathrm{tr}(\mathbf{P}_\lambda)/(\mathbf{P}_\lambda)_{jj}$ *bounded (choose* $\eta = c/\vartheta_0$ *and* $\varpi \to 0$ *slowly in (iii)).*

Proof. (i) Given $\mathcal{F}_1$, $(\mathbf{E}_2, \mathbf{f}_2)$ is Gaussian, centred and independent of (λ, Ψ) (Lemma 5.1(a)), and $\mathbf{E}_2 \perp \mathbf{f}_2$; so $\boldsymbol{\xi}$ is a centred Gaussian linear image with covariance $\mathbf{P}_\lambda(\sigma_2^2\mathbf{I} + \sigma_2^2\mathbf{K})\mathbf{P}_\lambda$. From $(\mathbf{G}_\lambda+\mathbf{E}_2)\tilde{\boldsymbol{\theta}} = \mathbf{c}+\mathbf{f}_2$ and $\mathbf{G}_\lambda\boldsymbol{\theta}_\lambda = \mathbf{c}$ we get $(\mathbf{G}_\lambda+\mathbf{E}_2)(\tilde{\boldsymbol{\theta}}-\boldsymbol{\theta}_\lambda) = \mathbf{f}_2-\mathbf{E}_2\boldsymbol{\theta}_\lambda$. (ii) Let $\mathbf{M} = \mathbf{G}_\lambda^{-1/2}\mathbf{E}_2\mathbf{G}_\lambda^{-1/2}$. Then $\mathbf{I} + \mathbf{P}_\lambda\mathbf{E}_2 = \mathbf{G}_\lambda^{-1/2}(\mathbf{I}+\mathbf{M})\mathbf{G}_\lambda^{1/2}$, so $\mathbf{G}_\lambda^{1/2}(\tilde{\boldsymbol{\theta}} - \boldsymbol{\theta}_\lambda - \boldsymbol{\xi}) = -(\mathbf{I}+\mathbf{M})^{-1}\mathbf{M}\mathbf{G}_\lambda^{1/2}\boldsymbol{\xi}$, whose norm is at most $\frac{q}{1-q}\|\mathbf{G}_\lambda^{1/2}\boldsymbol{\xi}\|$. Finally $|\mathbf{e}_j^\top\mathbf{v}| \le \|\mathbf{G}_\lambda^{-1/2}\mathbf{e}_j\|\|\mathbf{G}_\lambda^{1/2}\mathbf{v}\|$. (iii) $\Pr[|\xi_j + R_j| > z s_j] \le \Pr[|\xi_j| > (z-\varpi)s_j] + \Pr[|R_j| > \varpi s_j]$. The first term is $2Q(z-\varpi) \le \alpha + 2\varpi\phi(z-\varpi)$. For the second, $q \le \|\mathbf{E}_2\|_{\mathrm{op}}/\lambda_{\min}(\mathbf{G}_\lambda)$ and Lemma 4.1(c) give $\Pr[q > \eta] \le 2e^{-p(\eta\vartheta_0-1)^2}$. On $\{q \le \eta\}$, (ii) and Markov's inequality with $\mathbb{E}\|\mathbf{G}_\lambda^{1/2}\boldsymbol{\xi}\|^2 = \operatorname{tr}(\mathbf{G}_\lambda\mathbf{V}_0) = \sigma_2^2\operatorname{tr}((\mathbf{I}+\mathbf{K})\mathbf{P}_\lambda)$ give the last term. The bound on ω_j uses $s_j^2 \ge \sigma_2^2\|\mathbf{P}_\lambda\mathbf{e}_j\|^2 \ge \sigma_2^2(\mathbf{P}_\lambda)_{jj}^2$ (as $\mathbf{K} \succeq 0$) and $\|\mathbf{K}\| \le 2\|\boldsymbol{\theta}\|^2$. (v) On Ω_L, $\vartheta_0 \ge \vartheta^\star$ (Lemma 5.1(b)); integrate (iii)–(iv) over $\mathcal{F}_1$ and add $\Pr[\Omega_L^c] \le \alpha_L$. □

The bound in (iii) is a certificate that the error vanishes as $\vartheta_0 \to \infty$, but it is loose (it uses the worst direction of $\mathbf{M}$ and Markov's inequality). The *typical* size of the remainder relative to s_j is governed by

$$\sigma_2 \operatorname{tr}(\mathbf{P}_\lambda) \;\le\; \frac{p\,\sigma_2}{\lambda_{\min}(\mathbf{G}_\lambda)} = \frac{\sqrt{p}}{2\vartheta_0},$$

the ratio of the $O(\sigma_2^2)$ mechanism bias $\sigma_2^2\mathbf{P}_\lambda\mathbf{M}_\chi(\mathbf{P}_\lambda)\boldsymbol{\theta}_\lambda$ to the $O(\sigma_2)$ standard deviation. At $\vartheta_0 = \vartheta^\star = 2.5$ and $p = 11$ this can reach 0.66 when all eigenvalues of $\mathbf{G}_\lambda$ sit at the floor. "First order in $1/\vartheta$" therefore hides a factor $\sqrt{p}$ (Correction 7). In the simulations (Table 4) $\sigma_2\operatorname{tr}(\mathbf{P}_\lambda) \le 0.18$ in the realistic scenarios and 0.59 in the stress scenario.

Theorem 7.2 (S5(b): bias-aware interval for β_j)**.** *Assume (M1) the working model on the clipped scale: $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$ with (ε_i) conditionally independent given $\mathbf{X}$, $\mathbb{E}[\varepsilon_i \mid \mathbf{X}] = 0$ and $\operatorname{Var}(\varepsilon_i \mid \mathbf{X}) = s^2$; (M2) $\|\boldsymbol{\beta}\| \le B_\beta^{\mathrm{pub}}$ (possibly $+\infty$); (M3) $n > p$. Decompose*

$$\tilde{\theta}_j - \beta_j = b_j + S_j + \xi_j + R_j, \qquad b_j = -\lambda\,\mathbf{e}_j^\top\mathbf{P}_\lambda\Psi\boldsymbol{\beta}, \quad S_j = \mathbf{e}_j^\top\mathbf{P}_\lambda\mathbf{X}^\top\boldsymbol{\varepsilon}, \quad R_j = \tilde{\theta}_j - \theta_{\lambda,j} - \xi_j.$$

(i) [exact] $\mathbb{E}[\boldsymbol{\theta}_\lambda \mid \mathbf{X}, \mathbf{E}_1] = \boldsymbol{\beta} - \lambda\mathbf{P}_\lambda\Psi\boldsymbol{\beta}$; $|b_j| \le \lambda\|\Psi\mathbf{P}_\lambda\mathbf{e}_j\|\,\|\boldsymbol{\beta}\|$, with equality for $\boldsymbol{\beta} \parallel \Psi\mathbf{P}_\lambda\mathbf{e}_j$. $\operatorname{Var}(S_j \mid \mathbf{X}, \mathbf{E}_1) = s^2(\mathbf{P}_\lambda\mathbf{G}\mathbf{P}_\lambda)_{jj}$, and $\operatorname{Cov}(S_j, \xi_j) = 0$.

(ii) [exact] If $T \sim \mathcal{N}(b, v^2)$ and $\bar{B}$ is a random variable with $|b| \le \bar{B}$ on an event Ω, then $\Pr[|T| > v\,\mathrm{cv}_\alpha(\bar{B}/v)] \le \alpha + \Pr[\Omega^c]$, where $\mathrm{cv}_\alpha(t)$ is the $(1-\alpha)$-quantile of $|Z+t|$.

(iii) [exact on events] On $\Omega_L \cap \Omega_U$, $|b_j| \le \bar{b}_j^0 := \lambda\|\Psi\mathbf{P}_\lambda\mathbf{e}_j\|\hat{B}_\beta$.

(iv) [exact for the oracle interval] Let $v_j^2 = s^2(\mathbf{P}_\lambda\mathbf{G}\mathbf{P}_\lambda)_{jj} + (\mathbf{V}_0)_{jj}$. If $S_j + \xi_j \mid (\mathbf{X}, \mathbf{E}_1) \sim \mathcal{N}(0, v_j^2)$ and $R_j = 0$, then $\tilde{\theta}_j \pm v_j\,\mathrm{cv}_{\alpha_{\mathrm{CI}}}(\bar{b}_j^0/v_j)$ misses β_j with probability at most $\alpha_{\mathrm{CI}} + \alpha_L + \alpha_U'$, where $\alpha_U' = \alpha_U$ for the corrected U_{yy}^{c} and $2\alpha_U$ for the specification's U_{yy}.

(v) [first order] The reported interval replaces $\mathbf{P}_\lambda$ by $\tilde{\mathbf{P}}$ in $\bar{b}_j$ (relative error $\le \|\tilde{\mathbf{P}}\|\|\mathbf{E}_2\| = O_p(\vartheta_0^{-1})$, since $\mathbf{P}_\lambda - \tilde{\mathbf{P}} = \tilde{\mathbf{P}}\mathbf{E}_2\mathbf{P}_\lambda$), and v_j^2 by $V_{jj} = (\mathbf{V}_{\mathrm{mech}})_{jj} + \tilde{s}^2(\tilde{\mathbf{P}}\hat{\mathbf{G}}_2\tilde{\mathbf{P}})_{jj}$. It also drops R_j (Theorem 7.1(ii)) and ignores the dependence of $\mathbf{V}_0$ on $\boldsymbol{\theta}_\lambda$, hence on $\boldsymbol{\varepsilon}$, which is a relative $O(\|\boldsymbol{\theta}_\lambda - \mathbb{E}\boldsymbol{\theta}_\lambda\|/\|\boldsymbol{\theta}_\lambda\|)$ effect.

(vi) [asymptotic] Under (M1) the errors are bounded (clipping), so S_j is normal only asymptotically: by the Berry–Esseen bound for independent summands [11], $\sup_x|\Pr[S_j \le x\,\mathrm{sd}(S_j) \mid \mathbf{X}] - \Phi(x)| \le C_{\mathrm{BE}}\sum_i|a_{ij}|^3\mathbb{E}|\varepsilon_i|^3/(s^3\|\mathbf{a}_j\|^3)$ with $\mathbf{a}_j = \mathbf{X}\mathbf{P}_\lambda\mathbf{e}_j$. Together with (v): $\Pr[\beta_j \notin \mathrm{CI}_j,$ not refused$] \le \alpha_{\mathrm{CI}} + \alpha_L + \alpha_U' + \epsilon_j'$, with $\epsilon_j' \to 0$ when the Lyapunov ratio $\to 0$, $\vartheta^\star \to \infty$ and $B \to \infty$.

Proof. (i) $\mathbb{E}[\mathbf{c} \mid \mathbf{X}] = \mathbf{G}\boldsymbol{\beta}$ and $\mathbf{P}_\lambda\mathbf{G} = \mathbf{I} - \lambda\mathbf{P}_\lambda\Psi$; (λ, Ψ) is independent of $\boldsymbol{\varepsilon}$ given $\mathbf{X}$, because it depends on $\mathbf{G}$ and $\mathbf{E}_1$ only. Cauchy–Schwarz with $\Psi, \mathbf{P}_\lambda$ symmetric. $\mathbb{E}[\boldsymbol{\xi} \mid \mathbf{X}, \boldsymbol{\varepsilon}, \mathbf{E}_1] = 0$ gives the zero

covariance. (ii) $t \mapsto \Pr[|Z+t| \le c]$ is even and non-increasing in $|t|$ (Anderson's lemma [1]), so cv_α is non-decreasing in $|t|$. On Ω, $\mathrm{cv}_\alpha(\bar{B}/v) \ge \mathrm{cv}_\alpha(|b|/v)$, and $\Pr[|T| > v\,\mathrm{cv}_\alpha(|b|/v)] = \Pr[|Z + b/v| > \mathrm{cv}_\alpha(|b|/v)] = \alpha$. This is the construction of Armstrong and Kolesár [2], as used for regularized regression by Armstrong, Kolesár and Kwon [3]. (iii) Lemma 6.1(c) and (i). (iv) Apply (ii) with $\Omega = \Omega_L \cap \Omega_U$ and $\bar{B} = \bar{b}_j^0$. The dependence between $\bar{b}_j^0$ and T (through $\boldsymbol{\varepsilon}$ in U, and through $\mathbf{E}_1$) is harmless, because (ii) does not need $\bar{B}$ independent of T. (v), (vi) are as stated. □

Variance estimate. $\tilde{s}^2$ targets $\mathrm{RSS}_\lambda/(n-p)$, and $\mathbb{E}[\mathrm{RSS}_\lambda \mid \mathbf{X}] = s^2 \sum_i (1-h_i)^2 + \|\lambda \mathbf{X}\mathbf{P}_\lambda \Psi \boldsymbol{\beta}\|^2 \ge (n-p)s^2$, where $h_i \in [0,1)$ are the eigenvalues of $\mathbf{X}\mathbf{P}_\lambda\mathbf{X}^\top$ (at most p of them non-zero). So $\tilde{s}^2$ is conservative under the ridge (Table 5: 1.05, 1.48 and 1.79 against $s^2 = 1$ in three ridge scenarios). This conservativeness, and the worst-case nature of $\bar{b}_j$, explain the over-coverage of the β intervals.

Refusals and conditioning on the gate. All statements are unconditional, with a refusal producing no interval (equivalently, the trivial interval $\mathbb{R}$). Conditionally on passing, $\Pr[\text{miss} \mid \text{pass}] \le \Pr[\text{miss}, \text{pass}]/(1 - \Pr[\mathsf{refused}])$, with $\Pr[\mathsf{refused}]$ bounded in Lemma 5.1(d). The gate does not make the bootstrap moments finite (Lemma S6).

8 Lemma S6: the bootstrap RSS correction and covariance

Fix $\mathcal{F}_1$. Write $\mathrm{RSS}(\mathbf{G}', \mathbf{c}', yy') = yy' - \mathbf{c}'^\top\boldsymbol{\theta}' - \lambda\boldsymbol{\theta}'^\top\Psi\boldsymbol{\theta}'$ with $\boldsymbol{\theta}' = (\mathbf{G}' + \lambda\Psi)^{-1}\mathbf{c}'$; this equals $yy' - 2\mathbf{c}'^\top\boldsymbol{\theta}' + \boldsymbol{\theta}'^\top\mathbf{G}'\boldsymbol{\theta}'$, so $\mathrm{RSS}_\lambda := \mathrm{RSS}(\mathbf{G}, \mathbf{c}, \mathbf{y}^\top\mathbf{y}) = \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}_\lambda\|^2$ and $\widetilde{\mathrm{RSS}} = \mathrm{RSS}(\hat{\mathbf{G}}_2, \hat{\mathbf{c}}_2, \widehat{yy})$. Define the normalized quantities

$$\mathbf{M} = \mathbf{G}_\lambda^{-1/2}\mathbf{E}_2\mathbf{G}_\lambda^{-1/2}, \quad \mathbf{h} = \mathbf{G}_\lambda^{-1/2}\mathbf{f}_2, \quad \boldsymbol{\phi} = \mathbf{G}_\lambda^{1/2}\boldsymbol{\theta}_\lambda, \quad \mathbf{N} = \lambda\mathbf{G}_\lambda^{-1/2}\Psi\mathbf{G}_\lambda^{-1/2} \ (0 \preceq \mathbf{N} \preceq \mathbf{I}),$$

so that, exactly, $\widetilde{\mathrm{RSS}} = \mathbf{y}^\top\mathbf{y} + g - (\boldsymbol{\phi}+\mathbf{h})^\top(\mathbf{I}+\mathbf{M})^{-1}\big[\mathbf{I} + \mathbf{N}(\mathbf{I}+\mathbf{M})^{-1}\big](\boldsymbol{\phi}+\mathbf{h})$. For a symmetric truncation level $\eta \in (0,1)$ let $T_\eta = \{\|\mathbf{M}\|_{\mathrm{op}} \le \eta\}$. In the *truncated bootstrap* a draw is kept iff $\|(\hat{\mathbf{G}}_2 + \lambda\Psi)^{-1/2}\mathbf{E}^*(\hat{\mathbf{G}}_2 + \lambda\Psi)^{-1/2}\|_{\mathrm{op}} \le \eta$; this event is invariant under $\mathbf{E}^* \mapsto -\mathbf{E}^*$.

Lemma 8.1 (S6)**.** *Under Assumption 2.1, conditionally on $\mathcal{F}_1$:*

(a) ***Expansion.*** $\widetilde{\mathrm{RSS}} = \mathrm{RSS}_\lambda + R_1 + R_2 + R_3 + \dots$, *converging on T_η, where R_k is homogeneous of degree k in $(\mathbf{M}, \mathbf{h}, g)$:*

$$R_1 = g - 2\boldsymbol{\theta}_\lambda^\top\mathbf{f}_2 + \boldsymbol{\theta}_\lambda^\top\mathbf{E}_2\boldsymbol{\theta}_\lambda - 2\lambda\boldsymbol{\theta}_\lambda^\top\Psi\boldsymbol{\xi}, \qquad R_2 = -\boldsymbol{\xi}^\top(\mathbf{G} + 2\lambda\Psi)\boldsymbol{\xi} + 2\lambda\boldsymbol{\theta}_\lambda^\top\Psi\mathbf{P}_\lambda\mathbf{E}_2\boldsymbol{\xi},$$

with $\boldsymbol{\xi} = \mathbf{P}_\lambda(\mathbf{f}_2 - \mathbf{E}_2\boldsymbol{\theta}_\lambda)$, and $|R_k| \le c_k\|\mathbf{M}\|^{k-2}\big(\|\mathbf{M}\boldsymbol{\phi}\| + \|\mathbf{h}\| + \|\mathbf{N}\boldsymbol{\phi}\|\big)\big(\|\mathbf{M}\boldsymbol{\phi}\| + \|\mathbf{h}\|\big)$ for $k \ge 2$, with combinatorial constants $c_k \le 4(k+1)$.

(b) ***Second-order bias.*** $\mathbb{E}[R_1] = \mathbb{E}[R_3 \mid T_\eta] = 0$ *(odd degree, symmetric event), and*

$$\mathbb{E}[R_2] = b_{\mathrm{RSS}}(\mathbf{G}, \boldsymbol{\theta}_\lambda) := -\sigma_2^2\Big[\mathrm{tr}\big((\mathbf{G} + 2\lambda\Psi)\mathbf{P}_\lambda(\mathbf{I} + \mathbf{K}(\boldsymbol{\theta}_\lambda))\mathbf{P}_\lambda\big) + 2\lambda\,\boldsymbol{\theta}_\lambda^\top\Psi\mathbf{P}_\lambda\mathbf{M}_\chi(\mathbf{P}_\lambda)\boldsymbol{\theta}_\lambda\Big].$$

Hence $\mathbb{E}[\widetilde{\mathrm{RSS}} \mid T_\eta] - \mathrm{RSS}_\lambda = b_{\mathrm{RSS}}(\mathbf{G}, \boldsymbol{\theta}_\lambda) + \rho_1$, with $|\rho_1| \le C_\eta\,\vartheta_0^{-2}\,\mathcal{S}_2 + O(e^{-p(\eta\vartheta_0 - 1)^2})$, where $\mathcal{S}_2$ is the scale of the second-order terms, $\mathcal{S}_2 \asymp \sigma_2^2(1 + \|\boldsymbol{\theta}_\lambda\|^2)\,\mathrm{tr}(\mathbf{P}_\lambda) + \|\mathbf{N}\boldsymbol{\phi}\|\,\sigma_2(1 + \|\boldsymbol{\theta}_\lambda\|)\sqrt{\mathrm{tr}\,\mathbf{P}_\lambda}/\vartheta_0$.

(c) ***RSS correction.*** *Let $\mathrm{Corr}_\infty = \mathbb{E}^*[\mathrm{RSS}^* \mid T_\eta^*] - \widetilde{\mathrm{RSS}}$ be the $B \to \infty$ limit of* Corr *for the truncated bootstrap. Then $\mathrm{Corr}_\infty = b_{\mathrm{RSS}}(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}}) + \hat{\rho}_1$, and*

$$\mathbb{E}\big[\widetilde{\mathrm{RSS}} - \mathrm{Corr}_\infty \;\big|\; T_\eta\big] - \mathrm{RSS}_\lambda = O\big(\vartheta_0^{-2}\mathcal{S}_2\big).$$

The uncorrected bias $b_{\mathrm{RSS}} = O(\mathcal{S}_2)$ is removed up to a remainder smaller by a factor ϑ_0^{-2}: $O(\sigma_2^2)$ becomes $O(\sigma_2^4)$. The Monte Carlo error of Corr *with B draws is $O_p(\mathcal{S}_1/\sqrt{B})$ with $\mathcal{S}_1^2 = \mathrm{Var}(R_1) = \sigma_{yy}^2 + \sigma_2^2\, O\big((1+\|\boldsymbol{\theta}_\lambda\|^2)^2\big)$. Antithetic pairs $(\mathbf{E}^*, \mathbf{f}^*, g^*), (-\mathbf{E}^*, -\mathbf{f}^*, -g^*)$ cancel R_1 and R_3 exactly and reduce it to $O_p(\mathcal{S}_2/\sqrt{B})$.*

(d) ***Covariance.*** *$\mathrm{Var}(\tilde{\boldsymbol{\theta}} \mid T_\eta) = \mathbf{V}_0\,(1+O(\vartheta_0^{-2}))$ and $\mathrm{Cov}^*(\boldsymbol{\theta}^* \mid T_\eta^*) = \mathbf{V}_0(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}})(1+O(\hat{\vartheta}^{-2}))$, with $\mathbf{V}_0(\mathbf{G}', \boldsymbol{\theta}') = \sigma_2^2(\mathbf{G}'+\lambda\Psi)^{-1}(\mathbf{I}+\mathbf{K}(\boldsymbol{\theta}'))(\mathbf{G}'+\lambda\Psi)^{-1}$; and $\mathbf{V}_0(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}}) = \mathbf{V}_0(\mathbf{G}, \boldsymbol{\theta}_\lambda)(1+O_p(\vartheta_0^{-1}))$, in the sense of the diagonal entries relative to themselves. So $\mathbf{V}_{\mathrm{mech}}$ estimates the mechanism covariance to first order, with relative error $O_p(\vartheta_0^{-1}) + O(\vartheta_0^{-2}) + O_p(B^{-1/2})$.*

(e) ***Untruncated bootstrap.*** *Without truncation, $\mathbb{E}^*\|\boldsymbol{\theta}^*\| = \infty$ (already for $p = 1$, $\mathbb{E}|1/(a+e)| = \infty$ for Gaussian e), so $\mathrm{Cov}_b(\boldsymbol{\theta}_b^*)$ and $\overline{\mathrm{RSS}^*}$ do not converge as $B \to \infty$. For fixed B they coincide with the truncated versions except on an event of probability $\le B \cdot 2e^{-p(\eta\hat{\vartheta}-1)^2}$; the gate $\hat{\vartheta} \ge 1$ alone does not make this small.*

Proof [first order: constants c_k, C_η and the scale $\mathcal{S}_2$ are not made fully explicit]. (a) Expand $(\mathbf{I}+\mathbf{M})^{-1} = \sum_k (-\mathbf{M})^k$, which converges on T_η. Every term of degree $k \ge 2$ in $(\mathbf{M}, \mathbf{h})$ is a bilinear form whose two outer vectors are $\mathbf{h}$, $\mathbf{M}\boldsymbol{\phi}$ or $\mathbf{N}\boldsymbol{\phi}$, with at most one of them $\mathbf{N}\boldsymbol{\phi}$ and at most $k-2$ further factors $\mathbf{M}$ or $\mathbf{N}$ between them ($\|\mathbf{N}\| \le 1$); the number of such terms is at most $4(k+1)$ (outer pairs $(\boldsymbol{\phi}, \boldsymbol{\phi})$: $k+2$ terms; $(\mathbf{h}, \boldsymbol{\phi})$ and $(\boldsymbol{\phi}, \mathbf{h})$: $k+1$ each; $(\mathbf{h}, \mathbf{h})$: k), each with coefficient ± 1. The explicit R_1, R_2 follow by differentiation: with $\boldsymbol{\theta}' = \boldsymbol{\theta}_\lambda + \boldsymbol{\xi} + \boldsymbol{\zeta} + \ldots$, $\boldsymbol{\zeta} = -\mathbf{P}_\lambda \mathbf{E}_2 \boldsymbol{\xi}$, use $\mathbf{c} - \mathbf{G}\boldsymbol{\theta}_\lambda = \lambda\Psi\boldsymbol{\theta}_\lambda$ and $\mathbf{f}_2 - \mathbf{E}_2\boldsymbol{\theta}_\lambda = \mathbf{G}_\lambda\boldsymbol{\xi}$.

(b) Odd-degree terms have zero mean under the symmetric Gaussian law and on the symmetric event T_η. $\mathbb{E}[\boldsymbol{\xi}^\top \mathbf{A}\boldsymbol{\xi}] = \mathrm{tr}(\mathbf{A}\mathbf{V}_0)$ and $\mathbb{E}[\mathbf{E}_2\mathbf{P}_\lambda\mathbf{E}_2] = \sigma_2^2\mathbf{M}_\chi(\mathbf{P}_\lambda)$ give b_{RSS}; $\mathbb{E}[\mathbf{E}_2\mathbf{P}_\lambda\mathbf{f}_2] = 0$. The degree-4 terms are bounded via (a) by $c_4\mathbb{E}[\|\mathbf{M}\|^2(\cdot)(\cdot)]$, and $\mathbb{E}\|\mathbf{M}\|^2 \le (\mathbb{E}\|\mathbf{E}_2\|^2)/\lambda_{\min}(\mathbf{G}_\lambda)^2 = \vartheta_0^{-2}(1 + O(p^{-1/2}))$ by Lemma 4.1(d) and (b); also $\mathbb{E}\|\mathbf{h}\|^2 = \sigma_2^2\,\mathrm{tr}\,\mathbf{P}_\lambda$ and $\mathbb{E}\|\mathbf{M}\boldsymbol{\phi}\|^2 = \sigma_2^2\,\mathrm{tr}(\mathbf{P}_\lambda\mathbf{K}(\boldsymbol{\theta}_\lambda)) \le 2\sigma_2^2\|\boldsymbol{\theta}_\lambda\|^2\,\mathrm{tr}\,\mathbf{P}_\lambda$. Degrees ≥ 5 are summed geometrically on T_η. Removing the conditioning costs $\Pr[T_\eta^c] \le 2e^{-p(\eta\vartheta_0-1)^2}$ (Lemma 4.1(c)) times bounded conditional moments.

(c) Apply (b) in the bootstrap world centred at $(\hat{\mathbf{G}}_2, \hat{\mathbf{c}}_2, \widehat{yy})$, where the bootstrap law is exactly the mechanism law. Then $b_{\mathrm{RSS}}(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}}) - b_{\mathrm{RSS}}(\mathbf{G}, \boldsymbol{\theta}_\lambda)$ is a smooth function of $(\mathbf{E}_2, \mathbf{f}_2)$ whose linear part has zero mean; its quadratic part is $O(\vartheta_0^{-2}\mathcal{S}_2)$ in mean, since b_{RSS} is homogeneous of degree 2 in the noise scale and every derivative brings a factor of order ϑ_0^{-1}. The Monte Carlo statements follow from $\mathrm{Var}(\mathrm{RSS}^*) = \mathrm{Var}(R_1^*) + O(\mathcal{S}_2^2)$, and the antithetic average is even in the noise.

(d) $\mathrm{Var}(\tilde{\boldsymbol{\theta}}) = \mathrm{Var}(\boldsymbol{\xi}) + 2\,\mathrm{Cov}(\boldsymbol{\xi}, \boldsymbol{\zeta}) + \ldots$; $\mathrm{Cov}(\boldsymbol{\xi}, \boldsymbol{\zeta})$ is an odd (third) moment and vanishes, and the remaining terms are $O(\vartheta_0^{-2})$ relative by (a). The plug-in error comes from $\tilde{\mathbf{P}} - \mathbf{P}_\lambda = -\tilde{\mathbf{P}}\mathbf{E}_2\mathbf{P}_\lambda$ ($\|\cdot\| \le \|\tilde{\mathbf{P}}\|\|\mathbf{E}_2\|\|\mathbf{P}_\lambda\|$, relative $O_p(\vartheta_0^{-1})$) and $\tilde{\boldsymbol{\theta}} - \boldsymbol{\theta}_\lambda = O_p(\|\boldsymbol{\xi}\|)$.

(e) For $p = 1$, $\int |a+e|^{-1}\phi_\sigma(e)\,de = \infty$ because the density is positive at $e = -a$. For general p, $\|(\hat{\mathbf{G}}_2 + \mathbf{E}^* + \lambda\Psi)^{-1}\| \ge 1/|\lambda_{\min}(\cdot)|$, and the law of $\lambda_{\min}$ has a positive density at 0. Moreover $\hat{\vartheta} \ge 1$ only gives $\lambda_{\min}(\hat{\mathbf{G}}_2 + \lambda\Psi) \ge 2\sigma_2\sqrt{p}$, which for $p = 11$ is only $1.33\,\mathbb{E}\lambda_{\max}(\mathbf{E}^*)$, so near-singular draws are not rare when $\hat{\vartheta}$ is close to 1. □

Closed-form alternative. Because b_{RSS} and $\mathbf{V}_0$ are explicit, R may replace the Monte Carlo by $\mathrm{Corr}^{\mathrm{cf}} = b_{\mathrm{RSS}}(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}})$ and $\mathbf{V}_{\mathrm{mech}}^{\mathrm{cf}} = \mathbf{V}_0(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}})$. These have the same first-order guarantees as the truncated bootstrap, no Monte Carlo error, and no moment problem. They are post-processing, so they cost no privacy.

Numerical checks (`s3_s6_sim.py`, fixed design and data, 2×10^4 stage-2 releases for the "truth", 200 independent stage-2 releases each with $B = 400$ bootstrap draws). Table 4 shows that b_{RSS} matches the Monte Carlo bias within 2% when $\sigma_2\,\mathrm{tr}\,\mathbf{P}_\lambda \le 0.18$ and within 5% at 0.59. $\mathbf{V}_0$ and

$\mathbf{V}_{\text{mech}}$ match the true mechanism standard deviations within 2% in the realistic scenarios. The spread of the bootstrap standard errors across releases (3.5–4%) is essentially the Monte Carlo error $1/\sqrt{2B} = 3.5\%$. Antithetic pairs reduce the Monte Carlo s.d. of Corr by a factor of 5–100 at the same cost. In the stress scenario ($\vartheta_0 = 2.52$, outside the default design) $\mathbf{V}_0$ is 2–6% too small, the bootstrap is 3–10% too large, and Corr overshoots the bias by 12%. The corrected RSS still removes 87% of the bias.

scenario	ϑ_0	$\sigma_2 \operatorname{tr} \mathbf{P}_\lambda$	$\mathbb{E}[\widetilde{\mathrm{RSS}}] - \mathrm{RSS}_\lambda$ (MC)	b_{RSS}	Corr mean (sd)	antithetic (sd)	$\max_j \lvert \mathrm{sd}_{\mathbf{V}_{\text{mech}}}/\mathrm{sd} - 1 \rvert$
off ($\lambda = 0$)	48.7	0.017	-0.106 ± 0.062	-0.107	-0.08 (0.44)	-0.107 (0.004)	0.009
ridge $\mathbf{I}$	7.07	0.145	-5.19 ± 0.27	-5.29	-5.42 (1.90)	-5.33 (0.21)	0.017
ridge Π_O	8.05	0.111	-2.92 ± 0.13	-2.90	-3.03 (0.99)	-2.94 (0.12)	0.015
small n	8.17	0.183	-6.97 ± 0.16	-6.96	-6.98 (1.14)	-7.04 (0.23)	0.018
stress	2.52	0.588	-23.61 ± 0.17	-22.49	-26.48 (1.98)	-26.71 (1.58)	0.097

Table 4: Lemma S6 at fixed data and fixed (λ, Ψ). "±" is the Monte Carlo standard error; "sd" is the s.d. across 200 independent releases.

9 S7: the exact stage-1 radius and the transfer to the protocol

The exact radius of settings I and P replaces the union bound of Lemma S2(c) by a Monte Carlo quantile of $\lambda_{\max}$ of the unit stage-1 noise matrix. Let F_{d_A,d_O} be the distribution function of $\lambda_{\max}(\mathbf{E})$ for a masked symmetric noise matrix of unit scale with blocks (d_A, d_O), and $q_{1-\alpha} = F^{-1}_{d_A,d_O}(1-\alpha)$. Draw M independent copies, sort their largest eigenvalues $L_{(1)} \le \cdots \le L_{(M)}$, let j be the smallest integer with $\Pr[\mathrm{Bin}(M, 1-\alpha) \le j] \ge c$, and set $\hat{q} = L_{(j+1)}$, $r_1^{\mathrm{ex}} = \sigma_1 \hat{q}$ (the implementation's rule, with $c = 0.99$ and $M = \max\{2 \cdot 10^5, \lceil 2500/\alpha \rceil\}$). For release s let

$$\bar{w}_s = \frac{N}{2\sqrt{2}\,\Omega_f} + \sqrt{2\pi}\, 2^{-\kappa} \sigma_s + \bar{d}_s, \qquad \pi_s = \delta_B + N \,\mathrm{TV}_1(\sigma_{c,s}) + 2Ne^{-t_\eta^2/2},$$

with $\bar{d}_s$ the proved bound on the decoding offset (crypto proofs, Part A, Proposition on the decoding offset; it includes the rounding of the centre) and TV_1 bounded by Lemma S0(b).

Proposition 9.1 (S7: exact stage-1 radius)**.** *(a) The law of $\lambda_{\max}(\mathbf{E}_1)/\sigma_1$ is F_{d_A,d_O}, which depends only on (d_A, d_O) and is continuous.*

(b) Over the Monte Carlo draws, $\Pr[\hat{q} \ge q_{1-\alpha}] \ge c$.

(c) If the draws are independent of the release, $\Pr[\lambda_{\max}(\mathbf{E}_1) > r_1^{\mathrm{ex}}] = (M-j)/(M+1)$, at most α whenever $j + 1 \ge (1-\alpha)(M+1)$.

(d) Under the protocol, (c) holds with r_1^{ex} enlarged by $\sqrt{2}\,\bar{w}_1$ and the probability enlarged by π_1; so do Lemma S2(c) and the radius event of Lemma S3.

Proof. (a) $\mathbf{E}_1 = \sigma_1 \mathbf{E}$ with $\mathbf{E}$ of unit scale, whose law is fixed by the mask. $\lambda_{\max}(\mathbf{E}) \ge 0$ by interlacing (the R block is zero). If $\mathbf{C} \ne 0$, pick $\mathbf{a}, \mathbf{v}$ with $\mathbf{a}^\top \mathbf{C} \mathbf{v} > 0$; the vector $(\mathbf{a}, t\mathbf{v})$ gives $2t\mathbf{a}^\top \mathbf{C}\mathbf{v} + t^2 \mathbf{v}^\top \mathbf{B}\mathbf{v} > 0$ for small $t > 0$; so $\Pr[\lambda_{\max}(\mathbf{E}) = 0] \le \Pr[\mathbf{C} = 0] = 0$. For $t > 0$, $\{\lambda_{\max} \le t\}$ is a convex set of the Gaussian coordinates with non-empty interior (it contains a neighbourhood of $\mathbf{C} = 0$, $\mathbf{B} = -\mathbf{I}$), and $\{\lambda_{\max} = t\}$ lies in its boundary, which has Lebesgue measure zero. (b) $F(q_{1-\alpha}) = 1 - \alpha$ by continuity, the number K of draws below $q_{1-\alpha}$ is $\mathrm{Bin}(M, 1-\alpha)$, and $\hat{q} < q_{1-\alpha}$ iff $K \ge j+1$, of probability $1 - \Pr[K \le j] \le 1 - c$. (c) $F(L_i)$ are independent uniforms, so $F(\hat{q})$ is $\mathrm{Beta}(j+1, M-j)$ with mean $(j+1)/(M+1)$; conditionally on the draws the exceedance probability is $1 - F(\hat{q})$. (d) Lemma S0(b),(c) couple each release with Gaussian slot noise; on the coupling event

(probability $\geq 1-\pi_s$: coupling, certificate, no wraparound) the decoded entries differ from $\mathbf{z}_s + \boldsymbol{\xi}_s^G$ by a vector of norm at most $\bar{w}_s$ (the rounding of $\mathbf{Z}$, the flooded error, and the decoding offset), and the symmetric matrix assembled from it has operator norm at most $\sqrt{2}\,\bar{w}_s$; Weyl's inequality gives the enlargement. □

At $p = 11$, $d_A = 5$, $\alpha = 0.0025$: $M = 10^6$, $j = 997{,}615$, $\hat{q} = 7.415$ against 11.529 for the union bound, and the exact level is $(M-j)/(M+1) = 0.002385$. At the evaluation setting $N\,\mathrm{TV}_1(\sigma_{c,s}) \leq 2^{-122.2}$ ($\approx 10^{-37}$) and $\bar{w}_s$ is between $2\cdot 10^{-8}$ ($\bar{n}_O = 4{,}096$) and $5\cdot 10^{-6}$ ($\bar{n}_O = 10^6$, every row matchable), against $\sigma_1 \geq 128$. With the fixed public seed used in practice $\hat{q}$ is one number; (b) says that the procedure producing it errs with probability at most $1-c$, and (c) is a statement about the procedure with fresh draws. Parts (a)–(c) and the structure of (d) follow the companion book (Chapter 13, Proposition `inf.r1exact`); the book's caveat on (d), a numerically checked total-variation constant, is removed by Lemma S0(b) below, and its offset bound is proved in Part A.

10 S8: where the ridge switches off

Write $g_1^\star = g^\star + r_1$. The stage-1 rule returns $\lambda = 0$ iff $\lambda_{\min}(\hat{\mathbf{G}}_1) \geq g_1^\star$: then $\lambda_I = \lambda_O = 0$; conversely $\lambda_I > 0$ forces $\lambda_O > 0$, since $\lambda_O = 0$ means $\lambda_{\min}(\hat{\mathbf{G}}_1) \geq g_1^\star$.

Proposition 10.1 (S8: a bracket for the switch-off)**.** *Condition on the data, let $\alpha' \in (0,1)$, and let $q'_{1-\alpha'}$ be the $(1-\alpha')$ quantile of $\lambda_{\max}$ of the unit stage-1 noise matrix (or its union bound $2\sqrt{p} + 2\sqrt{\ln(1/\alpha')}$). Under Assumption 2.1:*

(a) if $\lambda_{\min}(\mathbf{G}) \geq g_1^\star + \sigma_1 q'_{1-\alpha'}$, the rule returns $\lambda = 0$ with probability at least $1-\alpha'$;

(b) if $\lambda_{\min}(\mathbf{G}) \leq g_1^\star - \sqrt{2}\,\sigma_1 z_{1-\alpha'}$, it returns $\lambda > 0$ with probability at least $1-\alpha'$;

(c) the threshold in (a) cannot be lowered in terms of $\lambda_{\min}(\mathbf{G})$ alone: for $\mathbf{G} = \lambda_{\min}(\mathbf{G})\mathbf{I}$ the rule returns $\lambda = 0$ with probability exactly $F_{d_A,d_O}((\lambda_{\min}(\mathbf{G}) - g_1^\star)/\sigma_1)$.

Under the protocol the thresholds move by $\sqrt{2}\,\bar{w}_1$ and the probabilities by π_1 (Proposition 9.1(d)).

Proof. (a) Weyl: $\lambda_{\min}(\hat{\mathbf{G}}_1) \geq \lambda_{\min}(\mathbf{G}) - \lambda_{\max}(-\mathbf{E}_1)$, and $\lambda_{\max}(-\mathbf{E}_1)$ has the law of $\sigma_1\lambda_{\max}(\mathbf{E})$. (b) For a unit bottom eigenvector $\mathbf{u}$ of $\mathbf{G}$, $\lambda_{\min}(\hat{\mathbf{G}}_1) \leq \lambda_{\min}(\mathbf{G}) + \mathbf{u}^\top\mathbf{E}_1\mathbf{u}$, a centred Gaussian with variance $\sigma_1^2(\sum_{i\in O} u_i^4 + 4\sum_{i<j,\chi_{ij}=1} u_i^2u_j^2) \leq 2\sigma_1^2$. (c) $\lambda_{\min}(c\mathbf{I} + \mathbf{E}_1) = c - \lambda_{\max}(-\mathbf{E}_1)$. □

With $\lambda_{\min}(\mathbf{G}) \approx n\ell$ the switch-off lies between $(1 - \sqrt{2}\sigma_1 z_{1-\alpha'}/g_1^\star)n_{\mathrm{off}}$ and $(1 + \sigma_1 q'_{1-\alpha'}/g_1^\star)n_{\mathrm{off}}$; on the simulation population at $\alpha' = 0.05$ ($q'_{0.95} = 6.303$, $\sigma_2/\sigma_1 = 0.645$) this is $[0.87, 1.35]\,n_{\mathrm{off}}$ under setting I and $[0.79, 1.58]\,n_{\mathrm{off}}$ under setting P, at every ε.

Lemma 10.2 (S8′: the size of the downward bias of $\lambda_{\min}$)**.** *For fixed symmetric $\mathbf{G}$ and a masked symmetric noise matrix $\mathbf{E}$ of scale σ (any variances $\leq \sigma^2$ off the diagonal and $\leq 2\sigma^2$ on it), with $\mathbf{g} \sim \mathcal{N}(0, \mathbf{I}_p)$:*

(a) $\mathbb{E}\lambda_{\min}(\mathbf{G}+\mathbf{E}) \leq \lambda_{\min}(\mathbf{G})$;

(b) $\mathbb{E}\lambda_{\min}(\mathbf{G}+\mathbf{E}) \geq m(\mathbf{G},\sigma) := \mathbb{E}\min_{\|\mathbf{u}\|=1}\big(\mathbf{u}^\top\mathbf{G}\mathbf{u} - 2\sigma\mathbf{g}^\top\mathbf{u}\big) \geq \lambda_{\min}(\mathbf{G}) - 2\sigma\mathbb{E}\|\mathbf{g}\|$; $m(\mathbf{G},\sigma)$ depends only on the eigenvalues $\mu_1 \leq \cdots \leq \mu_p$ of $\mathbf{G}$ and equals $\mathbb{E}[\theta - \sum_k \sigma^2 g_k^2/(\mu_k - \theta)]$, with $\theta < \mu_1$ the root of $\sum_k \sigma^2 g_k^2/(\mu_k-\theta)^2 = 1$;

(c) $\Pr[\lambda_{\min}(\mathbf{G}+\mathbf{E}) \leq \mathbb{E}\lambda_{\min}(\mathbf{G}+\mathbf{E}) - t] \leq e^{-t^2/(4\sigma^2)}$ for $t \geq 0$.

Proof. (a) $\lambda_{\min}(\mathbf{G}+\mathbf{E}) \leq \mathbf{u}^\top\mathbf{G}\mathbf{u} + \mathbf{u}^\top\mathbf{E}\mathbf{u}$ for a bottom eigenvector $\mathbf{u}$ of $\mathbf{G}$, and $\mathbb{E}\mathbf{u}^\top\mathbf{E}\mathbf{u} = 0$. (b) $-\lambda_{\min}(\mathbf{G}+\mathbf{E}) = \sup_{\mathbf{u}}(-\mathbf{u}^\top\mathbf{G}\mathbf{u} - X_{\mathbf{u}})$ with $X_{\mathbf{u}} = \mathbf{u}^\top\mathbf{E}\mathbf{u}$. The processes $(-\mathbf{u}^\top\mathbf{G}\mathbf{u} - X_{\mathbf{u}})$ and $(-\mathbf{u}^\top\mathbf{G}\mathbf{u} + Y_{\mathbf{u}})$, $Y_{\mathbf{u}} = 2\sigma\mathbf{g}^\top\mathbf{u}$, are Gaussian with equal means, and by the proof of Lemma S2(a)

$\mathbb{E}(X_{\mathbf{u}} - X_{\mathbf{v}})^2 \le \mathbb{E}(Y_{\mathbf{u}} - Y_{\mathbf{v}})^2$. The Sudakov–Fernique inequality holds for Gaussian vectors with equal means (its interpolation proof uses only the covariances; Chatterjee [7]), and it extends to the sphere by continuity; so $\mathbb{E}[-\lambda_{\min}(\mathbf{G}+\mathbf{E})] \le \mathbb{E}\sup_{\mathbf{u}}(-\mathbf{u}^\top\mathbf{G}\mathbf{u} + 2\sigma\mathbf{g}^\top\mathbf{u})$, which is (b). The minimum over the sphere of $\mathbf{u}^\top\mathbf{G}\mathbf{u} - 2\mathbf{h}^\top\mathbf{u}$ is a trust-region problem whose value is $\theta - \mathbf{h}^\top(\mathbf{G}-\theta\mathbf{I})^{-1}\mathbf{h}$ at the root $\theta < \mu_1$ of $\|(\mathbf{G}-\theta\mathbf{I})^{-1}\mathbf{h}\| = 1$ (the hard case has probability zero), and $\mathbf{g}$ is rotation invariant. The last inequality takes $\mathbf{u}^\top\mathbf{G}\mathbf{u} \ge \lambda_{\min}(\mathbf{G})$. (c) $\lambda_{\min}(\mathbf{G}+\mathbf{E})$ is an infimum of affine functions of the Gaussian coordinates with Lipschitz constants at most $\sqrt{2}\sigma$ (proof of Lemma S2(b)); Gaussian concentration. □

Remark 10.3 (What the lemma gives, and what remains open)**.** On the simulation population (Σ of the clipped, standardized design, eigenvalues $0.599, 0.621, 0.680, 0.754, \ldots$, $p = 11$), at $n = n_{\text{off}}$ under setting I and $\varepsilon = 1$, (b) bounds the downward bias by $3.19\,\sigma_1$, against $1.70\,\sigma_1$ in simulation, and $2\sigma_1\mathbb{E}\|\mathbf{g}\| \approx 6.5\,\sigma_1$ for $\mathbf{G} \propto \mathbf{I}$ (script `experiments/proofs_check/lmin_bias_bound.py`). Combined with (c), it guarantees the switch-off with probability 0.95 from $1.4\,n_{\text{off}}$ under setting I (above $1.5\,n_{\text{off}}$ under P), which is weaker than the bracket's 1.35 (and 1.58). By Proposition 10.1(c) no bound in terms of $\lambda_{\min}(\mathbf{G})$ alone can do better than the bracket; the empirical factors 1.2 (I) and 1.4 (P) reflect the spread of the population's spectrum. Turning them into bounds needs a tail bound for $\lambda_{\min}(\mathbf{G}+\mathbf{E})$ that uses the spectrum and is sharp to within much less than the factor 1.9 that the comparison of (b) loses: second-order perturbation, $\lambda_{\min}(\mathbf{G}+\mathbf{E}) \approx \mu_1 + \mathbf{u}_1^\top\mathbf{E}\mathbf{u}_1 - \sum_{k\ge2}(\mathbf{u}_k^\top\mathbf{E}\mathbf{u}_1)^2/(\mu_k - \mu_1)$, needs gaps $\mu_k - \mu_1$ above about $2\|\mathbf{E}\|$, and at n_{off} the three smallest eigenvalues of $n\Sigma$ lie within $0.1\,n\ell \approx 2.4\,\sigma_1$ of each other. This remains [open].

11 S9: mechanism bias with explicit constants

Fix (λ, Ψ) independently of the stage-2 noise, write $\mathbf{G}$ for $\mathbf{G}_\lambda$, $\mathbf{P} = \mathbf{G}^{-1}$, $\sigma = \sigma_2$, $\vartheta = \lambda_{\min}(\mathbf{G})/(2\sigma\sqrt{p})$, $\mathbf{H} = \mathbf{P}\mathbf{E}$, $\mathbf{M} = \mathbf{G}^{-1/2}\mathbf{E}\mathbf{G}^{-1/2}$, $q = \|\mathbf{M}\|_{\text{op}}$, and for $\bar\eta \in (0,1)$ let $T = \{q \le \bar\eta\}$; T is invariant under $\mathbf{E} \mapsto -\mathbf{E}$. Let $\kappa_p = (2\sqrt{p} + c_0 + 2\sqrt{2})/(2\sqrt{p})$ ($\kappa_p = 1.81$ at $p = 11$), so that $(\mathbb{E}\|\mathbf{E}\|_{\text{op}}^4)^{1/4} \le 2\sigma\sqrt{p}\,\kappa_p$ (Lemma S2(d)) and $\mathbb{E}q^4 \le \kappa_p^4\vartheta^{-4}$. The bias matrix of the paper is $\mathbf{P}\mathbf{M}_\chi(\mathbf{P})$, with $\mathbf{M}_\chi$ the contraction of Section 2.

Theorem 11.1 (S9: leading-order conditional bias, explicit remainder)**.** *If $\bar\eta\vartheta > 1$, then $\mathbb{E}[\tilde{\boldsymbol\theta} \mid T]$ exists and $\mathbb{E}[\tilde{\boldsymbol\theta} \mid T] - \boldsymbol\theta_\lambda = \sigma^2\mathbf{P}\mathbf{M}_\chi(\mathbf{P})\boldsymbol\theta_\lambda + \boldsymbol\rho$ with*

$$\|\boldsymbol\rho\|_{\mathbf{G}} \le \frac{1}{\Pr(T)}\Big[\frac{\kappa_p^4}{(1-\bar\eta)\vartheta^4}\big(\|\boldsymbol\theta_\lambda\|_{\mathbf{G}} + \sigma\sqrt{\operatorname{tr}\mathbf{P}}\big) + \frac{2\sqrt{2}\,\kappa_p^2}{\vartheta^2}e^{-p(\bar\eta\vartheta-1)^2/2}\|\boldsymbol\theta_\lambda\|_{\mathbf{G}}\Big],$$

where $\Pr(T^c) \le 2e^{-p(\bar\eta\vartheta-1)^2}$, and $\|\sigma^2\mathbf{P}\mathbf{M}_\chi(\mathbf{P})\boldsymbol\theta_\lambda\|_{\mathbf{G}} \le \kappa_p^2\vartheta^{-2}\|\boldsymbol\theta_\lambda\|_{\mathbf{G}}$. The moment noise $\mathbf{f}$ contributes nothing at order σ^2. For $\Psi = \mathbf{I}$, $\|\sigma^2\mathbf{P}\mathbf{M}_\chi(\mathbf{P})\boldsymbol\theta_\lambda\| \le (3 + p + d_O)\sigma^2\|\mathbf{c}\|/(\lambda_{\min}(\mathbf{G}) + \lambda)^3$ with $\mathbf{G}$ the unpenalized Gram; for $\Psi = \Pi_O$ and $\mathbf{A}$ invertible it is $O(\sigma^2/\lambda)$.

Proof. Expansion. On T, $(\mathbf{I}+\mathbf{H})^{-1} = \mathbf{I} - \mathbf{H} + \mathbf{H}^2 - \mathbf{H}^3 + \mathbf{H}^4(\mathbf{I}+\mathbf{H})^{-1}$, so $\tilde{\boldsymbol\theta} = (\mathbf{I} - \mathbf{H} + \mathbf{H}^2 - \mathbf{H}^3)(\boldsymbol\theta_\lambda + \mathbf{P}\mathbf{f}) + \mathbf{R}_4$ with $\|\mathbf{R}_4\|_{\mathbf{G}} \le \frac{q^4}{1-q}(\|\boldsymbol\theta_\lambda\|_{\mathbf{G}} + \|\mathbf{G}^{-1/2}\mathbf{f}\|)$, because $\mathbf{G}^{1/2}\mathbf{H}^k\mathbf{G}^{-1/2} = \mathbf{M}^k$; and $\|\tilde{\boldsymbol\theta}\|_{\mathbf{G}} \le (1-\bar\eta)^{-1}\|\mathbf{G}^{-1/2}(\mathbf{c}+\mathbf{f})\|$ on T, which is integrable. *Odd terms.* $\mathbb{E}[\mathbf{H}^m \mid T] = 0$ for odd m by symmetry, and every term containing $\mathbf{P}\mathbf{f}$ has zero conditional mean ($\mathbf{f}$ is centred and independent of $\mathbf{E}$). *Remainder.* $\mathbb{E}[q^4\mathbf{1}_T] \le \kappa_p^4\vartheta^{-4}$, and $\mathbb{E}\|\mathbf{G}^{-1/2}\mathbf{f}\| \le \sigma\sqrt{\operatorname{tr}\mathbf{P}}$, independently of q. *Conditioning.* $\mathbb{E}[\mathbf{H}^2 \mid T] - \mathbb{E}\mathbf{H}^2 = (\mathbb{E}[\mathbf{H}^2]\Pr(T^c) - \mathbb{E}[\mathbf{H}^2\mathbf{1}_{T^c}])/\Pr(T)$, of $\mathbf{G}$-operator norm at most $2(\mathbb{E}q^4)^{1/2}\Pr(T^c)^{1/2}/\Pr(T)$ by Cauchy–Schwarz; $\Pr(T^c) \le \Pr[\|\mathbf{E}\|_{\text{op}} > 2\bar\eta\vartheta\sigma\sqrt{p}] \le 2e^{-p(\bar\eta\vartheta-1)^2}$ by Lemma S2(c). *Contraction.* $\mathbb{E}\mathbf{H}^2 = \mathbf{P}\mathbb{E}[\mathbf{E}\mathbf{P}\mathbf{E}] = \sigma^2\mathbf{P}\mathbf{M}_\chi(\mathbf{P})$, of $\mathbf{G}$-norm at most $\mathbb{E}q^2 \le (\mathbb{E}q^4)^{1/2}$.

Decay. $\mathbf{M}_\chi(\mathbf{P}) = \Pi_O\mathbf{P} + \Pi_R\mathbf{P}\Pi_O + \mathrm{tr}(\mathbf{P})\Pi_O + \mathrm{tr}(\Pi_O\mathbf{P})\Pi_R - \Pi_O\,\mathrm{diag}(\mathbf{P})$, with term norms at most $\|\mathbf{P}\|$, $\|\mathbf{P}\|$, $p\|\mathbf{P}\|$, $d_O\|\mathbf{P}\|$, $\|\mathbf{P}\|$; with $\|\mathbf{P}_\lambda\| = 1/(\lambda_{\min}(\mathbf{G}) + \lambda)$ and $\|\boldsymbol{\theta}_\lambda\| \le \|\mathbf{P}_\lambda\|\|\mathbf{c}\|$ for $\Psi = \mathbf{I}$. For $\Psi = \Pi_O$ the O rows of $\mathbf{P}_\lambda$ are $(\mathbf{S}+\lambda\mathbf{I})^{-1}[-\mathbf{C}^\top\mathbf{A}^{-1},\ \mathbf{I}]$, of norm at most $(1+\|\mathbf{A}^{-1}\mathbf{C}\|)/(\lambda_{\min}(\mathbf{S})+\lambda)$, and every term of $\mathbf{M}_\chi(\mathbf{P}_\lambda)\boldsymbol{\theta}_\lambda$ contains a factor $\Pi_O\mathbf{P}_\lambda$ or $\mathbf{P}_\lambda\Pi_O$. □

This is the paper's Theorem on the leading-order bias and its ridge version, with the constants made explicit; it follows the companion book (Chapter 14, Theorems `bias.main` and `bias.ridge`, Proposition `bias.shrink`), whose proofs were checked step by step, including $c_0 = 2.512$ and the fourth moment $64\sigma^4$ of the centred norm. At $p = 11$, $\vartheta = 5$ and $\bar\eta = \frac{1}{2}$ the leading term is at most $0.13\|\boldsymbol{\theta}_\lambda\|_{\mathbf{G}}$ and the remainder at most $0.034(\|\boldsymbol{\theta}_\lambda\|_{\mathbf{G}} + \sigma\sqrt{\mathrm{tr}\,\mathbf{P}})$ (worst direction); in simulation the leading term matches the bias to within 2%. The correction $\tilde{\boldsymbol{\theta}} - \sigma^2\tilde{\mathbf{P}}\mathbf{M}_\chi(\tilde{\mathbf{P}})\tilde{\boldsymbol{\theta}}$ leaves a bias $O(\vartheta^{-4})$ and a fluctuation $O_p(\vartheta^{-3})$ relative to $\|\boldsymbol{\theta}_\lambda\|_{\mathbf{G}}$ (book, Proposition `bias.corr`; checked, with its constant not explicit).

12 S10: the FWS residual variance, binding clipping, and the full-record sensitivity

12.1 The plug-in residual variance of the FWS comparison

The paper's FWS comparison uses $\hat s^2 = \widehat{\mathrm{RSS}}/(n-p)$ with $\widehat{\mathrm{RSS}} = \widehat{yy} - \hat{\mathbf{c}}_2^\top\hat{\mathbf{G}}_2^{-1}\hat{\mathbf{c}}_2$, the case $\lambda = 0$ of Lemma S6. Let ϑ, T, κ_p be as in Section 11 with $\lambda = 0$, and $\hat{\boldsymbol{\beta}} = \mathbf{P}\mathbf{c}$.

Proposition 12.1 (S10a: bias of the plug-in RSS)**.** *If $\bar\eta\vartheta > 1$, then* $\mathbb{E}[\widehat{\mathrm{RSS}} \mid T] - \mathrm{RSS} = -\sigma^2\big(\mathrm{tr}\,\mathbf{P} + \hat{\boldsymbol{\beta}}^\top\mathbf{M}_\chi(\mathbf{P})\hat{\boldsymbol{\beta}}\big) + \rho$ *with*

$$|\rho| \le \frac{1}{\Pr(T)}\Big[\frac{\kappa_p^4}{(1-\bar\eta)\vartheta^4}\big(\|\hat{\boldsymbol{\beta}}\|_{\mathbf{G}}^2 + \sigma^2\,\mathrm{tr}\,\mathbf{P}\big) + \frac{\kappa_p^2}{\vartheta^2}\sigma^2\,\mathrm{tr}\,\mathbf{P} + \frac{2\sqrt{2}\,\kappa_p^2}{\vartheta^2}e^{-p(\bar\eta\vartheta-1)^2/2}\|\hat{\boldsymbol{\beta}}\|_{\mathbf{G}}^2\Big].$$

The leading term is negative: $\hat s^2$ underestimates the residual variance at order σ_2^2.

Proof. With $\boldsymbol{\phi} = \mathbf{G}^{1/2}\hat{\boldsymbol{\beta}}$ and $\mathbf{h} = \mathbf{G}^{-1/2}\mathbf{f}$, $\widehat{\mathrm{RSS}} - \mathrm{RSS} = g - [(\boldsymbol{\phi}+\mathbf{h})^\top(\mathbf{I}+\mathbf{M})^{-1}(\boldsymbol{\phi}+\mathbf{h}) - \|\boldsymbol{\phi}\|^2]$. Expand $(\mathbf{I}+\mathbf{M})^{-1}$ to $\mathbf{M}^4(\mathbf{I}+\mathbf{M})^{-1}$; on the symmetric event T the terms odd in $\mathbf{M}$ vanish, and so do the terms linear in $\mathbf{h}$; $\mathbb{E}\|\boldsymbol{\phi}+\mathbf{h}\|^2 = \|\boldsymbol{\phi}\|^2 + \sigma^2\,\mathrm{tr}\,\mathbf{P}$; $\boldsymbol{\phi}^\top\mathbb{E}[\mathbf{M}^2]\boldsymbol{\phi} = \sigma^2\hat{\boldsymbol{\beta}}^\top\mathbf{M}_\chi(\mathbf{P})\hat{\boldsymbol{\beta}}$. The remainder collects $\boldsymbol{\phi}^\top(\mathbb{E}[\mathbf{M}^2 \mid T] - \mathbb{E}\mathbf{M}^2)\boldsymbol{\phi}$ (bounded as in Theorem 11.1), $\mathbb{E}[\mathbf{h}^\top\mathbf{M}^2\mathbf{h} \mid T] \le \mathbb{E}[q^2]\sigma^2\,\mathrm{tr}\,\mathbf{P}/\Pr(T)$, and the quartic term, at most $\mathbb{E}[q^4\mathbf{1}_T](\|\boldsymbol{\phi}\|^2 + \sigma^2\,\mathrm{tr}\,\mathbf{P})/((1-\bar\eta)\Pr(T))$. $\mathbf{M}_\chi(\mathbf{P}) \succeq 0$ because it is $\sigma^{-2}\mathbb{E}[\mathbf{EPE}]$ with $\mathbf{P} \succ 0$. □

A Monte Carlo check with antithetic pairs ($p = 11$, $\vartheta = 7.0, 3.5, 1.8$) matches the leading term to 0.05–0.65% (`experiments/proofs_check/fws_s2_bias.py`). Rerunning the paper's FWS comparison on the synthetic population (setting I, the same cohorts, releases and FWS random numbers, 200 cohorts per cell) with $\widehat{\mathrm{RSS}}$ corrected by the leading term changes the coverage of the ten slopes by at most 0.0005 in each of the 41 ridge-on cells, and the median half-width by at most 3%: where the ridge is on, the mechanism noise dominates FWS's width, and $\hat s^2$ is floored at 0.1 in up to 49% of releases (`experiments/proofs_check/fws_s2_corrected.py`).

12.2 Binding clipping and the bias-aware interval

When clipping binds, the clipped response is heteroskedastic and $\mathbb{E}[\tilde{\mathbf{y}} \mid \mathbf{X}]$ need not be linear. Take as estimand the projection coefficient $\boldsymbol{\beta}^* = \mathbf{G}^{-1}\mathbf{X}^\top\mathbb{E}[\tilde{\mathbf{y}} \mid \mathbf{X}]$ of the clipped data (fixed design), with the $\tilde y_i$ conditionally independent given $\mathbf{X}$, and let $\bar\sigma^2 \ge \max_i \mathrm{Var}(\tilde y_i \mid \mathbf{X})$.

Proposition 12.2 (S10b: a valid interval under binding clipping)**.** $\mathbb{E}[\boldsymbol{\theta}_\lambda \mid \mathbf{X}] - \boldsymbol{\beta}^* = -\lambda \mathbf{P}_\lambda \Psi \boldsymbol{\beta}^*$, $\mathrm{Cov}(\boldsymbol{\theta}_\lambda \mid \mathbf{X}) \preceq \bar{\sigma}^2 \mathbf{P}_\lambda \mathbf{G} \mathbf{P}_\lambda$, *and the norm bound of Lemma S4 holds for* $\boldsymbol{\beta}^*$. *Consequently Theorem S5(b) holds for* $\boldsymbol{\beta}^*$, *with the same exact and first-order parts, when* $\tilde{s}^2$ *is replaced by* $\bar{\sigma}^2$ *in* $\mathbf{V} = \mathbf{V}_{\mathrm{mech}} + \tilde{s}^2 \tilde{\mathbf{P}} \hat{\mathbf{G}}_2 \tilde{\mathbf{P}}$. *Admissible public choices are:* $\bar{\sigma}^2 = B_y^2$, *with no model; and, when only the response is clipped and* $y_i = \mathbf{x}_i^\top \boldsymbol{\beta} + \varepsilon_i$ *before clipping with* $\mathrm{Var}(\varepsilon_i \mid \mathbf{X}) \le \sigma_0^2$, *any public* $\bar{\sigma}^2 \ge \sigma_0^2$.

Proof. $\mathbb{E}[\mathbf{c} \mid \mathbf{X}] = \mathbf{X}^\top \mathbb{E}[\tilde{\mathbf{y}} \mid \mathbf{X}] = \mathbf{G}\boldsymbol{\beta}^*$, so $\mathbb{E}[\boldsymbol{\theta}_\lambda \mid \mathbf{X}] = \mathbf{P}_\lambda \mathbf{G} \boldsymbol{\beta}^* = \boldsymbol{\beta}^* - \lambda \mathbf{P}_\lambda \Psi \boldsymbol{\beta}^*$; $\mathrm{Cov}(\boldsymbol{\theta}_\lambda \mid \mathbf{X}) = \mathbf{P}_\lambda \mathbf{X}^\top \boldsymbol{\Omega} \mathbf{X} \mathbf{P}_\lambda$ with $\boldsymbol{\Omega} = \mathrm{diag}(\mathrm{Var}(\tilde{y}_i \mid \mathbf{X})) \preceq \bar{\sigma}^2 \mathbf{I}$. Lemma S4 uses no model. The interval is $\tilde{\theta}_j \pm \mathrm{cv}_\alpha(\bar{b}_j / \sqrt{V_{jj}}) \sqrt{V_{jj}}$; if the true standard deviation is $s \le s' = \sqrt{V_{jj}}$ and $|b| \le \bar{b}_j$, then $c = \mathrm{cv}_\alpha(\bar{b}_j / s') s' > \bar{b}_j \ge |b|$ (since $\mathrm{cv}_\alpha(t) > t$ for $\alpha < \frac{1}{2}$), and $s \mapsto \Pr[|sZ + b| \le c] = \Phi((c-b)/s) - \Phi((-c-b)/s)$ has derivative $-[(c-b)\phi((c-b)/s) + (c+b)\phi((c+b)/s)]/s^2 < 0$; so coverage at s is at least coverage at s', which is the guarantee of Theorem S5(b). $|\tilde{y}_i| \le B_y$ gives $\mathrm{Var}(\tilde{y}_i \mid \mathbf{X}) \le B_y^2$. Clipping is 1-Lipschitz, and $\mathrm{Var}\, f(Y) = \frac{1}{2}\mathbb{E}(f(Y) - f(Y'))^2 \le \frac{1}{2}\mathbb{E}(Y - Y')^2 = \mathrm{Var}\, Y$ for 1-Lipschitz f and an independent copy Y', so $\mathrm{Var}(\mathrm{clip}(y_i) \mid \mathbf{X}) \le \sigma_0^2$. □

What remains [open] is a guarantee with the data-driven $\tilde{s}^2$: it estimates an average of the $\mathrm{Var}(\tilde{y}_i \mid \mathbf{X})$ plus a misspecification term, not their maximum, and a sandwich variance would need per-record residuals of O's data, that is, a further release. The public bound B_y^2 is conservative (on the synthetic population $B_y^2 = 8.35$ against a clipped residual variance of about 0.5, which widens the sampling part by about 4.1).

12.3 The full-record sensitivity in closed form

The paper's full-record baseline releases the upper triangle of $\sum_i \mathbf{w}_i \mathbf{w}_i^\top$, $\mathbf{w} = (\dot{\mathbf{x}}, \mathbf{x}_O, y)$, $\dot{\mathbf{x}} = (1, \mathbf{u})$ with an intercept ($\iota = 1$; $\dot{\mathbf{x}} = \mathbf{u}$ and $\iota = 0$ without), every entry with the same noise. Let $R_u = B_R^2 - \iota$, $o = B_O^2$, $w = B_y^2$, and $m_1 = \min(R_u, o)$, $m_2 = \max(R_u, o)$. Define

$$H(a) = \begin{cases} \frac{2}{3}a^2, & |a| \le 3m_1, \\ \frac{1}{2}a^2 + |a| m_1 - \frac{3}{2}m_1^2, & 3m_1 < |a| \le m_1 + 2m_2, \\ 2|a|(m_1 + m_2) - 2(m_1^2 + m_2^2 + m_1 m_2), & |a| > m_1 + 2m_2, \end{cases}$$

$K = \iota(2R_u + 2o) + 2R_u^2 + 2o^2 + 2R_u o$, and, for $t \in [-w, w]$,

$$\Psi_2(t) = K + (\iota + R_u + o)\Big(w + \frac{t^2}{w}\Big) - 2\iota t + \Big(w - \frac{t^2}{w}\Big)^2 + H(\iota + t),$$

and $\Psi_2^{\neg yy}$ the same without the term $(w - t^2/w)^2$.

Theorem 12.3 (S10c: full-record sensitivity)**.** *If* $\mathbf{u}$ *and* $\mathbf{x}_O$ *have dimension at least* 2, *the replace-one* ℓ_2 *sensitivities of the full-record release are* $\Delta^2_{\mathrm{full}} = \max_{|t| \le w} \Psi_2(t)$ *(joint mode),* $\Delta^2_{\mathrm{full},\neg yy} = \max_{|t| \le w} \Psi_2^{\neg yy}(t)$ *(separate mode) and* $\Delta^2_{G,\mathrm{full}} = K + H(\iota)$ *(stage 1).* Ψ_2 *is a piecewise quartic in* t, *so the maximum is attained at* $t = \pm w$, *at a point where* $|\iota + t| \in \{3m_1, m_1 + 2m_2\}$, *or at a real root of the cubic* Ψ_2' *on a piece: a finite enumeration, as for Proposition S1. If a block has dimension* 1, *the same values are valid, conservative upper bounds.*

Proof. The bracket. For a record $\mathbf{w}$ and its replacement $\mathbf{w}'$, write n_k, n_k' for the squared norms of the blocks $k \in \{\mathbf{u}, \mathbf{x}_O, y\}$ and g_k for their inner products. The change of the first moments has squared norm $\iota \sum_k (n_k + n_k' - 2g_k)$; that of the off-diagonal blocks of $\mathbf{w}\mathbf{w}^\top - \mathbf{w}'\mathbf{w}'^\top$ has squared Frobenius norm $\sum_{k<l}(n_k n_l + n_k' n_l' - 2g_k g_l)$; and $\|\operatorname{triu} \mathbf{S}_k\|^2 \le \|\mathbf{S}_k\|_F^2 = n_k^2 + n_k'^2 - 2g_k^2$ for each diagonal block. The supremum of the sum (the paper's expression for Δ^2_{full}) is attained: for a block of dimension at

least 2 the two vectors can be chosen with any norms and inner product so that $\mathbf{S}_k$ is diagonal, and for y the bound is exact with $g_y = yy'$.

Norms at the bounds. For fixed g and (y, y') every coefficient of n_k, n'_k ($k = \mathbf{u}, \mathbf{x}_O$) is non-negative, and the admissible set $|g_k| \le \sqrt{n_k n'_k}$ grows with the norms; so the supremum has $n_{\mathbf{u}} = n'_{\mathbf{u}} = R_u$ and $n_O = n'_O = o$. Then, with $s = y^2 + y'^2$ and $t = yy'$, the objective is

$$K + \iota(s - 2t) + (R_u + o)s + (s^2 - 4t^2) + h(g_{\mathbf{u}}, g_O), \qquad h = -2(\iota + t)(g_{\mathbf{u}} + g_O) - 2(g_{\mathbf{u}}^2 + g_O^2 + g_{\mathbf{u}} g_O),$$

over $|g_{\mathbf{u}}| \le R_u$, $|g_O| \le o$ (the term $s^2 - 4t^2 = (y^2 - y'^2)^2$ is absent in separate mode, and every term in y at stage 1).

The inner maximum. h is a concave quadratic; its unconstrained maximizer is $g_{\mathbf{u}} = g_O = -a/3$, $a = \iota + t$, with value $\frac{2}{3}a^2$. If $|a| > 3m_1$, the coordinate with the smaller box is clipped at $-\operatorname{sgn}(a)m_1$ and the other is $-\operatorname{sgn}(a)(|a| - m_1)/2$, feasible while $|a| \le m_1 + 2m_2$, with value $\frac{1}{2}a^2 + |a|m_1 - \frac{3}{2}m_1^2$; beyond, both are clipped, with value $2|a|(m_1 + m_2) - 2(m_1^2 + m_2^2 + m_1 m_2)$. Each candidate satisfies the KKT conditions (for the clipped coordinate the partial derivative $-2a - 4g_1 - 2g_2$ has the sign of $-a$ once $|a| > 3m_1$), so it is the maximum of the concave problem; this is $H(a)$, continuous at the breakpoints.

The outer maximum. For fixed t the objective is strictly increasing in s (its derivative is $\iota + R_u + o + 2s > 0$). Given t, the pairs (y, y') with $yy' = t$ and $y^2, y'^2 \le w$ have $s \le w + t^2/w$, with equality iff one of $|y|, |y'|$ equals B_y; and $(w + t^2/w)^2 - 4t^2 = (w - t^2/w)^2$. So the supremum is $\max_{|t| \le w} \Psi_2(t)$. A configuration with a block of dimension 1 embeds in dimension 2 (collinear vectors), so its supremum is at most the value for dimension 2. □

At unit bounds without an intercept, $\Psi_2(t) = 9 + t^4 + \frac{2}{3}t^2$, so $\Delta^2_{\text{full}} = 32/3$ (at $t = \pm 1$) and $\Delta^2_{G,\text{full}} = 6$, the paper's values. On the simulation population the enumeration gives $\Delta_{\text{full}} = 92.670$ and $\Delta_{G,\text{full}} = 85.816$ (the paper's 92.7 and 85.8). Over 19 bound vectors, with and without an intercept, it agrees with the paper's grid-and-local-search computation (`experiments/baselines/sens_full.py`) to the relative 10^{-12} of that computation's safeguard, in all three modes; with a block of dimension 1 it is above the grid value, as it must be (`experiments/proofs_check/sens_full_closed.py`). Of the 29 benchmark datasets only one (tamielectric, $d = 3$, so $\mathbf{u}$ has dimension 1) uses the dimension-1 case, for which the theorem gives a valid conservative bound and the grid computation remains a numerical refinement.

13 Simulation of Algorithm 6 end to end

`alg6.py` implements Algorithm 6 line by line (stage 1, r_1, $\hat{L}$, λ_I, closed-form λ_O checked against bisection, the κ_Ψ rule, stage 2, gate, ridge solve, bootstrap with $B = 300$, Corr, $\tilde{s}^2$, both intervals, U_{yy} as specified and corrected). `s3_s6_sim.py` runs 2000 replications per scenario: fixed clipped design ($d_R = 4$, $d_O = 6$, $p = 11$, correlated features clipped at the 99% norm quantile), bounded uniform errors ($s = 1$), B_y a true bound so that the working model holds exactly, $\alpha = 0.05$, $\alpha_L = \alpha_U = 0.0025$, $\vartheta^\star = 2.5$, $\kappa_\Psi = 2$, $\sigma_1 = 2.5\sigma_2$, joint mode.

The Wald interval for $\boldsymbol{\theta}_\lambda$ attains its nominal level in every realistic scenario and over-covers slightly in the stress scenario, because the bootstrap variance there is too large. The β interval over-covers whenever the ridge is active: $\bar{b}_j$ is a worst-case bound and $\tilde{s}^2$ is conservative. It is uninformative ($\hat{L} = 0$) without a public bound in two of the three ridge scenarios.

scenario	n	σ_2	med. λ	Ψ	refused	Wald cov. (range over j)	β cov.	$\overline{\widetilde{\mathrm{RSS}} - \mathrm{RSS}_\lambda}$ / after Corr
off	4000	4	0	–	0	0.946–0.958	0.945–0.968	-0.14 / -0.05 (± 0.20)
ridge $\mathbf{I}$	1500	20	477	$\mathbf{I}$	0	0.948–0.960	∞ ($\hat{L} = 0$)	-6.83 / -1.68 (± 0.86)
ridge $\mathbf{I}$, $B^{\mathrm{pub}} = 1.5\|\boldsymbol{\beta}\|$	1500	20	477	$\mathbf{I}$	0	0.948–0.960	1.000	as above
ridge $\mathbf{I}$, private $\hat{B}_\beta$	1500	14	167	$\mathbf{I}$	0	0.950–0.960	1.000†	-4.49 / -1.31 (± 0.66)
ridge Π_O	3000	10	439	Π_O	0	0.944–0.962	∞ ($\hat{L} = 0$)	-3.03 / -0.07 (± 0.42)
small n, B^{pub}	200	20	1021	$\mathbf{I}$	0	0.946–0.962	1.000	-7.43 / -0.41 (± 0.50)
stress, B^{pub}	60	20	319	$\mathbf{I}$	0	0.960–0.974	1.000	-23.35 / $+3.05$ (± 0.57)

Table 5: Algorithm 6, 2000 replications each; Wald target $1 - \alpha_{\mathrm{CI}} = 0.955$, β target ≥ 0.95. In every scenario $\Pr[\Omega_L]$, $\Pr[\vartheta_0 \geq \vartheta^\star]$, $\Pr[\hat{L} \leq \lambda_{\min}(\mathbf{G})]$, $\Pr[U_{yy} \geq \mathbb{E}\mathbf{y}^\top\mathbf{y}]$ (both versions) and the agreement of closed-form λ_O with bisection were all 1.000. "Stress" uses $\sigma_1 = 0.01\sigma_2$, so that $\vartheta_0 \approx \vartheta^\star$; it is outside the default design. †The private bound was finite in 96.1% of replications.

14 Corrections to the specification

Each item states what `protocol_v31.tex` says, what is wrong or imprecise, and the corrected statement.

Correction 1 (Sensitivity: $d_O = 1$, the regime-3 algorithm, and the $\mathbf{y}^\top\mathbf{y}$ mode)**.** DeltaRep's comment "for $d_O \geq 2$ these are exact suprema" is correct (Theorem 3.1). For $d_O = 1$ the same numbers are valid but conservative upper bounds; the exact values are in Proposition 3.5, e.g. $\Delta_G^{(1)} = 2\sqrt{ro}$ for $o \leq \frac{11+5\sqrt{5}}{8} r$. In regime 3, replace "one-dimensional maximization, rounded up" by the exact candidate enumeration of Theorem 3.1(iv) (three cubics). Rounding up may be kept as a safeguard. The joint/separate rule of Algorithm 1, line 5, is consistent with Theorem 3.1(ix): $\mathbf{y}^\top\mathbf{y}$ is free iff $\max_t G \leq 4w(r+o)$ in regime 3, always free if $w \leq 2(r+o)$, never free if $w > 3(r+o)$.

Correction 2 (Statements inherited from the earlier paper)**.** (i) "$\mathbf{y}^\top\mathbf{y}$ is free" is false in general: $(1, 1, 5)$ gives $\Delta = 26.07$ against 14.14. (ii) "Using the full Frobenius norm of $\Delta\mathbf{B}$ would be conservative" is false for $d_O \geq 2$: the supremum is the same (Lemma 3.2); the triu maximum is attained only when $\mathbf{D}$ is diagonal. (iii) The worst case is not "$(\mathbf{x}, y) \mapsto (-\mathbf{x}, -y)$" in general; that is regime 2 only. In regime 1 it is $y' = -y$ with $\cos\angle(\mathbf{x}, \mathbf{x}') = -(r-w)/(2o)$ and $\mathbf{D}$ diagonal, and in regime 3 $|y'| < B_y$ when $\mathbf{y}^\top\mathbf{y}$ costs. (iv) "At the optimum the norms are at their bounds" is true for $\|\dot{\mathbf{x}}\|, \|\mathbf{x}\|, \|\mathbf{x}'\|$ (Lemma 3.3), false for $|y'|$ in costly regime 3, and false for $d_O = 1$, where only $\max(|x|, |x'|) = B_O$.

Correction 3 (The radius r_1)**.** The comment "$\Pr[\|\mathbf{G} - \hat{\mathbf{G}}_1\|_{\mathrm{op}} > r_1] \leq \alpha_L$" (Algorithm 6, line 2) is not proved and would need $\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}} \leq 2\sigma\sqrt{p}$ (Remark 4.2). Replace it by

$$\Pr\big[\lambda_{\max}(\hat{\mathbf{G}}_1 - \mathbf{G}) > r_1\big] \leq \alpha_L,$$

which is proved (Lemma 4.1(c)) and is the only event Algorithm 6 uses. For a two-sided statement use $r_1 = \sigma_1(2\sqrt{p} + 2\sqrt{\ln(2/\alpha_L)})$. Optional improvement: since the law of $\lambda_{\max}(\mathbf{E}_1)/\sigma_1$ depends only on (p, d_A), r_1 may be set to its exact $(1-\alpha_L)$-quantile, computed offline by Monte Carlo with a certified error. For $p = 11$ this is $7.43\sigma_1$ instead of $11.53\sigma_1$, and it costs no privacy.

Correction 4 (λ_O and the Π_O ridge)**.** Replace the bisection by the closed form of Lemma 5.1(c). State explicitly that $\lambda_O < \infty$ iff $\lambda_{\min}(\mathbf{A}) > g^\star + r_1$ (almost surely), and that the Π_O ridge can never certify $\vartheta_0 > \lambda_{\min}(\mathbf{A})/(2\sigma_2\sqrt{p})$.

Correction 5 (U_{yy})**.** As written, U_{yy} bounds $\mathbb{E}[\mathbf{y}^\top\mathbf{y} \mid \mathbf{X}]$ with probability $\geq 1-2\alpha_U$, not $1-\alpha_U$ (two one-sided events at level α_U each). The Hoeffding term $B_y^2\sqrt{n\ln(1/\alpha_U)/2}$ is itself correct, but it requires the y_i to be conditionally independent given $\mathbf{X}$. Corrected:

$$U_{yy} = \min\Big\{\widehat{yy} + z_{1-\alpha_U/2}\,\sigma_{yy} + B_y^2\sqrt{n\ln(2/\alpha_U)/2},\ nB_y^2\Big\} \qquad \text{or}$$
$$U_{yy} = \min\Big\{\widehat{yy} + \sqrt{2\ln(1/\alpha_U)(nB_y^4/4+\sigma_{yy}^2)},\ nB_y^2\Big\}.$$

Correction 6 (Estimand under clipping)**.** "CI for β_j" must name its target. Either (A) assume the working linear model, conditionally independent and homoskedastic, *for the clipped data*; this holds if the bounds are true bounds, so that clipping never binds. Or (B) target the projection coefficient $\boldsymbol{\beta}^* = \mathbf{G}^{-1}\tilde{\mathbf{X}}^\top\mathbb{E}[\tilde{\mathbf{y}} \mid \tilde{\mathbf{X}}]$, for which the bias bound remains valid but the homoskedastic variance is not, because clipping induces heteroskedasticity, so coverage is not guaranteed. Exactly Gaussian errors cannot hold with $|y| \leq B_y$; normality is asymptotic (Theorem 7.2(vi)).

Correction 7 (First-order accuracy is governed by $\sigma_2\,\mathrm{tr}\,\mathbf{P}_\lambda$, not by $1/\vartheta$ alone)**.** The validity ratio ϑ_0 controls convergence of the Neumann series. The relative size of the neglected mechanism bias in the Wald interval is of order $\sigma_2\,\mathrm{tr}(\mathbf{P}_\lambda) \leq \sqrt{p}/(2\vartheta_0)$, which at $\vartheta^\star = 2.5$, $p = 11$ may be 0.66. Coverage claims are asymptotic as $\vartheta_0 \to \infty$ at fixed p. Report $\sigma_2\,\mathrm{tr}(\tilde{\mathbf{P}})$ as a diagnostic, and consider centring at the bootstrap bias-corrected $2\tilde{\boldsymbol{\theta}} - \overline{\boldsymbol{\theta}^*}$.

Correction 8 (Bootstrap)**.** Without truncation the bootstrap mean and covariance of $\boldsymbol{\theta}^*$ do not exist (Lemma 8.1(e)); the gate $\hat{\vartheta} \geq 1$ does not fix this. Use the symmetric truncation of Lemma 8.1 (keep a draw iff $\|(\hat{\mathbf{G}}_2+\lambda\Psi)^{-1/2}\mathbf{E}^*(\hat{\mathbf{G}}_2+\lambda\Psi)^{-1/2}\|_{\mathrm{op}} \leq \eta$, e.g. $\eta = 1/2$), or the closed forms $b_{\mathrm{RSS}}(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}})$ and $\mathbf{V}_0(\hat{\mathbf{G}}_2, \tilde{\boldsymbol{\theta}})$. Use antithetic pairs if Monte Carlo is kept. "$\mathbf{V}_{\mathrm{mech}}$ is the mechanism covariance" holds to first order only (relative $O_p(\vartheta_0^{-1})$).

Correction 9 (The bias bound $\bar{b}_j$)**.** $\bar{b}_j = \lambda\|\Psi\tilde{\mathbf{P}}\mathbf{e}_j\|\hat{B}_\beta$ uses $\tilde{\mathbf{P}}$; the rigorous bound uses $\mathbf{P}_\lambda$. Either state the interval as first order, or inflate: on $\{\|\mathbf{E}_2\|_{\mathrm{op}} \leq r_2\}$ with $r_2 = \sigma_2(2\sqrt{p} + 2\sqrt{\ln(2/\alpha_2)})$ and $\|\tilde{\mathbf{P}}\|r_2 < 1$, $\|\Psi\mathbf{P}_\lambda\mathbf{e}_j\| \leq \|\Psi\tilde{\mathbf{P}}\mathbf{e}_j\| + \|\tilde{\mathbf{P}}\|r_2\|\tilde{\mathbf{P}}\mathbf{e}_j\|/(1-\|\tilde{\mathbf{P}}\|r_2)$, at the price of a further α_2 in the split.

Correction 10 (The private bound on $\|\boldsymbol{\beta}\|$)**.** $\sqrt{U_{yy}/\hat{L}} < \infty$ iff $\lambda_{\min}(\hat{\mathbf{G}}_1) > r_1$. When the ridge is active, $\lambda_{\min}(\hat{\mathbf{G}}_1) < g^\star + r_1$, and in our scenarios $\hat{L} = 0$ in every replication of two of three ridge settings. Even when finite, the bound is of order $\sqrt{(s^2 + \boldsymbol{\beta}^\top\Sigma\boldsymbol{\beta})/\lambda_{\min}(\Sigma)}$, far above $\|\boldsymbol{\beta}\|$. The specification should say that the β interval is informative essentially only with a public B_β^{pub} when the ridge binds.

Correction 11 (Gate and refusals)**.** State whether coverage is unconditional (refusal = no claim) or conditional on passing; the latter loses a factor $1/(1-\Pr[\mathsf{refused}])$. The rigorous refusal bound at the defaults is weak, $\alpha_L + e^{-p(\vartheta^\star-2)^2} = \alpha_L + 0.064$ for $p = 11$, although Monte Carlo gives essentially zero. The notation table should distinguish ϑ_0 (guaranteed $\geq \vartheta^\star$ on Ω_L) from the gate statistic $\hat{\vartheta}$.

Correction 12 (Minor)**.** (i) The reported "shrinkage $-\lambda\tilde{\mathbf{P}}\Psi\tilde{\boldsymbol{\theta}}$" is not $\boldsymbol{\theta}_\lambda - \hat{\boldsymbol{\beta}}_{\mathrm{OLS}}$, which equals $-\lambda\mathbf{G}^{-1}\Psi\boldsymbol{\theta}_\lambda$ exactly; it agrees with it only to first order in $\lambda\|\mathbf{P}_\lambda\Psi\|$, which is not small when the ridge binds. Rename it or report the exact form when $\mathbf{G}$ is well conditioned. (ii) $\tilde{s}^2$ is conservative under the ridge (Section 7); this is a feature, but it should be stated. (iii) The printed $\Delta = 86.18$, $\Delta_G = 77.20$ correspond to unrounded bounds; with $(5.59, 6.24, 2.89)$ the values are 86.15, 77.16. (iv) σ_{yy} is used in Algorithm 6 but not defined: it is σ_2 in joint mode and σ_y in separate mode.

15 Steps not made fully rigorous

(1) $\mathbb{E}\|\mathbf{E}\|_{\mathrm{op}} \le 2\sigma\sqrt{p}$ (two-sided) is not proved (Remark 4.2); it is not needed by the algorithm.
(2) (Resolved.) The constant of TV_1 is now proved, with an explicit σ^{-3} term (Lemma 2.2(b)).
(3) The boundary $w^*(r, o)$ of the free set in regime 3 is characterized exactly as $\max_t G = 4w(r+o)$, but its interval structure and the table of Remark 3.4 are numerical. Only $2(r + o) \le w^* \le 3(r + o)$ and $w^*(r, 0) = \frac{11+5\sqrt{5}}{8} r$ are proved.
(4) For $d_O = 1$, $\Delta^{(1)}$ (with $\mathbf{y}$) has no closed form; it is an exact two-variable maximization.
(5) Lemma S6: the constants c_k, C_η and the scale $\mathcal{S}_2$ are given in order of magnitude only; the remainder orders are rigorous in form but not in constants.
(6) Proposition S5: replacing the oracle quantities ($\mathbf{V}_0$, $\mathbf{P}_\lambda$ in $\bar{b}_j$, v_j) by plug-ins is first order. Normality of the sampling term is asymptotic. The variation of $\mathbf{V}_0$ with $\boldsymbol{\theta}_\lambda$ is neglected.
(7) The homoskedastic variance is not valid under binding clipping (Correction 6). A public variance bound gives a valid interval (Proposition 12.2); a guarantee with the data-driven $\tilde{s}^2$ is open.
(8) The size of the downward bias of $\lambda_{\min}(\hat{\mathbf{G}}_1)$ is bounded only by the comparison of Lemma 10.2, which is about a factor 1.9 loose on the simulation population; the empirical switch-off factors 1.2 and 1.4 are not bounds (Remark 10.3).

Scripts

All in `v31/proofs/scratch_stats/` (Python 3, NumPy, SciPy; `sim_core.py` imported read-only): `s1_common.py` (exact Δ, pieces, worst-case configuration), `s1_symbolic.py` (exact verification of the polynomial identities of the S1 proof), `s1_check.py` (Table 2), `s1_freeset.py` and `s1_band.py` (free set), `s2_check.py` (Table 3, discrete-Gaussian and TV checks), `s2_goe_ratio.py` (Remark 4.2), `alg6.py` (Algorithm 6), `s3_s6_sim.py` (Tables 4–5); outputs in `*.out`.

References

[1] T. W. Anderson, "The integral of a symmetric unimodal function over a symmetric convex set and some probability inequalities," *Proceedings of the American Mathematical Society* 6(2):170–176, 1955.

[2] T. B. Armstrong and M. Kolesár, "Simple and honest confidence intervals in nonparametric regression," *Quantitative Economics* 11(1):1–39, 2020.

[3] T. B. Armstrong, M. Kolesár and S. Kwon, "Bias-aware inference in regularized regression models," arXiv:2012.14823, 2020 (revised 2023). *[arXiv number verified; published version not checked.]*

[4] A. S. Bandeira and R. van Handel, "Sharp nonasymptotic bounds on the norm of random matrices with independent entries," *Annals of Probability* 44(4):2479–2506, 2016.

[5] S. Boucheron, G. Lugosi and P. Massart, *Concentration Inequalities: A Nonasymptotic Theory of Independence*, Oxford University Press, 2013. *[Theorem 5.6 is the Gaussian concentration inequality; numbering from memory.]*

[6] C. L. Canonne, G. Kamath and T. Steinke, "The discrete Gaussian for differential privacy," in *Advances in Neural Information Processing Systems 33 (NeurIPS)*, pp. 15676–15688, 2020. *[The*

bounds $\mathrm{Var} \le \sigma^2$ *and* $\mathbb{E}e^{tX} \le e^{t^2\sigma^2/2}$ *are proved there; a secondary source places the subgaussian bound at Corollary 17 of the arXiv version (arXiv:2004.00010); numbering not checked against the paper.]*

[7] S. Chatterjee, "An error bound in the Sudakov–Fernique inequality," arXiv:math/0510424, 2005. *[The comparison is stated for Gaussian vectors with equal means; the theorem number was not checked.]*

[8] J. M. Danskin, "The theory of max-min, with applications," *SIAM Journal on Applied Mathematics* 14(4):641–664, 1966.

[9] X. Fernique, "Régularité des trajectoires des fonctions aléatoires gaussiennes," in *École d'Été de Probabilités de Saint-Flour IV—1974*, Lecture Notes in Mathematics 480, Springer, 1975, pp. 1–96.

[10] W. Hoeffding, "Probability inequalities for sums of bounded random variables," *Journal of the American Statistical Association* 58(301):13–30, 1963.

[11] V. V. Petrov, *Sums of Independent Random Variables*, Springer, 1975. *[Berry–Esseen for non-identically distributed summands, Ch. V.]*

[12] V. N. Sudakov, "Gaussian random processes and measures of solid angles in Hilbert space," *Soviet Mathematics Doklady* 12:412–415, 1971.

[13] B. S. Tsirelson, I. A. Ibragimov and V. N. Sudakov, "Norms of Gaussian sample functions," in *Proceedings of the Third Japan–USSR Symposium on Probability Theory*, Lecture Notes in Mathematics 550, Springer, 1976, pp. 20–41.

[14] R. Vershynin, *High-Dimensional Probability: An Introduction with Applications in Data Science*, Cambridge University Press, 2018. *[Thm. 7.2.11 (Sudakov–Fernique), Cor. 4.2.13 (nets), Ex. 4.4.3 (quadratic forms on nets); numbering from memory.]*

Protocol v3.1: composition across runs

Proof of Proposition `prop:compose` for runs fixed in advance and for adaptively chosen runs under the ledger's filter, and a hockey-stick treatment of the smudging slack

Draft for review

Abstract

This note proves the composition statement of protocol v3.1 (Proposition `prop:compose` of the paper) and closes the check item "the adaptive hybrid of Proposition `prop:compose` is argued here but not yet written out in the proofs". We model all researchers as one adaptive coalition that chooses each run's parameters and inputs from the joint view of earlier runs, follows the protocol inside each run, may interleave and abort runs, and faces O's ledger. We prove that the joint view is $(\varepsilon_L, \delta_L)$-IND-CDP for adjacency on the base register, under each admission rule (Theorem 4.2). The proof has four parts: one statistical hybrid that removes every encryption error at once, a computational hybrid run by run, a self-contained Rényi filter for adaptively parameterized lattice Gaussian mechanisms, and the conversion. Part 2 measures the smudging step in hockey-stick divergence (Lemma 5.1, Theorem 5.2). With a reserved share $\xi_L = 10^{-3}$ of ε_L, the ledger's flooding requirement κ_L at $(\varepsilon_L, \delta_L, \zeta_L) = (10, 10^{-6}, 0.1)$ falls from $48, 46, 44, 42, 40$ to $22, 21, 19, 17, 15$ for $\varepsilon_{\text{run}} = 0.5, 1, 2, 4, 8$, so it no longer binds against the floor $\kappa_{\min} = 40$.

References to other documents. This note is Part C of the proofs appendix (`protocol_v31_proofs.tex`). [CP] is `crypto_proofs.tex`, the proofs of the cryptographic and DP claims (Part A of the proofs appendix); result numbers are those of its current build, and labels are given in typewriter font. [P] is the paper, `paper/main.tex`. Notation follows [CP].

Status (2026-10-25). The paper, the implementation and [CP] now follow this note: the paper states Theorem 4.2 as Proposition `prop:compose` with base-register adjacency (`def:baseadj`); its ledger is the ledger of Definition 5.3 with $\xi_L = 10^{-3}$ and $\delta' = 2^{-100}$, charged in full at admission; its κ_L is that of Proposition 5.5; and the floor is 2^{42} in the paper, in [CP] (Corollary 2.4, corrected) and in the implementation. The issues of Section 6 are resolved accordingly, except that rule (N) with adaptively chosen requests remains unproved (Remark 4.4).

Summary of results.

- **Theorem 4.2** replaces Proposition `prop:compose`. For every PPT coalition of researchers, with arbitrary collusion, adaptivity, interleaving and aborts, the joint view is simulated by an ideal experiment with slack $\mathbb{E}[\sum_{k \text{ admitted}} \gamma_k] + \text{negl}(\lambda)$. The ideal experiment is ρ^*-zCDP whenever the admitted runs satisfy $\sum_k \rho_k \le \rho^*$ pointwise. Under the filter rule the joint view is $(\varepsilon_L, \delta_L)$-IND-CDP for every adaptive coalition. Under the non-adaptive rule the same holds when the run requests are fixed in advance.
- **The adaptive case goes through as the paper's sketch claims**, with three precisions.

(i) All smudging steps of all runs are made in one statistical hybrid, before any computational step. The slack is then a single expectation under the real experiment, which the pointwise filter bounds. Taken step by step, the sketch's "expectation of γ_k" would be an expectation under a different hybrid for each step.

(ii) The δ_B part of a run's slack must be charged for *every planned release of every admitted run*, aborted or not. A researcher's decision to continue may depend on s_R and on its own encryption errors, which enter e_{tot}. The simplest rule is to charge a run's full cost at admission, with no refund.

(iii) The RLWE assumptions are needed for adaptively chosen parameter sets (a finite menu suffices) and against non-uniform adversaries, since the reduction knows the data.

- **The Rényi filter** (Theorem 3.2) is proved here from scratch by the supermartingale argument. It covers lattice Gaussian steps whose dimension, centre and σ depend on the history, together with refusals and other public steps.
- **The δ budget adds up exactly**: $\delta(\rho_L, \varepsilon_L) + (1 + e^{\varepsilon_L})\Gamma_L = (1 - \zeta_L)\delta_L + \zeta_L\delta_L = \delta_L$. Numerically the sum is $9.99999999999 \cdot 10^{-7}$. The paper's κ_L values $48, 46, 44, 42, 40$ are reproduced.
- **Part 2.** A shifted lattice Gaussian with $\|\mathbf{e}\| \le B_{\text{tot}}$ and $\sigma \ge 2^{\kappa} B_{\text{tot}}/\sqrt{2\pi}$ is (ε', δ')-indistinguishable from the unshifted one with $\varepsilon' \le 2\sqrt{\pi \ln(1/\delta')}\, 2^{-\kappa} + \pi 4^{-\kappa}$. This bound needs no smoothing condition. Composing with an (ε', δ') smudging step gives a three-budget ledger and a new rule for κ_L (Section 5).
- **Issues in the paper** are collected in Section 6. None invalidates the stated guarantee once the ledger charges at admission. They include a constant in the smudging lemma: the factor $1 + 2^{-80}$ needs $\sigma_c \ge 2^{42}$, not 2^{20}.

Contents

1 Setting **2**
- 1.1 The base register and its adjacency . . . 2
- 1.2 Runs, costs and the ledger . . . 3
- 1.3 The multi-run experiment . . . 4
- 1.4 Assumptions . . . 4

2 Divergences on countable spaces **5**

3 Adaptive processes: a Rényi filter and two hybrid lemmas **7**

4 Composition across runs **9**
- 4.1 The ideal experiment . . . 9
- 4.2 Proof of Theorem 4.2 . . . 11
- 4.3 Corner cases . . . 14

5 Part 2: the smudging step in hockey-stick divergence **16**
- 5.1 The shifted lattice Gaussian . . . 16
- 5.2 Composition with an (ε', δ') smudging step . . . 16
- 5.3 A ledger with a hockey-stick slack, and the new κ_L . . . 18

6 Issues found in the paper's current text **19**

7 What is used and what is proved here 20

1 Setting

1.1 The base register and its adjacency

The paper's adjacency ([P], Definition `def:adj`) is stated for the padded scope register of one run, with the cohort fixed. Runs on one base register may use different scopes, models and cohorts, so the joint statement needs an adjacency on the base register itself.

Definition 1.1 (Base register; base adjacency)**.** A *base register* is $\mathfrak{D} = (\mathrm{id}_j, \mathbf{a}_j, \mathbf{x}_j, y_j)_{j \le n_\mathrm{b}}$. Here id_j is an identifier, $\mathbf{a}_j$ holds the public attributes that scopes may use, $\mathbf{x}_j$ holds all of O's feature columns and y_j is the response. The *fixed part* of $\mathfrak{D}$ is $\mathrm{fix}(\mathfrak{D}) = (n_\mathrm{b}, (\mathrm{id}_j, \mathbf{a}_j)_j)$. Two base registers are adjacent, $\mathfrak{D} \sim_\mathrm{b} \mathfrak{D}'$, if $\mathrm{fix}(\mathfrak{D}) = \mathrm{fix}(\mathfrak{D}')$ and $(\mathbf{x}_j, y_j) \neq (\mathbf{x}'_j, y'_j)$ for exactly one j.

Definition 1.2 (Dataset of a run)**.** A run's public parameters P fix four things:

- a scope predicate Φ on public attributes;
- a row map g_P from $(\mathbf{x}, y)$ to O's preprocessed and clipped columns, which selects the model's columns and applies public constants;
- a padded size $\bar{n}_O \ge |\{j : \Phi(\mathbf{a}_j)\}|$;
- the number of dummy rows.

Given dummy identifiers $\mathbf{d}$, the run's dataset is the scope register $D_P(\mathfrak{D}; \mathbf{d})$. It consists of the rows $(\mathrm{id}_j, g_P(\mathbf{x}_j, y_j))$ with $\Phi(\mathbf{a}_j)$, padded with rows $(\mathrm{d}_i, \mathbf{0}, 0)$ to $\bar{n}_O$ rows.

Lemma 1.3 (Base adjacency induces scope adjacency)**.** *Let $\mathfrak{D} \sim_\mathrm{b} \mathfrak{D}'$. Then for every P and every $\mathbf{d}$, the datasets $D_P(\mathfrak{D}; \mathbf{d})$ and $D_P(\mathfrak{D}'; \mathbf{d})$ either coincide or are adjacent in the sense of* [P]*, Definition* `def:adj`*. Adjacent here means that they have the same identifiers and the same scope, and differ in $(\mathbf{x}_O, y)$ of exactly one row.*

Proof. The selected rows depend on the public attributes, which are equal. The map g_P acts row by row with public constants, so only the row j in which $\mathfrak{D}$ and $\mathfrak{D}'$ differ can change, and only if $\Phi(\mathbf{a}_j)$ holds. The dummy rows are equal. □

The lemma needs row-wise preprocessing with public constants. A data-dependent transformation, such as standardizing O's columns by register means, would let one row change every row, and is excluded by Assumption 1.9(A4).

1.2 Runs, costs and the ledger

Definition 1.4 (Run request; costs)**.** A *run request* is a pair (P, χ):

- $P \in \mathcal{P}_\lambda$ is a public parameter vector of the paper's Section `sec:params`. It comprises the scope, the model and its preprocessing, the bounds, the run budget $(\varepsilon_\mathrm{run}, \delta_\mathrm{run})$, the splits, ζ, the policy for κ, δ_B, and an HE parameter set $p \in \Pi_\lambda$ (ring dimension, modulus chain, distributions).
- χ is the researcher's private input: at most $n_R^{\max}$ cohort identifiers, with features $\mathbf{X}_R$.

Calibrate(P) is deterministic. It either refuses or returns κ, n_rel and $(\sigma_{c,s}, \rho_s, B_{\mathrm{tot},s})_{s \le n_\mathrm{rel}}$. The run's *costs* are

$$\rho(P) = \sum_{s \le n_\mathrm{rel}} \rho_s, \qquad \gamma(P) = \sum_{s \le n_\mathrm{rel}} \gamma_s, \qquad \gamma_s = \delta_B + \gamma'(\kappa), \qquad \gamma'(\kappa) = 2^{-\kappa}(1 + 2^{-80}) + 2^{-97}.$$

The costs are public, because they are functions of P.

Throughout, $\delta(\rho,\varepsilon)$ is the tight conversion of Lemma 2.3, which is the function of [CP], Proposition 4.4 (`prop:conv`).

Definition 1.5 (Ledger)**.** The ledger of a base register has public parameters $(\varepsilon_L, \delta_L, \zeta_L)$. It derives

$$\rho_L = \max\{\rho \geq 0 : \delta(\rho, \varepsilon_L) \leq (1-\zeta_L)\delta_L\}, \qquad \Gamma_L = \frac{\zeta_L \delta_L}{1+e^{\varepsilon_L}}.$$

Its state is $(\rho_\Sigma, \Gamma_\Sigma)$, initially $(0,0)$. It processes requests one at a time. Let a request have costs (ρ_k, γ_k) and let CALIBRATE not refuse it. The ledger uses one of two rules, fixed when the ledger is created:

(F) *filter:* admit iff $\rho_\Sigma + \rho_k \leq \rho_L$ and $\Gamma_\Sigma + \gamma_k \leq \Gamma_L$;
(N) *non-adaptive:* admit iff $\delta(\rho_\Sigma + \rho_k, \varepsilon_L) + (1+e^{\varepsilon_L})(\Gamma_\Sigma + \gamma_k) \leq \delta_L$.

On admission the ledger sets $(\rho_\Sigma, \Gamma_\Sigma) \leftarrow (\rho_\Sigma + \rho_k, \Gamma_\Sigma + \gamma_k)$ at once. The full planned cost is charged before any message of the run that depends on a secret, and it is not refunded if the run is aborted.

Remark 1.6 (Differences from the paper's ledger)**.** The paper's earlier text defined the state as the sum of the costs "of every release made" and admitted a run against its full cost. The implementation (`vfr/dp.py`, class `Ledger`, and `vfr/party_o.py`) checked a run's full cost at setup but charged each release only when it was made. This was a reservation-free scheme with implicit refunds on abort. Both now charge the full cost at admission, atomically (the implementation under an exclusive file lock), with no refund. Remark 4.3 shows which refunds are sound: those of ρ_s and of $\gamma'(\kappa)$ for releases not made. The δ_B part of an aborted run's unmade releases must stay charged. Without reservation, two runs admitted against the same state can together exceed the cap. Definition 1.5 avoids both problems.

1.3 The multi-run experiment

Definition 1.7 (The real experiment $\mathsf{Real}_{\mathcal{A}}(\mathfrak{D})$)**.** Let $\mathcal{A}$ be an interactive PPT machine, the *coalition*. It receives 1^λ and draws its random tape $\omega_{\mathcal{A}}$. Until it halts it repeats one of the following actions, each chosen as a function of $\omega_{\mathcal{A}}$ and of everything it has received so far.

(a) *Request.* $\mathcal{A}$ submits a run request (P_k, χ_k); requests are numbered in order of submission. O computes CALIBRATE(P_k) and answers $\perp_{\text{cal}}$ if it refuses. Otherwise O submits the costs to the ledger and answers $\perp_{\text{led}}$ if the ledger refuses. Otherwise run k is *admitted*, the ledger charges it, and the run becomes *open*. O draws fresh coins $\omega_{O,k}$, and the honest researcher machine $\mathsf{R}_k = \mathsf{R}(P_k, \chi_k; \omega_{R,k})$ gets fresh uniform coins $\omega_{R,k}$, which are given to $\mathcal{A}$. The coins $\omega_{O,k}$ comprise:
- O's key shares s_O and their errors;
- the common random strings;
- the OPRF key, the shuffle τ_k, the dummy identifiers and the PSI sanitization coins;
- for each release $s \leq n_{\text{rel}}$, the re-randomization coins (u, e_0, e_1) and the noise $\eta \sim D_{\mathbb{Z}^N, \sigma_{c,s}}$.

(b) *Step.* $\mathcal{A}$ selects an open run. The next message of that run is computed by the honest party whose turn it is, either R_k or $\mathsf{O}(\mathfrak{D}, P_k; \omega_{O,k})$, and delivered. $\mathcal{A}$ receives every message that R_k sends or receives.

(c) *Abort.* $\mathcal{A}$ closes an open run, and no further message of that run is sent. In Algorithm 5 the researcher requests each release, so not requesting a release is an abort.

(d) *Halt.* $\mathcal{A}$ outputs a bit, which is the output of $\mathsf{Real}_{\mathcal{A}}(\mathfrak{D})$.

$\text{view}_{\mathcal{A}}(\mathfrak{D})$ consists of $\omega_{\mathcal{A}}$ and every answer, coin $\omega_{R,k}$ and message that $\mathcal{A}$ received.

The model makes the researchers semi-honest within each run: the researcher side of every run is the honest machine. Across runs they are fully adaptive in what they request, when, how they

interleave open runs, and when they abort. One machine stands for all researchers, so arbitrary collusion is covered. Non-colluding researchers are the special case in which $\mathcal{A}$'s actions for researcher i depend only on i's part of the view. The view of any subset of researchers is a function of $\text{view}_{\mathcal{A}}$.

Definition 1.8 (IND-CDP of the multi-run protocol)**.** The protocol is (ε, δ)-*IND-CDP for* $\sim_{\mathrm{b}}$ if, for every PPT coalition $\mathcal{A}$, there is a negligible ν such that for all $\mathfrak{D} \sim_{\mathrm{b}} \mathfrak{D}'$

$$\Pr[\mathsf{Real}_{\mathcal{A}}(\mathfrak{D}) = 1] \leq e^{\varepsilon} \Pr[\mathsf{Real}_{\mathcal{A}}(\mathfrak{D}') = 1] + \delta + \nu(\lambda).$$

$\mathcal{A}$'s output may be any PPT function of its view, so this is the IND-CDP of [P], Definition `def:simcdp`, applied to the joint view. The coalition's adaptive strategy is built into the experiment.

1.4 Assumptions

Assumption 1.9. (A1) *Honest data holder and honest, atomic ledger.* O follows the protocol. Its ledger behaves as in Definition 1.5, with one rule. O produces decryption shares only within admitted runs. The answers $\perp_{\text{cal}}$ and $\perp_{\text{led}}$ depend only on P_k and the ledger state. O aborts a run only on events determined by $\text{fix}(\mathfrak{D})$, P_k, its coins other than (s_O, its errors, u, e_0, e_1, η), and the messages it has received. An overflow of the PSI tables is an example of such an event.

(A2) *Semi-honest researchers with arbitrary collusion*, as in Definition 1.7. In particular each R_k clips its features, and its coins are uniform.

(A3) *Fresh coins.* Each admitted run draws $\omega_{O,k}$ and $\omega_{R,k}$ independently of each other and of everything before, with the protocol's distributions. The CKKS key material of O and R (s_O, s_R, their errors and the common random strings) is fresh in every run. (u, e_0, e_1, η) is fresh in every release, and η is *exactly* $D_{\mathbb{Z}^N, \sigma_{c,s}}$.

(A4) *Static register, public scopes, row-wise public preprocessing.* $\mathfrak{D}$ is the same in all runs. Every scope uses public attributes only, and every g_P is row-wise with public constants. Lemma 1.3 therefore holds.

(A5) *Per-run hypotheses.* Every $P \in \mathcal{P}_\lambda$ that Calibrate does not refuse satisfies the hypotheses of [CP], Theorem 4.11 (`thm:output`), for every release s:

- (M1)–(M5);
- $\sigma_{c,s} \geq \max\{2^{\kappa} B_{\text{tot},s}/\sqrt{2\pi},\ 2^{42}\}$ (see Remark 2.6 for the floor);
- $\rho_s \geq \widehat{\Delta}^2_{\text{enc},s}/(2\sigma^2_{c,s})$, where $\widehat{\Delta}_{\text{enc},s}$ bounds the sensitivity of the centre f_s under scope adjacency, for every arrangement and every researcher input in the clipped domain ([CP], Corollary 3.13, `cor:sensZ`);
- $B_{\text{tot},s}$ satisfies [CP], Theorem 3.12 (`thm:01`): $\Pr[\|e_{\text{tot},s}\|_2 > B_{\text{tot},s}] \leq \delta_B$ for every dataset in the clipped domain, every researcher input and every value of the public randomness. The probability is over the run's key shares, errors and fresh coins.

(A6) *RLWE for every parameter set in a finite menu.* [CP], Assumptions 4.9 (`as:rlweu`) and 4.10 (`as:rlweO`) hold for every $p \in \Pi_\lambda$, and $|\Pi_\lambda| \leq \text{poly}(\lambda)$. They hold against non-uniform adversaries: for every polynomial t, the best advantages $\text{Adv}_{1,p}(\lambda, t(\lambda))$ and $\text{Adv}_{2,p}(\lambda, t(\lambda))$ of circuits of size t are negligible.

(A7) *Efficiency.* Calibrate and the ledger tests run in deterministic polynomial time. The centre $f_s(D; \pi)$ is computable in time polynomial in λ and $|\mathfrak{D}|$. The exact sampler runs in expected polynomial time; a reduction may cut it off at a polynomial time bound, which changes its output with probability $2^{-\lambda}$, so it suffices.

(A8) *PPT coalition.* $\mathcal{A}$ runs in time $t_{\mathcal{A}}(\lambda) \leq \text{poly}(\lambda)$, so it makes at most $t_{\mathcal{A}}(\lambda)$ requests and actions.

(A3) is used only for the CKKS key material and for the release coins. Fresh OPRF keys and shuffles matter for cohort and linkage privacy (G2, G4) across runs, not for Theorem 4.2: the theorem never replaces the Link phase, and the sensitivity bound holds for every arrangement.

2 Divergences on countable spaces

Every object below is a probability mass function (pmf) on a countable set. This covers every transcript of the protocol, which consists of bit strings, ring elements and lattice points.

Definition 2.1 (Divergences)**.** Let P, Q be pmfs on a countable set $\mathcal{Y}$. Let $\phi_a(p,q) = p^a q^{1-a}$ for $q > 0$, with $\phi_a(0,0) = 0$ and $\phi_a(p,0) = \infty$ for $p > 0$.

- For $a > 1$, $\mathsf{M}_a(P\|Q) = \sum_y \phi_a(P(y), Q(y))$, and the Rényi divergence is $D_a(P\|Q) = \frac{1}{a-1} \ln \mathsf{M}_a(P\|Q)$.
- For $\varepsilon \geq 0$, the hockey-stick divergence is

$$\mathsf{E}_\varepsilon(P\|Q) = \sum_y (P(y) - e^\varepsilon Q(y))_+ = \sup_{\varphi:\mathcal{Y}\to[0,1]} \big(\mathbb{E}_P\varphi - e^\varepsilon \mathbb{E}_Q\varphi\big).$$

- $\mathrm{SD} = \mathsf{E}_0$ is the statistical distance.

P and Q are (ε, δ)-*indistinguishable* if $\mathsf{E}_\varepsilon(P\|Q) \leq \delta$ and $\mathsf{E}_\varepsilon(Q\|P) \leq \delta$. A finite $D_a(P\|Q)$ implies $P \ll Q$.

The supremum form of E_ε is attained at $\varphi = \mathbf{1}\{P > e^\varepsilon Q\}$. It covers randomized tests: a distinguisher $\mathcal{B}$ with coins defines $\varphi(y) = \Pr[\mathcal{B}(y) = 1] \in [0,1]$.

Lemma 2.2 (Post-processing; mixtures with common weights)**.** *Let K be a Markov kernel from $\mathcal{Y}$ to a countable $\mathcal{Z}$, and write $PK(z) = \sum_y P(y)K(z \mid y)$. Then $\mathsf{M}_a(PK\|QK) \leq \mathsf{M}_a(P\|Q)$ and $\mathsf{E}_\varepsilon(PK\|QK) \leq \mathsf{E}_\varepsilon(P\|Q)$. In particular, for pmfs μ on a countable index set and families (P_r), (Q_r),*

$$\mathsf{M}_a\Big(\textstyle\sum_r \mu_r P_r \,\Big\|\, \sum_r \mu_r Q_r\Big) \leq \sum_r \mu_r \mathsf{M}_a(P_r\|Q_r), \qquad \mathsf{E}_\varepsilon\Big(\textstyle\sum_r \mu_r P_r \,\Big\|\, \sum_r \mu_r Q_r\Big) \leq \sum_r \mu_r \mathsf{E}_\varepsilon(P_r\|Q_r).$$

Proof. Both ϕ_a and $\psi_\varepsilon(p,q) = (p - e^\varepsilon q)_+$ are convex, lower semicontinuous and positively homogeneous of degree one on $[0,\infty)^2$. Here ϕ_a is the closed perspective of $x \mapsto x^a$. Such a function f is subadditive, since $f(x + x') = 2f(\frac{x+x'}{2}) \leq f(x) + f(x')$. By lower semicontinuity this extends to countable sums: $f(\sum_i x_i) \leq \liminf_n \sum_{i\leq n} f(x_i) \leq \sum_i f(x_i)$. Hence

$$\mathsf{M}_a(PK\|QK) = \sum_z \phi_a\Big(\sum_y K(z \mid y)(P(y), Q(y))\Big) \leq \sum_z \sum_y K(z \mid y)\, \phi_a(P(y), Q(y)) = \mathsf{M}_a(P\|Q),$$

and likewise for ψ_ε. For the mixture statement, apply this to the joint pmfs $\mu_r P_r(y)$ and $\mu_r Q_r(y)$ of (r, y), whose M_a and E_ε equal the right-hand sides, and to the projection $(r, y) \mapsto y$. □

Lemma 2.3 (Conversion; [2], [CP], Proposition 4.4)**.** *Suppose $D_a(P\|Q) \leq a\rho$ for every $a > 1$. Then for every $\varepsilon \geq 0$,*

$$\mathsf{E}_\varepsilon(P\|Q) \leq \delta(\rho,\varepsilon) = \inf_{a>1} e^{(a-1)(a\rho-\varepsilon)} \frac{(a-1)^{a-1}}{a^a}.$$

Moreover $\delta(\rho,\varepsilon) \leq e^{-(\varepsilon-\rho)^2/(4\rho)}$ for $\varepsilon \geq \rho$. Hence $\mathsf{E}_\varepsilon(P\|Q) \leq \delta'$ whenever $\varepsilon \geq \rho + 2\sqrt{\rho \ln(1/\delta')}$. The function $\delta(\cdot,\varepsilon)$ is continuous, and strictly increasing on the set where it is below 1; for $\varepsilon > 0$, $\delta(0,\varepsilon) = 0$.

Proof. For $x \ge 0$ and $a > 1$ we have $(x - e^{\varepsilon})_+ \le c_a x^a$, where

$$c_a = \sup_{x>0} \frac{(x-e^{\varepsilon})_+}{x^a} = e^{-(a-1)\varepsilon} \frac{(a-1)^{a-1}}{a^a}.$$

The supremum is attained at $x = ae^{\varepsilon}/(a-1)$. Since $P \ll Q$,

$$\mathsf{E}_{\varepsilon}(P\|Q) = \sum_{Q(y)>0} Q(y)\Big(\frac{P(y)}{Q(y)} - e^{\varepsilon}\Big)_+ \le c_a \mathsf{M}_a(P\|Q) \le c_a e^{(a-1)a\rho}.$$

Take the infimum over a. For the simple form, note that $(a-1)^{a-1}/a^a = \frac{1}{a}(1-\frac{1}{a})^{a-1} \le 1$, and take $a = (\varepsilon+\rho)/(2\rho)$, which exceeds 1 when $\varepsilon > \rho$. The exponent is then $(a-1)(a\rho-\varepsilon) = -(\varepsilon-\rho)^2/(4\rho)$.

Continuity is [CP], Proposition 4.4. For strict monotonicity, let $\rho < \rho'$ with $\delta(\rho',\varepsilon) < 1$. The objective is continuous in a, tends to 1 as $a \to 1$ and to $+\infty$ as $a \to \infty$, so its infimum $\delta(\rho',\varepsilon) < 1$ is attained at some a^*. Then $\delta(\rho,\varepsilon) \le e^{(a^*-1)(a^*\rho-\varepsilon)}c' < e^{(a^*-1)(a^*\rho'-\varepsilon)}c' = \delta(\rho',\varepsilon)$, where $c' = (a^*-1)^{a^*-1}/a^{*a^*}$. □

Lemma 2.4 (Integer shifts of the lattice Gaussian; [CP], Lemma 2.2)**.** *Let $\eta \sim D_{\mathbb{Z}^n,\sigma}$ and $\mathbf{e} \in \mathbb{Z}^n$. Then $D_a(\mathbf{e}+\eta\|\eta) \le a\|\mathbf{e}\|_2^2/(2\sigma^2)$ for every $a > 1$.*

Proof. Let $Z = \rho_\sigma(\mathbb{Z}^n)$. Since $\mathbf{e}$ is integral, $\mathbf{e}+\eta$ has pmf $\rho_\sigma(y-\mathbf{e})/Z$ on $\mathbb{Z}^n$. By the identity $a\|y-\mathbf{e}\|^2 + (1-a)\|y\|^2 = \|y-a\mathbf{e}\|^2 + a(1-a)\|\mathbf{e}\|^2$,

$$\mathsf{M}_a = e^{a(a-1)\|\mathbf{e}\|^2/(2\sigma^2)} \frac{\rho_\sigma(\mathbb{Z}^n - a\mathbf{e})}{Z} \le e^{a(a-1)\|\mathbf{e}\|^2/(2\sigma^2)}.$$

The inequality is [CP], Lemma 2.1 (`lem:gsum`), proved there by Poisson summation. □

The same bound holds for $D_a(\eta\|\mathbf{e}+\eta)$: shift both laws by $-\mathbf{e}$, which is a bijection of $\mathbb{Z}^n$.

Lemma 2.5 (Smudging in statistical distance; [CP], Lemma 2.3(b) and Corollary 2.4)**.** *Let $\eta \sim D_{\mathbb{Z}^N,\sigma}$ with $N \le 2^{20}$, let $\mathbf{e} \in \mathbb{Z}^N$ with $\|\mathbf{e}\|_2 \le B$, and let $\sigma \ge \max\{2^\kappa B/\sqrt{2\pi}, 2^{42}\}$. Then $\mathrm{SD}(\mathbf{e}+\eta,\eta) \le \gamma'(\kappa) = 2^{-\kappa}(1+2^{-80}) + 2^{-97}$.*

Proof. [CP], Lemma 2.3(b), with smoothing error $\epsilon = 2^{-100}$ and $\sigma_2 = 2.1$, gives

$$\mathrm{SD} \le \frac{\|\mathbf{e}\|}{\sqrt{2\pi}\,\sigma_1} + \frac{4\epsilon}{1-\epsilon} \le 2^{-\kappa}\frac{\sigma}{\sigma_1} + 2^{-97.99},$$

where $\sigma_1 = \sqrt{\sigma^2 - \sigma_2^2}$. The choice $\sigma_2 = 2.1$ is admissible for $N \le 2^{20}$, because by [4, Lemma 3.3]

$$\frac{\eta_\epsilon(\mathbb{Z}^N)}{\sqrt{2\pi}} \le \frac{\sqrt{\ln(2N(1+2^{100}))/\pi}}{\sqrt{2\pi}} \le 2.07.$$

Finally, $\sigma/\sigma_1 = (1-\sigma_2^2/\sigma^2)^{-1/2} \le 1 + \sigma_2^2/\sigma^2 \le 1 + 4.41\cdot 2^{-84} < 1 + 2^{-81}$. □

Remark 2.6 (The floor on σ_c)**.** [CP] and [P] stated the bound $\gamma'(\kappa)$ under $\sigma_c \ge \max\{2^\kappa B_{\mathrm{tot}}/\sqrt{2\pi}, 2^{20}\}$; both now use 2^{42}. At $\sigma = 2^{20}$, however, σ/σ_1 is $1 + 2^{-38.9}$, and the proof's bound $1 + \sigma_2^2/\sigma^2$ is $1 + 2^{-37.9}$. Either way the factor is not $1 + 2^{-80}$. Correspondingly, the last step of [CP], Corollary 2.4 ($2^{-\kappa}\sigma/\sigma_1 + 2^{-97.9} \le 2^{-\kappa} + 2^{-97}$ "for every $\sigma \ge 2^{20}$") fails for $\kappa < 60$. The factor $1 + 2^{-80}$ needs $\sigma \ge 2.1 \cdot 2^{40} \approx 2^{41.1}$. This has no practical effect: at every evaluation setting $\kappa \ge 40$ and $B_{\mathrm{tot}} \ge 2^{17}$, so $\sigma_c \ge 2^\kappa B_{\mathrm{tot}}/\sqrt{2\pi} \ge 2^{55} \gg 2^{42}$. The fix is to write 2^{42} for 2^{20}, or $1 + 2^{-37}$ for $1 + 2^{-80}$. Nothing below depends on the exact form of $\gamma'(\kappa)$, only on Lemma 2.5 holding for it.

3 Adaptive processes: a Rényi filter and two hybrid lemmas

Definition 3.1 (Adaptive process)**.** Let $\mathcal{Y}$ be countable and $M \in \mathbb{N}$. An *adaptive process* P of horizon M on $\mathcal{Y}$ is a family of kernels: for each $t \le M$ and each history $h \in \mathcal{Y}^{t-1}$, a pmf $P_t(\cdot \mid h)$ on $\mathcal{Y}$. Its law is $P(y) = \prod_{t \le M} P_t(y_t \mid y_{<t})$, also written P. A process that stops early outputs a fixed symbol $\perp$ with probability one from then on. A family $c = (c_t)$ of functions $c_t : \mathcal{Y}^{t-1} \to [0, \infty)$ is *predictable*: the value at step t depends only on the history before t.

Theorem 3.2 (Rényi filter for zCDP)**.** *Let P and Q be adaptive processes of horizon M on $\mathcal{Y}$, and let c be predictable. Suppose*

(i) $D_a\big(P_t(\cdot \mid h) \| Q_t(\cdot \mid h)\big) \le a\, c_t(h)$ *for all* $t \le M$*, all* $h \in \mathcal{Y}^{t-1}$ *and all* $a > 1$*;*

(ii) $\sum_{t \le M} c_t(y_{<t}) \le \rho^*$ *for every* $y \in \mathcal{Y}^M$*.*

Then $D_a(P\|Q) \le a\rho^*$ *for every* $a > 1$*.*

Proof. Fix $a > 1$ and abbreviate $P_t(y_t \mid y_{<t})$ to P_t, and similarly Q_t and c_t. By (i), $P_t(\cdot \mid h) \ll Q_t(\cdot \mid h)$ for every h, so the products below contain no term $\phi_a(p, 0)$ with $p > 0$. For $0 \le t \le M$ put

$$S_t = \sum_{y_{\le t} \in \mathcal{Y}^t} \prod_{j \le t} \phi_a(P_j, Q_j)\, e^{-a(a-1)c_j}, \qquad S_0 = 1.$$

Sum first over y_t, with $y_{<t}$ fixed. Since c_t depends only on $y_{<t}$,

$$S_t = \sum_{y_{<t}} \Big(\prod_{j<t} \phi_a(P_j, Q_j) e^{-a(a-1)c_j} \Big)\, e^{-a(a-1)c_t}\, \mathsf{M}_a\big(P_t(\cdot \mid y_{<t}) \| Q_t(\cdot \mid y_{<t})\big) \le S_{t-1},$$

because $\mathsf{M}_a(P_t\|Q_t) = e^{(a-1)D_a(P_t\|Q_t)} \le e^{a(a-1)c_t}$ by (i). Hence $S_M \le 1$. The function ϕ_a is multiplicative, so $\phi_a(P(y), Q(y)) = \prod_t \phi_a(P_t, Q_t)$, and by (ii)

$$\mathsf{M}_a(P\|Q) = \sum_y \prod_t \phi_a(P_t, Q_t) \le e^{a(a-1)\rho^*} \sum_y \prod_t \phi_a(P_t, Q_t) e^{-a(a-1)c_t} = e^{a(a-1)\rho^*} S_M \le e^{a(a-1)\rho^*}.$$

Take $\frac{1}{a-1} \ln$. □

Remark 3.3 (What the proof is, and attribution)**.** S_t is $\mathbb{E}_Q[Z_t]$ for $Z_t = \prod_{j \le t} (P_j/Q_j)^a e^{-a(a-1)c_j}$, and the inequality $S_t \le S_{t-1}$ is the Q-supermartingale property of Z_t. This is the standard argument for Rényi filters. Condition (ii) holds for every history, which is exactly what a filter guarantees. A filter stops the process, or makes it output $\perp$ at cost 0, as soon as the next cost would exceed the remaining budget. That decision depends only on the accumulated public costs, so it is a stopping time.

The horizon is finite because the coalition is PPT (Assumption 1.9(A8)), so no optional-stopping theorem is needed. An unbounded horizon would need a passage to the limit $M \to \infty$, which is not needed here.

The cost $a\, c_t$ is linear in a with one predictable c_t for all orders, so one stopping rule serves every order and the conclusion is zCDP. Privacy filters were introduced by Rogers, Roth, Ullman and Vadhan [6]. The Rényi version, with adaptively chosen per-step Rényi costs and a fixed total, is due to Feldman and Zrnic [3], who prove a more general individual form [*citation detail unverified:* their main theorem is quoted as Theorem 1.1 in the arXiv version; numbering in the NeurIPS version not checked]. Whitehouse et al. [7] give fully adaptive filters for approximate DP and approximate zCDP. Nothing below depends on the exact statements of these papers.

Corollary 3.4 (Steps of the protocol)**.** *The following steps satisfy condition (i) of Theorem 3.2 for the processes P (dataset $\mathfrak{D}$) and Q (dataset $\mathfrak{D}'$).*

(a) A public step*, whose kernel is the same under $\mathfrak{D}$ and $\mathfrak{D}'$, with $c_t = 0$. This includes deterministic steps (a function of h), refusals, and draws of coins that do not depend on $(\mathbf{x}_j, y_j)_j$.*

(b) A lattice Gaussian step*, $P_t(\cdot \mid h) = \mathrm{law}(F(h) + \eta)$ and $Q_t(\cdot \mid h) = \mathrm{law}(F'(h) + \eta)$, with $\eta \sim D_{\mathbb{Z}^{n(h)},\sigma(h)}$, $F(h) - F'(h) \in \mathbb{Z}^{n(h)}$ and $\|F(h) - F'(h)\|_2 \le \Delta(h)$, with $c_t(h) = \Delta(h)^2/(2\sigma(h)^2)$. The dimension, the centre, the bound and σ may all depend on the history.*

Deterministic post-processing of the transcript preserves the conclusion of Theorem 3.2 (Lemma 2.2).

Proof. (a) is immediate. For (b), D_a is invariant under the bijection $y \mapsto y - F'(h)$ of $\mathbb{Z}^{n(h)}$; apply Lemma 2.4 with $\mathbf{e} = F(h) - F'(h)$. □

Lemma 3.5 (Statistical hybrid lemma with predictable distances)**.** *Let P and Q be adaptive processes of horizon M on $\mathcal{Y}$, and let β be predictable with $\mathrm{SD}(P_t(\cdot \mid h), Q_t(\cdot \mid h)) \le \beta_t(h)$ for all t and h. Then*

$$\mathrm{SD}(P, Q) \le \mathbb{E}_{Y\sim P}\Big[\sum_{t\le M} \beta_t(Y_{<t})\Big].$$

Proof. For $0 \le t \le M$, let $H^{(t)}$ be the process with kernels P_j for $j \le t$ and Q_j for $j > t$, so that $H^{(M)} = P$ and $H^{(0)} = Q$. The processes $H^{(t)}$ and $H^{(t-1)}$ have the same law of $Y_{<t}$, namely P's, and the same kernels after step t; they differ only in the kernel at step t. By Lemma 2.2 (a mixture over $y_{<t}$ with common weights, followed by the common continuation kernel),

$$\mathrm{SD}(H^{(t)}, H^{(t-1)}) \le \sum_{y_{<t}} P(y_{<t})\, \mathrm{SD}\big(P_t(\cdot \mid y_{<t}), Q_t(\cdot \mid y_{<t})\big) \le \mathbb{E}_P[\beta_t(Y_{<t})].$$

Sum over t and use the triangle inequality. □

The point of the lemma is that every term is an expectation under the *same* law P, because each hybrid keeps P's kernels on the prefix. Applied with P the real experiment, it gives a slack that is one expectation under the real experiment.

Lemma 3.6 (Adaptive composition of hockey-stick closeness)**.** *Let P and Q be adaptive processes of horizon M on $\mathcal{Y}$. Let $\varepsilon_t \ge 0$ and $\delta_t \ge 0$ be predictable with $\mathsf{E}_{\varepsilon_t(h)}\big(P_t(\cdot \mid h) \| Q_t(\cdot \mid h)\big) \le \delta_t(h)$ for all t and h, and suppose $\sum_t \varepsilon_t(y_{<t}) \le E$ for every y. Then*

$$\mathsf{E}_E(P\|Q) \le \mathbb{E}_{Y\sim P}\Big[\sum_{t\le M} \delta_t(Y_{<t})\Big].$$

Proof. Put $P'_t(y_t \mid h) = \min\{P_t(y_t \mid h), e^{\varepsilon_t(h)} Q_t(y_t \mid h)\}$. Then $\sum_{y_t}(P_t - P'_t)(y_t \mid h) = \mathsf{E}_{\varepsilon_t(h)}(P_t\|Q_t) \le \delta_t(h)$. Let $P'(y) = \prod_t P'_t(y_t \mid y_{<t})$. Then $P'(y) \le e^{\sum_t \varepsilon_t} Q(y) \le e^E Q(y)$. Telescoping, with $0 \le P'_j \le P_j$,

$$P(y) - P'(y) = \sum_{t\le M} \Big(\prod_{j<t} P'_j\Big)(P_t - P'_t)\Big(\prod_{j>t} P_j\Big) \ \ge 0.$$

Summing over y, the factor $\prod_{j>t} P_j$ sums to 1 over $y_{>t}$. Bounding $P'_j \le P_j$ for $j < t$ gives

$$\sum_y (P - P')(y) \le \sum_t \sum_{y_{<t}} P(y_{<t})\, \delta_t(y_{<t}) = \mathbb{E}_P\Big[\sum_t \delta_t\Big].$$

For any $\varphi : \mathcal{Y} \to [0,1]$, $\mathbb{E}_P\varphi \le \mathbb{E}_{P'}\varphi + \sum_y (P - P')(y) \le e^E \mathbb{E}_Q\varphi + \mathbb{E}_P[\sum_t \delta_t]$. □

The ε_t must sum to at most E pointwise, because they enter as a factor e^E. The δ_t enter only through their expectation under P.

Lemma 3.7 (Adaptively chosen parameter sets)**.** *For each p in a finite set Π, let G_p be a decision problem: a challenger draws $b \in \{0,1\}$ and hands out a challenge X_p^b. Let $\mathcal{B}$ be an adversary that first computes, possibly at random, a choice $p \in \Pi$, then receives X_p^b, and outputs a bit. Let $\mathcal{B}_p$ run $\mathcal{B}$ until it chooses some p', output 0 if $p' \neq p$, and otherwise forward its own challenge to $\mathcal{B}$ and output $\mathcal{B}$'s bit. Then $\mathrm{Adv}(\mathcal{B}) \leq \sum_{p\in\Pi} \mathrm{Adv}_p(\mathcal{B}_p)$, and each $\mathcal{B}_p$ runs in the time of $\mathcal{B}$.*

Proof. $\mathcal{B}$'s choice is made before it sees the challenge, so it is independent of b. Hence $\Pr[\mathcal{B}=1 \mid b] = \sum_p \Pr[\mathcal{B} \text{ chooses } p]\Pr[\mathcal{B}=1 \mid p,b] = \sum_p \Pr[\mathcal{B}_p = 1 \mid b]$. Take the difference over b and apply the triangle inequality. □

4 Composition across runs

4.1 The ideal experiment

Definition 4.1 (The ideal experiment $\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D})$)**.** $\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D})$ is $\mathsf{Real}_{\mathcal{A}}(\mathfrak{D})$ with the same coalition, CALIBRATE, ledger, refusals and honest machines R_k, with one change: in every admitted run k, O is replaced by the ideal party $\tilde{\mathsf{O}}_k$. It sends its messages at the same points of the protocol as O.

(I1) It sends the public parameters, as O does.

(I2) It replaces O's key-generation shares (the public-key share and the rotation-key shares) by independent uniform elements of the corresponding rings.

(I3) It runs O's Link phase (OPRF, labeled PSI and SANITIZEBFV) honestly on the identifiers of $D_{P_k}(\mathfrak{D};\mathbf{d}_k)$, with fresh coins (OPRF key, τ_k, $\mathbf{d}_k$, sanitization coins). This uses only $\mathrm{fix}(\mathfrak{D})$ and P_k.

(I4) It receives R_k's uploads and ignores them; it performs no aggregation.

(I5) At each release s that $\mathcal{A}$ requests, it draws $c^* \leftarrow \mathrm{U}(\mathcal{R}_{q_0})$ and $\eta \sim D_{\mathbb{Z}^N,\sigma_{c,s}}$, sets

$$w_{k,s} = f_{k,s}(\mathfrak{D}) + \eta, \qquad f_{k,s}(\mathfrak{D}) = f_s\big(D_{P_k}(\mathfrak{D};\mathbf{d}_k);\pi_k,\chi_k\big) \in \mathbb{Z}^N,$$

and sends $(c^*,\ w_{k,s} - c^* s_R \bmod q_0)$. Here f_s is the centre of [CP], Definition 3.7 (`def:centre`), π_k is the arrangement defined by τ_k, and s_R is read from $\omega_{R,k}$.

$\tilde{\mathsf{O}}_k$ uses $(\mathbf{x}_j, y_j)_j$ only through the values $w_{k,s}$. Run by run, $\tilde{\mathsf{O}}_k$ is the simulator of [CP], Theorem 4.11, applied to that theorem's mechanism $\mathcal{M}_k = (\omega_{R,k}, \Lambda_k, (f_{k,s} + \eta_{k,s})_s)$.

Theorem 4.2 (Composition across runs; replaces Proposition `prop:compose`)**.** *Assume Assumption 1.9, and let $\mathcal{A}$ be a PPT coalition. Let $n_{\mathrm{rel}} \leq n_{\max}$ for every $P \in \mathcal{P}_\lambda$. Let $t_{\mathcal{B}}(\lambda)$ be the polynomial running time of the reductions in Step 3 of the proof. Put*

$$\nu_{\mathcal{A}}(\lambda) = t_{\mathcal{A}}(\lambda) \sum_{p\in\Pi_\lambda} \big(n_{\max}\mathrm{Adv}_{1,p}(\lambda, t_{\mathcal{B}}(\lambda)) + \mathrm{Adv}_{2,p}(\lambda, t_{\mathcal{B}}(\lambda))\big),$$

which is negligible.

(a) Simulation. *For every $\mathfrak{D}$,*

$$\big|\Pr[\mathsf{Real}_{\mathcal{A}}(\mathfrak{D})=1] - \Pr[\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D})=1]\big| \leq \Gamma_{\mathcal{A}}(\mathfrak{D}) + \nu_{\mathcal{A}}(\lambda),$$

where $\Gamma_{\mathcal{A}}(\mathfrak{D}) = \mathbb{E}_{\mathsf{Real}_{\mathcal{A}}(\mathfrak{D})}\big[\sum_{k \text{ admitted}} \gamma_k\big]$.

(b) Ideal composition. *Suppose that in every execution the admitted runs satisfy* $\sum_k \rho_k \le \rho^*$*, for a constant* ρ^**. Then* $D_a\big(\mathrm{view}_{\mathcal{A}}^{\mathsf{Ideal}}(\mathfrak{D}) \| \mathrm{view}_{\mathcal{A}}^{\mathsf{Ideal}}(\mathfrak{D}')\big) \le a\rho^*$ *for all* $a > 1$ *and* $\mathfrak{D} \sim_{\mathrm{b}} \mathfrak{D}'$*. Hence, for every* $\varepsilon \ge 0$*,*

$$\Pr[\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D}) = 1] \le e^{\varepsilon} \Pr[\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D}') = 1] + \delta(\rho^*, \varepsilon).$$

(c) Accounting. *Suppose in addition that* $\sum_k \gamma_k \le \Gamma^*$ *over the admitted runs in every execution. Then for every* $\varepsilon \ge 0$ *and* $\mathfrak{D} \sim_{\mathrm{b}} \mathfrak{D}'$*,*

$$\Pr[\mathsf{Real}_{\mathcal{A}}(\mathfrak{D}) = 1] \le e^{\varepsilon} \Pr[\mathsf{Real}_{\mathcal{A}}(\mathfrak{D}') = 1] + \delta(\rho^*, \varepsilon) + (1 + e^{\varepsilon})\Gamma^* + (1 + e^{\varepsilon})\nu_{\mathcal{A}}(\lambda).$$

(d) Rule (F), any adaptive coalition. *Under the filter rule, (b) and (c) hold with* $(\rho^*, \Gamma^*) = (\rho_L, \Gamma_L)$*, and the protocol is* $(\varepsilon_L, \delta_L)$*-IND-CDP for* $\sim_{\mathrm{b}}$*. The* δ *budget is used exactly:* $\delta(\rho_L, \varepsilon_L) + (1 + e^{\varepsilon_L})\Gamma_L = \delta_L$*.*

(e) Rule (N), requests fixed in advance. *Suppose the ledger uses the non-adaptive rule and* $\mathcal{A}$*'s requests are a fixed list* $(P_k, \chi_k)_{k \le K}$*, submitted in this order and chosen independently of* $\omega_{\mathcal{A}}$ *and of every message.* $\mathcal{A}$ *may still schedule, abort and stop adaptively. Let* $(\rho^{\mathrm{fin}}, \Gamma^{\mathrm{fin}})$ *be the ledger's final state when the whole list is submitted. Then (b) and (c) hold with* $(\rho^*, \Gamma^*) = (\rho^{\mathrm{fin}}, \Gamma^{\mathrm{fin}})$*,* $\delta(\rho^{\mathrm{fin}}, \varepsilon_L) + (1 + e^{\varepsilon_L})\Gamma^{\mathrm{fin}} \le \delta_L$*, and the protocol is* $(\varepsilon_L, \delta_L)$*-IND-CDP for* $\sim_{\mathrm{b}}$ *against such coalitions.*

Parts (a) and (b) together say that the joint view is SIM-CDP with an *interactive* mechanism: $\mathcal{A}$ interacting with the parties $\tilde{\mathsf{O}}_k$ is a post-processing of an adaptive ρ^*-zCDP mechanism, with statistical slack $\Gamma_{\mathcal{A}}(\mathfrak{D}) \le \Gamma^*$.

4.2 Proof of Theorem 4.2

Step 0: the experiments as adaptive processes. Make three changes of presentation, none of which changes the law of $\mathrm{view}_{\mathcal{A}}$.

(i) *Draw a run's coins at admission.* When run k is admitted, draw $\omega_{R,k}$ and all of $\omega_{O,k}$ except the noises $\eta_{k,s}$. Each $\eta_{k,s}$ is drawn at release s. By (A3) these coins are independent of everything earlier, and drawing them earlier changes nothing.

(ii) *Rewrite the share* as in hybrid H1 of [CP], Theorem 4.11. O's share $h_O = c'_0 + c'_1 s_O + \eta$ equals $v - c'_1 s_R$ with $v = c'_0 + c'_1 s + \eta$, and by [CP], Proposition 3.8 (`prop:decomp`),

$$v \equiv f_{k,s}(\mathfrak{D}) + e_{k,s} + \eta \pmod{q_0}, \qquad e_{k,s} = e_{\mathrm{tot},s} \text{ of run } k \in \mathbb{Z}^N.$$

(iii) *Honest completions.* The honest machine R_k depends only on $(P_k, \chi_k, \omega_{R,k})$ and on O's messages in run k, and those depend only on $(\mathfrak{D}, P_k, \omega_{O,k})$ and R_k's messages. So each admitted run has an *honest completion*: the transcript it would have if $\mathcal{A}$ never aborted it. The honest completion is a function of $(\mathfrak{D}, P_k, \chi_k, \omega_{O,k}, \omega_{R,k})$ and does not depend on other runs or on the scheduling. Let $e_{k,s}$ be the error of release s in the honest completion, for every $s \le n_{\mathrm{rel},k}$. When release s is made, it is the error of the actual release. It does not involve any η, because M_1, M_2 (and M_y) are computed before any release. By (A5) and (A3), for every history up to the admission of run k,

$$\Pr\big[\|e_{k,s}\|_2 > B_{\mathrm{tot},k,s} \;\big|\; \text{history up to the admission of run } k\big] \le \delta_{B,k}. \tag{1}$$

Every experiment below is then an adaptive process on a countable alphabet. Its steps are the draw of $\omega_{\mathcal{A}}$, the actions of $\mathcal{A}$ (deterministic given the history), O's answers, the draws of coins at admission, and the messages of open runs. The horizon is $M \le \mathrm{poly}(\lambda)$, by (A8). $\mathcal{A}$'s output is a function of the process. The ledger's state is a function of the history, and every refusal is a deterministic function of the history.

Step 1: one statistical hybrid for all runs. Let $\mathsf{Hyb}^{\mathrm{st}}_{\mathcal{A}}(\mathfrak{D})$ be $\mathsf{Real}_{\mathcal{A}}(\mathfrak{D})$ with v replaced by $f_{k,s}(\mathfrak{D})+\eta$ in every release that is made. The share sent is then $f_{k,s}(\mathfrak{D})+\eta-c'_1 s_R \bmod q_0$. This is H2 of [CP], Theorem 4.11, done in every run at once.

The two processes have the same kernels at every step except the release steps. At a release step (k,s) with history h, c'_1, s_R, $f_{k,s}(\mathfrak{D})$ and $e_{k,s}$ are functions of h, and η is fresh. The real kernel is the law of $(c'_1,\ f_{k,s}+e_{k,s}+\eta-c'_1 s_R \bmod q_0)$, and the hybrid kernel is the same with $e_{k,s}$ removed. By Lemma 2.2, which covers the shift and the reduction modulo q_0, and Lemma 2.5, their distance is at most

$$\beta_{(k,s)}(h)=\mathrm{SD}(e_{k,s}+\eta,\eta)\le \mathbf{1}\{\|e_{k,s}\|>B_{\mathrm{tot},k,s}\}+\gamma'(\kappa_k).$$

By Lemma 3.5, with the real experiment as P,

$$\begin{aligned}\mathrm{SD}\big(\mathsf{Real}_{\mathcal{A}}(\mathfrak{D}),\mathsf{Hyb}^{\mathrm{st}}_{\mathcal{A}}(\mathfrak{D})\big)&\le \mathbb{E}\Big[\sum_{(k,s)\ \mathrm{made}}\mathbf{1}\{\|e_{k,s}\|>B_{\mathrm{tot},k,s}\}\Big]+\mathbb{E}\Big[\sum_{(k,s)\ \mathrm{made}}\gamma'(\kappa_k)\Big]\\&\le\sum_k\sum_{s\le n_{\mathrm{rel},k}}\mathbb{E}\big[\mathbf{1}\{k\ \mathrm{adm}\}\,\mathbf{1}\{\|e_{k,s}\|>B_{\mathrm{tot},k,s}\}\big]+\mathbb{E}\Big[\sum_{k\ \mathrm{adm}}n_{\mathrm{rel},k}\gamma'(\kappa_k)\Big]\\&\le\mathbb{E}\Big[\sum_{k\ \mathrm{adm}}n_{\mathrm{rel},k}\big(\delta_{B,k}+\gamma'(\kappa_k)\big)\Big]=\Gamma_{\mathcal{A}}(\mathfrak{D}).\end{aligned}$$

Expectations are under $\mathsf{Real}_{\mathcal{A}}(\mathfrak{D})$. The second line uses $\{(k,s)\text{ made}\}\subseteq\{k\text{ adm}\}$. The third line conditions on the history up to the admission of run k, of which $\{k\text{ adm}\}$ is a function, and applies (1).

Why every planned release is charged δ_B. The only usable bound on $\mathbf{1}\{\text{made}\}\mathbf{1}\{\text{bad}\}$ is $\mathbf{1}\{\text{adm}\}\mathbf{1}\{\text{bad}\}$. $\mathcal{A}$'s decision to request release s may depend on its view, which contains s_R and R's encryption errors. Both enter $e_{k,s}$, through T_1, T_2 and T_6 of [CP], Proposition 3.8. So $\Pr[\text{made}\wedge\text{bad}]$ need not be at most $\delta_B\Pr[\text{made}]$. The term $\gamma'(\kappa_k)$, by contrast, bounds the distance for every value of $e_{k,s}$ of norm at most B_{tot}, so it is charged only for releases made.

Expectation versus pointwise bound. Step 1 needs only the expectation $\Gamma_{\mathcal{A}}(\mathfrak{D})$ under the real experiment, for $\mathfrak{D}$ and for $\mathfrak{D}'$. Lemma 3.5 makes it a single expectation; doing the smudging steps one at a time, interleaved with computational steps, would give expectations under different hybrids. The pointwise bound $\sum_{\mathrm{adm}}\gamma_k\le\Gamma^*$ that the ledger enforces implies $\Gamma_{\mathcal{A}}(\mathfrak{D})\le\Gamma^*$ for every $\mathfrak{D}$ and every $\mathcal{A}$. The ledger cannot compute expectations, which depend on $\mathfrak{D}$ and $\mathcal{A}$, so the pointwise rule is the implementable one. The zCDP part (Step 4) genuinely needs the pointwise bound on $\sum\rho_k$.

Step 2: the order of hybrids. Order the admitted runs by admission time, $k=1,2,\dots,T$ with $T\le t_{\mathcal{A}}(\lambda)$. For $1\le k\le t_{\mathcal{A}}(\lambda)$ and $0\le j\le n_{\max}+1$, let $\mathsf{Hyb}_{k,j}$ be the following experiment.

- Admitted runs $k'<k$ are played by the ideal parties $\tilde{\mathsf{O}}_{k'}$.
- Run k is played as in $\mathsf{Hyb}^{\mathrm{st}}$, except that $c'_{1,k,s}$ is uniform in $\mathcal{R}_{q_0}$ for $s\le\min(j,n_{\mathrm{rel},k})$, and, if $j=n_{\max}+1$, O's key-generation shares are uniform.
- Runs $k'>k$ are played as in $\mathsf{Hyb}^{\mathrm{st}}$.

Then $\mathsf{Hyb}_{1,0}=\mathsf{Hyb}^{\mathrm{st}}$, $\mathsf{Hyb}_{k,n_{\max}+1}=\mathsf{Hyb}_{k+1,0}$, and $\mathsf{Hyb}_{t_{\mathcal{A}},n_{\max}+1}=\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D})$.

The middle identity holds because, in $\mathsf{Hyb}_{k,n_{\max}+1}$, run k is already ideal. Every share sent is $f_{k,s}+\eta-c^*s_R$ with c^* uniform, and the key shares are uniform. The Link phase is O's honest code. The ciphertexts M_s no longer enter the view, so neither do the conversion, the aggregation and the key switching that produce them. This is exactly $\tilde{\mathsf{O}}_k$. If fewer than k runs are admitted, $\mathsf{Hyb}_{k,j}$ does not depend on j.

Step 3: the computational steps. *Re-randomizer steps,* $1 \le j \le n_{\max}$. The reduction $\mathcal{B}_{k,j}$ plays the adaptive-parameter game of Lemma 3.7 for [CP], Assumption 4.9. It has $\mathfrak{D}$ as non-uniform advice. It draws $\omega_{\mathcal{A}}$ and runs $\mathcal{A}$, playing every other party and the ledger itself, exactly as in $\mathsf{Hyb}_{k,j-1}$.

- *Runs admitted before run k* are played by $\tilde{\mathsf{O}}_{k'}$. $\mathcal{B}_{k,j}$ can do this: the Link phase needs $\mathrm{fix}(\mathfrak{D})$; a release needs $f_{k',s}(\mathfrak{D}) + \eta$, computable in polynomial time by (A7); the key shares and c^* are uniform.
- *Run k.* When run k is admitted with parameter set p_k, $\mathcal{B}_{k,j}$ announces p_k and receives a challenge (a^*, b^*) over $\mathcal{R}_{q_0}$. It draws all of run k's coins itself, with two exceptions. First, the public-key CRS a is a uniform lift of a^*: $a \equiv a^* \pmod{q_0}$, uniform modulo the full modulus. Second, $(u_{k,j}, e_{1,k,j})$ is not drawn.
  - It sets $c'_{1,k,s}$ uniform for $s < j$.
  - It sets $c'_{1,k,j} = c_{1,k,j} + b^* \bmod q_0$. Here $c_{1,k,j}$ is computable from the public randomness and the data, by [CP], Lemma 3.6 (`lem:c1indep`).
  - It sets $c'_{1,k,s} = c_{1,k,s} + a u_{k,s} + e_{1,k,s}$ for $s > j$, which it can compute since it knows a.

  Each release that is made sends $(c'_{1,k,s},\ f_{k,s}(\mathfrak{D}) + \eta - c'_{1,k,s} s_R \bmod q_0)$. In $\mathsf{Hyb}^{\mathrm{st}}$ nothing else in run k uses $(u_{k,j}, e_{1,k,j})$: c'_0 is never sent, and it is no longer used, because v no longer contains $e_{k,j}$, whose term T_6 was the only other occurrence. If $j > n_{\mathrm{rel},k}$, $\mathcal{B}_{k,j}$ ignores the challenge.
- *Runs admitted after run k* are played as in $\mathsf{Hyb}^{\mathrm{st}}$. $\mathcal{B}_{k,j}$ draws their keys and coins and computes everything from $\mathfrak{D}$.
- *Public computations.* $\mathcal{B}_{k,j}$ computes Calibrate, the ledger and every refusal from public data (A7). It obtains the adaptively chosen requests by running $\mathcal{A}$ on the view it has generated, which has exactly the law of $\mathcal{A}$'s view in the hybrid. This is why the adaptively chosen parameters are efficiently computable by the reduction: they are the outputs of a PPT machine on an input the reduction has produced.
- $\mathcal{B}_{k,j}$ outputs $\mathcal{A}$'s output.

If $b^* = a^* u + e_1$ with fresh $u \leftarrow \chi_{\mathrm{key}}$ and $e_1 \leftarrow \chi_{\mathrm{err}}$, the simulation has the law of $\mathsf{Hyb}_{k,j-1}$. If b^* is uniform, $c'_{1,k,j}$ is uniform and the simulation has the law of $\mathsf{Hyb}_{k,j}$. If run k is never admitted, or its release j is never made, the two hybrids coincide, and $\mathcal{B}_{k,j}$'s output does not depend on the challenge.

$\mathcal{B}_{k,j}$ runs in time $t_{\mathcal{B}}(\lambda) = O(t_{\mathcal{A}}(\lambda)) + t_{\mathcal{A}}(\lambda) \cdot \mathrm{poly}(\lambda)$, because it plays at most $t_{\mathcal{A}}$ runs of polynomial cost. By Lemma 3.7,

$$\left|\Pr[\mathsf{Hyb}_{k,j-1} = 1] - \Pr[\mathsf{Hyb}_{k,j} = 1]\right| \le \sum_{p \in \Pi_\lambda} \mathrm{Adv}_{1,p}(\lambda, t_{\mathcal{B}}(\lambda)).$$

Key-share step, $j = n_{\max} + 1$. $\mathcal{B}_{k,n_{\max}+1}$ plays the game for [CP], Assumption 4.10, in the same way. It embeds the challenge as O's key-generation shares in run k, uses the challenge's a and $a_{g,j}$ as run k's common random strings, and computes the joint keys from the challenge and R_k's shares; it knows $\omega_{R,k}$. In $\mathsf{Hyb}_{k,n_{\max}}$ the view depends on run k's $(s_O, e_O^{\mathrm{pk}}, e_O^{\mathrm{rk}})$ only through these shares:

- every release share is $f + \eta - c'_1 s_R$ with c'_1 uniform;
- c'_0 is unused;
- the conversion $c_0 - a s_O$, the aggregation and the key switching produce only the ciphertexts M_s, which no longer enter the view.

So $\mathcal{B}_{k,n_{\max}+1}$ never needs s_O. A real challenge gives $\mathsf{Hyb}_{k,n_{\max}}$, and a uniform one gives $\mathsf{Hyb}_{k,n_{\max}+1}$. The loss is at most $\sum_p \mathrm{Adv}_{2,p}(\lambda, t_{\mathcal{B}}(\lambda))$.

Number of runs and total loss. There are at most $t_{\mathcal{A}}(\lambda)(n_{\max} + 1)$ steps, by (A8), so

$$\left|\Pr[\mathsf{Hyb}^{\mathrm{st}} = 1] - \Pr[\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D}) = 1]\right| \le \nu_{\mathcal{A}}(\lambda).$$

$\nu_{\mathcal{A}}$ is negligible by (A6), since $t_{\mathcal{A}}$, $n_{\max}$ and $|\Pi_\lambda|$ are polynomial and $t_{\mathcal{B}}$ is polynomial. Together with Step 1 this proves (a).

Step 4: the ideal experiment is ρ^*-zCDP. Write $\mathsf{Ideal}_{\mathcal{A}}(\mathfrak{D})$ as an adaptive process Y with three kinds of random step.

(a) The draw of $\omega_{\mathcal{A}}$.

(b) For each admitted run k, one step that draws r_k. It consists of $\omega_{R,k}$, $\tilde{\mathsf{O}}_k$'s Link coins (OPRF key, τ_k, $\mathbf{d}_k$, sanitization coins), its uniform key shares and $(c^*_{k,s})_{s\le n_{\mathrm{rel},k}}$.

(c) For each release (k,s) that $\mathcal{A}$ requests, one step that outputs $w_{k,s}\in\mathbb{Z}^N$.

Everything else is a deterministic function of the history: $\mathcal{A}$'s actions, the refusals, the ledger state, the PSI messages, and the release messages $(c^*, w - c^* s_R \bmod q_0)$. The transcript Y contains τ_k even though $\mathcal{A}$ does not see it. $\mathcal{A}$'s view is a deterministic function of Y.

Compare Y under $\mathfrak{D}$ and $\mathfrak{D}'$, with $\mathfrak{D}\sim_{\mathrm{b}}\mathfrak{D}'$. Steps (a) and (b) have kernels that do not depend on $(\mathbf{x}_j,y_j)_j$, so they are public steps with cost 0. A step (c) with history h is a lattice Gaussian step in the sense of Corollary 3.4(b), with $F(h)=f_{k,s}(\mathfrak{D})$ and $F'(h)=f_{k,s}(\mathfrak{D}')$. Its dimension N and $\sigma_{c,k,s}$ are given by P_k, and π_k and χ_k are coordinates of h. By Lemma 1.3, the two scope registers coincide or are adjacent, so by (A5) $\|F(h)-F'(h)\|_2\le\widehat{\Delta}_{\mathrm{enc},k,s}$ for every arrangement and every researcher input. Both centres are integral, so the step's cost is $c(h)=\widehat{\Delta}^2_{\mathrm{enc},k,s}/(2\sigma^2_{c,k,s})\le\rho_{k,s}$.

For every history, $\sum_t c_t\le\sum_{k\ \mathrm{adm}}\sum_s\rho_{k,s}=\sum_{k\ \mathrm{adm}}\rho_k\le\rho^*$. This holds for every history, not only almost surely, because the ledger state is a function of the history and enforces the bound. Theorem 3.2 gives $D_a(Y(\mathfrak{D})\|Y(\mathfrak{D}'))\le a\rho^*$, and Lemma 2.2 transfers this to $\mathcal{A}$'s view and output. Lemma 2.3 gives the hockey-stick bound. The roles of $\mathfrak{D}$ and $\mathfrak{D}'$ are symmetric. This proves (b).

Refusals and the gate. The refusals are deterministic functions of the history, that is, public steps of cost 0. The stage-2 gate is part of $\mathcal{A}$'s computation. The theorem bounds the unconditional joint law of the view, of which the refusal pattern and the gate's outcome are coordinates. We never condition on them. A guarantee conditional on an event such as "run k was admitted" or "the gate passed" does not follow and is not claimed.

Step 5: conversion and the budget. Write $\Gamma=\Gamma_{\mathcal{A}}(\mathfrak{D})$, $\Gamma'=\Gamma_{\mathcal{A}}(\mathfrak{D}')$ and $\nu=\nu_{\mathcal{A}}(\lambda)$. By (a) for $\mathfrak{D}$, then (b), then (a) for $\mathfrak{D}'$,

$$\begin{aligned}\Pr[\mathsf{Real}(\mathfrak{D})=1]&\le\Pr[\mathsf{Ideal}(\mathfrak{D})=1]+\Gamma+\nu\le e^{\varepsilon}\Pr[\mathsf{Ideal}(\mathfrak{D}')=1]+\delta(\rho^*,\varepsilon)+\Gamma+\nu\\&\le e^{\varepsilon}\big(\Pr[\mathsf{Real}(\mathfrak{D}')=1]+\Gamma'+\nu\big)+\delta(\rho^*,\varepsilon)+\Gamma+\nu.\end{aligned}$$

With $\Gamma,\Gamma'\le\Gamma^*$ this is (c). This is [P], Lemma `lem:sim2ind`, with the factor $1+e^{\varepsilon}$ at the final ε.

(d) Under rule (F), every admission keeps both sums within their caps, and the sums are the ledger's state. So in every execution $\sum_{\mathrm{adm}}\rho_k\le\rho_L$ and $\sum_{\mathrm{adm}}\gamma_k\le\Gamma_L$. Take $\varepsilon=\varepsilon_L$. By the definition of ρ_L and the continuity and strict monotonicity of $\delta(\cdot,\varepsilon_L)$ (Lemma 2.3), $\delta(\rho_L,\varepsilon_L)=(1-\zeta_L)\delta_L$. Also $(1+e^{\varepsilon_L})\Gamma_L=\zeta_L\delta_L$, so the sum is δ_L exactly. The term $(1+e^{\varepsilon_L})\nu$ is negligible because ε_L is a constant. At $(10,10^{-6},0.1)$ the sum evaluates to $9.99999999999\cdot10^{-7}$ (script `scripts/kappa_numbers.py`).

(e) The requests are fixed, so their costs are fixed. The full cost is charged at admission, so the ledger's decisions depend only on the cost sequence. The set $\mathcal{K}$ of runs admitted when the whole list is submitted is therefore deterministic. In every execution the admitted runs form a subset of $\mathcal{K}$, since $\mathcal{A}$ may stop early, so the sums are at most $(\rho^{\mathrm{fin}},\Gamma^{\mathrm{fin}})$. If $\mathcal{K}\neq\emptyset$, the test at the last admission in $\mathcal{K}$ reads exactly $\delta(\rho^{\mathrm{fin}},\varepsilon_L)+(1+e^{\varepsilon_L})\Gamma^{\mathrm{fin}}\le\delta_L$. If $\mathcal{K}=\emptyset$, no run is admitted and the view does not depend on $(\mathbf{x}_j,y_j)_j$ at all. □

4.3 Corner cases

Refused runs. A run refused by CALIBRATE or by the ledger produces one message, $\perp_{\mathrm{cal}}$ or $\perp_{\mathrm{led}}$. That message is a function of P_k and the ledger state, both public. No key is generated and no share is sent. In the proof such a run is a public step of cost 0 (Step 4). The refusal pattern can depend on earlier outputs, through $\mathcal{A}$'s requests, but it is part of the view, which the theorem covers unconditionally.

The stage-2 gate. A refused gate means that R withholds its estimate after the stage-2 release has been made and charged. This is post-processing of $\mathcal{A}$'s view and changes nothing. The paper's sentence "a refused stage-2 gate still spends the budget" is therefore correct and necessary.

Aborted runs. $\mathcal{A}$ may abort any open run at any time, adaptively. The proof handles this through the honest completion (Step 0(iii)): the ideal experiment releases only the prefix that $\mathcal{A}$ requests, and the costs of unmade releases are simply not used in Step 4. An abort by O is covered by (A1). Such an abort depends only on public and data-independent quantities, such as a PSI table overflow, which depends on the identifiers and the OPRF key. The accounting must keep the δ_B part of every planned release of an admitted run (Step 1 and Remark 4.3).

The PSI transcript. The identifiers are part of $\mathrm{fix}(\mathfrak{D})$, so the Link phase of every run is identical under $\mathfrak{D}$ and $\mathfrak{D}'$. The ideal experiment runs O's real Link code (I3), so no PSI security is used for this theorem. What the Link phase discloses $(I,\ \pi|_I,\ n)$ is outside the DP guarantee, and over many runs this disclosure accumulates: [P], Remark `rem:revealed`, and G4. The OPRF quota limits it as a matter of policy, not as part of this theorem.

Different researchers' cohorts. The cohorts are inputs χ_k of the coalition, chosen adaptively, possibly overlapping, and not protected. Step 4 compares the two datasets for a *fixed* history, which fixes χ_k. The sensitivity bound holds for every cohort in the clipped domain. So the cohorts may depend on earlier outputs in arbitrary ways. A row matched in many runs is charged in every run, through $\sum \rho_k$.

One base register, several scopes and models. By Lemma 1.3, base adjacency induces adjacency or equality of every run's scope register. This holds for every scope, model, column subset and public row-wise preprocessing. The ledger must therefore be shared across all scopes and models of a base register, as the paper specifies; a ledger per scope would not give Theorem 4.2. The guarantee is per base register and per row. If O holds several base registers that contain the same people, protecting a person across them needs one ledger over their union, which is outside this theorem.

Concurrency. Several runs may be open at once and interleaved arbitrarily (Definition 1.7(b)). The proof indexes hybrids by admission order and orders the statistical and filter steps by global time. Both are insensitive to interleaving, provided admission is atomic and charges the full cost.

The number of runs. By (A8) the coalition makes at most $t_{\mathcal{A}}(\lambda)$ requests, which bounds the number of computational hybrids. Under rule (F) the ledger gives a concrete bound as well. Every admitted run has $\gamma_k \geq n_{\mathrm{rel},k}(\delta_{B,k} + 2^{-97})$, so at most $\Gamma_L / \min_k \gamma_k$ runs are admitted. At $(10, 10^{-6}, 0.1)$ with $n_{\mathrm{rel}} = 2$ and $\delta_B = 2^{-50}$ this is 2555. In $\nu_{\mathcal{A}}$ one may therefore replace $t_{\mathcal{A}}$ by this

constant, which makes the computational loss independent of the coalition's running time. Refused requests involve no key and add nothing.

Not covered. The following are outside Theorem 4.2:

- a register that changes between runs;
- researchers who deviate within a run, for example by uploading noisier ciphertexts, which would void B_{tot};
- a dishonest or non-atomic ledger;
- side channels such as the time O takes to refuse.

Remark 4.3 (Which refunds are sound)**.** Suppose the ledger reserves the full cost at admission and refunds part of it when a run ends early. The proof uses only two bounds. Step 4 uses $\sum_{\mathrm{made}} \rho_{k,s} \le \rho^*$. Step 1 uses $\sum_{\mathrm{made}} \gamma'(\kappa_k) + \sum_{\mathrm{adm}} n_{\mathrm{rel},k}\delta_{B,k} \le \Gamma^*$ in expectation. Refunding $\rho_{k,s}$ and $\gamma'(\kappa_k)$ of the releases not made is therefore sound. Refunding their $\delta_{B,k}$ is not covered. The paper's state, "the sum of the costs of every release made", refunds all three. The simplest correct rule charges the whole run at admission. Since $\delta_B \le \gamma_s$, the difference is a fraction $\delta_B/\gamma_s \approx \delta_B 2^{\kappa}$ of the slack of unmade releases: about 1.5% at $\kappa = 44$.

Remark 4.4 (Rule (N) with adaptive requests)**.** With adaptively chosen requests, the final state $(\rho_\Sigma, \Gamma_\Sigma)$ of rule (N) is random. Each trajectory satisfies the rule's inequality, but the trade-off between ρ and Γ varies from one trajectory to another. Theorem 4.2(c) then applies only with $\rho^* = \sup \rho_\Sigma$ and $\Gamma^* = \sup \Gamma_\Sigma$ over executions. If some trajectories spend almost all of δ_L on ρ and others on Γ, this can approach $2\delta_L$. We do not claim that rule (N) fails adaptively, only that the proof does not cover it. The paper's prescription is consistent with this: use the filter unless every run is fixed in advance.

Remark 4.5 (Comparison with the paper's proof sketch)**.** The sketch's three parts (simulation run by run, ideal composition by a Rényi filter, conversion at ε_L) are those of the proof above, with four differences:

(i) the smudging steps of all runs are done together, first, so that the slack is a single expectation under Real (Lemma 3.5);
(ii) the slack counts δ_B for all planned releases of admitted runs;
(iii) the reductions need RLWE for adaptively chosen parameter sets and against non-uniform adversaries;
(iv) the Rényi filter is proved here (Theorem 3.2), so the citation of [3] becomes an attribution.

5 Part 2: the smudging step in hockey-stick divergence

5.1 The shifted lattice Gaussian

Lemma 5.1 (Hockey-stick smudging, explicit constants)**.** *Let $\eta \sim D_{\mathbb{Z}^N,\sigma}$ and $\mathbf{e} \in \mathbb{Z}^N$ with $\|\mathbf{e}\|_2 \le B$ and $\sigma \ge 2^{\kappa}B/\sqrt{2\pi}$. For $\delta' \in (0,1)$ let*

$$\varepsilon'(\kappa,\delta') = \min\{\varepsilon \ge 0 : \delta(\pi 4^{-\kappa}, \varepsilon) \le \delta'\} \le \pi 4^{-\kappa} + 2\sqrt{\pi \ln(1/\delta')}\, 2^{-\kappa}.$$

(a) $\mathbf{e} + \eta$ and η are $(\varepsilon'(\kappa,\delta'),\delta')$-indistinguishable.
(b) Side information. *Let $(Z,\mathbf{e})$ have any joint law with $\mathbf{e} \in \mathbb{Z}^N$, and let η be independent of $(Z,\mathbf{e})$. Then both $\mathsf{E}_{\varepsilon'}\big((Z,\mathbf{e}+\eta)\|(Z,\eta)\big)$ and $\mathsf{E}_{\varepsilon'}\big((Z,\eta)\|(Z,\mathbf{e}+\eta)\big)$ are at most $\delta' + \Pr[\|\mathbf{e}\| > B]$.*

No smoothing condition and no floor on σ are needed, and N is arbitrary. For $\delta' = 2^{-100}$, $\varepsilon' \le 29.52 \cdot 2^{-\kappa} + \pi 4^{-\kappa}$; at $\kappa = 40$ this is $2.7 \cdot 10^{-11}$, and the tight value is $2.0 \cdot 10^{-11}$.

Proof. (a) Let $\rho' = \|\mathbf{e}\|^2/(2\sigma^2) \le B^2 \cdot 2\pi/(2 \cdot 4^\kappa B^2) = \pi 4^{-\kappa}$. By Lemma 2.4, and the remark after it, $D_a(\mathbf{e}+\eta\|\eta) \le a\rho'$ and $D_a(\eta\|\mathbf{e}+\eta) \le a\rho'$ for all $a > 1$. By Lemma 2.3 and the monotonicity of $\delta(\cdot,\varepsilon)$, both hockey-stick divergences at $\varepsilon'(\kappa,\delta')$ are at most $\delta(\rho',\varepsilon') \le \delta(\pi 4^{-\kappa},\varepsilon') \le \delta'$. The closed-form bound on ε' is the simple form in Lemma 2.3.

(b) Let μ be the law of $(Z,\mathbf{e})$. Both laws are mixtures over $(z,\mathbf{e})$ with the common weights $\mu(z,\mathbf{e})$: $(Z,\mathbf{e}+\eta)$ mixes $\delta_z \otimes \mathrm{law}(\mathbf{e}+\eta)$, and (Z,η) mixes $\delta_z \otimes \mathrm{law}(\eta)$. By Lemma 2.2, each divergence is at most $\sum \mu(z,\mathbf{e})\mathsf{E}_{\varepsilon'}(\cdot\|\cdot)$. Each term is at most δ' when $\|\mathbf{e}\| \le B$, by (a), and at most 1 otherwise. □

Table 1: Exact check of Lemma 5.1 in one dimension, worst case $e = B$, $\sigma = 2^\kappa B/\sqrt{2\pi}$. "Exact ε" is the smallest ε whose exact hockey-stick divergence is at most δ', summed over $\pm 40\sigma$. The last column is the exact divergence at the tight bound, in both directions. The Rényi bound of Lemma 2.4 was also checked exactly ($\sigma \in \{3,10,50\}$, $a \in \{1.5,2,8\}$). Script `scripts/kappa_numbers.py`.

κ	B	σ	δ'	ε' tight	ε' simple	exact ε	E at tight bound
6	50	1276.6	10^{-6}	0.15886	0.20664	0.14574	$2.36 \cdot 10^{-7}$
6	50	1276.6	10^{-9}	0.21364	0.25291	0.20188	$1.85 \cdot 10^{-10}$
8	40	4085.2	10^{-6}	0.036278	0.051517	0.032917	$2.53 \cdot 10^{-7}$
8	40	4085.2	10^{-9}	0.050799	0.063085	0.047807	$1.93 \cdot 10^{-10}$
10	30	12255.5	10^{-6}	0.0081637	0.012870	0.0072998	$2.74 \cdot 10^{-7}$
10	30	12255.5	10^{-9}	0.012044	0.015762	0.011281	$2.02 \cdot 10^{-10}$

5.2 Composition with an (ε',δ') smudging step

For an admitted run k, let $\delta'_k \in (0,1)$ be a public constant (a protocol constant such as 2^{-100}, or part of P_k), and put

$$\varepsilon'_k = n_{\mathrm{rel},k}\,\varepsilon'(\kappa_k,\delta'_k), \qquad \hat{\gamma}_k = n_{\mathrm{rel},k}(\delta_{B,k} + \delta'_k).$$

Theorem 5.2 (Composition with the hockey-stick slack)**.** *Assume Assumption 1.9; the floor 2^{42} in (A5) is not needed here. Let $\mathcal{A}$ be a PPT coalition, and suppose that in every execution the admitted runs satisfy $\sum_k \rho_k \le \rho^*$, $\sum_k \varepsilon'_k \le E^*$ and $\sum_k \hat{\gamma}_k \le \Gamma^*$. Then for every $\varepsilon \ge 0$ and $\mathfrak{D} \sim_{\mathrm{b}} \mathfrak{D}'$,*

$$\Pr[\mathsf{Real}_{\mathcal{A}}(\mathfrak{D}) = 1] \le e^{\varepsilon+2E^*}\Pr[\mathsf{Real}_{\mathcal{A}}(\mathfrak{D}') = 1] + e^{E^*}\delta(\rho^*,\varepsilon) + \big(1 + e^{\varepsilon+E^*}\big)\Gamma^* + \mathrm{negl}(\lambda).$$

Proof. Step 1′, replacing Step 1. Apply Lemma 3.6 twice: once to $(P,Q) = (\mathsf{Real},\mathsf{Hyb}^{\mathrm{st}})$ and once to $(\mathsf{Hyb}^{\mathrm{st}},\mathsf{Real})$, both under $\mathfrak{D}$. At a release step (k,s) with history h, take $\varepsilon_t(h) = \varepsilon'(\kappa_k,\delta'_k)$ and $\delta_t(h) = \delta'_k + \mathbf{1}\{\|e_{k,s}(h)\| > B_{\mathrm{tot},k,s}\}$; at every other step take $(0,0)$. The two kernels at a release step differ by the shift $e_{k,s}(h)$ of $f + \eta - c'_1 s_R$. By Lemma 2.2 and Lemma 5.1(a), their hockey-stick divergence at ε' is at most δ'_k in both directions when $\|e_{k,s}(h)\| \le B_{\mathrm{tot},k,s}$, and at most 1 trivially otherwise.

In every execution $\sum_t \varepsilon_t \le \sum_{\mathrm{adm}} \varepsilon'_k \le E^*$. As in Step 1, $\mathbb{E}_P[\sum_t \delta_t] \le \mathbb{E}[\sum_{\mathrm{adm}} \hat{\gamma}_k] \le \Gamma^*$. This holds for $P = \mathsf{Real}$, and also for $P = \mathsf{Hyb}^{\mathrm{st}}$, because there the coins of each run are drawn as in Real and (1) holds unchanged. Hence for every test φ, in particular $\mathcal{A}$'s output,

$$\mathbb{E}_{\mathsf{Real}(\mathfrak{D})}\varphi \le e^{E^*}\mathbb{E}_{\mathsf{Hyb}^{\mathrm{st}}(\mathfrak{D})}\varphi + \Gamma^*, \qquad \mathbb{E}_{\mathsf{Hyb}^{\mathrm{st}}(\mathfrak{D})}\varphi \le e^{E^*}\mathbb{E}_{\mathsf{Real}(\mathfrak{D})}\varphi + \Gamma^*,$$

and likewise for $\mathfrak{D}'$.

Steps 2–4 are unchanged: $|\Pr[\mathsf{Hyb}^{\mathrm{st}}(\mathfrak{D}) = 1] - \Pr[\mathsf{Ideal}(\mathfrak{D}) = 1]| \le \nu$, and the ideal experiment is ρ^*-zCDP.

Chain. Write $p_X(\mathfrak{D}) = \Pr[X(\mathfrak{D}) = 1]$. Then

$$\begin{aligned} p_{\mathsf{Real}}(\mathfrak{D}) &\le e^{E^*} p_{\mathsf{Hyb}^{\mathrm{st}}}(\mathfrak{D}) + \Gamma^* \le e^{E^*}\big(p_{\mathsf{Ideal}}(\mathfrak{D}) + \nu\big) + \Gamma^* \\ &\le e^{E^*}\big(e^{\varepsilon} p_{\mathsf{Ideal}}(\mathfrak{D}') + \delta(\rho^*, \varepsilon) + \nu\big) + \Gamma^* \\ &\le e^{E^*}\big(e^{\varepsilon}(p_{\mathsf{Hyb}^{\mathrm{st}}}(\mathfrak{D}') + \nu) + \delta(\rho^*, \varepsilon) + \nu\big) + \Gamma^* \\ &\le e^{E^*}\big(e^{\varepsilon}(e^{E^*} p_{\mathsf{Real}}(\mathfrak{D}') + \Gamma^* + \nu) + \delta(\rho^*, \varepsilon) + \nu\big) + \Gamma^*, \end{aligned}$$

which is the claim, with $\mathrm{negl} = (e^{E^*} + e^{\varepsilon + E^*})\nu$. □

How computational and (ε', δ') steps are mixed. The hockey-stick step is information-theoretic and holds for every $[0,1]$-valued test, so in particular for the event that a PPT distinguisher outputs 1. It is applied only between Real and $\mathsf{Hyb}^{\mathrm{st}}$, before any computational replacement, where $e_{k,s}$ is defined and Lemma 5.1 applies. The computational steps are applied between $\mathsf{Hyb}^{\mathrm{st}}$ and Ideal, and they are additive. In the chain the multiplicative factors e^{E^*} and e^{ε} multiply the negligible ν, which stays negligible because ε and E^* are constants. No hockey-stick bound is ever applied across a computational step. That would need a separate argument, and it is not needed here.

5.3 A ledger with a hockey-stick slack, and the new κ_L

Definition 5.3 (Ledger with an ε-share for smudging)**.** The parameters are $(\varepsilon_L, \delta_L, \zeta_L)$, a share $\xi_L \in (0,1)$ of ε_L reserved for smudging, and a per-release δ', a protocol constant such as 2^{-100}. Put

$$E'_L = \frac{\xi_L \varepsilon_L}{2}, \qquad \rho_L = \max\big\{\rho : \delta\big(\rho, (1-\xi_L)\varepsilon_L\big) \le e^{-E'_L}(1-\zeta_L)\delta_L\big\}, \qquad \Gamma'_L = \frac{\zeta_L \delta_L}{1 + e^{\varepsilon_L - E'_L}}.$$

The state is $(\rho_\Sigma, E_\Sigma, \Gamma_\Sigma)$. Rule (F′) admits a run iff $\rho_\Sigma + \rho_k \le \rho_L$, $E_\Sigma + \varepsilon'_k \le E'_L$ and $\Gamma_\Sigma + \hat{\gamma}_k \le \Gamma'_L$, and charges all three in full at admission.

Corollary 5.4. *Under rule (F′) the protocol is $(\varepsilon_L, \delta_L)$-IND-CDP for $\sim_{\mathrm{b}}$, against every PPT coalition.*

Proof. Apply Theorem 5.2 with $\varepsilon = (1-\xi_L)\varepsilon_L$, $(\rho^*, E^*, \Gamma^*) = (\rho_L, E'_L, \Gamma'_L)$. Then:

- $\varepsilon + 2E'_L = \varepsilon_L$;
- $e^{E'_L}\delta(\rho_L, \varepsilon) \le (1-\zeta_L)\delta_L$;
- $\varepsilon + E'_L = \varepsilon_L - E'_L$, so $(1 + e^{\varepsilon + E'_L})\Gamma'_L = \zeta_L \delta_L$.

□

For runs fixed in advance, the analogue of rule (N) admits a run iff

$$e^{E_\Sigma}\delta(\rho_\Sigma, \varepsilon_L - 2E_\Sigma) + (1 + e^{\varepsilon_L - E_\Sigma})\Gamma_\Sigma \le \delta_L,$$

with the sums taken after adding the run. Correctness follows from Theorem 5.2 as in Theorem 4.2(e).

Proposition 5.5 (The new κ_L rule)**.** *Let $\tilde{\rho}_{\mathrm{run}}$ be as in* [P]*, equation* `eq:kappa`*: the budget at $\delta_{\mathrm{DP}} = (1-\zeta)\delta_{\mathrm{run}}$, a lower bound on ρ_{run} once $\kappa \ge \kappa_{\mathrm{run}}$. Let ρ_L and E'_L be as in Definition 5.3. Define*

$$\kappa_L = \min\Big\{\kappa \in \mathbb{Z} :\ n_{\mathrm{rel}}\, \varepsilon'(\kappa, \delta') \le \frac{\tilde{\rho}_{\mathrm{run}}}{\rho_L} E'_L\Big\}.$$

Suppose every run on the register uses $\kappa_k \ge \kappa_L$. Then the condition $E_\Sigma + \varepsilon'_k \le E'_L$ holds whenever $\rho_\Sigma + \rho_k \le \rho_L$. The slack condition of rule (F′) holds whenever the zCDP condition does, provided δ_B satisfies the admissibility constraint $n_{\mathrm{rel}}(\delta_B + \delta') \le (\tilde{\rho}_{\mathrm{run}}/\rho_L)\Gamma'_L$, which involves δ_B alone and not κ.

Proof. $\varepsilon'_k \le (\tilde\rho_{\mathrm{run},k}/\rho_L)E'_L \le (\rho_k/\rho_L)E'_L$, so $\sum_k \varepsilon'_k \le (\sum_k \rho_k/\rho_L)E'_L \le E'_L$. The same argument gives the bound on $\sum_k \hat\gamma_k$. □

Table 2: The ledger's flooding requirement κ_L at $(\varepsilon_L, \delta_L, \zeta_L) = (10, 10^{-6}, 0.1)$, with $n_{\mathrm{rel}} = 2$, $\delta_B = 2^{-50}$, $\delta_{\mathrm{run}} = 10^{-9}$ and $\zeta = 0.1$. "SD" is the paper's rule: the smallest κ with $\gamma \le (\tilde\rho_{\mathrm{run}}/\rho_L)\Gamma_L$. The hockey-stick columns use Proposition 5.5 with $\delta' = 2^{-100}$ and the tight ε'; the simple closed form gives the same values except 23 instead of 22 at $\xi_L = 10^{-3}$, $\varepsilon_{\mathrm{run}} = 0.5$, and 14 instead of 13 at $\xi_L = 10^{-2}$, $\varepsilon_{\mathrm{run}} = 4$. "Loss in ρ_L" is the relative decrease of ρ_L from reserving ξ_L. The δ_B admissibility constraint of Proposition 5.5 holds in every cell, with a margin of 6.6× at $\varepsilon_{\mathrm{run}} = 0.5$ up to 1300× at $\varepsilon_{\mathrm{run}} = 8$.

		κ_L for $\varepsilon_{\mathrm{run}} =$				
rule	loss in ρ_L	0.5	1	2	4	8
SD (paper), $\rho_L = 1.52971$	–	48	46	44	42	40
hockey-stick, $\xi_L = 10^{-2}$ ($E'_L = 0.05$)	1.98%	19	17	15	13	12
hockey-stick, $\xi_L = 10^{-3}$ ($E'_L = 0.005$)	0.198%	22	21	19	17	15
hockey-stick, $\xi_L = 10^{-4}$ ($E'_L = 0.0005$)	0.020%	26	24	22	20	18

Effect on κ. With $\kappa = \max\{\kappa_{\min}, \kappa_{\mathrm{run}}, \kappa_L\}$, $\kappa_{\min} = 40$ and the paper's $\kappa_{\mathrm{run}} = 36, 37, 38, 41, 46$, the ledger cap $(10, 10^{-6})$ now gives $\kappa = 40, 40, 40, 41, 46$, compared with $48, 46, 44, 42, 46$ under the statistical-distance rule. This is the κ of a single-run ledger, so the right-hand columns of [P], Table `tab:kappa`, become equal to its single-run columns. The run budget is unchanged, so $\log_2 \Omega_f$ at $\bar n_O = 10^6$ becomes 61.78, 62.75, 63.69, 65.66 and 71.94, compared with 69.88, 68.78, 67.72, 66.66 and 71.94. That saves 8.1, 6.0, 4.0, 1.0 and 0 bits of scale and of q_0. The cost is a relative decrease of 0.2% in ρ_L at $\xi_L = 10^{-3}$, and one more ledger counter.

Remark 5.6 (The per-run statement)**.** The same lemma applies to a single run ([CP], Theorem 4.11), with the share ξ taken from $\varepsilon_{\mathrm{run}}$. At $\xi = 10^{-3}$ and $\delta' = 2^{-100}$ the per-run requirement becomes $\kappa_{\mathrm{run}} = 18, 17, 16, 15, 14$, and the slack is at most 0.53% of δ_{run}. Then $\kappa = \kappa_{\min} = 40$ for every row. This changes Theorem `thm:output`'s statement, so it is not part of the proposed paper text.

Remark 5.7 (A Rényi-aggregated variant)**.** One can also account the smudging steps in Rényi divergence. Truncate the error to the event $\|e_{k,s}\| \le B$, which costs $\sum_{\mathrm{adm}} n_{\mathrm{rel}}\delta_B$ in statistical distance. The truncated real process is then Rényi-close to $\mathsf{Hyb}^{\mathrm{st}}$ with cost $\pi 4^{-\kappa_k}$ per release, and these costs compose by Theorem 3.2. Converting once at (E'_L, δ'') makes the smudging ε grow like the square root of the number of releases rather than linearly. At $\xi_L = 10^{-3}$ and $\delta'' = 2^{-100}$ this gives $\kappa_L = 18, 17, 16, 15, 14$. Since $\kappa_{\min}$ binds either way, the simpler per-release form is proposed.

Remark 5.8 (δ_B is not helped)**.** δ_B remains multiplied by $1 + e^{\varepsilon_L - E'_L}$. At $\varepsilon_L = 20$, $n_{\mathrm{rel}}(1 + e^{20})\delta_B \approx 8.6 \cdot 10^{-7}$ exceeds $\zeta_L\delta_L = 10^{-7}$ for $\delta_L = 10^{-6}$, so the paper's remark that δ_B must then be lowered stands.

6 Issues found in the paper's current text

(I1) *The sketch's slack.* "Step k costs the expectation of γ_k … slack $\Gamma = \mathbb{E}\sum_k \gamma_k$." When the smudging steps are interleaved with the computational steps run by run, each expectation is taken under a different hybrid, so the sum is not a single expectation. The fix is to do all smudging first (Lemma 3.5); the conclusion stands.

(I2) *Aborts and the δ_B part.* The ledger state "of every release made", together with admission against the full cost, refunds unmade releases. The δ_B part of an admitted run must stay charged for every planned release (Step 1; Remark 4.3). The fix is to charge a run's full cost at admission. The numerical effect is tiny, but the statement as written is not covered by the proof.

(I3) *Atomicity.* The implementation checks a run's full cost at setup but charges at each release, so two concurrently admitted runs can together exceed the cap. Admission must reserve the full cost atomically.

(I4) *Adjacency of the joint statement.* G3 and Proposition `prop:compose` speak of the joint view over a base register, but [P], Definition `def:adj`, defines adjacency on one scope register with the cohort fixed. A base-register adjacency is needed (Definition 1.1). So are the facts that it induces scope adjacency (Lemma 1.3, which needs public scopes and row-wise public preprocessing) and that cohorts may be chosen adaptively.

(I5) *RLWE for adaptively chosen parameters.* The run's HE parameter set may depend on earlier outputs. The reduction needs the assumption for a finite menu of parameter sets, or an adaptive-parameter form; it loses a factor $|\Pi_\lambda|$ (Lemma 3.7). It also needs security against non-uniform adversaries, because the reduction knows the data.

(I6) *The smudging constant.* The factor $1 + 2^{-80}$ in γ_s needs $\sigma_c \geq 2^{41.1}$, not 2^{20} (Remark 2.6). With 2^{20} the factor is $1 + 2^{-37.9}$, and [CP], Corollary 2.4's last inequality fails for $\kappa < 60$. This affects [P], Lemma `lem:smudge` and Theorem `thm:output`(i), and [CP], Corollary 2.4. It has no numerical effect at any admissible setting.

(I7) *Wording of the κ_L guarantee.* "κ_L makes the ledger run out of zCDP budget before it runs out of slack, whatever the order and the sizes of the runs." This holds only if every run on the register uses $\kappa \geq \kappa_L$ computed with the same (ρ_L, Γ_L). A run with a fixed, smaller κ can exhaust the slack first. Privacy is unaffected, because the filter checks both conditions.

(I8) *Rule (N).* "Runs fixed in advance" should say that the requests (parameters and inputs) are fixed before the first run and do not depend on any output. Rule (N) is not shown to be valid for adaptive requests (Remark 4.4).

(I9) *Implementation details.* The filter test in `Ledger.admits` allows a relative excess of 10^{-12} over the caps. `rho_tight` returns the root of a bracket, which may lie slightly above the true ρ_L. A rigorous ledger should round ρ_L and Γ_L down and test without tolerance. The effect on δ is a relative 10^{-11}.

No inequality in the budget is reversed. The definitions of ρ_L and Γ_L make the δ budget add up exactly (Theorem 4.2(d)). The paper's κ_L values $48, 46, 44, 42, 40$ and κ_{run} values $36, 37, 38, 41, 46$ are reproduced exactly.

7 What is used and what is proved here

Proved here, self-contained.

- Lemmas 2.2 (post-processing, common mixtures) and 2.3 (conversion), with proofs.
- Theorem 3.2 (Rényi filter for zCDP) and Corollary 3.4.

- Lemma 3.5 (statistical hybrid lemma), Lemma 3.6 (adaptive composition of hockey-stick closeness) and Lemma 3.7 (adaptive parameter choice).
- Lemma 1.3, Theorem 4.2, Lemma 5.1, Theorem 5.2, Corollary 5.4 and Proposition 5.5.
- All numbers, in `scripts/kappa_numbers.py`, with output in `scripts/kappa_numbers.out`. The script also computes exactly a two-step adaptive composition of one-dimensional lattice Gaussians, in which the second step's σ depends on the first output. With the filter, $D_a \le a\rho^*$ holds at every tested order; a rule that overspends on some branches violates it.

Used from [CP] (proved there, not re-verified here).

- Lemma 2.1 (Gaussian sums maximal at the lattice) and Lemma 2.2 (integer shifts), restated with proof as Lemma 2.4.
- Lemma 2.3(b) (smudging), with the constant corrected in Remark 2.6.
- Lemma 3.6 (c_1 independent of secrets), Proposition 3.8 (error decomposition), Theorem 3.12 (O1) and Corollary 3.13 (sensitivity of the centre).
- The per-run hybrids H1–H4 of Theorem 4.11, which are re-run here inside the multi-run experiment.

External results.

- [4, Lemma 3.3], the bound on the smoothing parameter of $\mathbb{Z}^N$, through Lemma 2.5.
- Decision-RLWE (Assumption 4.9) and RLWE with circular security for O's key shares (Assumption 4.10), as assumptions.
- The definitions of zCDP [1] and of SIM-/IND-CDP [5].
- [6, 3, 7] are cited for attribution only. No step depends on their exact statements, which were not checked against the full texts.
- The CKS conversion [2] is reproved in Lemma 2.3.

References

[1] M. Bun and T. Steinke. Concentrated differential privacy: Simplifications, extensions, and lower bounds. In *Theory of Cryptography (TCC 2016-B)*, LNCS 9985, pages 635–658. Springer, 2016.

[2] C. L. Canonne, G. Kamath, and T. Steinke. The discrete Gaussian for differential privacy. In *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*, 2020.

[3] V. Feldman and T. Zrnic. Individual privacy accounting via a Rényi filter. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021. arXiv:2008.11193.

[4] D. Micciancio and O. Regev. Worst-case to average-case reductions based on Gaussian measures. *SIAM Journal on Computing*, 37(1):267–302, 2007.

[5] I. Mironov, O. Pandey, O. Reingold, and S. Vadhan. Computational differential privacy. In *Advances in Cryptology – CRYPTO 2009*, LNCS 5677, pages 126–142. Springer, 2009.

[6] R. Rogers, A. Roth, J. Ullman, and S. Vadhan. Privacy odometers and filters: Pay-as-you-go composition. In *Advances in Neural Information Processing Systems 29 (NIPS 2016)*, 2016.

[7] J. Whitehouse, A. Ramdas, R. Rogers, and Z. S. Wu. Fully-adaptive composition in differential privacy. In *Proceedings of ICML*, PMLR vol. 202, 2023.