


Privacy-Preserving Vertical Federated Linear Regression

When the organization holds the response: protocol,
privacy analysis and evaluation





1

The problem

Joint regression across data that cannot be pooled

Two datasets, one question, no pooling



Researcher R

- Smaller cohort with rich domain variables: health, education, surveys
- Wants to know whether these predict an administrative outcome
- Holds features X_R and identifiers



Organization O (e.g. StatCan)


- Population register: tax filings, income, benefits
- Holds the outcome y and features X_O
- Cannot release records; law and governance bar direct sharing

**6-8
weeks**

of vetting and approvals today before a researcher can link a cohort to agency records in a Research Data Centre. Neither dataset alone answers the question.

Same people, different variables

Vertically partitioned data: overlapping rows, disjoint columns




Goal

- Fit y on $[X_R | X_O]$ over the n matched individuals
- Neither party sees the other's raw data
- O must not learn which of its records are in R's cohort
- The released model must not expose individuals

What has to be computed

OLS on the matched join needs a Gram matrix that spans both parties

$$\beta = (Z^T Z)^{-1} Z^T y, \quad Z = [X_R^I \mid X_O^I]$$



Only A is purely R's. Everything that touches O's data is computed under R's encryption key and noised by O before R sees it.

Three guarantees, honest-but-curious parties



Input privacy

Each party's view can be simulated from its own input and the agreed output. Neither sees the other's raw data.



Linkage privacy

O learns nothing about which of its records match R's cohort (only public sizes).



Output privacy

What R receives is differentially private for each matched person's O-side data (x_O, y) , with R's features held fixed (replace-one).



Parties follow the protocol but analyse everything they receive. R holds all decryption keys and learns the match set by design; O controls the privacy budget.

Where this work sits

Approach	Where the data sit	Record linkage	Output privacy	Gap for our setting
Central DP regression (AdaSSP, Wang 2018)	One trusted curator holds the joined table	Not needed	DP	Assumes the pooled data exist
VFL with HE (e.g. Yang et al. 2019)	Split across parties	Assumed or separate PSI	Usually none	Model can leak individuals
Private set intersection (Chen-Laine-Rindal)	Split across parties	Yes	Not applicable	Linkage only, no model
This work	Split; never pooled	Labeled PSI, private to O	DP, budget controlled by O	Narrower guarantee: protects O's side; R learns membership

Same statistical engine as AdaSSP (perturb the sufficient statistics, then solve a ridge), delivered without pooling, and with R's own block kept exact.

2

What we did

Reviewed the draft, found what breaks, and rebuilt it



1. Review

Draft paper and code checked line by line

2. Run

End-to-end over sockets: plaintext and real OpenFHE

3. Attack

Three privacy reviews plus an independent critique

4. Redesign

Revised protocol, TikZ figure, rewritten paper

5. Simulate

Accuracy and bias-correction studies

6. Benchmark

AdaSSP on 29 of Wang's UCI datasets

Critical finding: the noise did not protect the ciphertext

100%

of trials in which R identified the true dataset from O's reply: matched and unmatched records alike (real OpenFHE, 20/20).

73%

the most that $\epsilon = 1$ differential privacy allows.

Why

- The draft added the noise as a plaintext scalar, which changes only c_0
- c_1 is a deterministic function of R's own ciphertexts, keys and O's data
- R, who holds the key, recomputes c_1 for each candidate dataset and compares

Fix

- Add a fresh encryption of the noise, which re-randomizes c_1
- Switch to the last modulus level, then flood the error term
- BFV linkage replies: statistical flooding; CKKS: DP-calibrated flooding

Review scorecard

Finding	Severity	Status in the revised protocol
Plaintext-scalar noise leaks through the ciphertext	Critical	Fixed: fresh encryption plus flooding
Calibration step shares raw-data statistics	Critical	Fixed: per-party commitments, public constants only
Add/remove sensitivity vs replace-one adjacency (true ϵ up to $\sim 2\times$ the claim)	High	Fixed: replace-one Δ , closed form
No privacy ledger across runs	High	Fixed: zCDP ledger at O
PSI reveals O 's row order and bin loads	High	Fixed: secret permutation, public padding
Adaptive λ cancels the bias correction	Validity	Fixed: λ set from public inputs
CKKS slot leak; HE parameter security	Checked	Not an issue with OpenFHE
R learns which of its cohort are in the register	By design	Stated in Definition 4

Contributions

1

Corrected data flow

R's Gram block A is computed locally, is never sent and never noised. The sensitivity loses its B_R^4 term.

2

Revised protocol

Sanitized replies, replace-one DP, zCDP ledger, per-party calibration, hardened PSI, ridge fixed from public inputs.

3

Theory

Finite-noise bias as a conditional expectation, a zero-budget correction, and a closed-form replace-one sensitivity.

4

Evaluation

End-to-end runs, simulation studies of accuracy and of the correction, and a benchmark against AdaSSP.



3


Methodology

The revised protocol, phase by phase

The protocol at a glance

- 0 Setup and commitments**
Only public constants cross
- 1 Private linkage**
BFV labeled PSI; R learns the match set
- 2 Encrypted aggregation**
CKKS, depth 1; DP noise added by O
- 3 Solve at R**
Fixed ridge, bias correction; post-processing

Orange in the figure marks what changed from the draft.



Phases 0-1: commitments and private linkage

0 Setup and commitments

- R: standardizes with its own constants (kept), clips rows to B_R , generates BFV and CKKS keys
- O: standardizes with public constants, clips to B_O and B_y
- O computes Δ_{rep} and σ , and refuses the run if the ledger would exceed its cap
- **Only public constants cross: no quantiles, clip rates or sample moments**

1 Labeled PSI with alignment

- O applies a secret permutation τ , so labels are random positions
- Keyed hashing; bins padded to a public degree D and count α
- R cuckoo-hashes its IDs and sends encrypted windowed powers
- O evaluates membership and label polynomials, masks, and sanitizes
- R learns I and the aligned positions; O learns only public sizes

Phase 2: encrypted aggregation

O multiplies R's encrypted vectors by its own plaintext columns (depth 1)

$$\text{Enc}(B_{jk}) = \sum_i \text{Enc}(b_i) \cdot X_{O,ij} \cdot X_{O,ik}$$

$$\text{Enc}(C_{jk}) = \sum_i \text{Enc}(\dot{X}_R, ij) \cdot X_{O,ik}$$

$$\text{Enc}(c_{O,k}) = \sum_i \text{Enc}(b_i) \cdot X_{O,ik} \cdot y_i$$

$$\text{Enc}(c_{R,j}) = \sum_i \text{Enc}(\dot{X}_R, ij) \cdot y_i$$

b: selection vector; $\dot{X}_R$: R's features in O's order (zero off the match set)



$A = \dot{X}_R^T \dot{X}_R$ stays with R

Exact, never transmitted, never noised.



One noise draw per entry

z is the same in every CKKS slot.
Independent noise per slot could be averaged away.

Each released entry: $ct \leftarrow \text{Sanitize_CKKS}(ct \boxplus \text{Enc_pk}(z))$, $z \sim \text{discrete Gaussian } N_{\mathbb{Z}}(0, \sigma^2)$ from a CSPRNG

Calibrating the noise, and keeping count

Replace-one sensitivity

The adjacency replaces one matched person's (x_O, y) . The draft used the add/remove bound, which understates Δ .

(B_R, B_O, B_y)	Δ_{add}	Δ_{rep}	Ratio
(1, 1, 1)	2.24	3.16	1.41
(3.16, 3, 3)	20.6	29.4	1.43
$(\sqrt{5}, \sqrt{5}, 5)$	30.4	31.6	1.04
Simulation population	58.2	86.2	1.48

$\Delta_{\text{add}} \leq \Delta_{\text{rep}} \leq 2\Delta_{\text{add}}$. Releasing $y^T y$ costs nothing under replace-one.



σ from the analytic Gaussian

Smallest σ meeting (ϵ, δ) for the whole stacked release (Balle-Wang). Discrete Gaussian at the CKKS resolution.



zCDP ledger at O

$\rho_{\text{run}} = \Delta^2/2\sigma^2 + \rho_{\text{err}}$. O refuses a run once the cumulative spend would pass its cap. Re-runs, retraining and inference queries all count.

Sanitizing O's replies (circuit privacy)

R holds the key, so every reply must reveal nothing beyond its decrypted value

1 Fresh encryption

$ct \boxplus \text{Enc}_{pk}(z)$ makes c_1 computationally uniform: it no longer depends on O's data.

2 Modulus switch

Drop to the last level, which scales down the data-dependent error.

3 Flood the error

BFV: statistical flooding with $N \cdot 2^{-\kappa} \leq 2^{-40}$. CKKS: DP-calibrated flooding with width $\varsigma = S_{err} / \sqrt{(2\rho_{err})}$.

Why DP-calibrated flooding for CKKS

Classical 2^{40} flooding would swamp CKKS precision. A DP-sized flood needs only a few bits above the error; its cost joins the ledger.

Open proof obligation

An explicit bound on S_{err} , how much one record can change the ciphertext error, for the exact encoding and key switching.

Phase 3: fixed ridge and zero-budget bias correction

Ridge, fixed before decryption

$$\lambda = \max\{ 0, 2\rho\sigma\sqrt{p} - n\ell \}, \quad \Psi = I$$

Release gate

$$\hat{\rho} = \lambda_{\min}(\tilde{G} + \lambda I) / (2\sigma\sqrt{p}) \geq 1$$

Private ridge estimate

$$\beta_{\lambda} = (\tilde{G} + \lambda I)^{-1} \tilde{c}$$

Bias-corrected estimate

$$\beta_{bc} = \beta_{\lambda} - \sigma^2 M(\tilde{P}_{\lambda}) \beta_{\lambda}$$

Theorem (conditional bias)

$$E[\beta \mid \mathcal{E}_{\eta}] - \beta = \sigma^2 M(P) \beta + O(\sigma^4) + O(\exp(-p(\eta\rho-1)^2/2))$$

$$M(P) = (P^2 + \text{tr}(P)P - P \text{diag}(P)) \Pi_O + (P \Pi_O P + \text{tr}(\Pi_O P) P) \Pi_R$$

The plain mean does not exist, hence the conditioning on $\mathcal{E}_{\eta}$.

Why λ must not depend on the noise

Chosen from $\tilde{G}$, λ moves with the noise. The correction then removed none of the bias (1.0-1.6× left); with fixed λ it removes 87-98%.

How we evaluate



Simulation population

Features on real-world scales (income, age, binary indicators, counts), correlated. Public standardization constants; committed 99.5% bounds; $R^2 \in \{0.3, 0.5, 0.9\}$.



Monte Carlo

n from 100 to 100k matched records; $\epsilon \in \{0.5, 1, 2, 4, 8\}$, $\delta = 10^{-5}$; 200 replicates, each with new data and a new DP release.



Metrics

RMSE against the true β (standardized); per-coefficient efficiency ratio; interval coverage; MSE ratio corrected/uncorrected.



Benchmark

29 UCI datasets from Wang's published run, with his preprocessing; AdaSSP as published and under our guarantee; 10-fold test MSE.



4

Results

What we have so far, and what is still to run

The cryptographic layer is exact

Draft protocol, run end-to-end over sockets with plaintext and real OpenFHE backends

0

misaligned rows

PSI recovered exactly the true match set in all 6 cohorts, plaintext and OpenFHE

$\approx 10^{-10}$

$\max |\beta - \beta_{\text{OLS}}|$

with the noise switched off:
encryption adds no error

13.8 s

end-to-end on OpenFHE

$n_O = 4,000$ records, $n = 400$ matched (plaintext backend: 3.2 s)



All error comes from the DP noise and the ridge it forces. The revised protocol's sanitization has not been implemented in OpenFHE yet (see the placeholder slide).

Accuracy against intersection size

Coefficient error vs intersection size, $R^2 = 0.5$

Line: median of per-run RMSE over 200 replicates (new data + new DP release each run); $\delta = 1e-5$, replace-one σ , $p^* = 2$

Fixed ridge active ($\lambda = 2\rho^*\sigma\sqrt{p} - n\ell > 0$)



Within 2× of non-private OLS at

$\epsilon = 8$ **$n \approx 3.2k$**

$\epsilon = 4$ **$n \approx 9.4k$**

$\epsilon = 2$ **$n \approx 34k$**

$\epsilon = 1$ **$n \approx 105k$**

- Small n : the ridge shrinks $\hat{\beta}$ towards 0
- Then the DP error falls like $1/n$ and merges into the OLS floor
- Run-to-run spread is about $\times 2.5-3$

Key numbers

Median RMSE of the standardized slopes against the true β (ratio to the OLS floor), $R^2 = 0.5$

n	$\varepsilon = 0.5$	$\varepsilon = 1$	$\varepsilon = 2$	$\varepsilon = 4$	$\varepsilon = 8$	OLS floor
1,000	0.231 (9.9×)	0.213 (9.1×)	0.181 (7.7×)	0.138 (5.9×)	0.078 (3.3×)	0.023
10,000	0.092 (12×)	0.044 (5.8×)	0.025 (3.4×)	0.014 (1.9×)	0.010 (1.3×)	0.0075
100,000	0.0087 (3.4×)	0.0051 (2.0×)	0.0034 (1.4×)	0.0029 (1.2×)	0.0026 (1.0×)	0.0025

n at which VFL comes within 2× of the OLS floor


R^2	$\varepsilon = 0.5$	$\varepsilon = 1$	$\varepsilon = 2$	$\varepsilon = 4$	$\varepsilon = 8$
0.3	≈240k*	75k	21k	7.1k	1.9k
0.5	≈380k*	≈105k*	34k	9.4k	3.2k
0.9	>100k	>100k	66k	30k	13k

** extrapolated beyond $n = 100k$*

Comparing coefficients on different scales

Efficiency ratio per coefficient: DP error in units of that coefficient's own OLS standard error, $R^2 = 0.5$

SE_OLS,j = empirical sd of non-private OLS over the same replicates. 1 = as good as OLS. Black line = OLS's own RMSE/SE (above 1 only where clipping bias shows).



Raw scale

Coefficients span 5.6 orders of magnitude ($\approx \$0.2$ per \$ income vs $\approx \$10k$ per binary). A pooled raw RMSE is 87% driven by two binary coefficients.


Efficiency ratio

$RMSE_DP,j / SE_OLS,j$: unit-free. A ratio r means the release is about as accurate as OLS on n/r^2 records.

Result


Once the ridge is off, DP costs about the same for every coefficient ($\pm 15\%$); O-side coefficients run $\sim 10\%$ higher.

Does the bias correction matter?




-7 to -12%

MSE vs the protocol target when the ridge is on. The correction also shrinks noise; it never raised this MSE (90 settings).



+0.3 to +4.3%


MSE vs OLS under a moderate ridge: the correction restores shrinkage that noise had offset.



≥ 130

averaged releases before the bias alone would matter; one release: at most 0.75% of MSE.

Where the error comes from when σ is large



Error anatomy

- Heavy ridge: 72–90% of the error is shrinkage, 11–31% noise variance
- Mechanism bias²: at most 0.75% anywhere

Inference

- $\lambda = 0$: 95% intervals cover 0.94–0.96; SEs within 1–3% of the truth
- Ridge on: coverage collapses. Clipping shifts the target at large n (every DP method pays this)

Benchmark against AdaSSP (Wang 2018)


Test MSE at $\varepsilon = 1$, $\delta = 10^{-6}$, 10-fold CV. Ratio to non-private in brackets; bold = best private method

Dataset	n	d	Non-private	AdaSSP (published)	AdaSSP (matched)	VFL, ours	VFL, uncorrected
housing	506	13	0.0396	0.0713 (1.80×)	0.0819 (2.07×)	0.0779 (1.97×)	0.0773 (1.95×)
wine	1,599	11	0.0201	0.0338 (1.68×)	0.0391 (1.95×)	0.0370 (1.84×)	0.0368 (1.83×)
airfoil	1,503	5	0.0533	0.0579 (1.09×)	0.0625 (1.17×)	0.0617 (1.16×)	0.0616 (1.16×)
concrete	1,030	8	0.0444	0.0657 (1.48×)	0.0774 (1.74×)	0.0751 (1.69×)	0.0745 (1.68×)
bike	17,379	17	0.0279	0.0288 (1.03×)	0.0299 (1.07×)	0.0295 (1.06×)	0.0295 (1.06×)
elevators	16,599	18	0.00868	0.0131 (1.51×)	0.0154 (1.78×)	0.0149 (1.71×)	0.0147 (1.70×)
pol	15,000	26	0.135	0.140 (1.04×)	0.146 (1.08×)	0.144 (1.06×)	0.143 (1.06×)
kin40k	40,000	8	0.0632	0.0633 (1.00×)	0.0633 (1.00×)	0.0633 (1.00×)	0.0633 (1.00×)
protein	45,730	9	0.119	0.120 (1.01×)	0.120 (1.01×)	0.120 (1.01×)	0.120 (1.01×)

Published AdaSSP protects the whole record, a stronger guarantee than ours, and uses Wang's noise constants: at $\varepsilon = 1$ its per-entry noise on X^TX (11.9) is below ours (13.4), which is still to be checked. Matched = AdaSSP's adaptive λ under our guarantee; idealized, since its $\lambda_{\min}$ step needs X^TX in the clear.

Across privacy levels

Test MSE ÷ non-private, three datasets of increasing size; bars: 95% CI of the mean over the 10 folds



Ours better / worse (of 29)	$\epsilon = 0.1$	$\epsilon = 1$	$\epsilon = 10$
vs AdaSSP (published)	14 / 8	6 / 19	1 / 21
vs AdaSSP (matched)	27 / 0	27 / 0	21 / 3
Bias-corrected vs uncorrected	14 / 8	6 / 19	3 / 19

At small ϵ the fixed ridge wins: AdaSSP spends a third of its budget on $\lambda_{\min}$. From $\epsilon \approx 0.5$ AdaSSP's adaptive ridge wins (median 4% lower MSE at $\epsilon = 1$); the gap shrinks as n grows.

Paired over folds, 95% CI; the rest show no significant difference.

Still to run



Intersection-size sweep on real data

Graphic to be done: test MSE vs matched cohort size n (log) for bike, elevators, pol, kin40k and protein, one panel per dataset, lines for $\epsilon \in \{0.5, 1, 2, 4\}$; AdaSSP and non-private for reference.

Cost of the revised protocol in OpenFHE

n_O	d	PSI (s)	Agg. (s)	Traffic (MB)	Flood cost
4,000	10	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>
32,000	10	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>
65,536	20	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>
10^6 (chunked)	20	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>	<i>tbd</i>

Also pending

- Proof of the CKKS error-sensitivity bound S_{err}
- Finite-precision accounting for the discrete Gaussian
- Standard errors for the ridge regime

Takeaways and next steps

1

The revised protocol closes the critical and high gaps under honest-but-curious parties

2

Accuracy reaches 2× of non-private OLS at $n \approx 3k-100k$, depending on ϵ

3

Against AdaSSP: ahead at small ϵ , behind at large ϵ ; the gap shrinks as n grows

4

The bias correction is a small effect (under 3% of test MSE); ridge shrinkage dominates when σ is large

Next

- Real-data intersection-size sweeps
- Check Wang's noise constants; use the joint row bound to tighten our Δ
- Implement the revised protocol in OpenFHE and measure its cost
- Prove the S_{err} bound; complete finite-precision accounting
- Standard errors; membership-private variant for sensitive registers